

# Introduction to Computer Vision

---

**Kaveh Fathian**

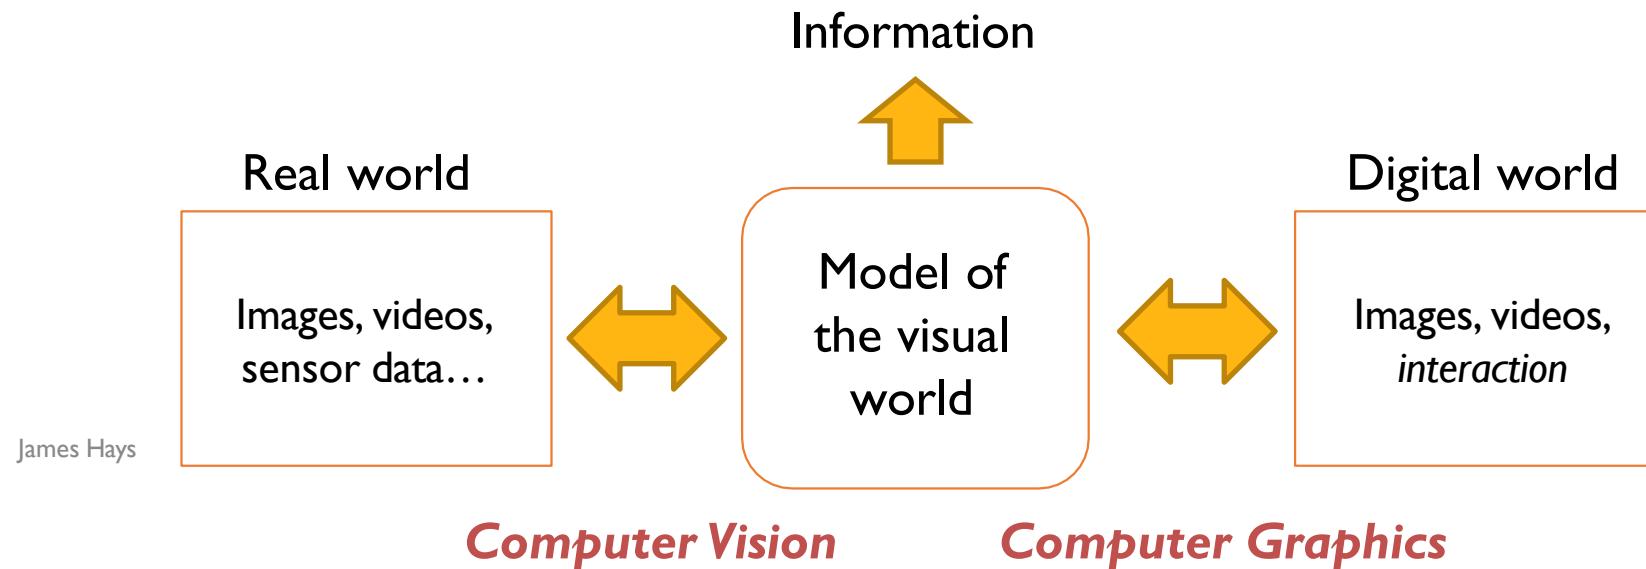
Assistant Professor

Computer Science Department

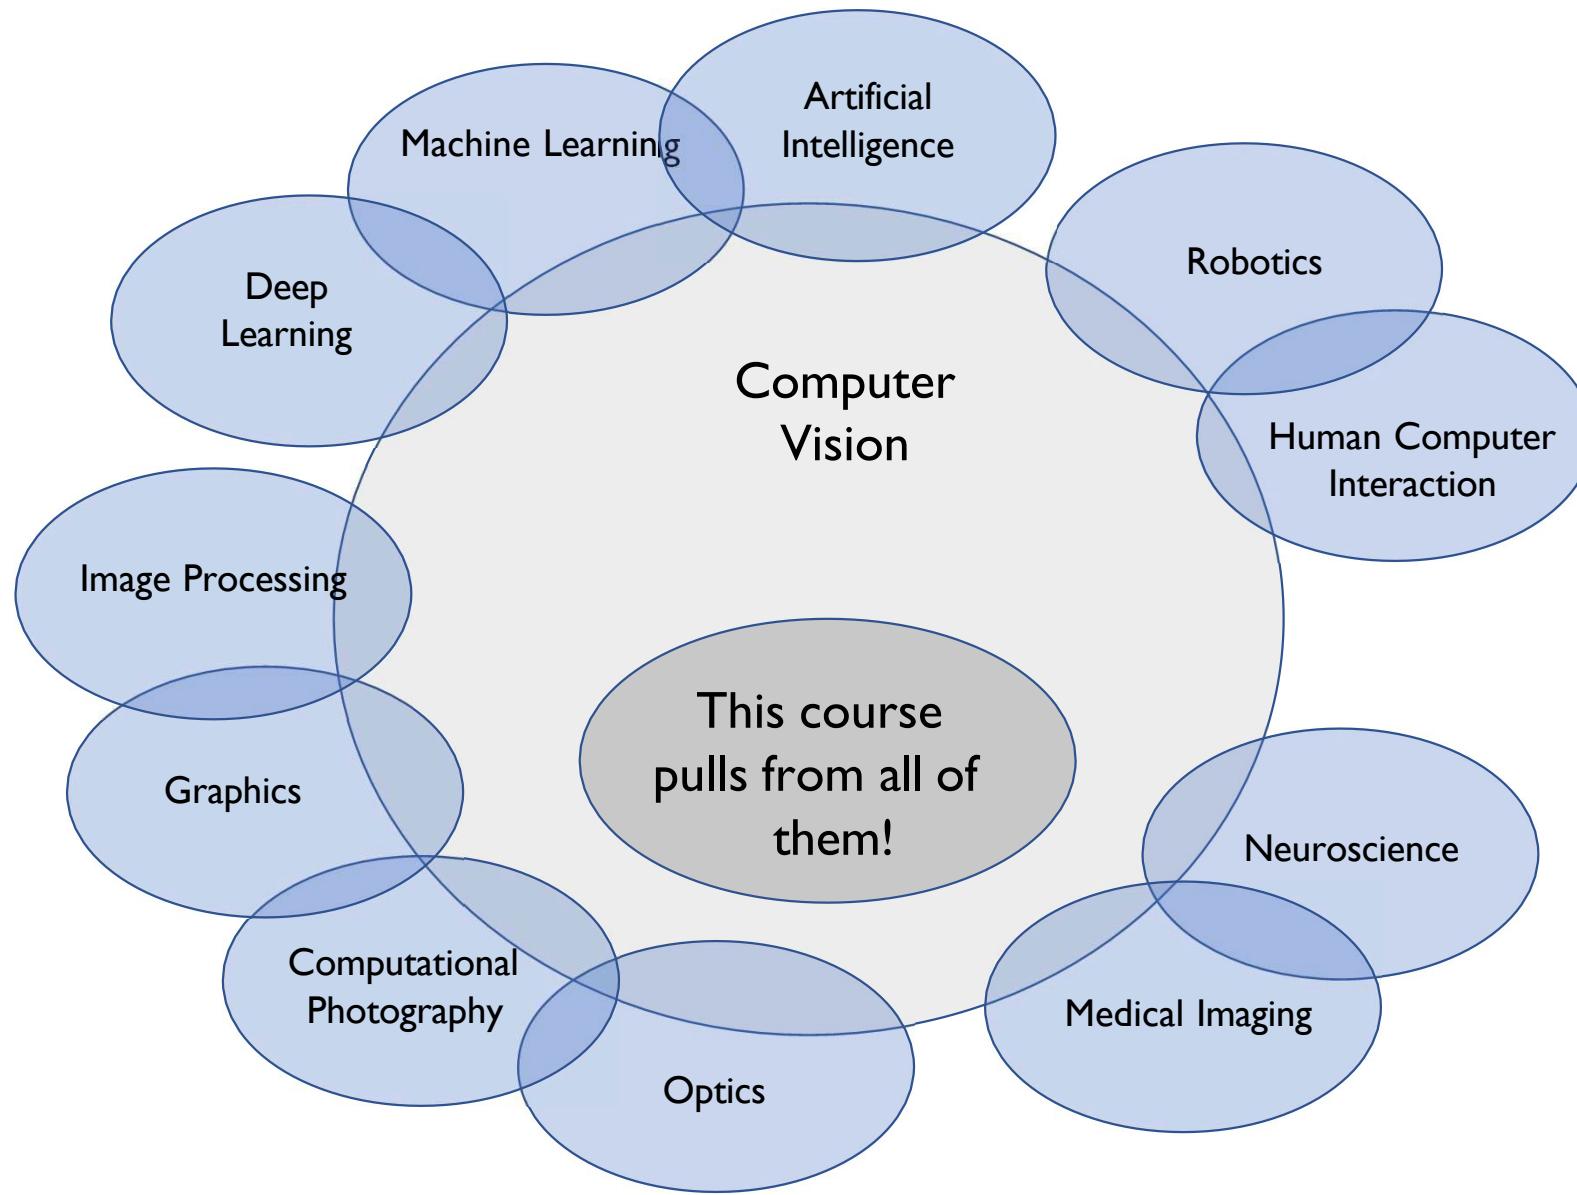
Colorado School of Mines

**Lecture 2**

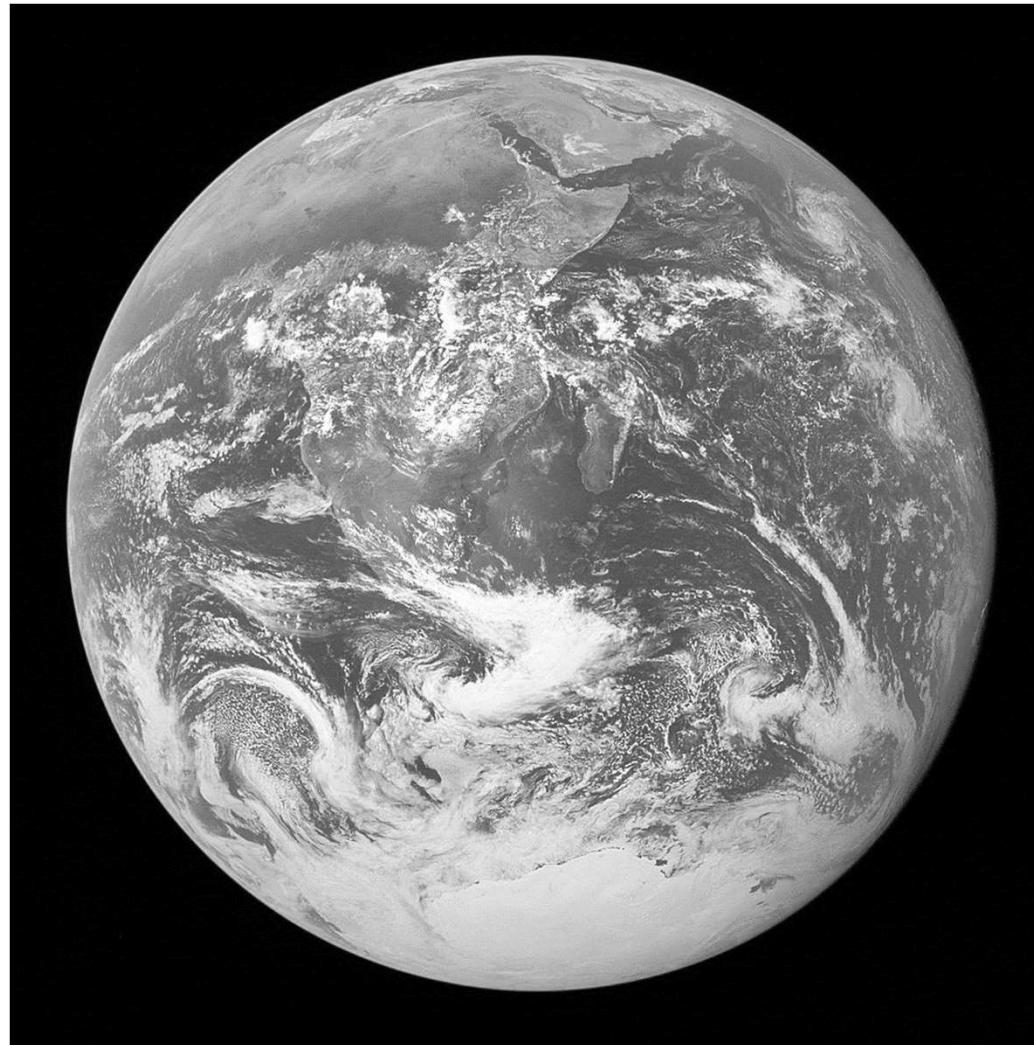
# Computer Vision and Nearby Fields



# Scope of Computer Vision

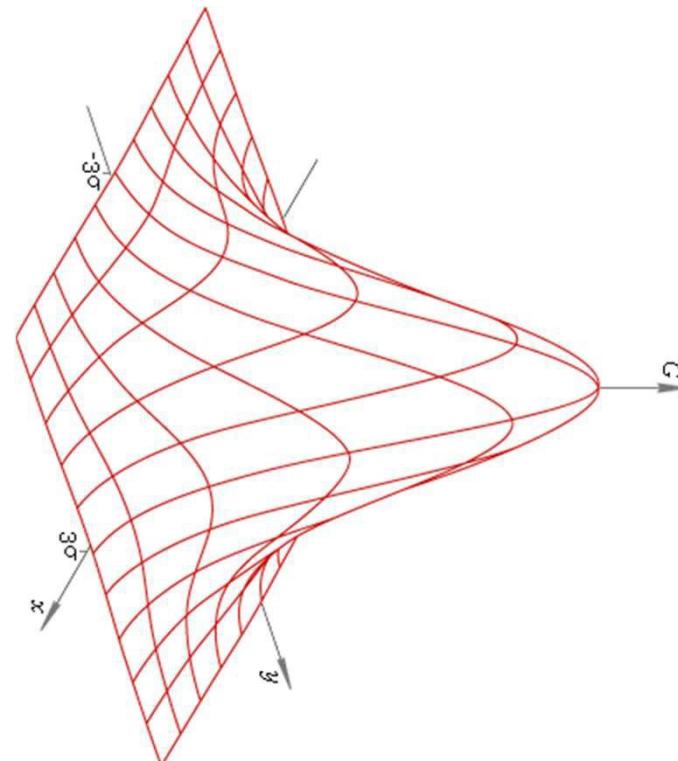
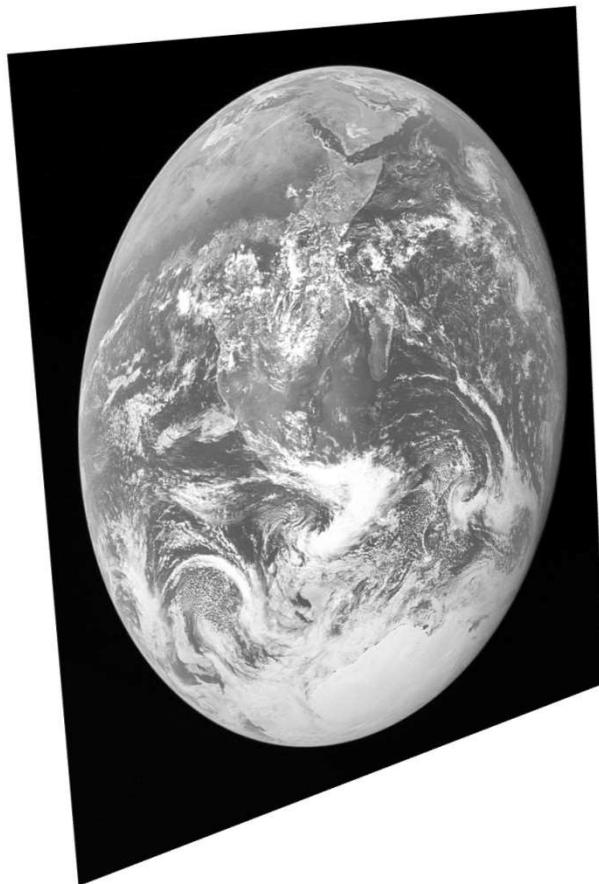


# Image



Blue Marble. NASA | Apollo 17

# Image

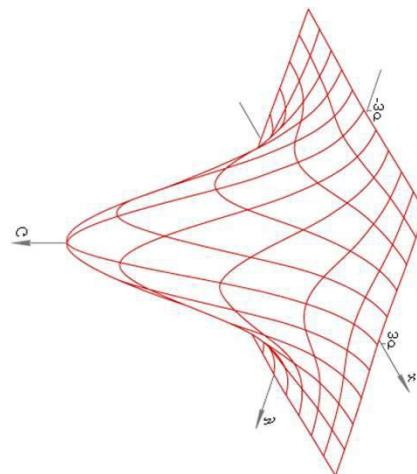
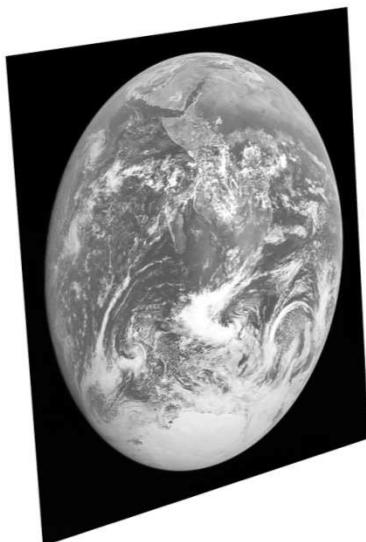


Blue Marble. NASA | Apollo 17

# Signal

## Definition

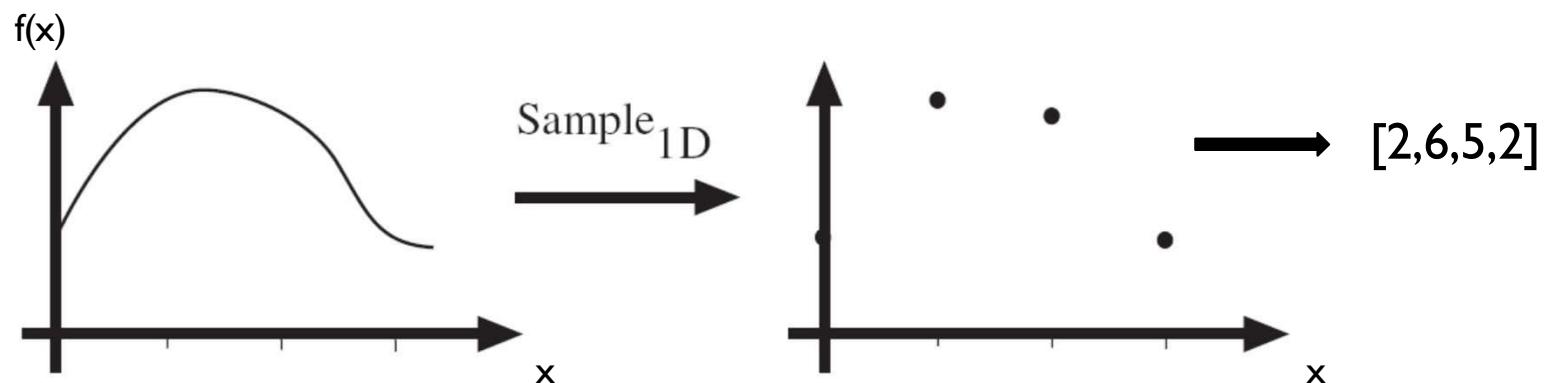
**Signal:** A (multi-dimensional) function that contains information about a phenomenon.



- Signals can be
  - Continuous: light
  - Discrete: measurement of a light
  - Sampling: reduction of continuous signal to a discrete signal
- Any phenomenon
  - Light
  - Heat
  - Gravity

# Sampling

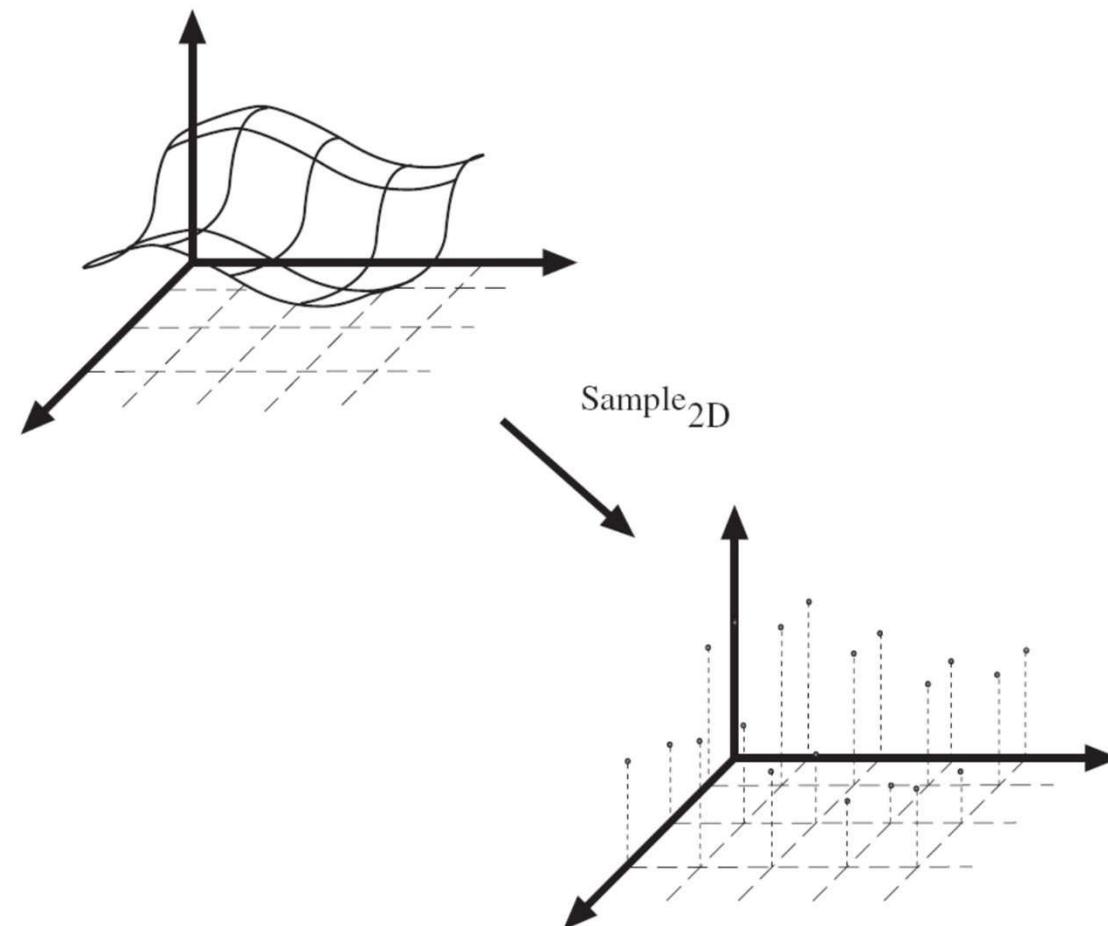
- Sampling in 1D takes a function and returns a vector whose elements are values of that function at the sample points.



Danny Alexander

# Sampling

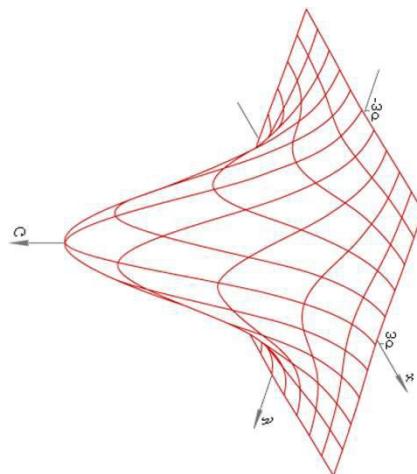
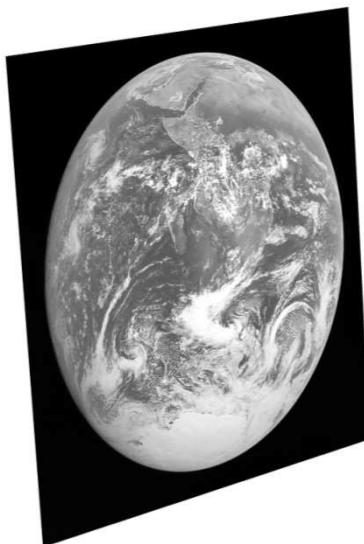
- Sampling in 2D takes a function and returns a matrix



# 2D Image

## Definition

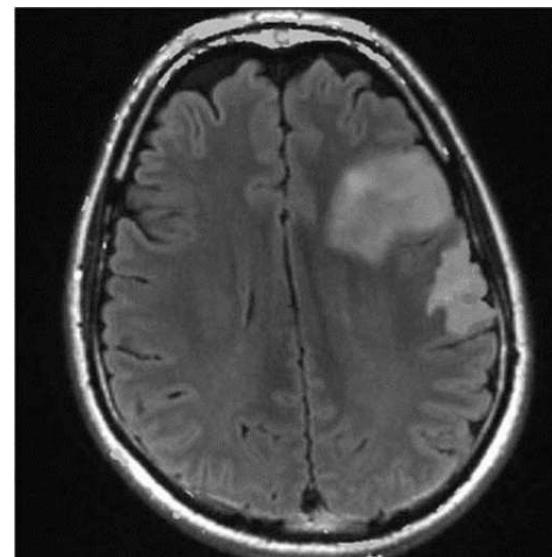
**Image:** A sampling of a function that contains information about a 2D\* signal.



\* or a 2D projection of a multi-dimensional signal

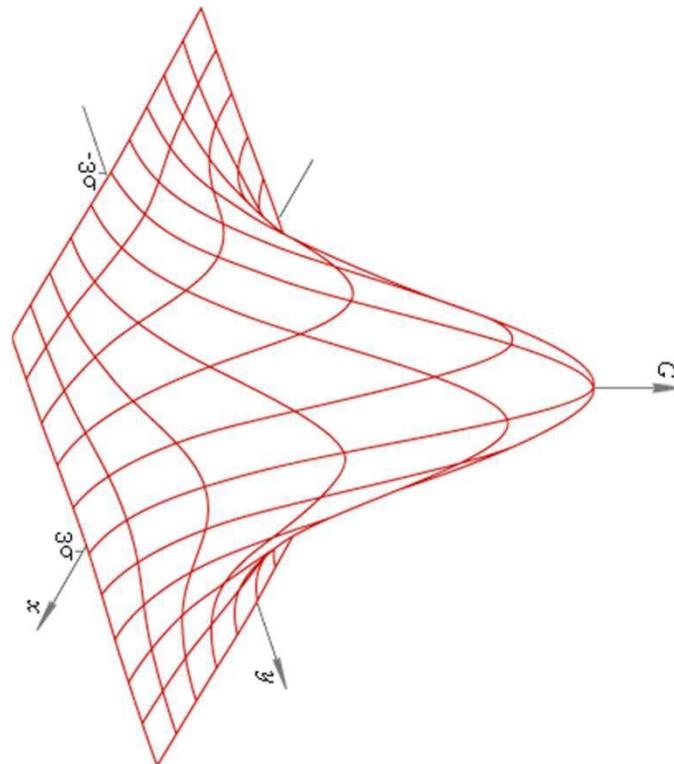
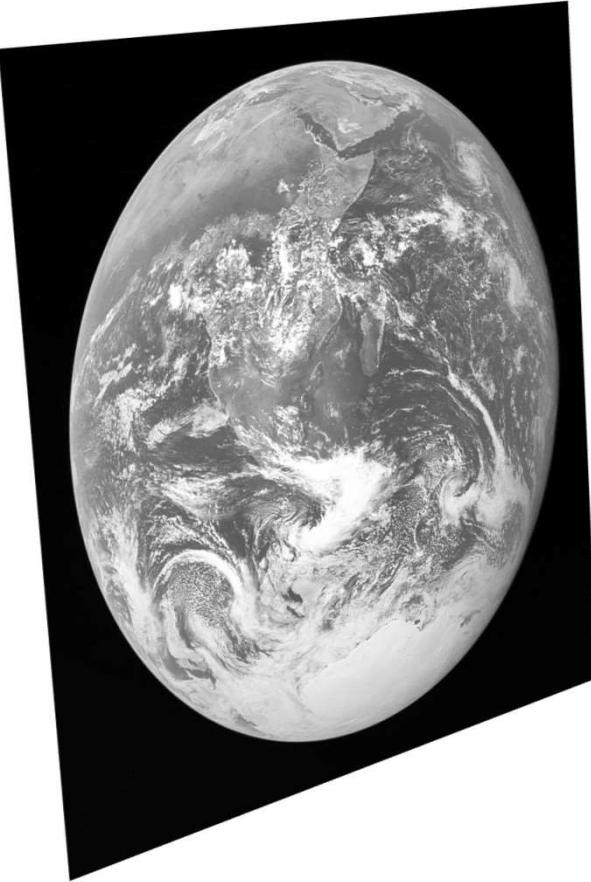
- Function stores ‘brightness’
- 2D signals are special for us
  - Brightness along x and y dimensions
- Video: xy-coordinates + time
  - Time-varying 2D signal

# Example of 2D Images

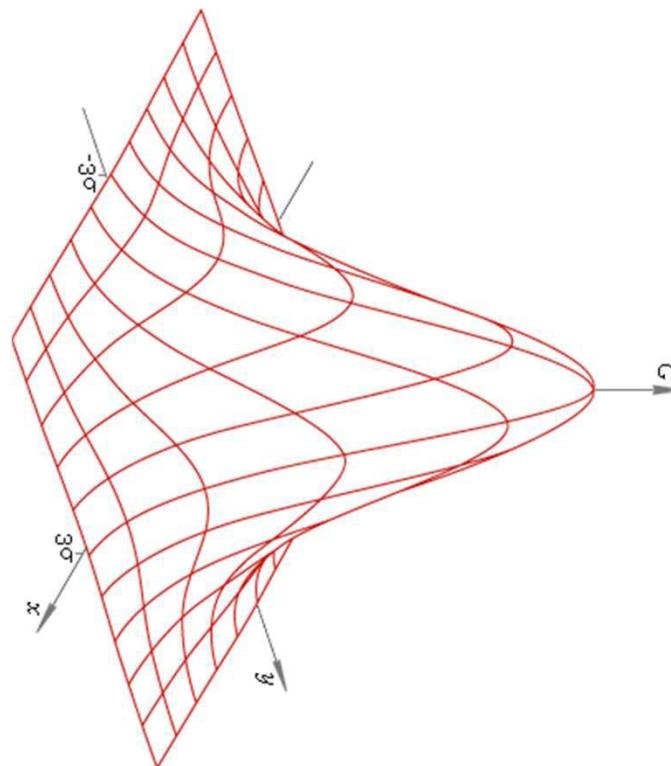
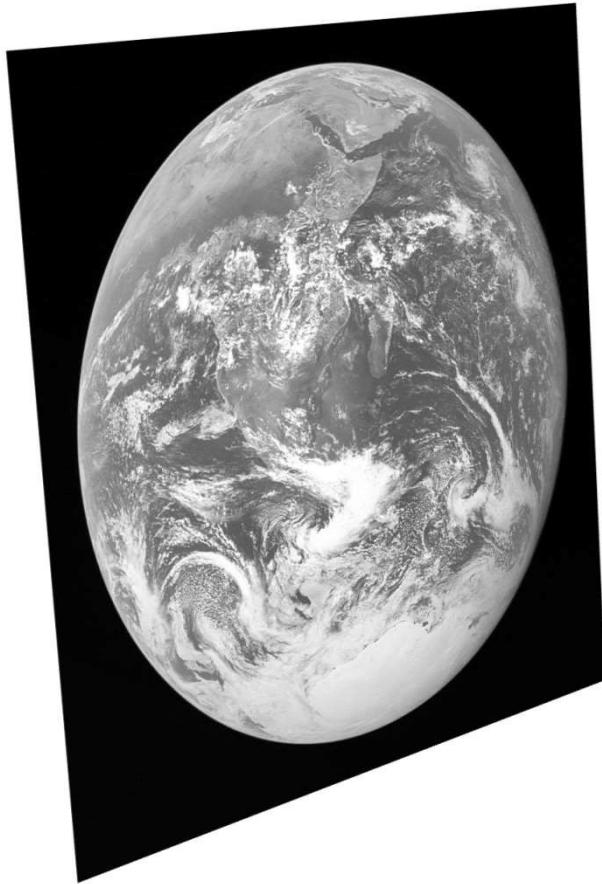


[http://radiologykey.com/wp-content/uploads/2016/01/9781604063325\\_c014\\_f004.jpg](http://radiologykey.com/wp-content/uploads/2016/01/9781604063325_c014_f004.jpg)

# Sampling in Practice

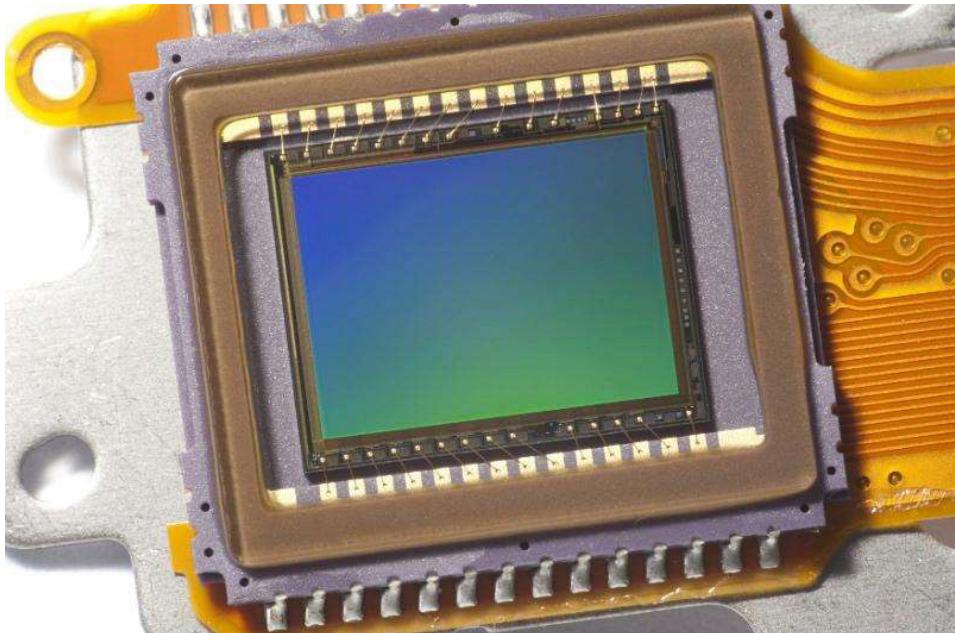


# Sampling in Practice: Digital Image

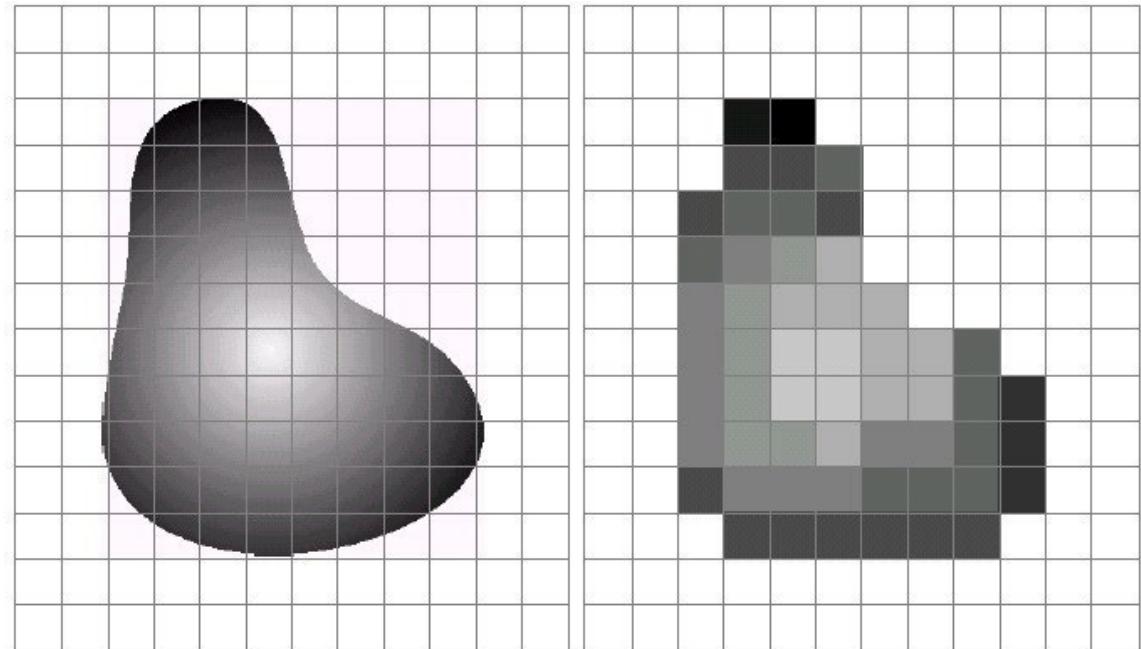


CMOS / CCD

# Sensor Array



CMOS sensor



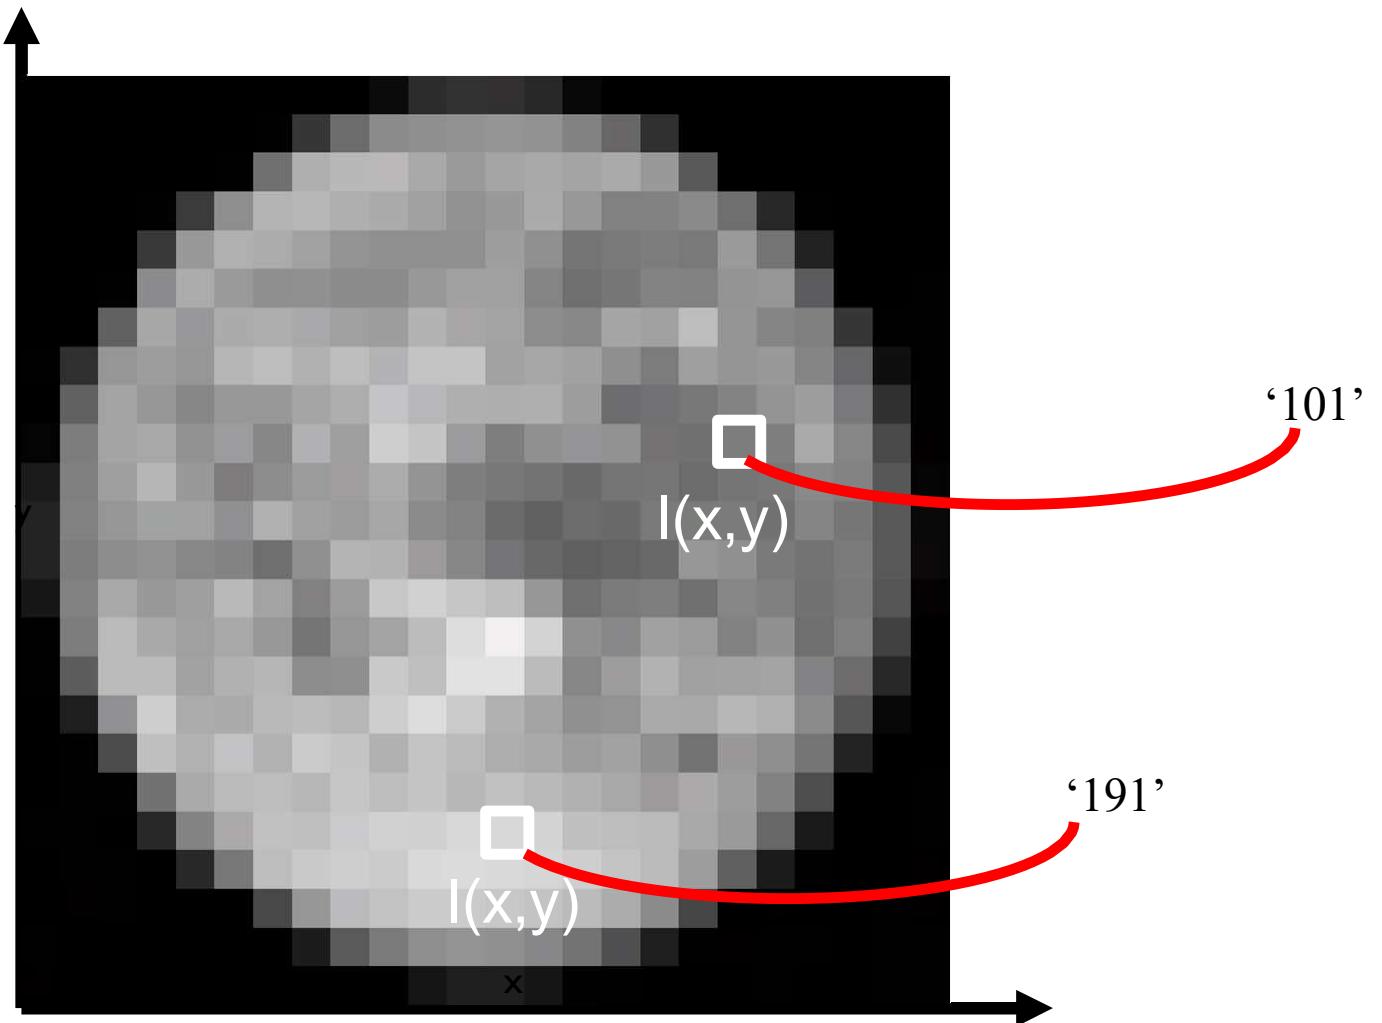
a b

**FIGURE 2.17** (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.

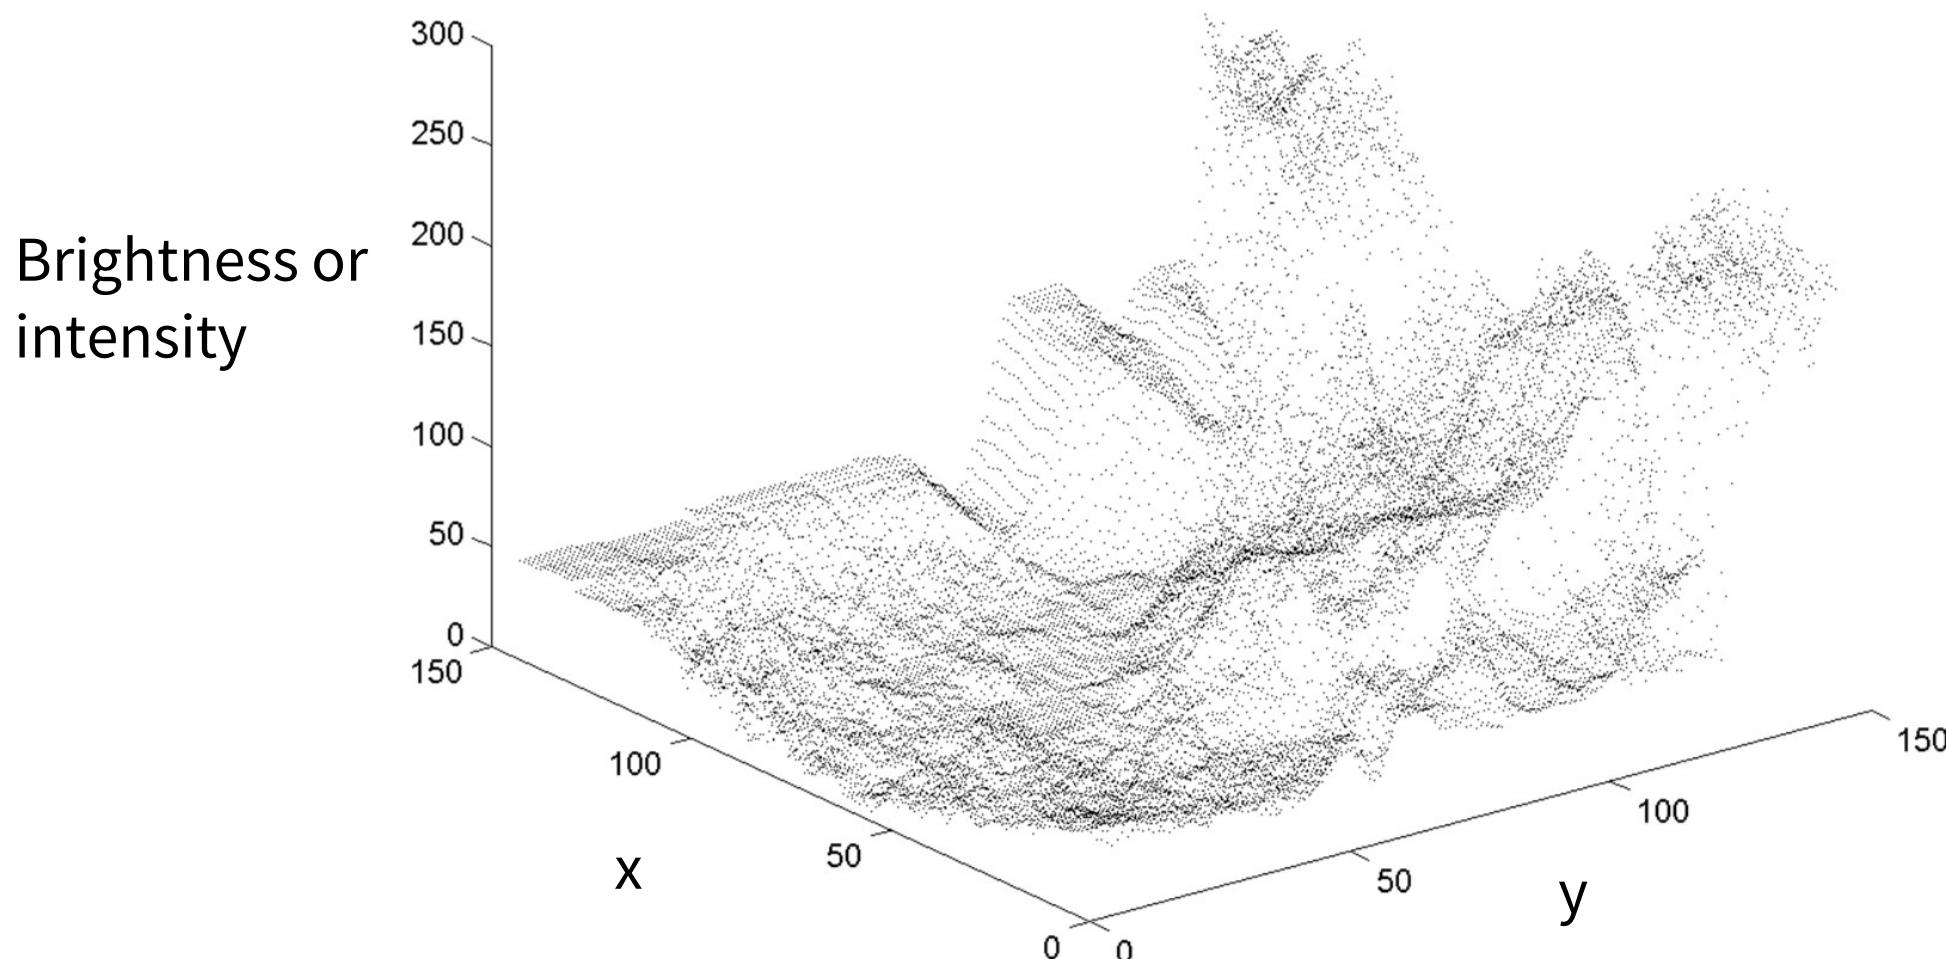
James Hays

# Elements of a Digital Image

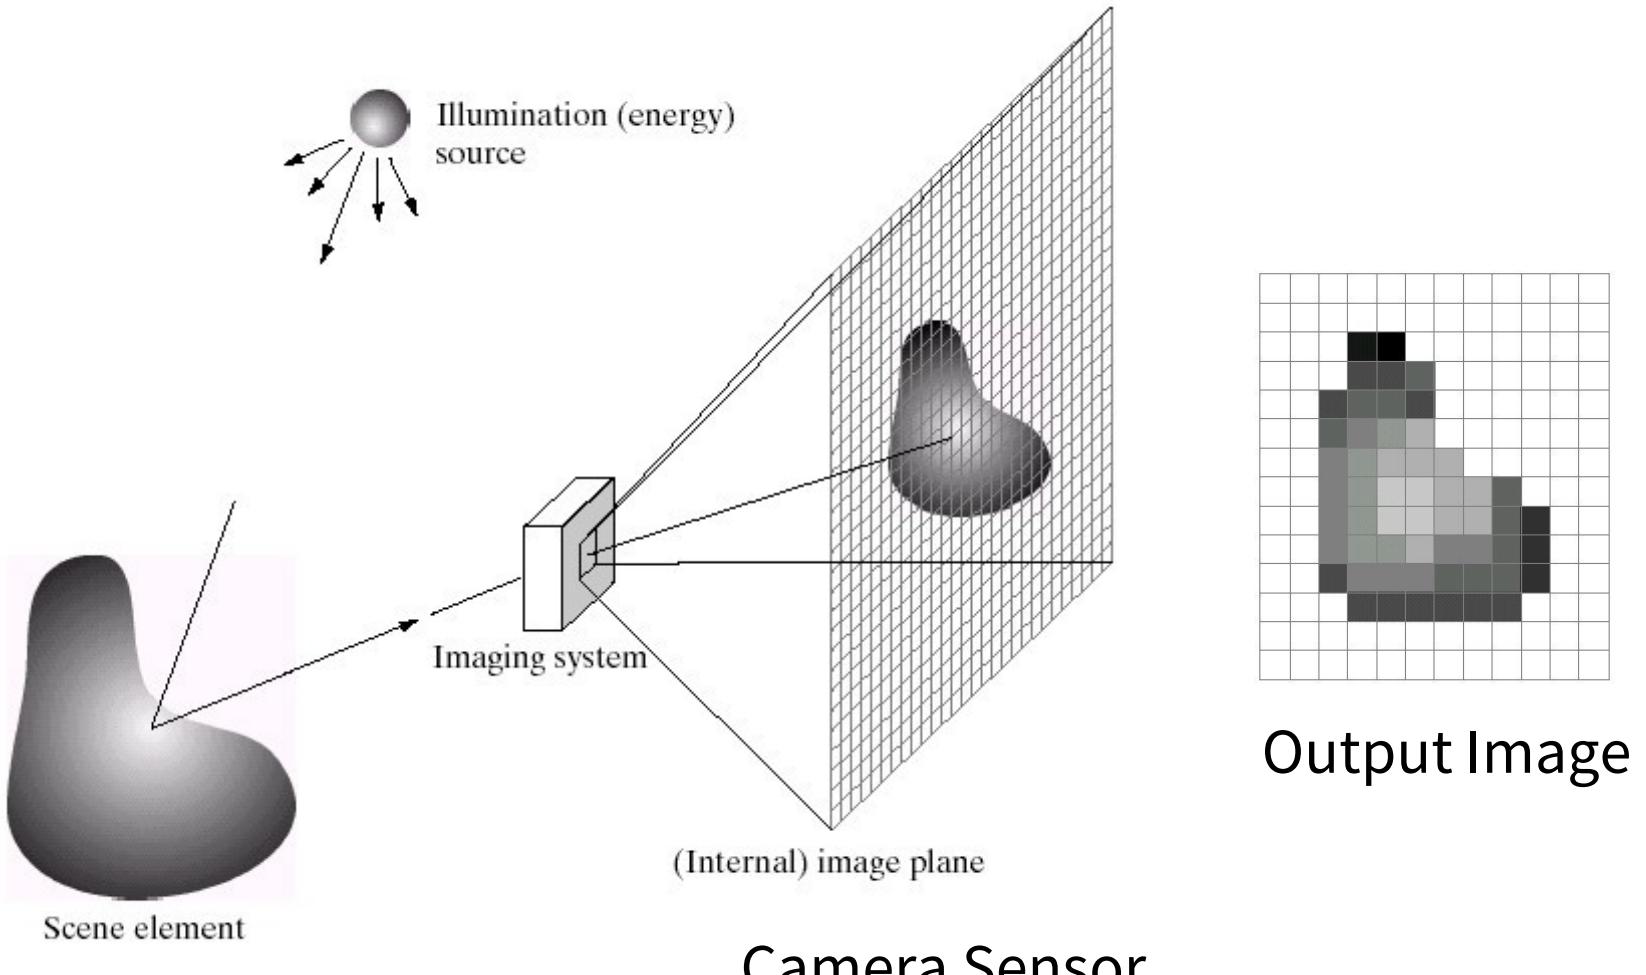
**Pixel:** picture element



# Digital Image is a 2D Signal

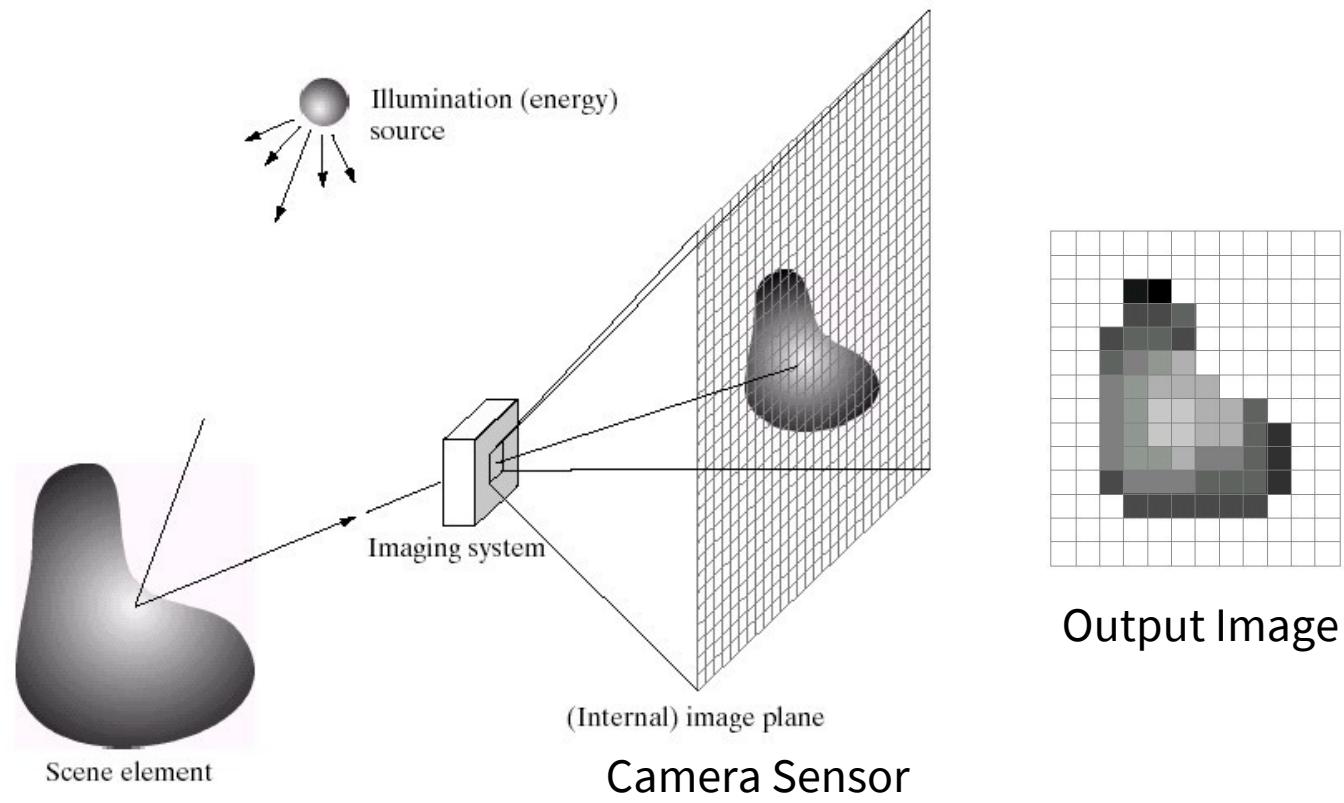


# Light Integration Over the “Frustum”

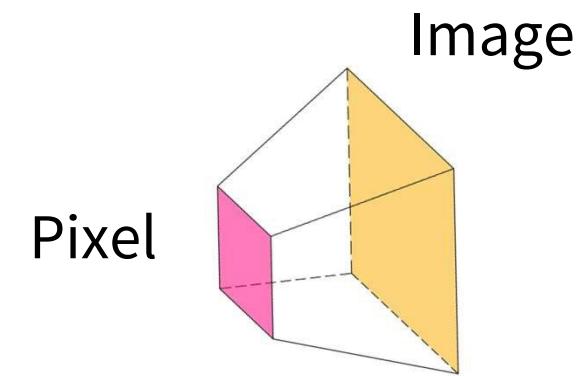


James Hays

# Light Integration Over the “Frustum”



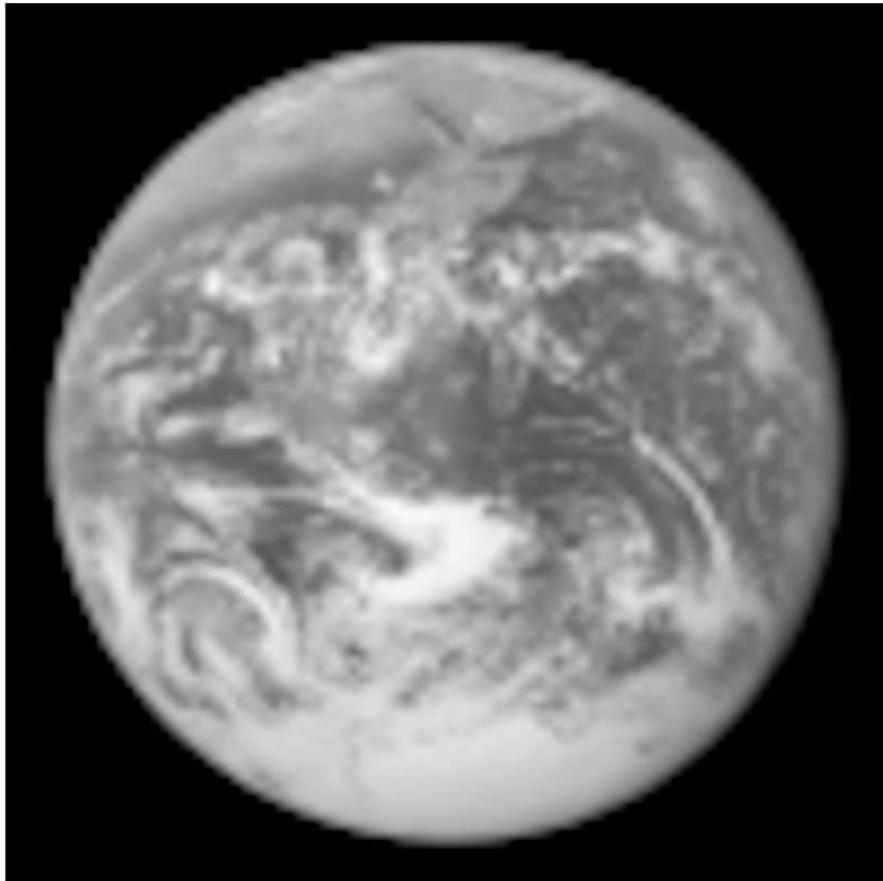
James Hays



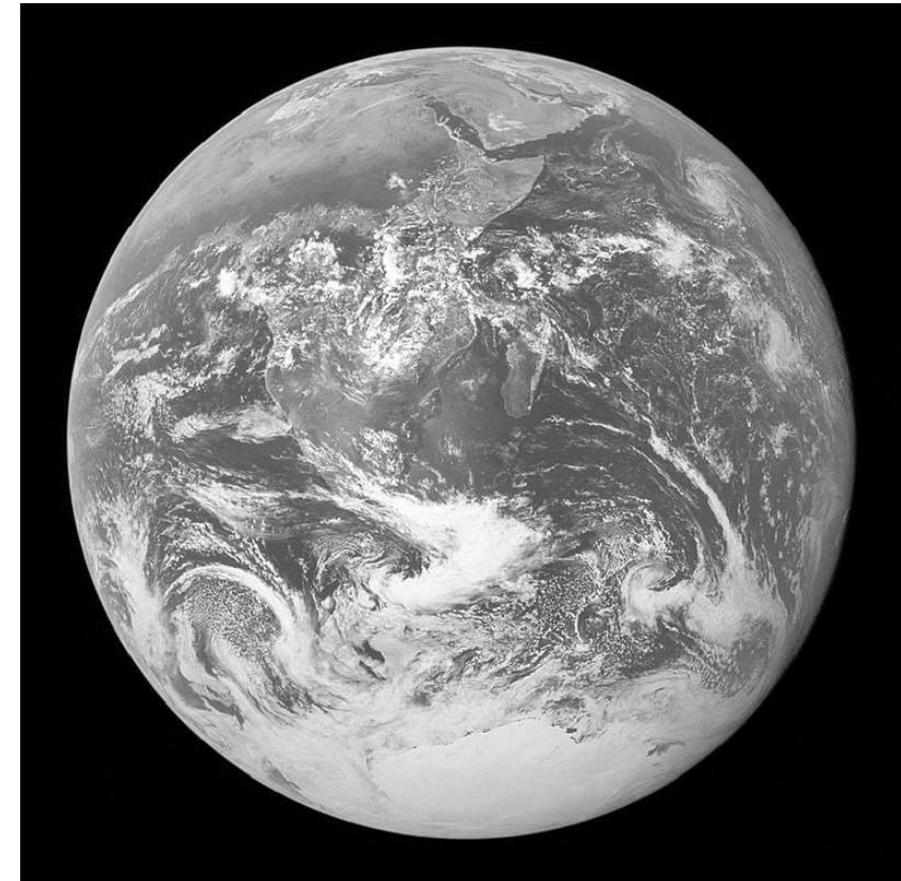
A frustum approximates the ray space that a pixel samples

# Resolution: geometric vs. spatial

Both images are 1000x1000 pixels—same spatial resolution:

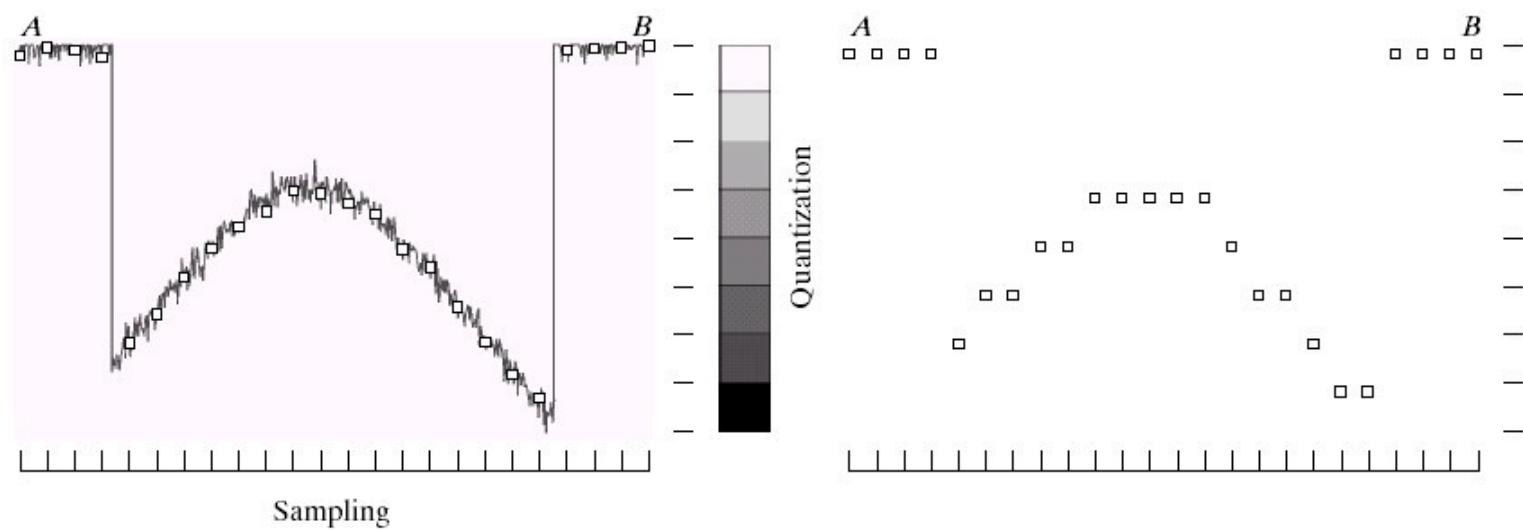
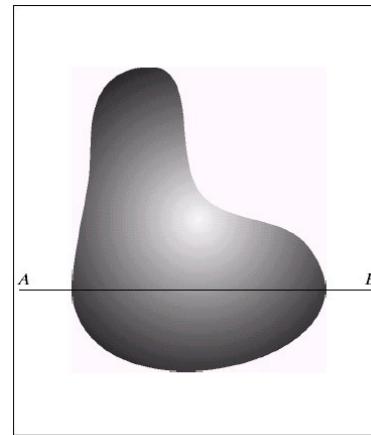


Low geometric resolution



High geometric resolution

# Quantization



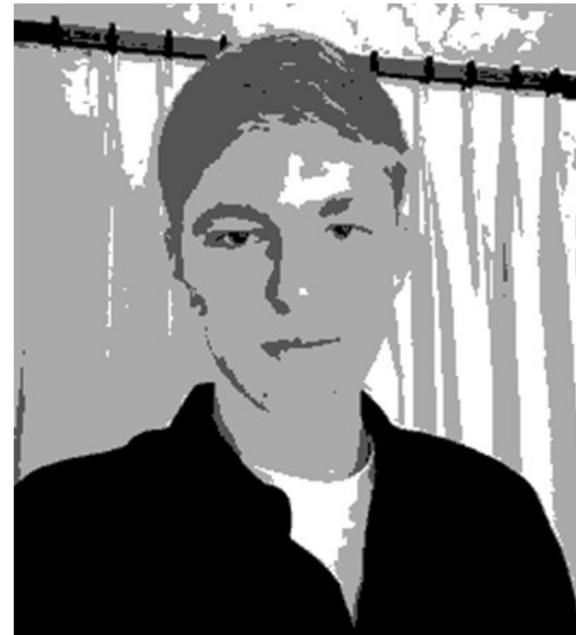
# Quantization Effects – Radiometric Resolution



8 bit – 256 levels



4 bit – 16 levels



2 bit – 4 levels



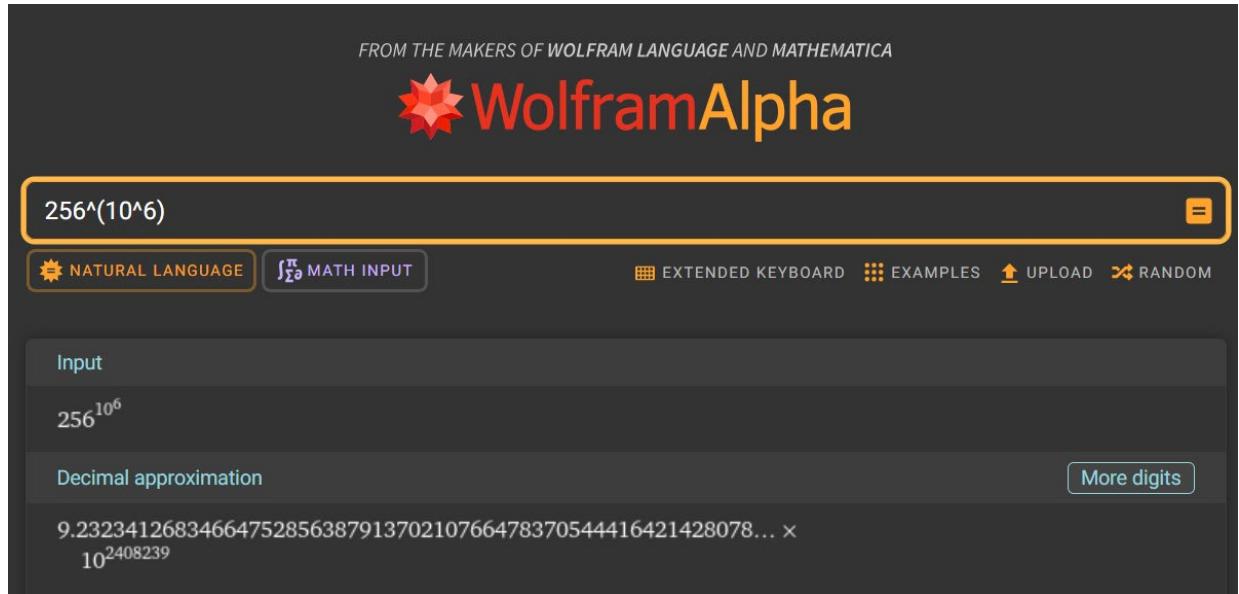
1 bit – 2 levels

We often call this ***bit depth***.

For photography, this is also related to ***dynamic range***.

# Dimensionality of an Image

- An image of size 1000x1000 with 8-bit quantization per pixel
  - $= 256 \text{ values}^{\wedge}(1000 \times 1000)$
- All images ever created
  - of size 1000x1000



\* it is estimated that there may be  $10^{78}$  to  $10^{82}$  atoms in the known universe!

- Computer vision as making sense of an extremely high-dimensional data
  - ‘natural’ images are a subspace

# Images in Python (import numpy)

$N \times M$  grayscale image “im”

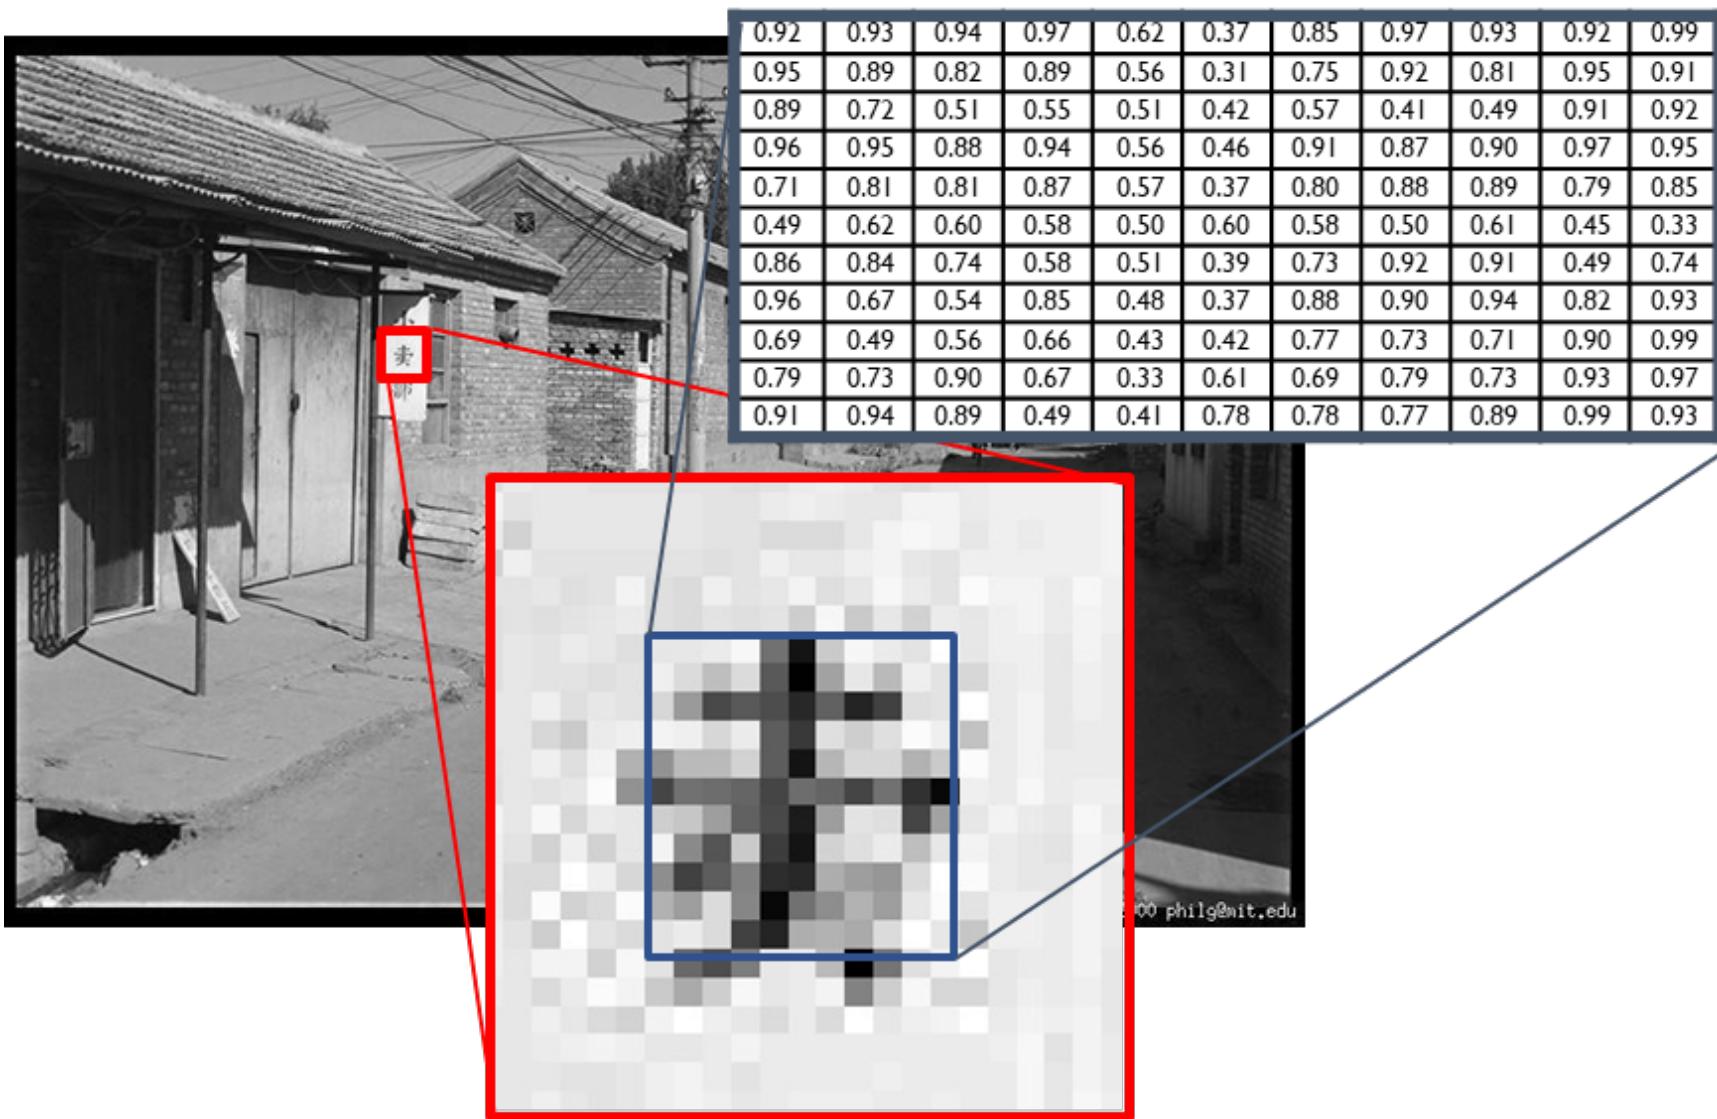
- $im[0, 0]$  = top-left pixel value
- $im[y, x]$  =  $y$  pixels down,  $x$  pixels to right
- $im[N-1, M-1]$  = bottom-right pixel

Row ↓      Column →

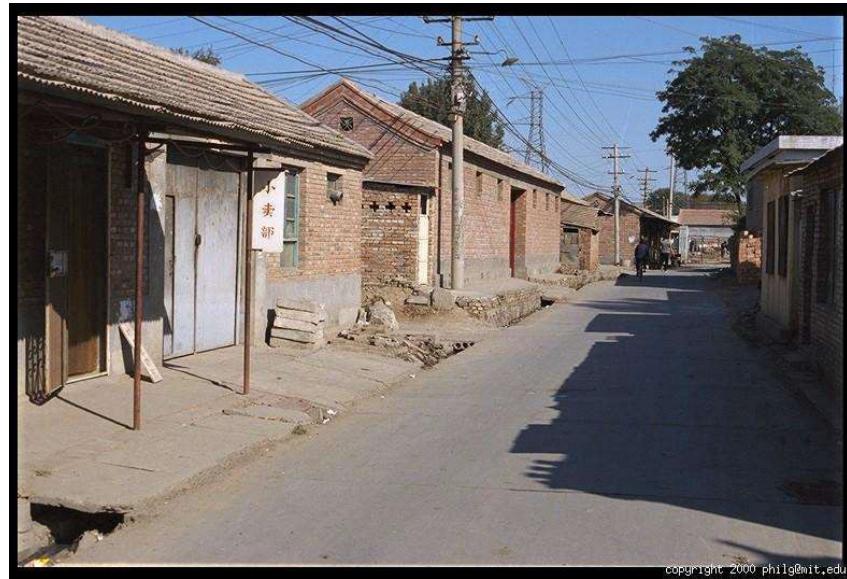
0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93

James Hays

# Grayscale Intensity



# Color



Red intensity



Green

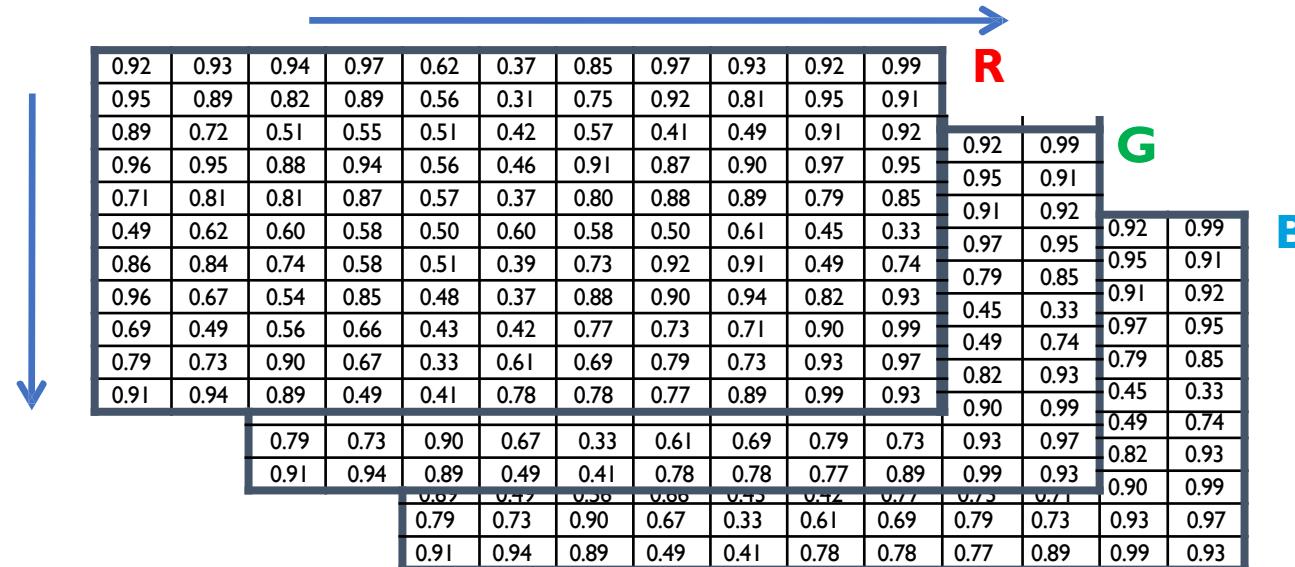


Blue

# Images in Python (import numpy)

$N \times M$  grayscale image “im”

- $im[0, 0, 0]$  = top-left pixel value, **red channel**
- $im[y, x, 1]$  =  $y$  pixels down,  $x$  pixels to right, **green channel**
- $im[N-1, M-1, 2]$  = bottom-right pixel, **blue channel**

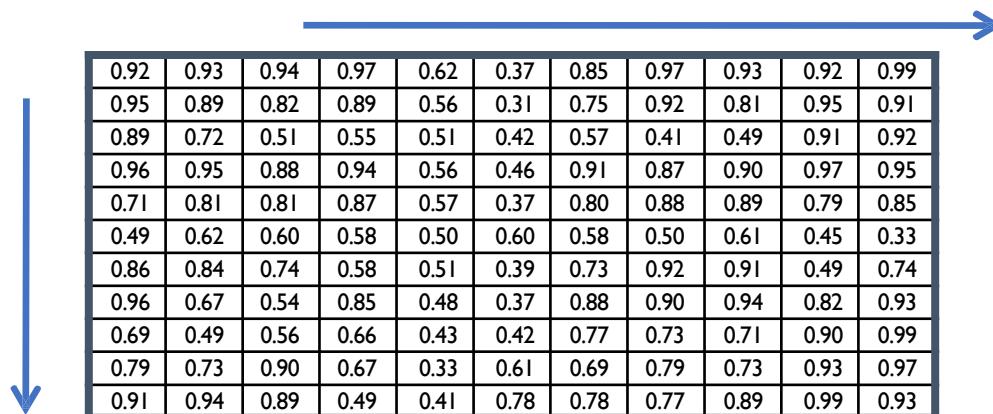


James Hays

# Images in Python (import numpy, scikit)

Take care of types!

- uint8 (values 0 to 255)      - `io.imread("file.jpg")`
- float32 (values 0 to 255)      - `io.imread("file.jpg").astype(np.float32)`
- float32 (values 0 to 1)      - `img_as_float32(io.imread("file.jpg"))`



James Hays

# Images in Python

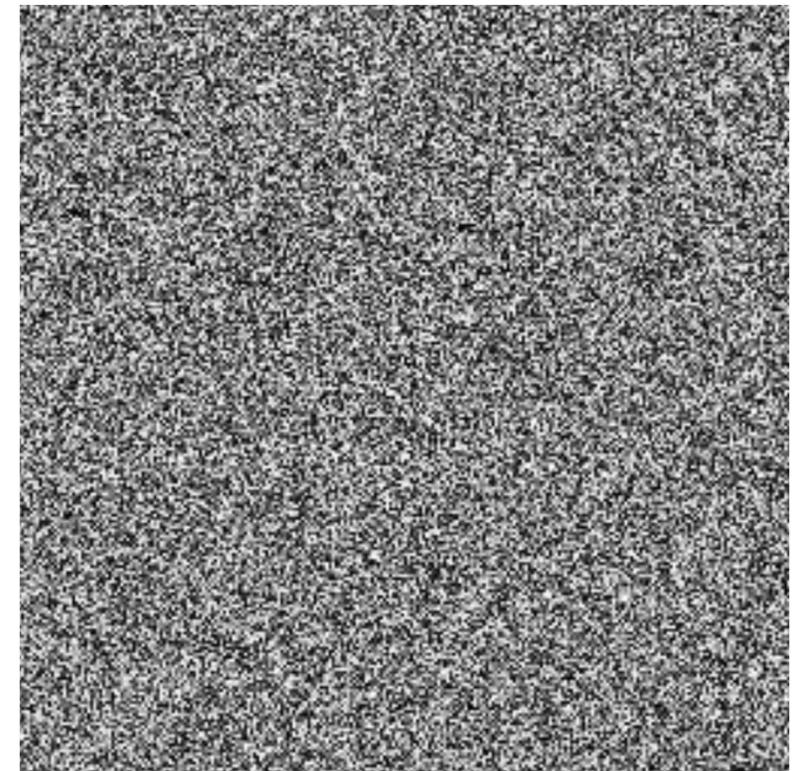
```
>>> from numpy import random as r  
>>> I = r.rand(256,256)
```

- What is this? What does it look like?
- Which values does it have?
- How many are there?

# Images in Python

```
>>> from matplotlib import pyplot as p  
>>> I = r.rand(256,256)  
>>> p.imshow(I, cmap='gray')  
>>> p.show()
```

Is it an image?





# Coding In This Course

## We use

- Package manager [conda](#)
- [Python](#) and its virtual environments
- `numpy` numerics library
- [Visual Studio Code](#) (VSCode) debugger



## GenAI (e.g. ChatGPT)

- Do **NOT** use! Otherwise, you won't learn to code!



## Be kind & supportive

- Students have diverse background (CS/EE/ME/...)
- Python proficiency varies among students
- Share your skills & knowledge with others on Ed

# Python Refresher

1. Download **tutorials**:

<https://github.com/ariarobotics/cv/tree/main/code/tutorials>

2. **Install Git:** follow instructions in “Github Guide”
3. **Install Python/Conda:** follow instructions in “Python Setup”
4. **Install VSCode:** follow instructions in “Setting up VSCode”
5. Run & review the numpy tutorial (on [Google Colab](#))
  1. Finish all 10 exercises at the end of the tutorial