

Introduction to Computer Vision

Kaveh Fathian

Assistant Professor

Computer Science Department

Colorado School of Mines

Lecture 3



Images in Python (import numpy)

$N \times M$ grayscale image “im”

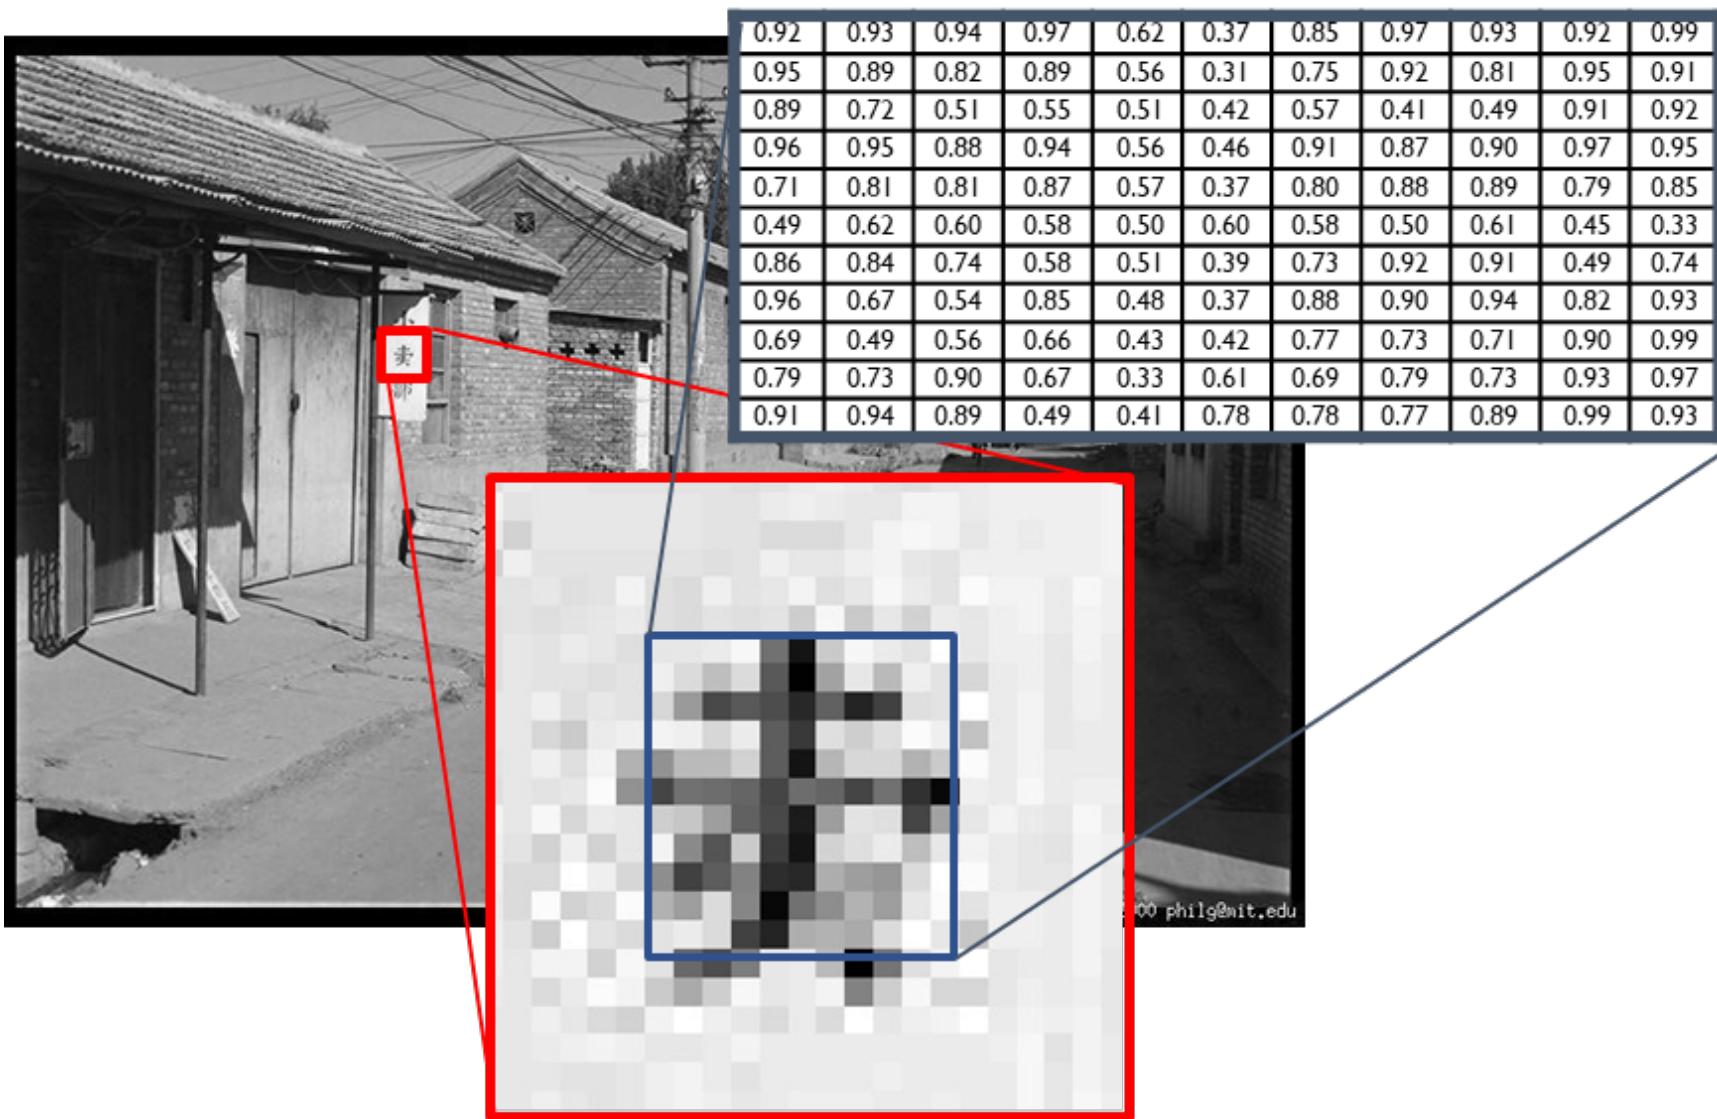
- $im[0, 0]$ = top-left pixel value
- $im[y, x]$ = y pixels down, x pixels to right
- $im[N-1, M-1]$ = bottom-right pixel

The diagram shows a 12x12 grid of numerical values representing a grayscale image. The columns are labeled from 0 to 11, and the rows are labeled from 0 to 11. A blue arrow labeled 'y' points downwards, indicating the row index. A blue arrow labeled 'x' points to the right, indicating the column index. The matrix values range from 0.37 to 0.99.

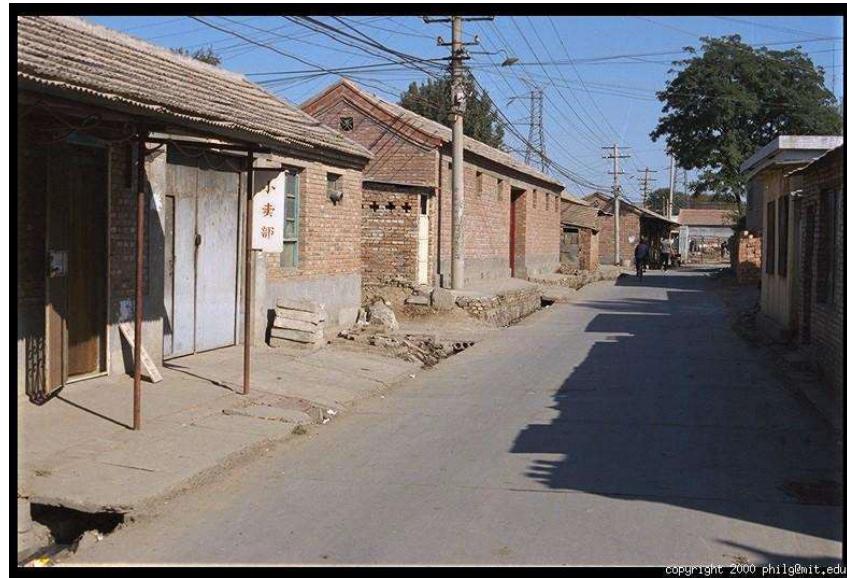
Row	Column	0	1	2	3	4	5	..	-	-	-	-
0	0	0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
1	1	0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
2	2	0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
3	3	0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
4	4	0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
.	5	0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
.	6	0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
.	7	0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
.	8	0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
.	9	0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
.	10	0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93
.	11											

James Hays

Grayscale Intensity



Color



Red intensity



Green

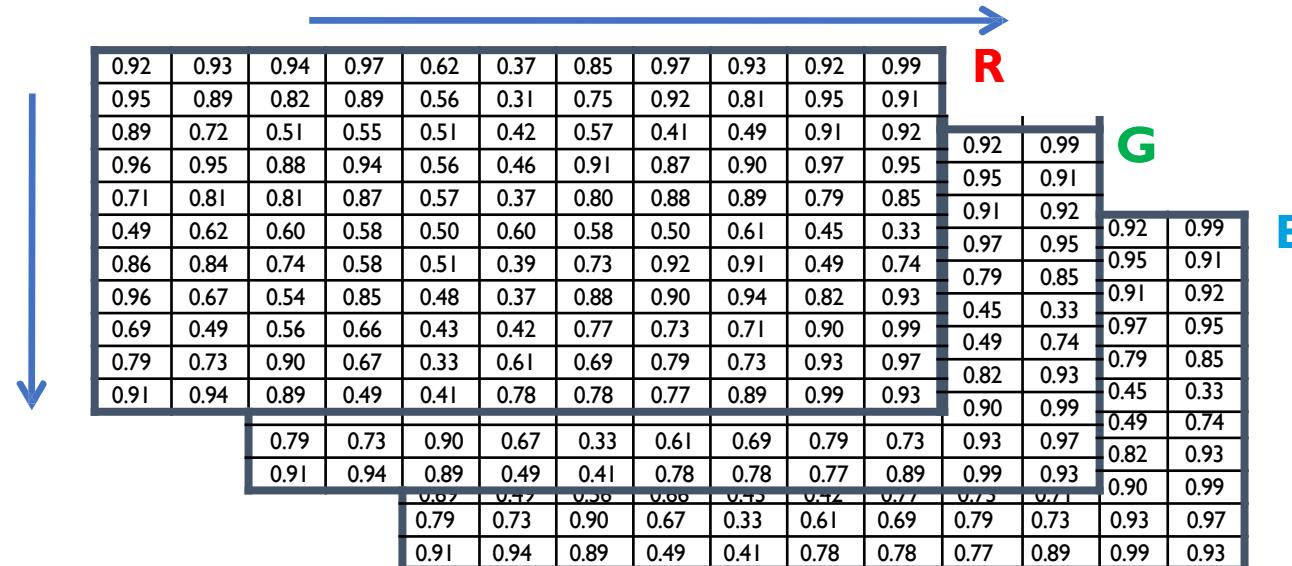


Blue

Images in Python (import numpy)

$N \times M$ grayscale image “im”

- $im[0, 0, 0]$ = top-left pixel value, **red channel**
- $im[y, x, 1]$ = y pixels down, x pixels to right, **green channel**
- $im[N-1, M-1, 2]$ = bottom-right pixel, **blue channel**



James Hays

Image Filtering



Image Filtering

Learning outcomes:


- Filter
- Image Filtering
- Low-pass Filter
- High-pass Filter
- Convolution
- Correlation
- Kernel
- Linearity and Shift Invariance
- Separability
- Sobel Filter
- Convolution
- Correlation & Template Matching
- Non-Linear Filters



Filtering

Definition

Filtering: An operation that modifies a (measured) signal.



- Modifications include
 - Removing undesirable components
 - Transforming signal in a desirable way
 - Extract specific components of a signal
- Can be analog or discrete signals
- Can be in any dimension
 - mainly 2D for images (gray)
 - 3D for colored

1D Filtering: Moving Average

$$I \in \mathcal{R}^{m \times 1}$$

↑ filter ↑ Signal

$$h[n] = \frac{1}{k} \sum_{i=n-k+1}^n (1) I[i]$$


Window size 'k'

$\underbrace{h[n]}_{1 \times 1} = \frac{1}{k} \underbrace{\begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}}_{1 \times k}^T \underbrace{I[n - k + 1 : n]}_{k \times 1}$

$\text{# of previous values}$

$\begin{bmatrix} I(n-k+1) \\ \vdots \\ I(n-1) \\ I(n) \end{bmatrix}_{k \times 1}$

https://en.wikipedia.org/wiki/Moving_average

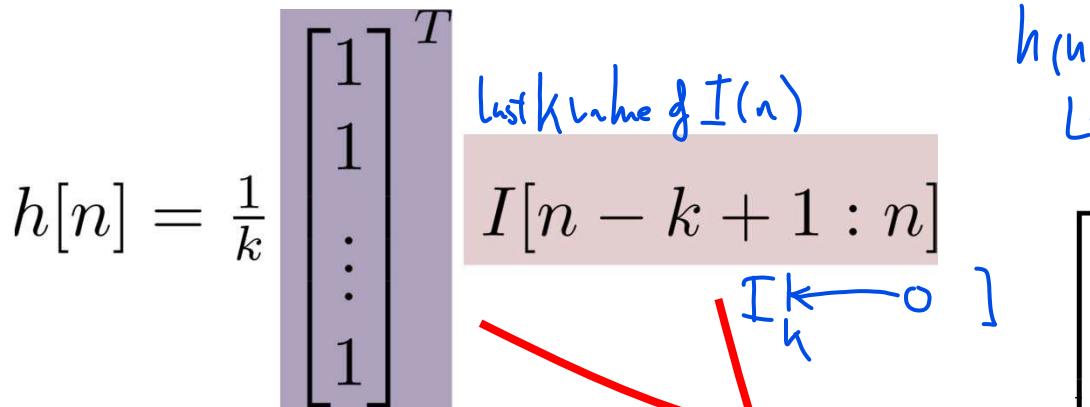


1D Filtering: Moving Average

$$I \in \mathcal{R}^{m \times 1}$$

$$h[n] = \frac{1}{k} \sum_{i=n-k+1}^n (1) I[i]$$

Window size 'k'



https://en.wikipedia.org/wiki/Moving_average

$$\text{ex: } I(n) = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ \vdots \end{bmatrix}$$

$$h(n) = [0\ 0\ 0\ 1]$$

$$\hookrightarrow I'(n) = \begin{bmatrix} 4 \\ 5 \\ \vdots \end{bmatrix}$$

(Identity)

$$h(n) = [1\ 0\ 0\ 0]$$

$$\hookrightarrow I''(n) = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

(delay)


$$\begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$

What is this filter? Identity

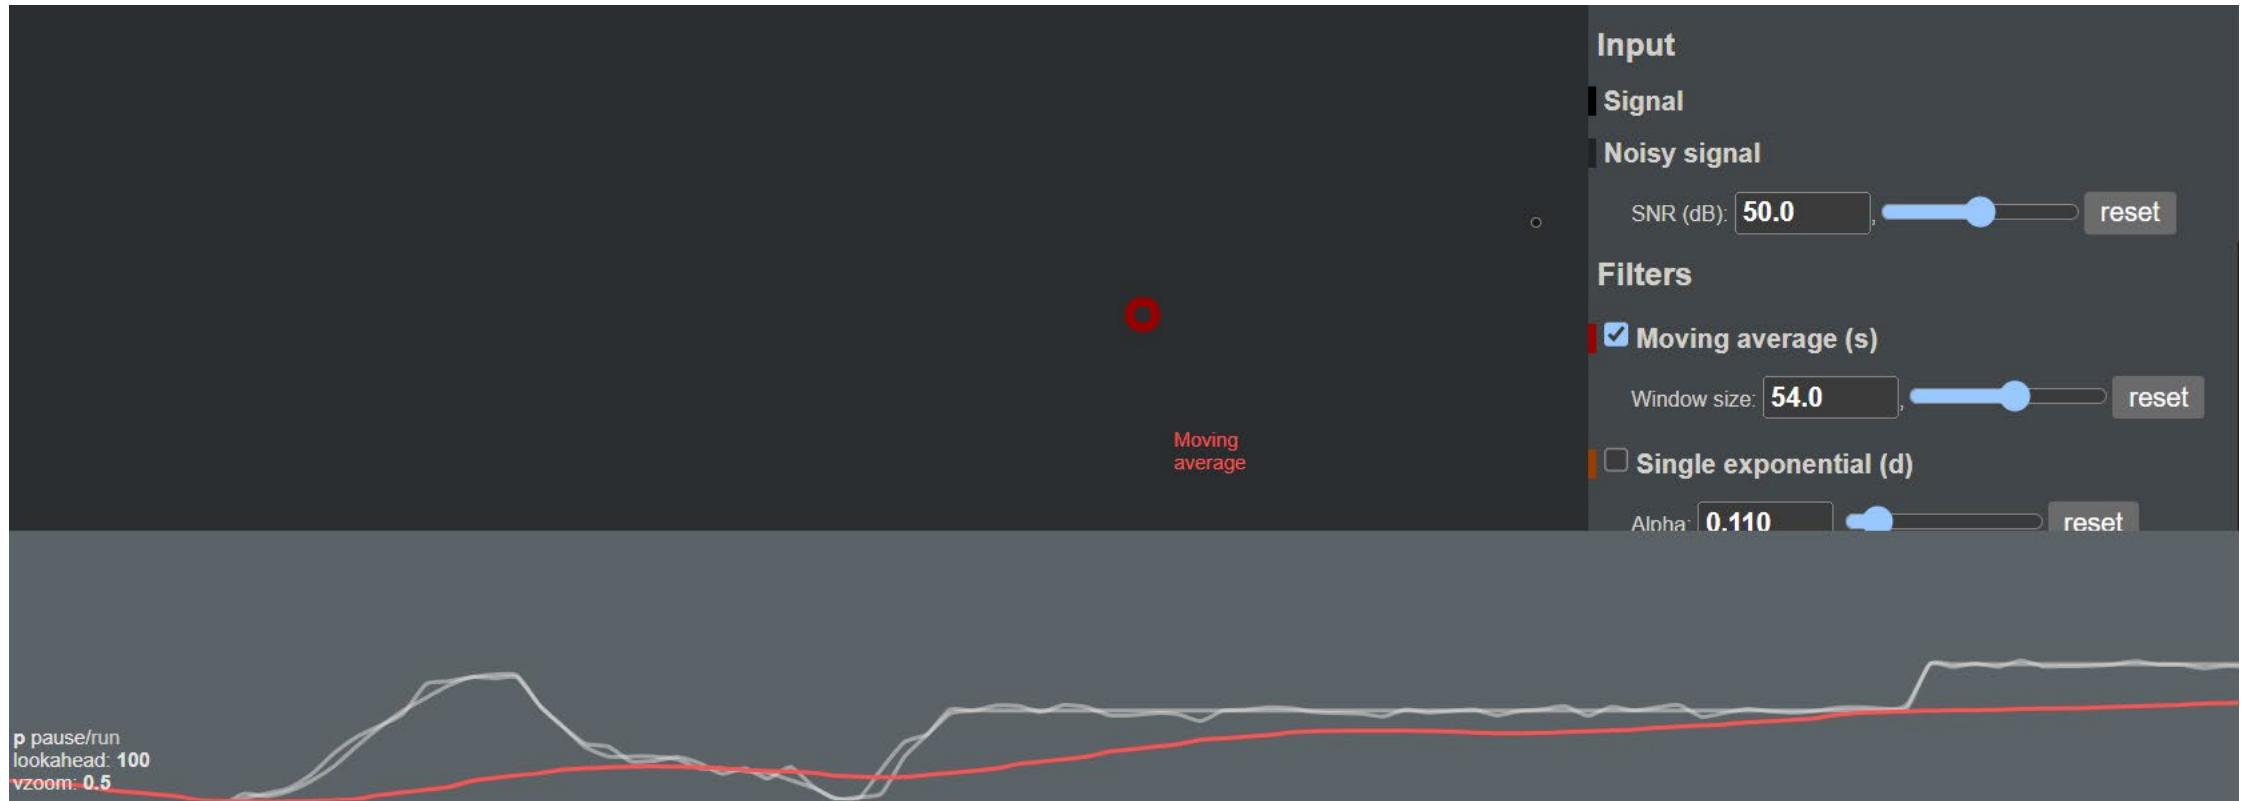
What is the difference?

$I[n - k/2 : n + k/2]$ delay filtered

$\left[\begin{array}{c} \leftarrow \circ \rightarrow \\ k/2 \quad k/2 \end{array} \right]$



Filtering Demo



<https://gery.casiez.net/1euro/InteractiveDemo/>

Image Filtering

Compute function of local neighborhood at each position:

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

filter *image*
 ↑ ↑
 y x

Window size 'k' and 'l'

Image Filtering

Compute function of local neighborhood at each position:

h =output

I =image

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

2d coords=k, l 2d coords=m, n

$[]$ $[]$ $[]$

Fundamental Equations of Computer Vision

1. Image Filtering

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l]$$

2. Optical Flow

$$I(x, y, t) = I(x + \Delta x, y + \Delta y, t + \Delta t)$$


3. Camera Geometry

$$\mathbf{x} = \mathbf{K}[\mathbf{R} \ \mathbf{t}] \mathbf{X} \quad x^T F x' = 0$$

4. Machine Learning

$$\arg \min_{\mathbf{S}} \sum_{i=1}^k \sum_{\mathbf{x} \in S_i} \|\mathbf{x} - \boldsymbol{\mu}_i\|^2. \quad y = \varphi(\sum_{i=1}^n w_i x_i + b) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

Example: Box Filter / Kernel



- every values are 1

Step:

1. time by $\frac{1}{k}$ | 2. sum all value in box | 3. put sum into filter of

Image Filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$I[.,.]$

10	10	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0


$h[.,.]$

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l] \quad m = 1, n = 1 \\ k, l = [-1, 0, 1]$$

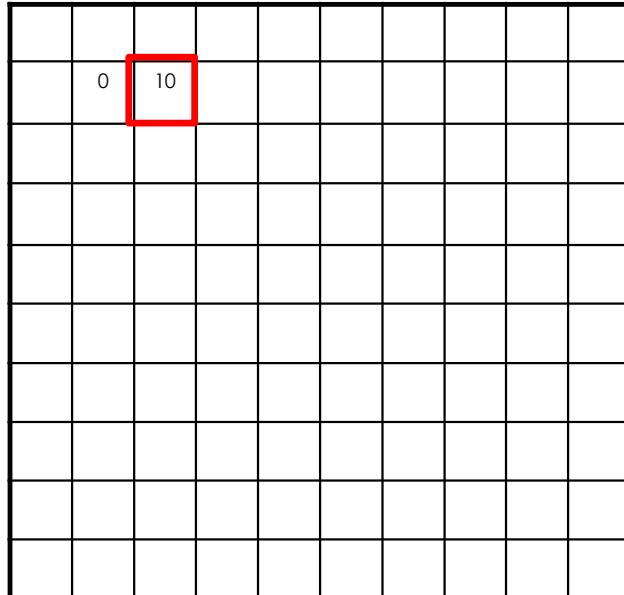
Image Filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$I[.,.]$



$h[.,.]$

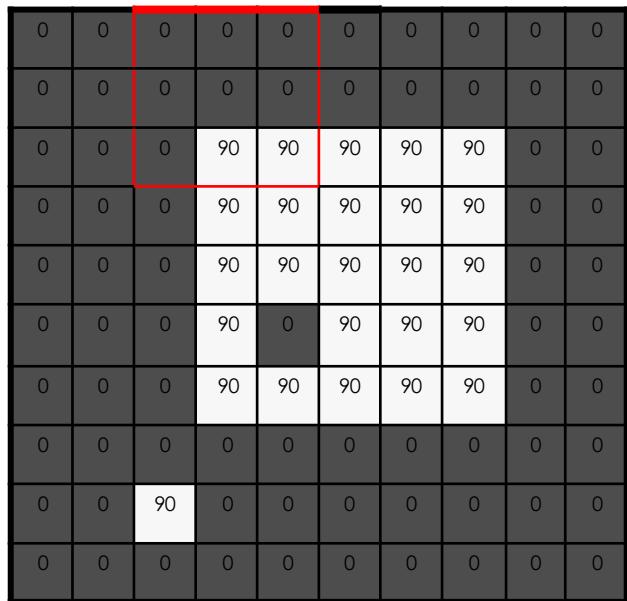


$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$
$$\begin{aligned} m &= 2, n = 1 \\ k, l &= [-1, 0, 1] \end{aligned}$$

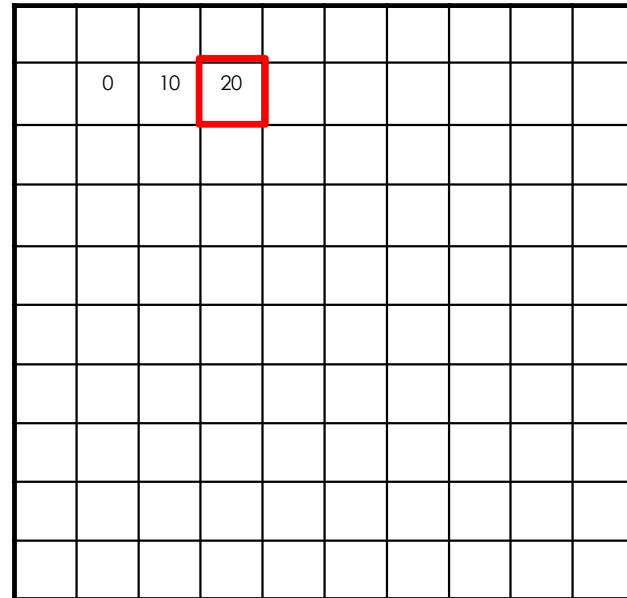
Image Filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$I[.,.]$



$h[.,.]$

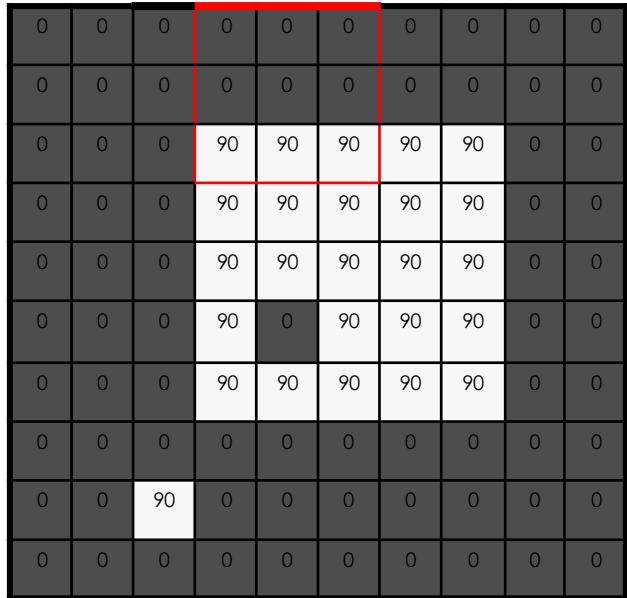


$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l] \quad m = 3, n = 1 \\ k, l = [-1, 0, 1]$$


Image Filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$I[.,.]$



$h[.,.]$



$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l] \quad m = 4, n = 1 \\ k, l = [-1, 0, 1]$$

Image Filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$I[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

$$\begin{aligned} m &= 5, n = 1 \\ k, l &= [-1, 0, 1] \end{aligned}$$

Image Filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

$$I[\cdot, \cdot]$$

$$h[.,.]$$

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l] \quad \begin{matrix} m = 4, n = 6 \\ k, l = [-1, 0, 1] \end{matrix}$$

Image Filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$I[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

	0	10	20	30	30					

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

$$\begin{aligned} m &= 6, n = 4 \\ k, l &= [-1, 0, 1] \end{aligned}$$

Image Filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

→ 2D moving average

Box Filter / Kernel

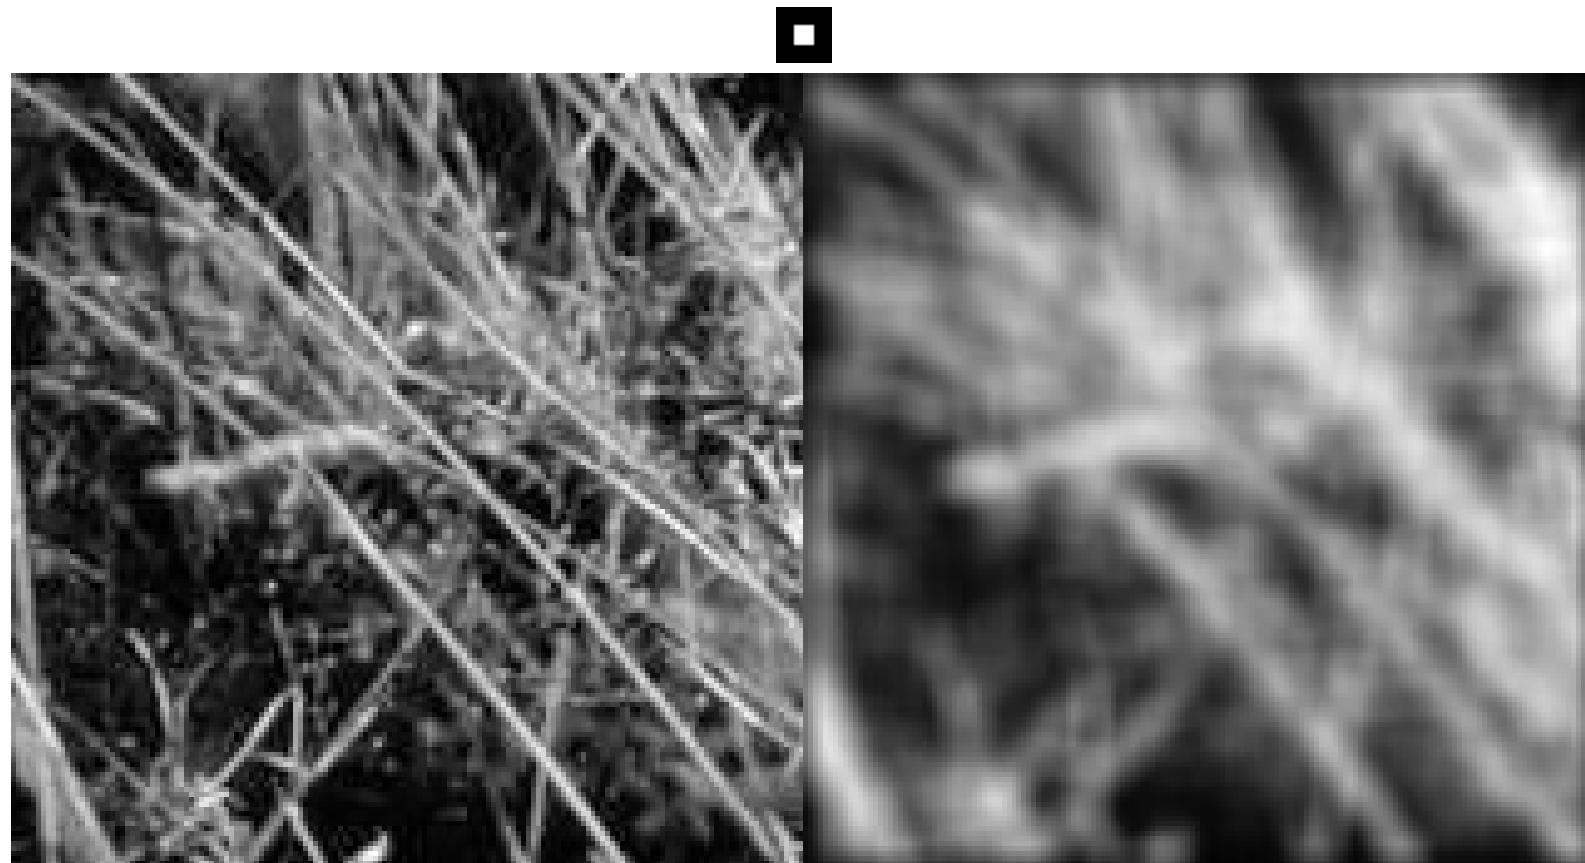
What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove ‘sharp’ features)
- Why does it sum to one?

$$f[\cdot, \cdot]$$
$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

filter size ↑ → blur ↑
 ↓
 image size ↓

Smoothing with Box Filter

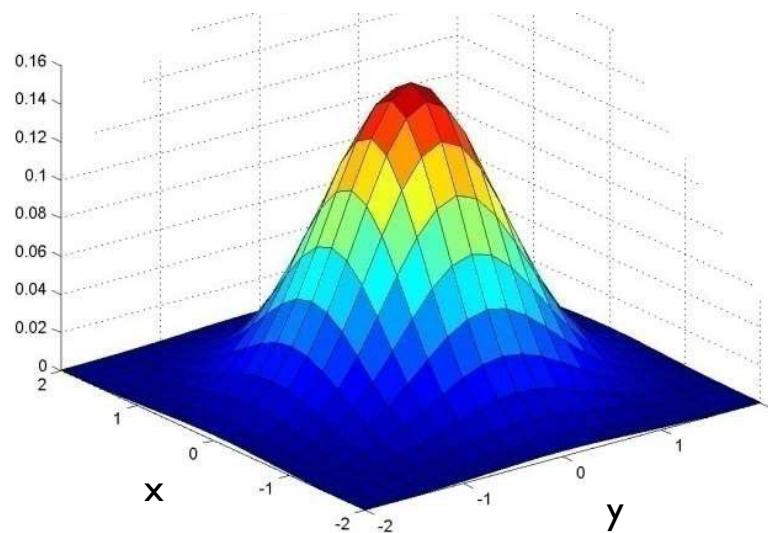


$$f[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$


David Lowe

Kaveh Fathian

Gaussian Filter / Kernel



Viewed
from top

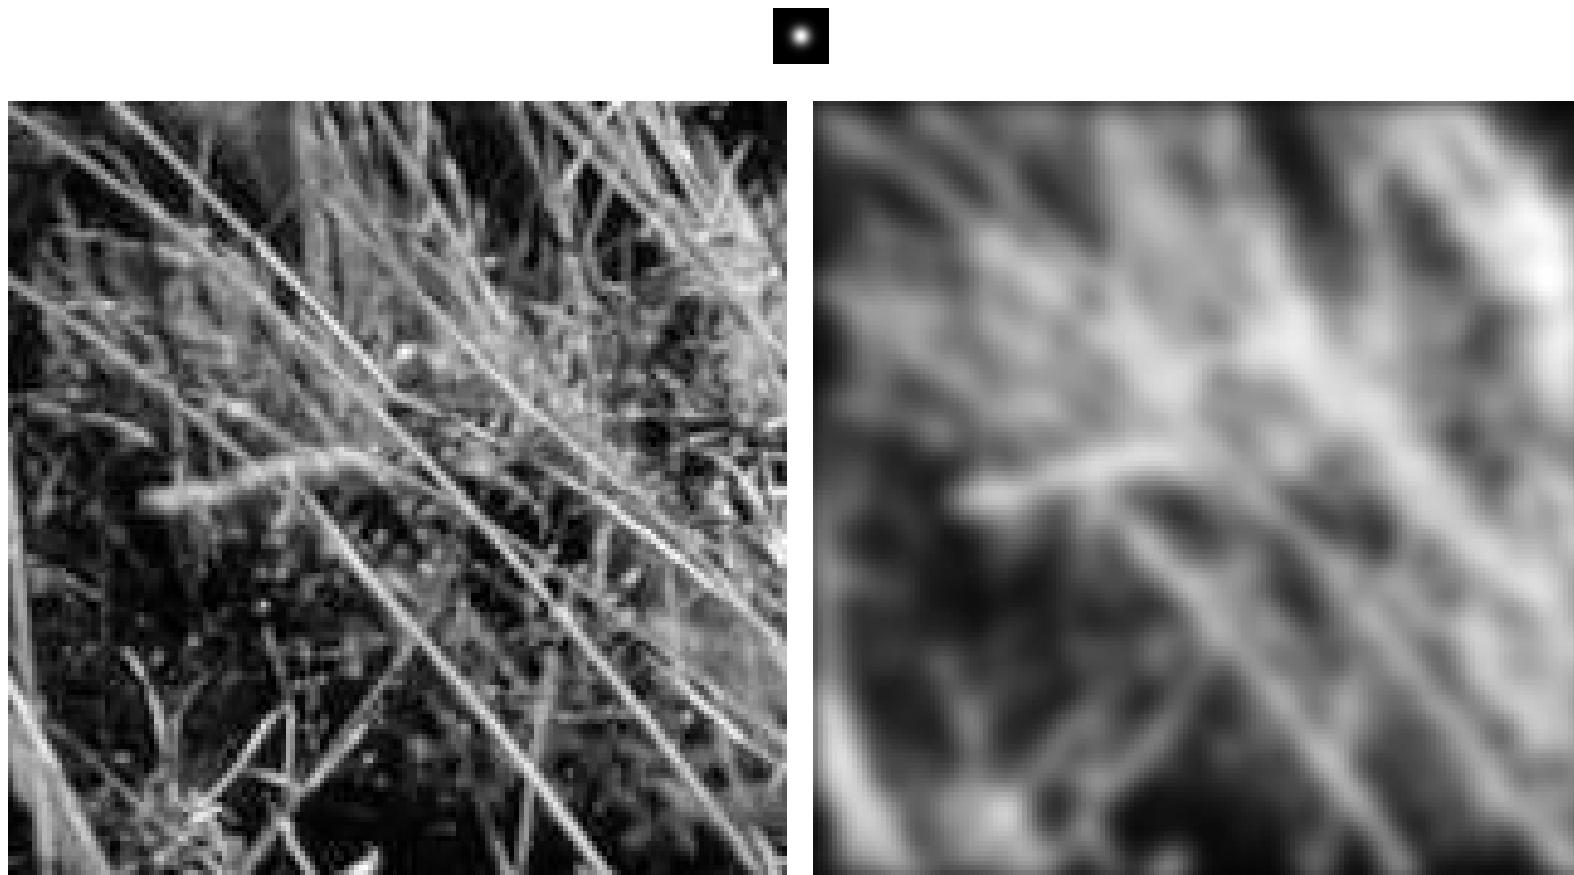


$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

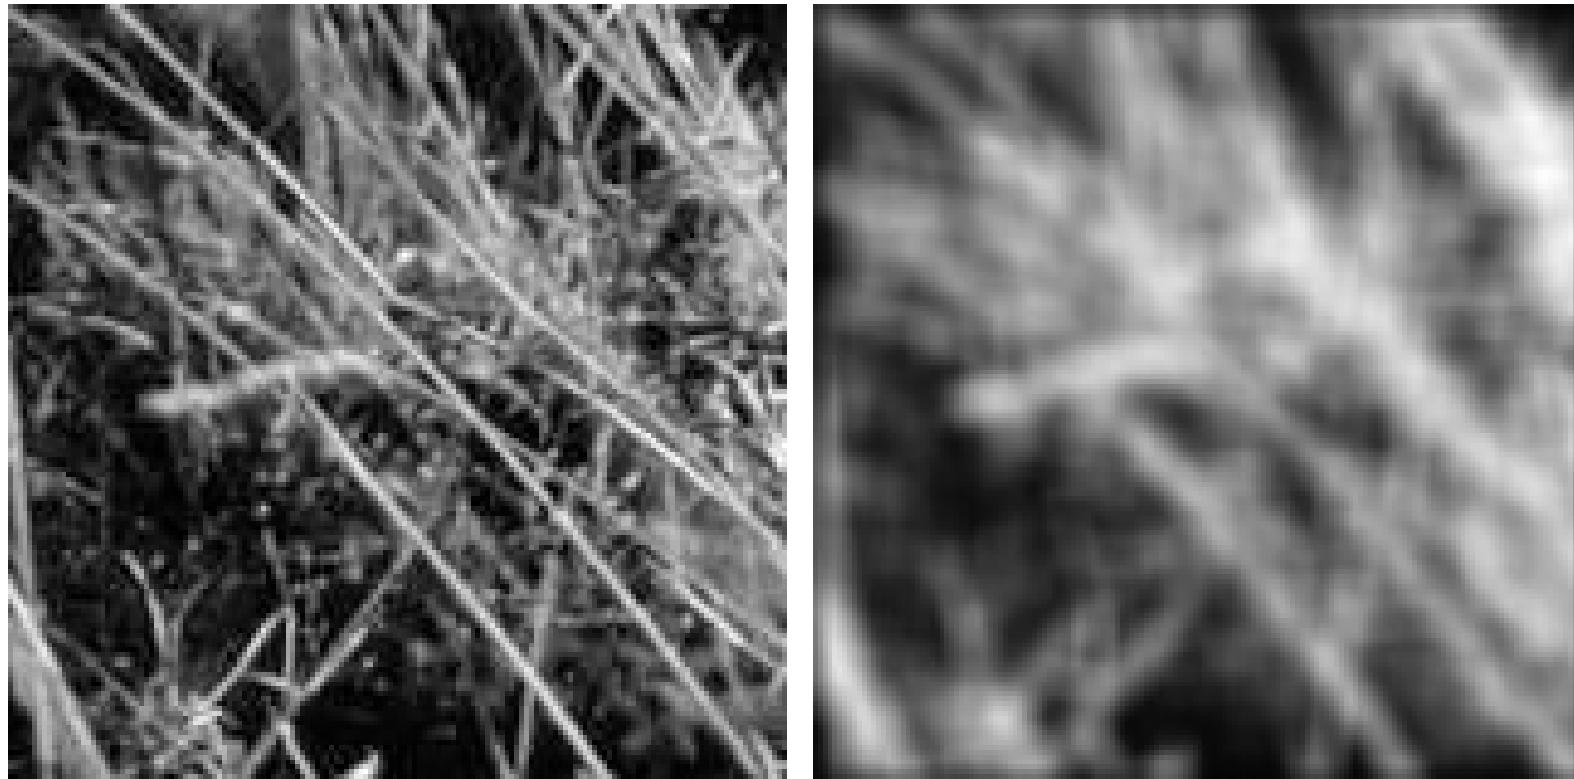
x	0.003	0.013	0.022	0.013	0.003
y	0.013	0.059	0.097	0.059	0.013
x	0.022	0.097	0.159	0.097	0.022
y	0.013	0.059	0.097	0.059	0.013
x	0.003	0.013	0.022	0.013	0.003

Filter/Kernel size 5×5 ,
Standard deviation $\sigma = 1$

Smoothing with Gaussian Filter



Smoothing with Box Filter



Linear Filter Properties

Linearity:

$$\text{imfilter}(I, f_1 + f_2) =$$

$$\text{imfilter}(I, f_1) + \text{imfilter}(I, f_2)$$

Shift/translation invariance:

Same behavior given intensities regardless of pixel location m,n

$$\text{imfilter}(I, \text{shift}(f)) = \text{shift}(\text{imfilter}(I, f))$$

Gaussian Filter Properties

Gaussian convolved with Gaussian...
...is another **Gaussian**

- So can smooth with small-width kernel, repeat, and get same result as larger-width kernel
- Convoluting twice with Gaussian kernel of width σ is same as convoluting once with kernel of width $\sigma\sqrt{2}$

Separable kernel

- Factors into product of two 1D Gaussians

Image Filtering Applications

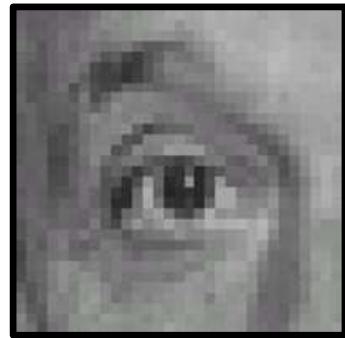
Compute function of local neighborhood at each position:

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

Really important!

- **Enhance images**
 - Denoise, resize, increase contrast, etc.
- **Extract information** from images
 - Texture, edges, distinctive points, etc.
- **Detect patterns**
 - Template matching

Practice with linear filters



1.

0	0	0
0	1	0
0	0	0

2.

0	0	0
0	0	1
0	0	0

3.

1	0	-1
2	0	-2
1	0	-1

4.

0	0	0
0	2	0
0	0	0

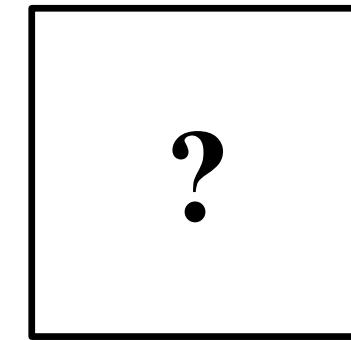
$$- \frac{1}{9} \begin{array}{|ccc|} \hline 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \\ \hline \end{array}$$

Practice with linear filters




Original

0	0	0
0	1	0
0	0	0

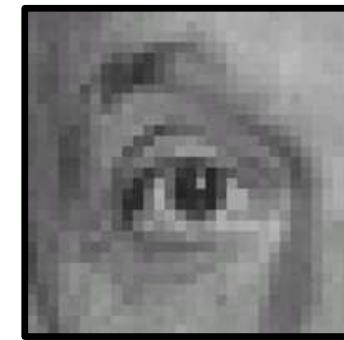


Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



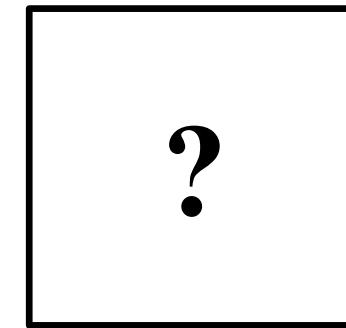
Filtered
(no change)

Practice with linear filters

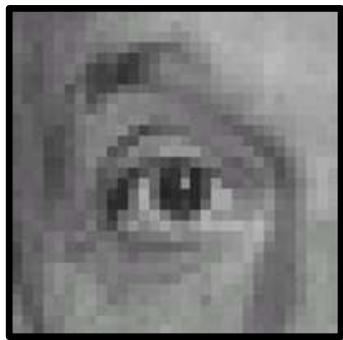


Original

0	0	0
0	0	1
0	0	0

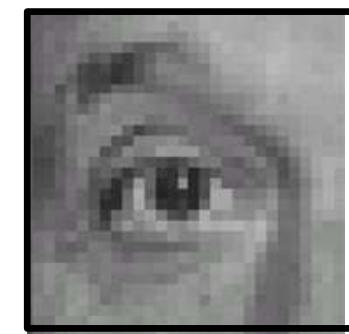


Practice with linear filters



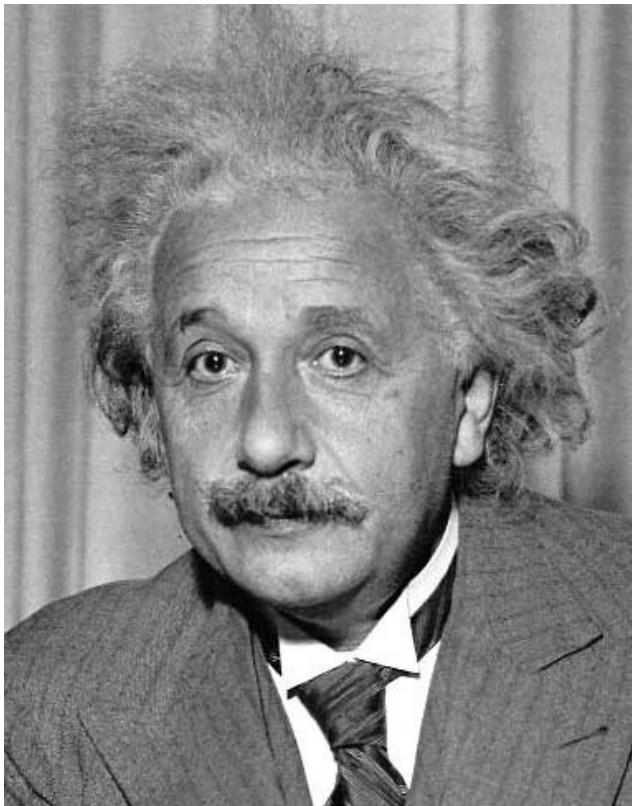
Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Practice with linear filters



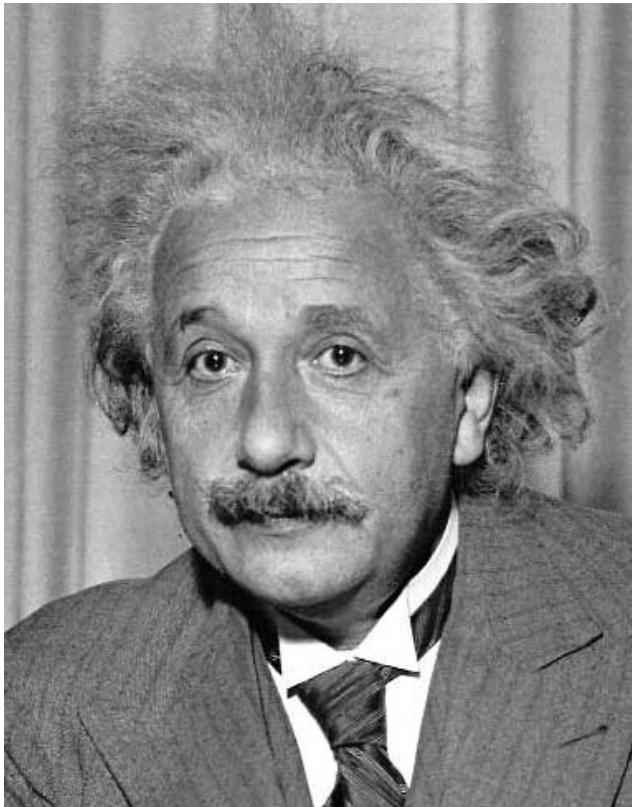
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

Practice with linear filters



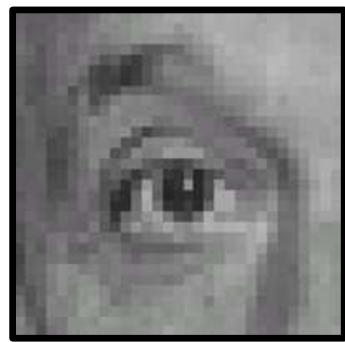
1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

Practice with linear filters



Original

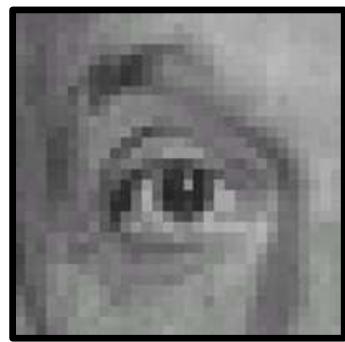
$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

$$- \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

(Note that filter sums to 1)

Practice with linear filters



Original

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

-

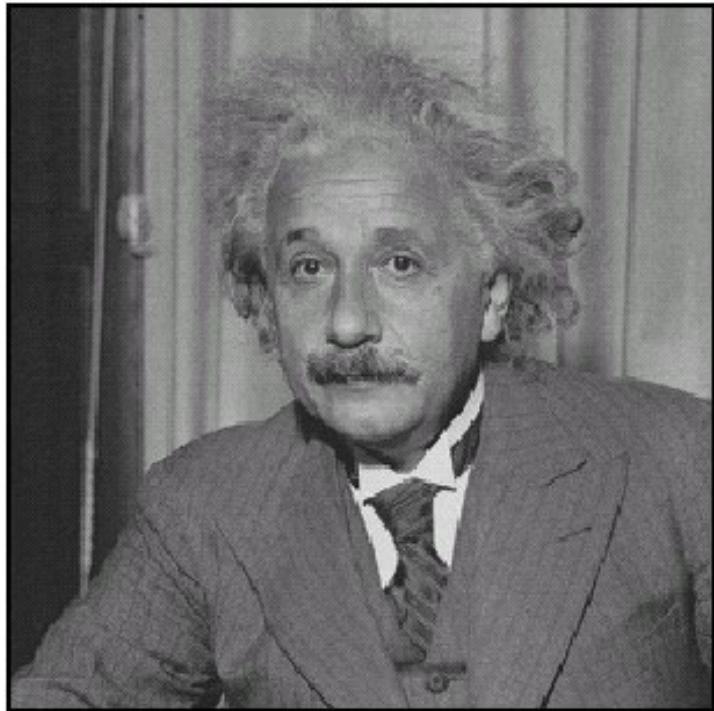
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$




Sharpening filter

- Accentuates differences with local average

Practice with linear filters



before


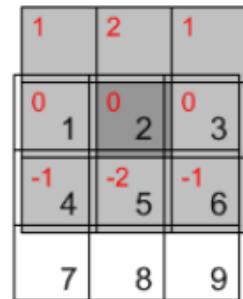
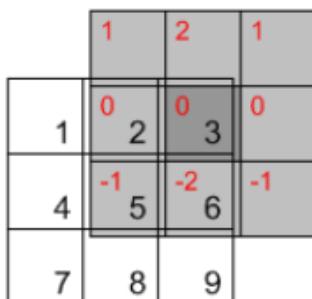





after

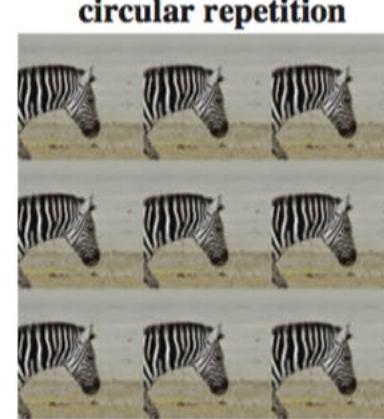

Filters in Practice

What about near the edge?

- The filter window falls off the edge of the image
- Need to extrapolate (padding)
- Methods:
 - clip filter (black / zero padding)
 - wrap around
 - copy edge
 - reflect across edge

Input



Correlation and Convolution

Definition

- **2D correlation**

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

* Symmetric in the geometric sense, not in the matrix linear algebra sense.

Filtering

Definition

- 2D correlation

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

- 2D convolution

$$h[m, n] = \sum_{k,l} f[k, l] I[m - k, n - l]$$

e.g., `h = scipy.signal.convolve2d(f, I)`

Convolution is the same as correlation with a 180° rotated filter kernel.
Correlation and convolution are identical when the filter kernel is symmetric*.

* Symmetric in the geometric sense, not in the matrix linear algebra sense.

Convolution Properties

Commutative: $a * b = b * a$

- Conceptually no difference between filter and signal
- But filtering implementations might break this equality, e.g., image edges

Associative: $a * (b * c) = (a * b) * c$

- Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Why important?
- Correlation is **NOT** associative (rotation effect)

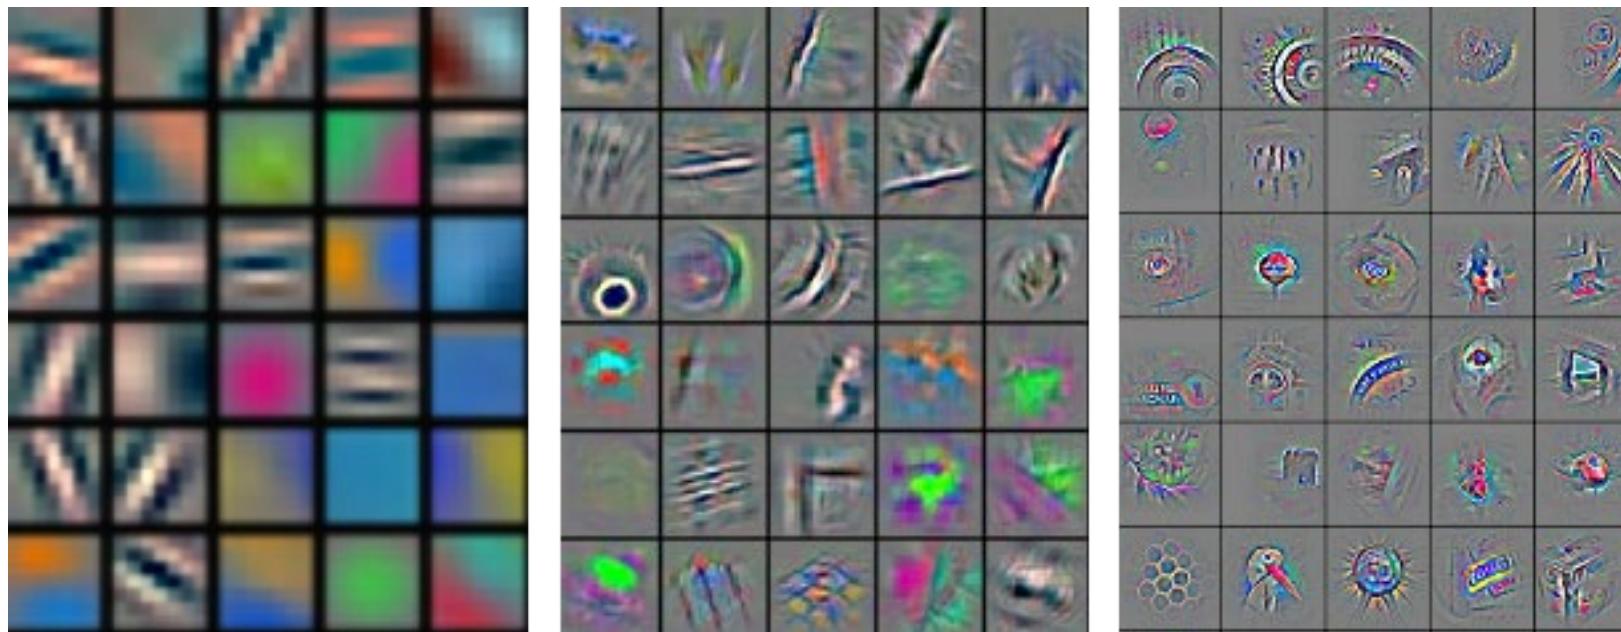
Distributes over addition: $a * (b + c) = (a * b) + (a * c)$

Scalars factor out: $ka * b = a * kb = k(a * b)$

Identity: unit impulse $e = [0, 0, 1, 0, 0]$, $a * e = a$

Convolution in Convolutional Neural Networks

- Convolution is the basic operation in CNNs
- Learning convolution kernels allows us to learn which ‘features’ provide useful information in images.



Linear Filters

Linearity:

$$\text{imfilter}(I, f_1 + f_2) = \text{imfilter}(I, f_1) + \text{imfilter}(I, f_2)$$

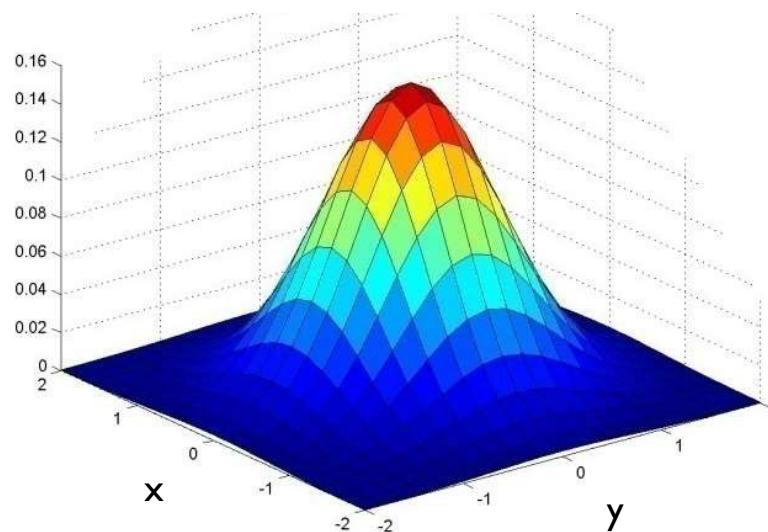
Shift/translation invariance:

Same behavior given intensities regardless of pixel location m,n

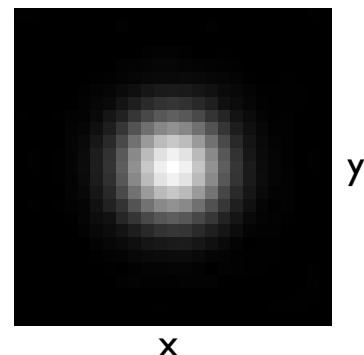
$$\text{imfilter}(I, \text{shift}(f)) = \text{shift}(\text{imfilter}(I, f))$$

Any linear, shift-invariant operator can be represented as a convolution

Gaussian Filter / Kernel



Viewed
from top



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

x	0.003	0.013	0.022	0.013	0.003
y	0.013	0.059	0.097	0.059	0.013
x	0.022	0.097	0.159	0.097	0.022
y	0.013	0.059	0.097	0.059	0.013
x	0.003	0.013	0.022	0.013	0.003

Filter/Kernel size 5×5 ,
Standard deviation $\sigma = 1$

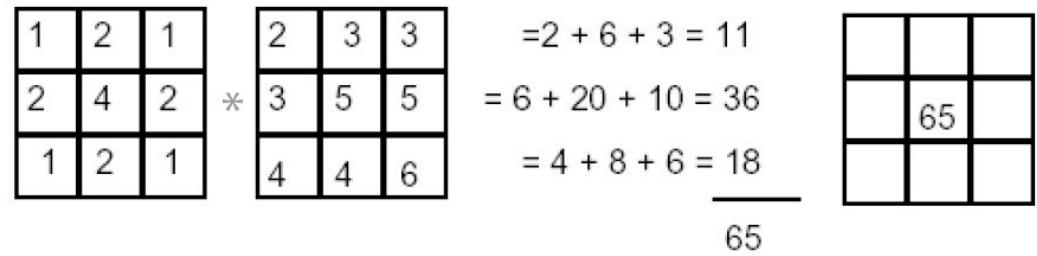
Separability of the Gaussian Filter

- The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y
- In this case, the two functions are the (identical) 1D Gaussian:

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

Separability Example

2D convolution (center location only):

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{r} = 2 + 6 + 3 = 11 \\ = 6 + 20 + 10 = 36 \\ = 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$


The filter factors into a product of 1D filters:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Perform convolution along rows:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 11 & & \\ \hline 18 & & \\ \hline 18 & & \\ \hline \end{array}$$

Followed by convolution along the remaining column:

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 11 & & \\ \hline 18 & & \\ \hline 18 & & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & 65 & \\ \hline \end{array}$$

Separability Use

MxN image, PxQ filter

- 2D convolution: $\sim MNPQ$ multiply-adds
- Separable 2D: $\sim MN(P+Q)$ multiply-adds

Speed up = $PQ/(P+Q)$ 9x9 filter = $\sim 4.5x$ faster

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Sobel Filter

```
>> I = img_to_float32( io.imread('luke.jpg') )  
>> h = convolve2d( I, sobelKernel )  
>> plt.imshow( h )
```

1	2	1
0	0	0
-1	-2	-1

Sobel



Sobel Filter

- What happens to negative numbers?

1	2	1
0	0	0
-1	-2	-1

Sobel



Sobel Filter

$h(:,:,1) < 0$



$h(:,:,1) > 0$



Sobel Filter

- For visualization:
 - Shift image + 0.5
 - If gradients are small, scale edge response

```
>> I = img_to_float32( io.imread('luke.jpg') )  
>> h = convolve2d( I, sobelKernel )  
  
>> plt.imshow( h )  
  
>> plt.imshow( h + 0.5 )
```

1	2	1
0	0	0
-1	-2	-1

Sobel



Practice

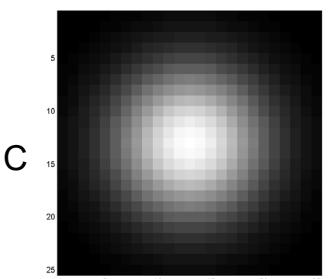
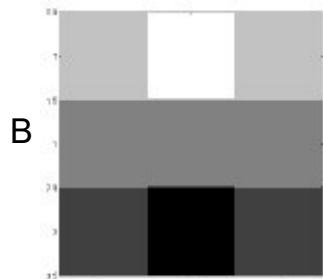
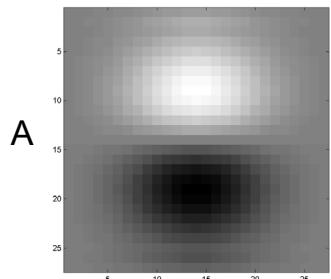
* = Convolution operator

$$1) \underline{\quad} = D * B$$


$$2) A = \underline{\quad} *$$

$$3) F = D * \underline{\quad}$$

$$4) \underline{\quad} = D * D$$



Derek Hoiem



Kaveh Fathian

Practice

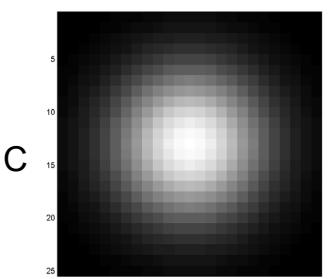
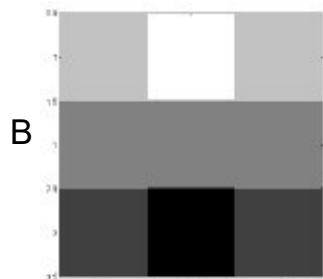
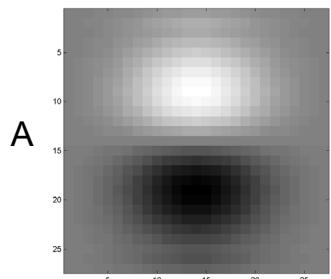
* = Convolution operator

$$1) G = D * B$$



$$2) A = B * C$$

$$3) F = D * E$$

$$4) I = D * D$$



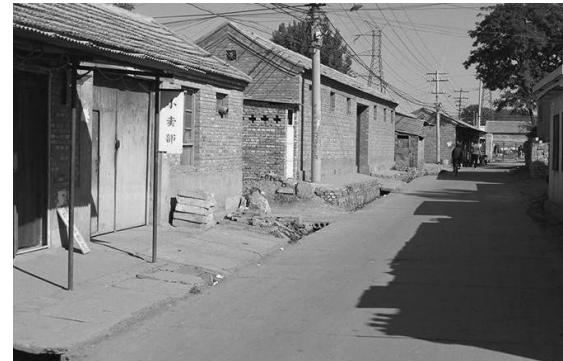
Derek Hoiem



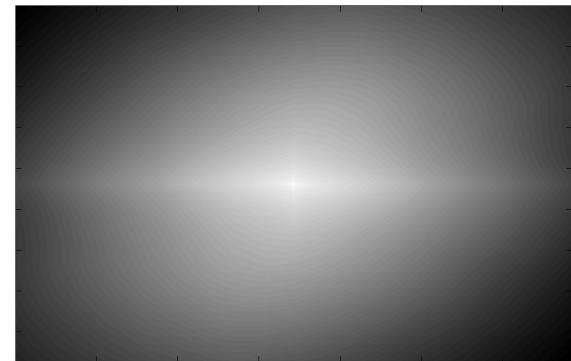
Kaveh Fathian

Convolution

D (275 x 175 pixels)



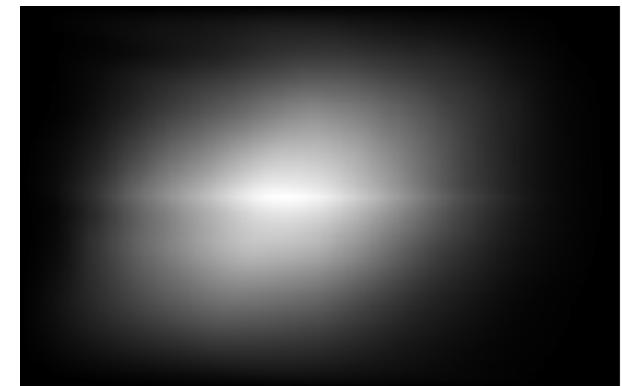
I (275 x 175 pixels)



```
>> D = img_to_float32( io.imread( 'convexample.png' ) )
>> I = convolve2d( D, D )
>> np.max(I) 1.1021e+04
```

I_norm

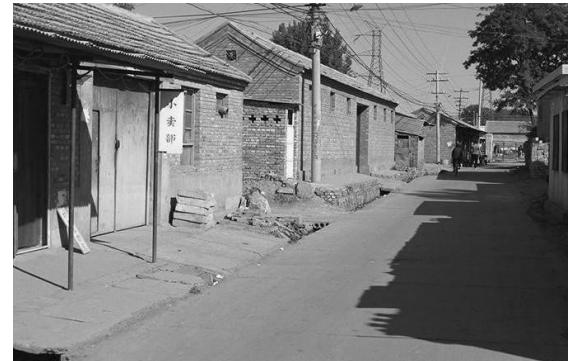
```
# Normalize for visualization
>> I_norm = (I - np.min(I)) / (np.max(I) - np.min(I))
>> plt.imshow( I_norm )
```



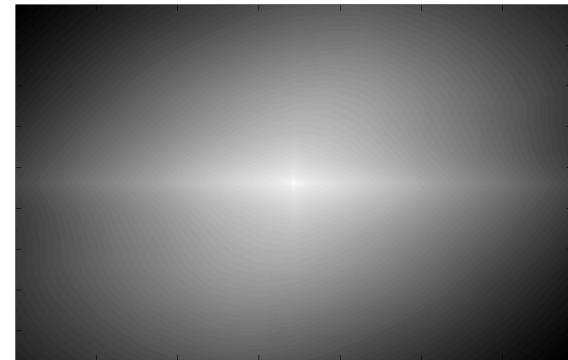
Zero padding causes black regions

Convolution

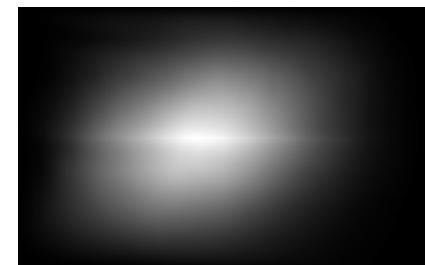
D (275 x 175 pixels)



I (275 x 175 pixels)



```
>> I = convolve2d( D, D, mode='full' )  
(Default; pad with zeros)
```



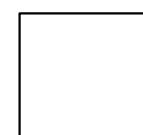
549 x 349

```
>> I = convolve2d( D, D, mode='same' )  
(Return same size as D)
```



275 x 175

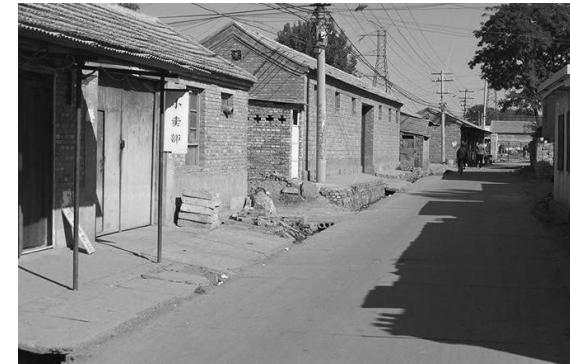
```
>> I = convolve2d( D, D, mode='valid' )  
(No padding)
```




Value = 10528.3 1x1

Convolution

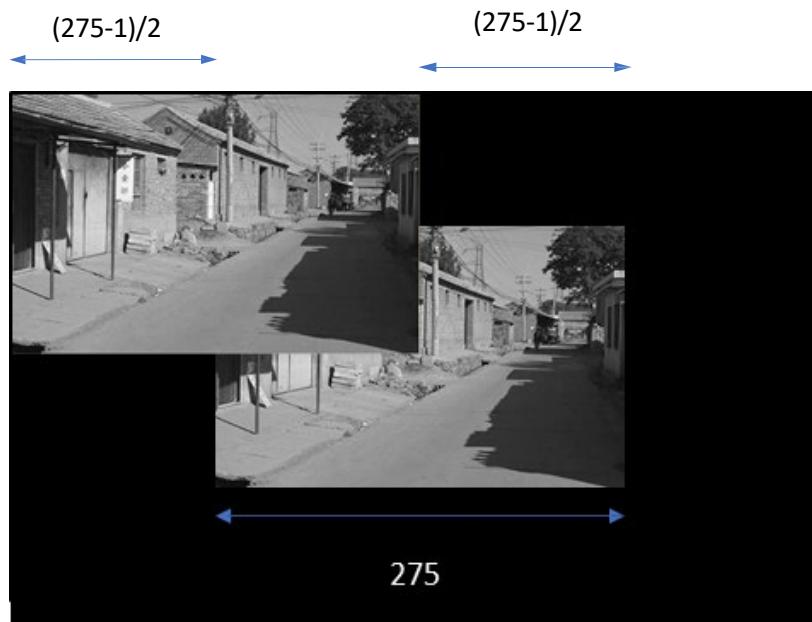
D (275 x 175 pixels)



I (275 x 175 pixels)



```
>> I = convolve2d( D, D, mode='full' )  
(Default; pad with zeros)
```



$$\text{For } x: 275 + (275-1)/2 + (275-1)/2 = 549$$

I_norm (549 x 349 pixels)



Correlation

- When the filter ‘looks like’ the image = ‘template matching’
- Filtering viewed as comparing an image of what you want to find against all image regions.
- For **symmetric** filters: use either convolution or correlation.
- For **nonsymmetric** filters: correlation is template matching.

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

e.g., `h = scipy.signal.correlate2d(f, I)`

As brightness in I increases, the response in h will increase, as long as f is positive.



Correlation

Let's see if we can use correlation to 'find' the parts of the image that look like the filter

D (275×175 pixels).



```
>> f = D[ 57:117, 107:167 ]
```

Expect response 'peak' in middle of I

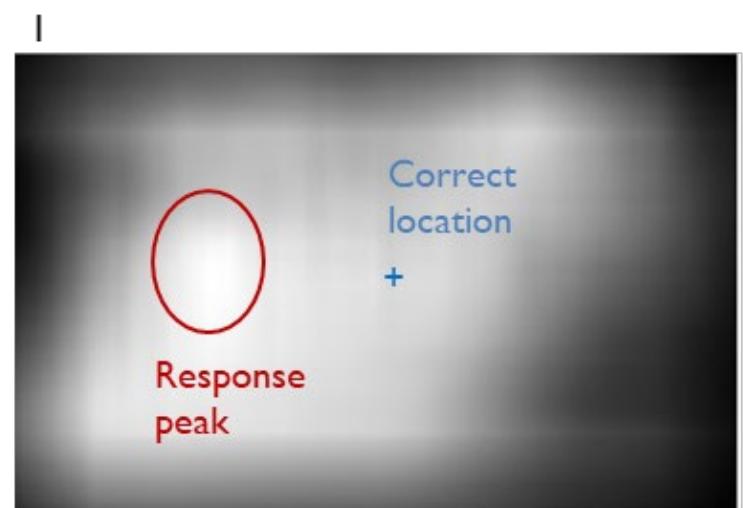
```
>> I = correlate2d( D, f, 'same' )
```

Hmm...

That didn't work – why not?



f
 61×61



Correlation

Let's see if we can use correlation to 'find' the parts of the image that look like the filter

D (275 x 175 pixels).



```
>> f = D[ 57:117, 107:167 ]
```

Expect response 'peak' in middle of I



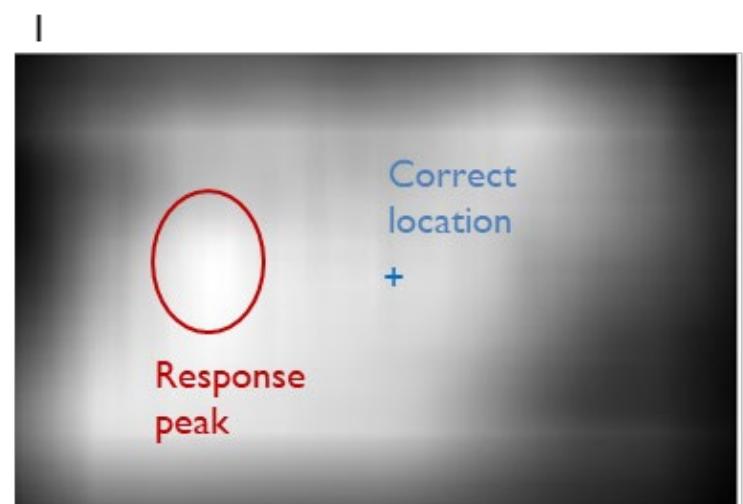
f
61 x 61

```
>> I = correlate2d( D, f, 'same' )
```

Hmm...

That didn't work – why not?

Overall brighter regions will give higher correlation response
-> not useful!



Correlation

OK, so let's subtract the mean

```
>> f = D[ 57:117, 107:167 ]  
>> f2 = f - np.mean(f)  
>> D2 = D - np.mean(D)
```

Now zero centered.

Score is higher only when dark parts match and when light parts match.

```
>> I2 = correlate2d( D2, f2, 'same' )
```

D2 (275 x 175 pixels)



f2
61 x 61

Remember –
float data type –
goes negative!

I2



Correlation

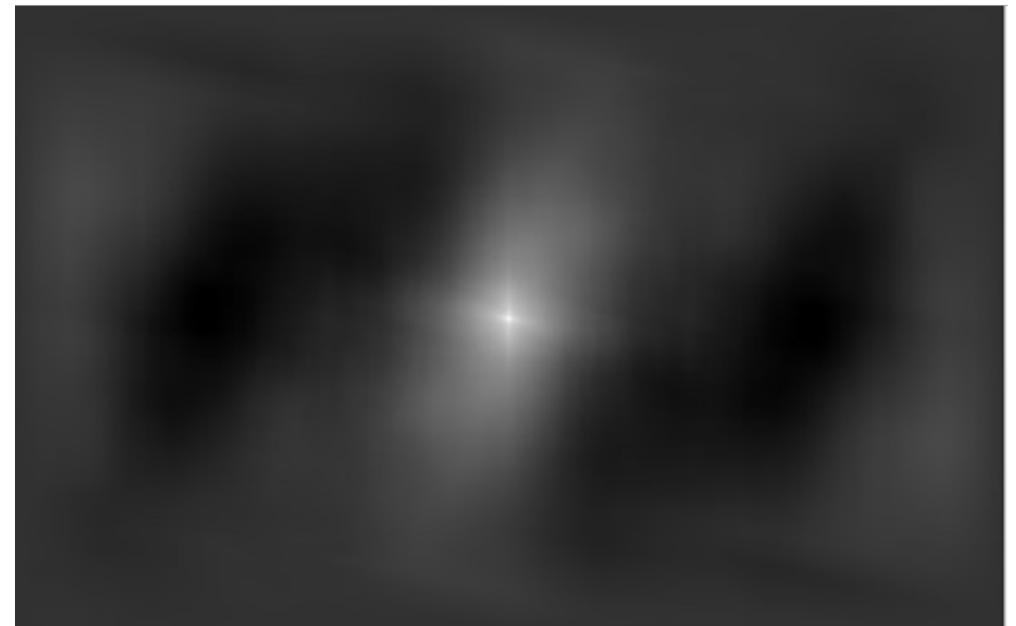
Or our original example: D2 correlated with itself

```
>> I3 = correlate2d( D2, D2, 'full' )
```

D2 (275 x 175 pixels)



I3



Correlation vs. Convolution

What happens with convolution?

```
>> f = D[ 57:117, 107:167 ]  
>> f2 = f - np.mean(f)  
>> D2 = D - np.mean(D)           I2  
>> I2 = convolve2d( D2, f2, 'same' )
```

D2 (275 x 175 pixels)



f2
61 x 61



Non-Linear Filters

- **Rank filters** are non-linear filters based on ordering of gray levels
 - e.g., median, min, max, range filters
- **Median filter** operates over a window by selecting the median intensity in the window

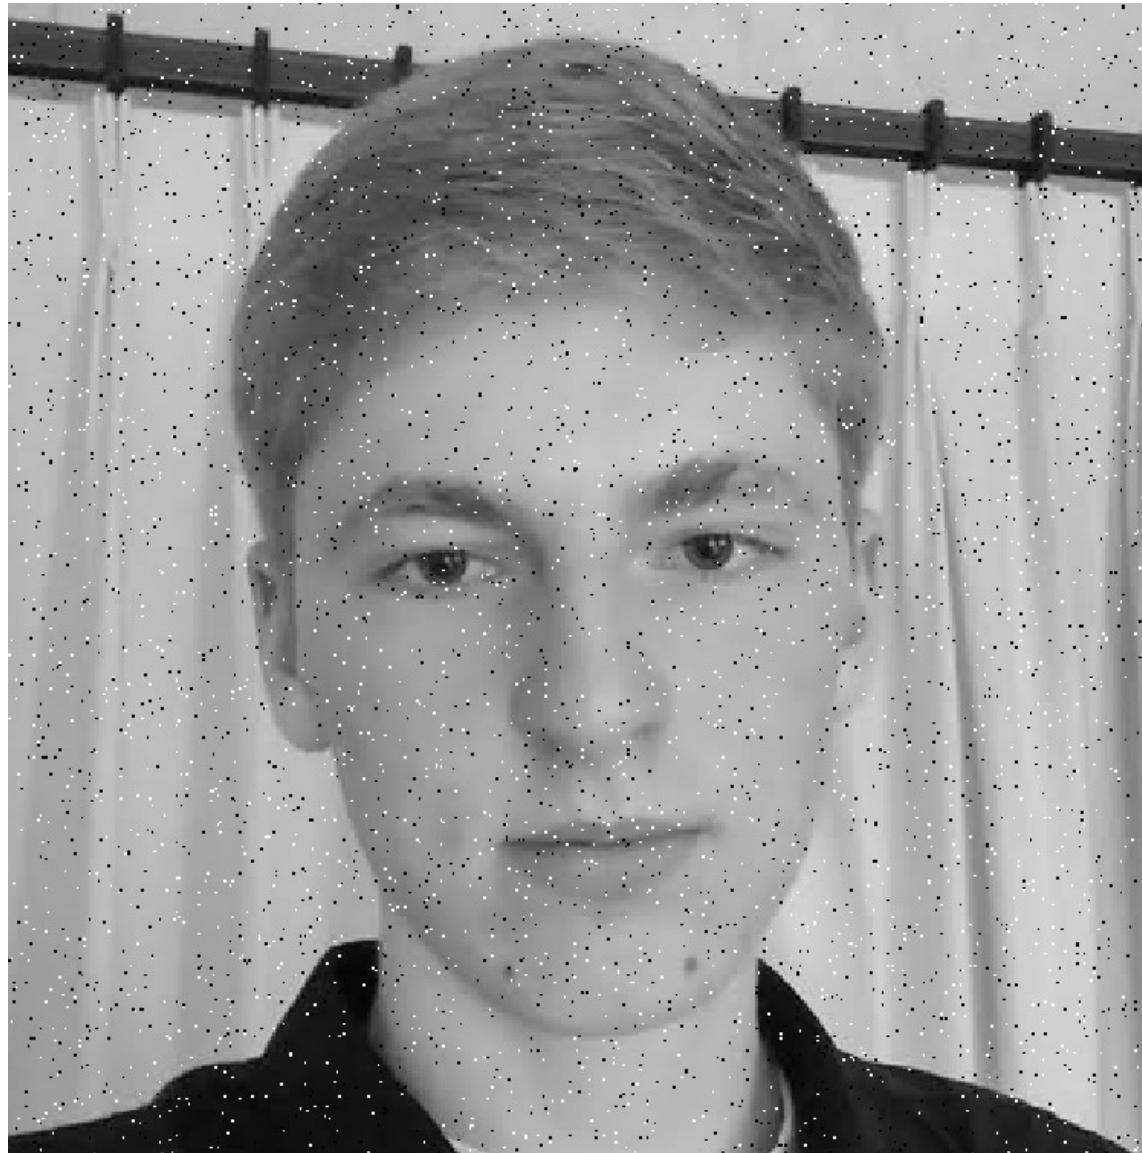
 $I[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $h[.,.]$

Steve Seitz, Steve Marschner

Salt and Pepper Noise



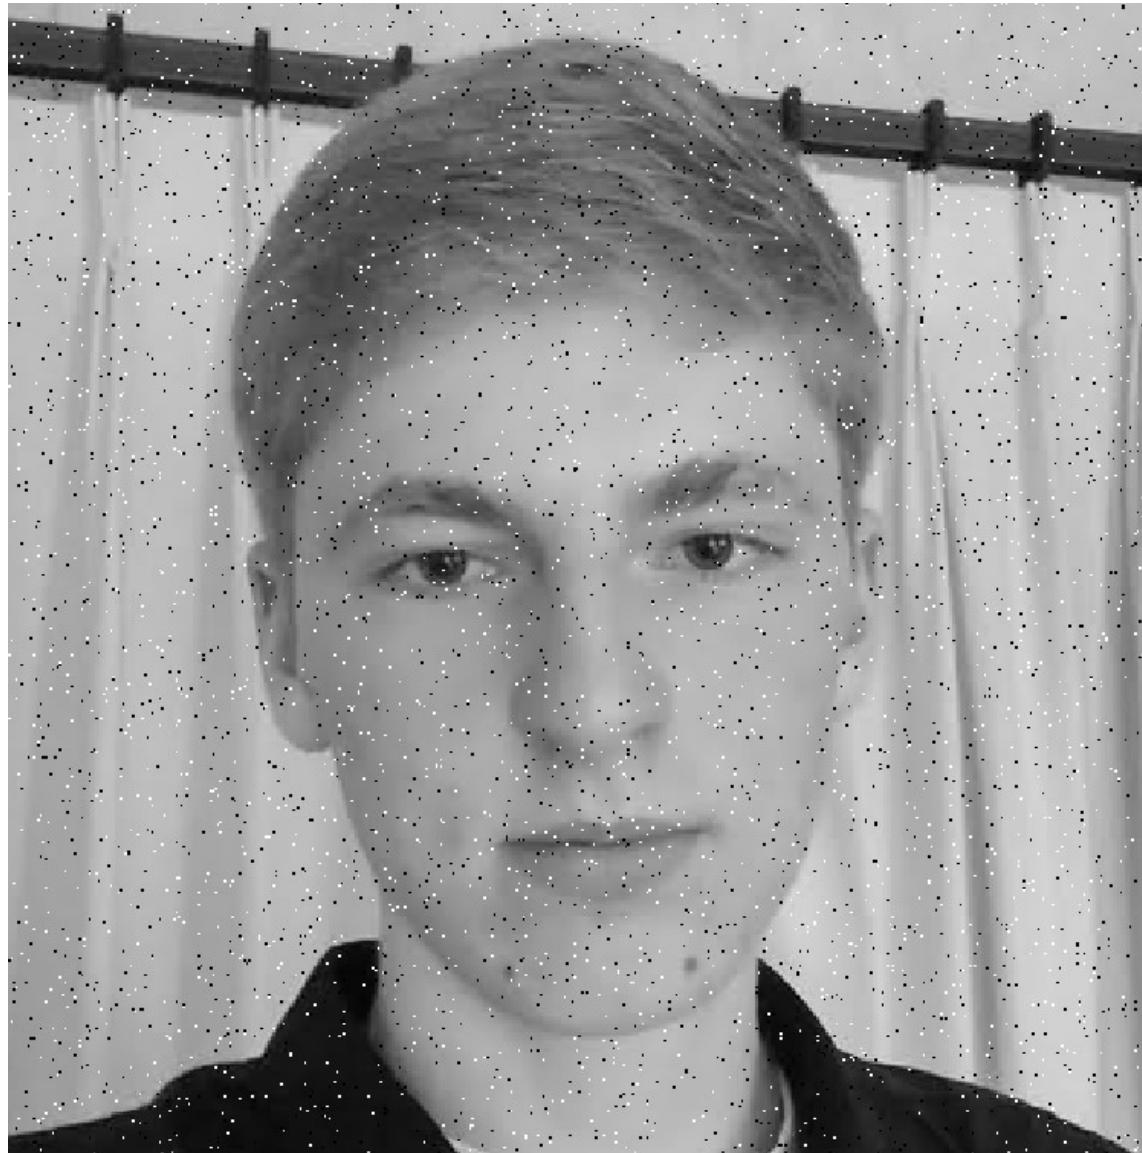
3 x 3 Mean Filter



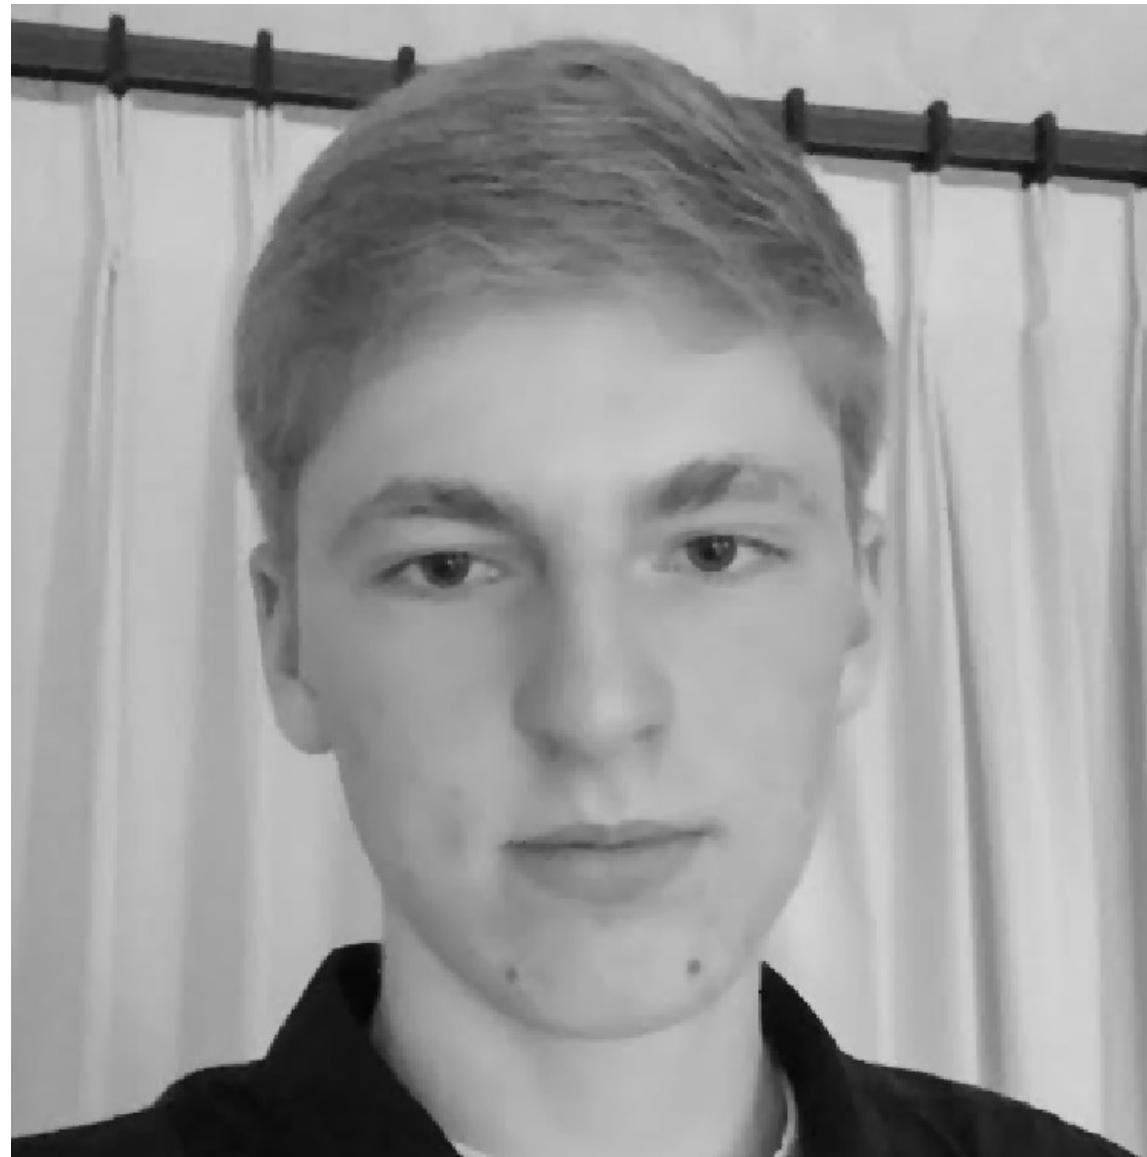
11 x 11 Mean Filter



Salt and Pepper Noise



3 x 3 Median Filter



11 x 11 Median Filter



Median filters

- Operates over a window by selecting the median intensity in the window
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?
- Interpretation: Median filtering is sorting

