

Introduction to Computer Vision

Kaveh Fathian

Assistant Professor

Computer Science Department

Colorado School of Mines

Lecture 4

Fourier Transforms & Images

Learning outcomes:

- Image downsampling
- Aliasing
- Gaussian image pyramid
- Laplacian image pyramid
- Fourier series
- Frequency domain
- Fourier transform
- Frequency-domain filtering
- Revisiting sampling



Fourier Transforms & Images

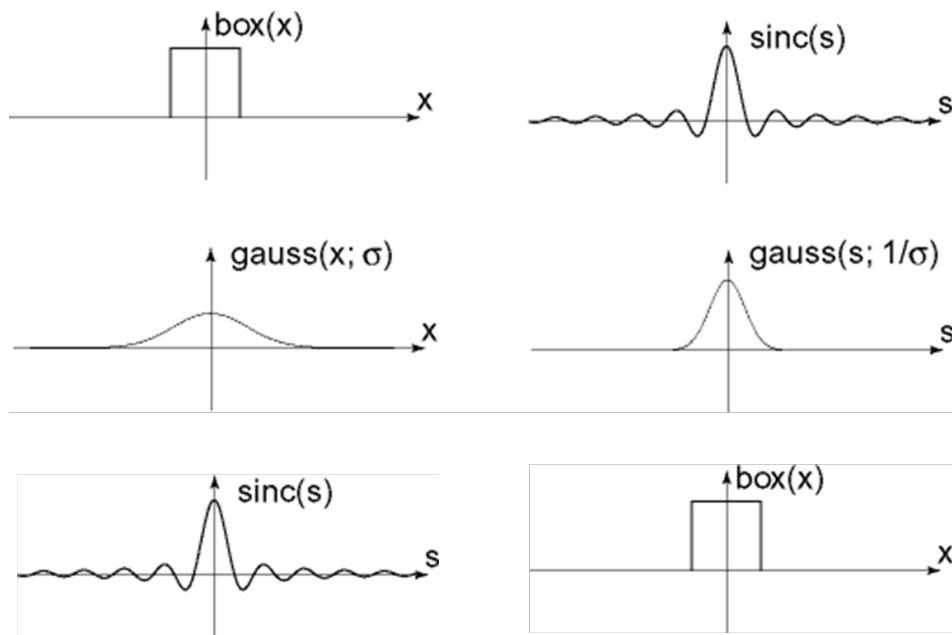
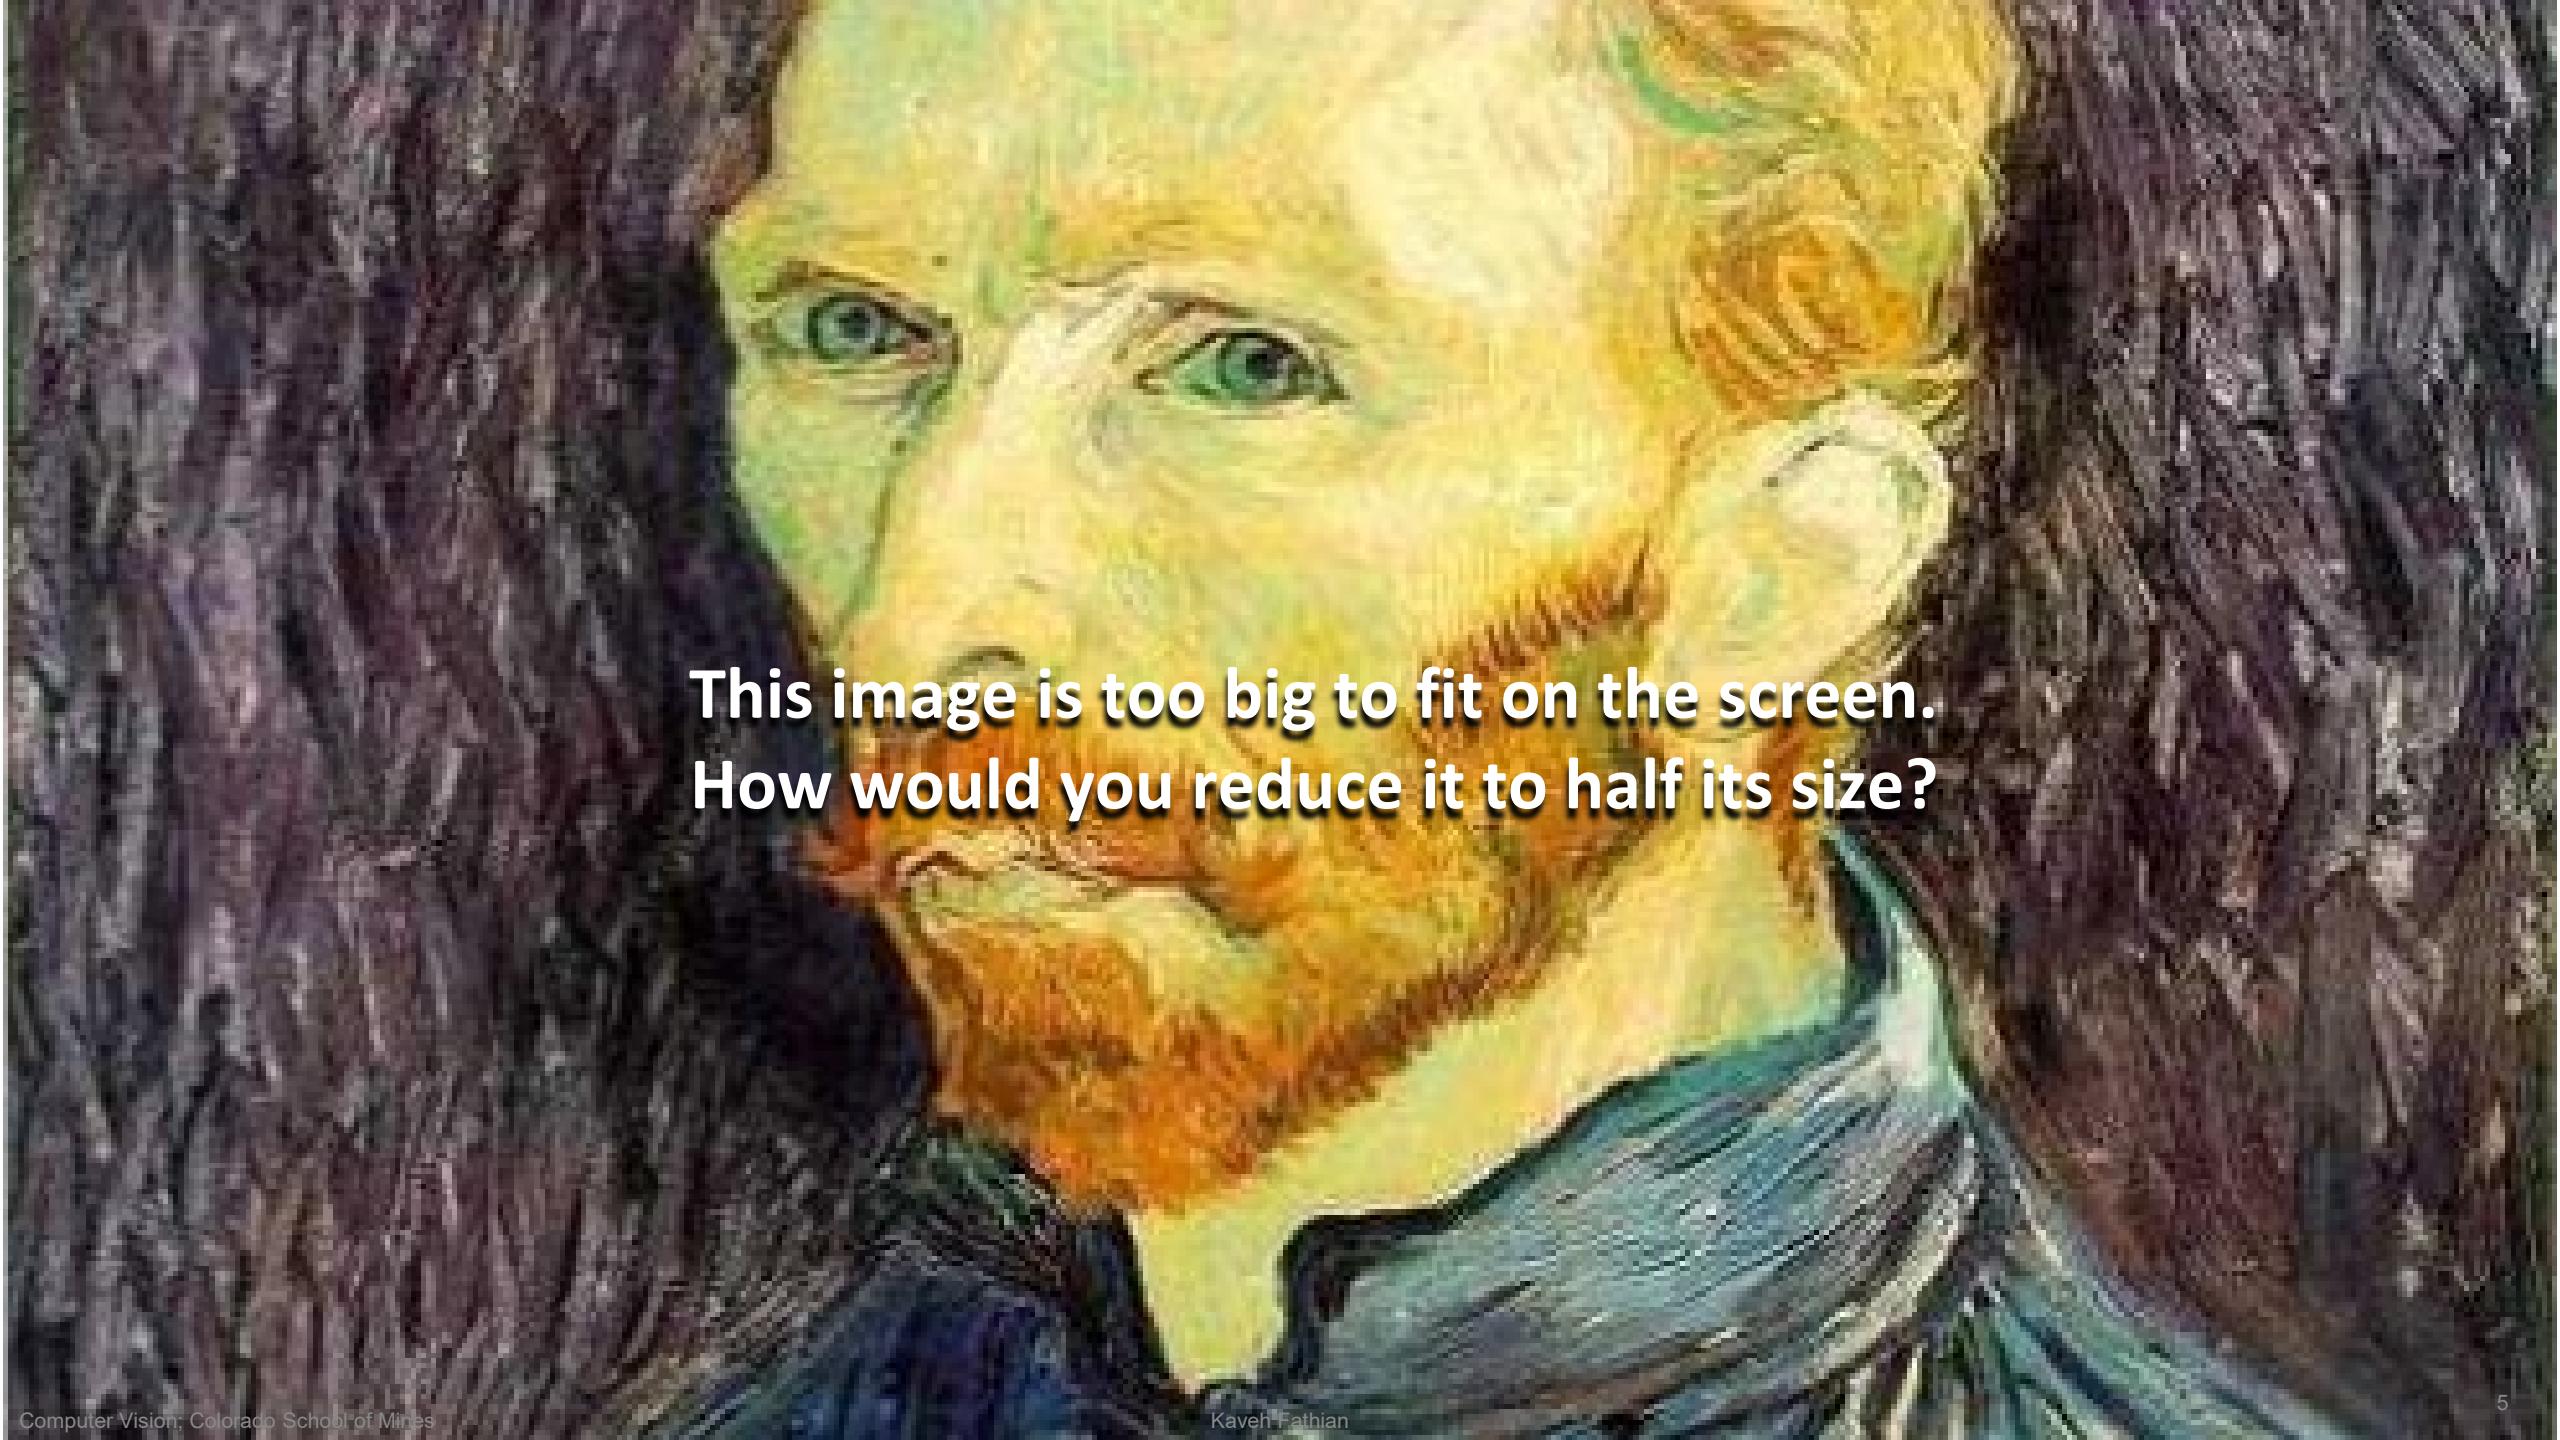
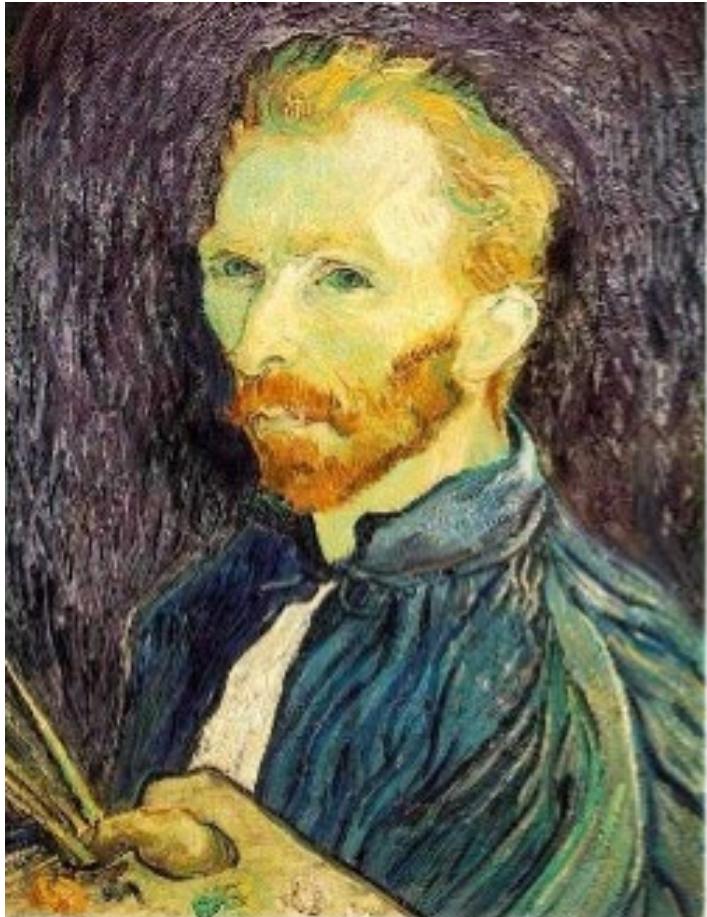


Image downsampling



**This image is too big to fit on the screen.
How would you reduce it to half its size?**

Naïve image downsampling



1/2

Throw away half the rows and columns

delete even rows
delete even columns



1/4

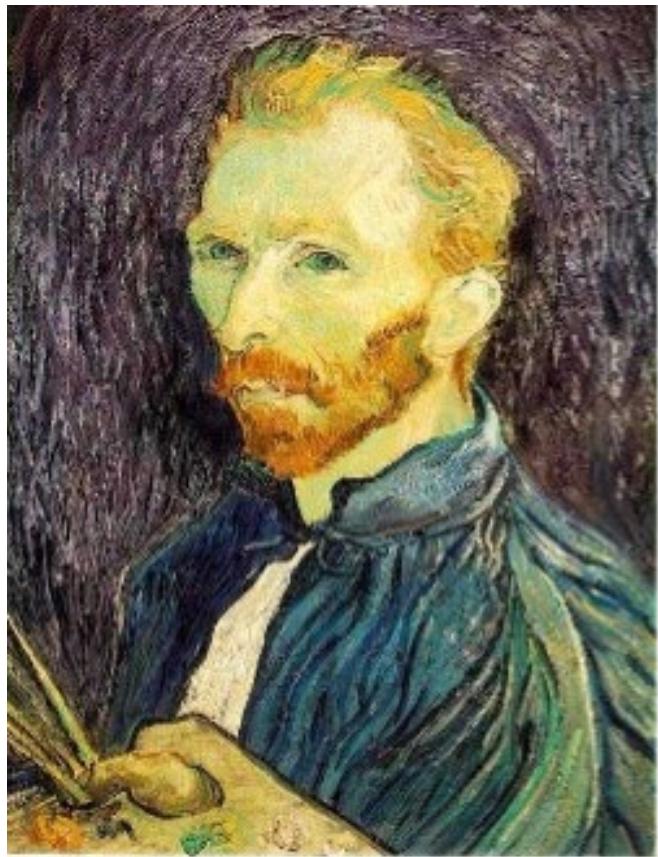
delete even rows
delete even columns



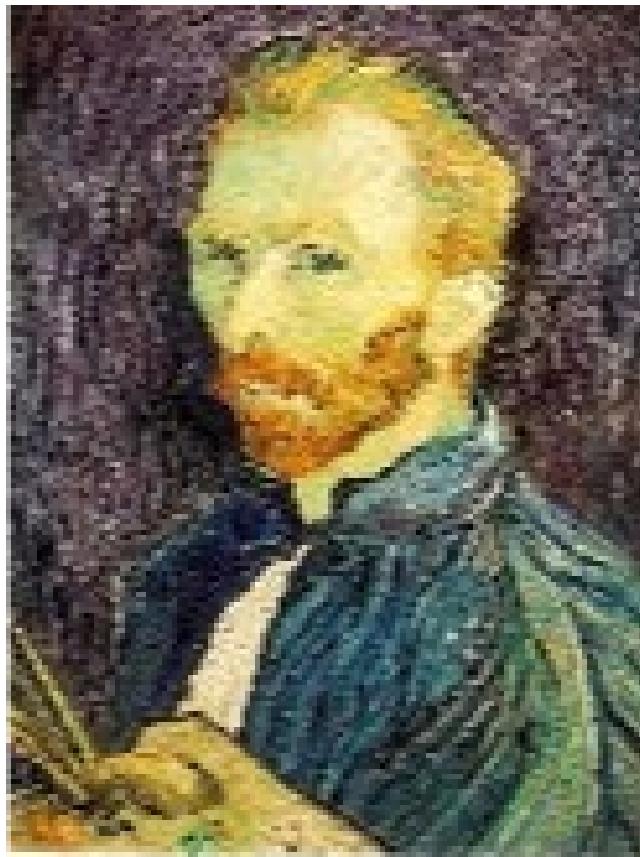
1/8

What is the problem with this approach?

Naïve image downsampling



1/2



1/4 (2x zoom)

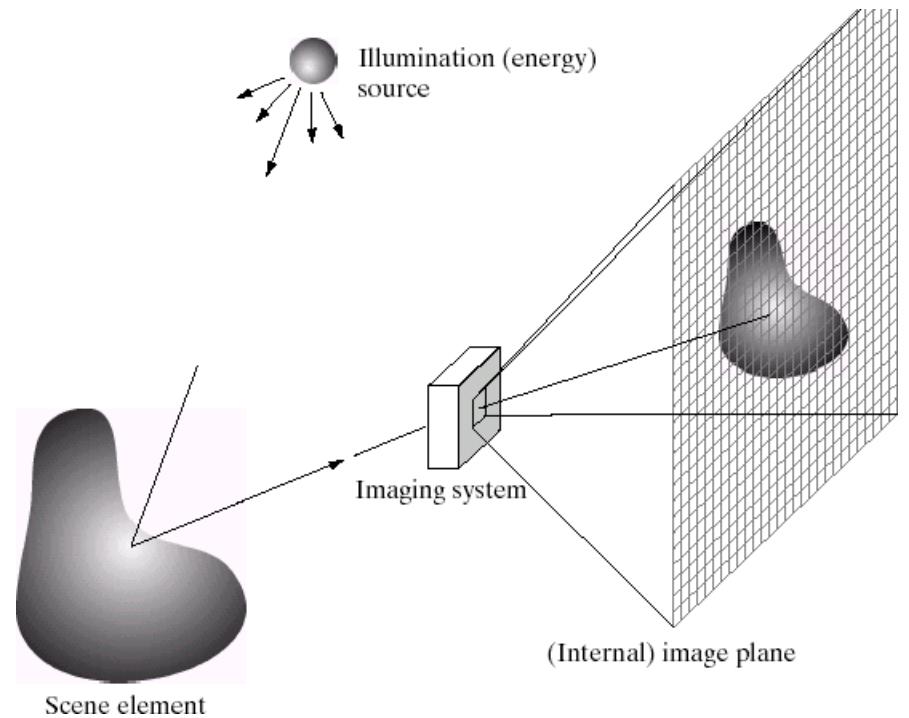


1/8 (4x zoom)

Why is the 1/8 image so pixelated (and do you know what this effect is called)?

Aliasing

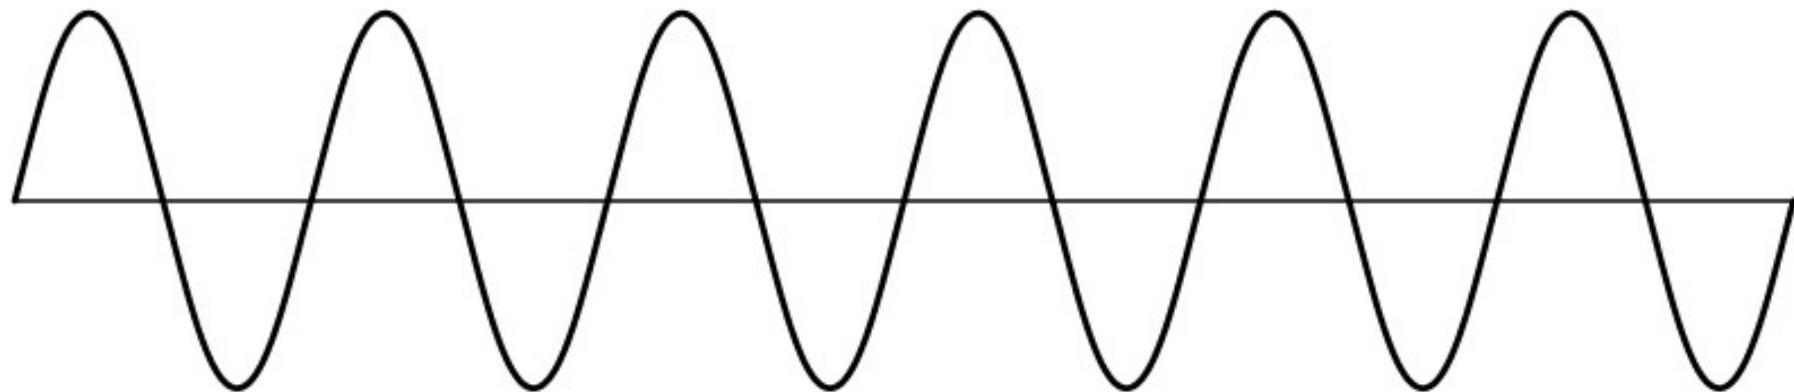
Reminder



Images are a *discrete*, or *sampled*, representation of a *continuous* world

Sampling

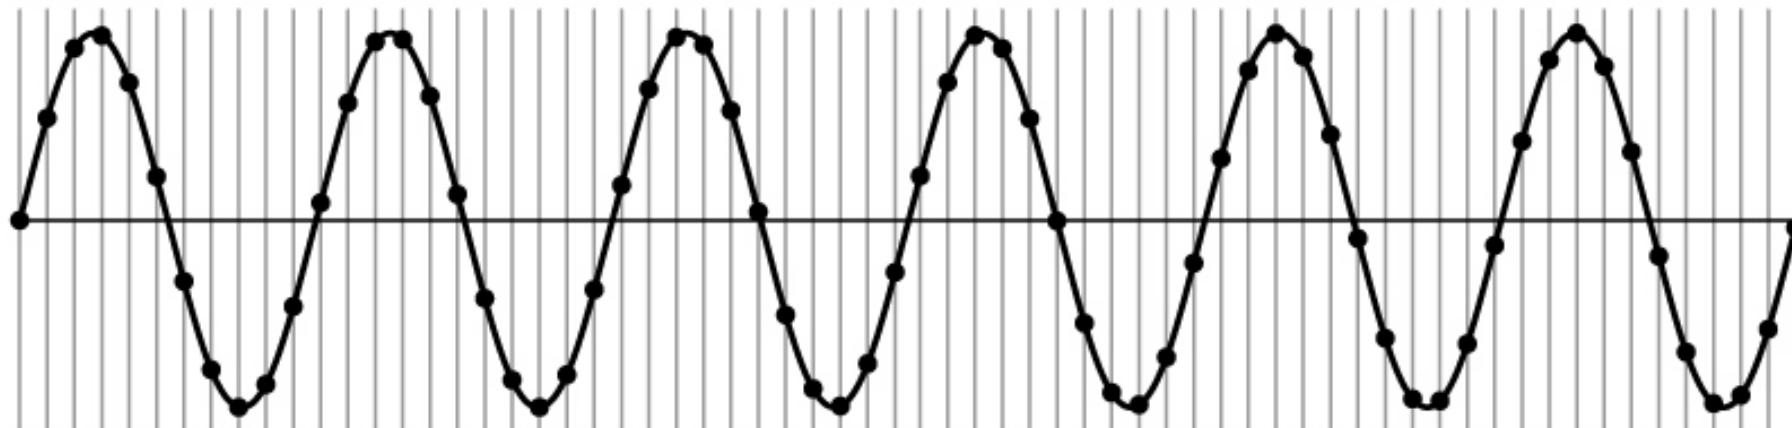
Very simple example: a sine wave



How would you discretize this signal?

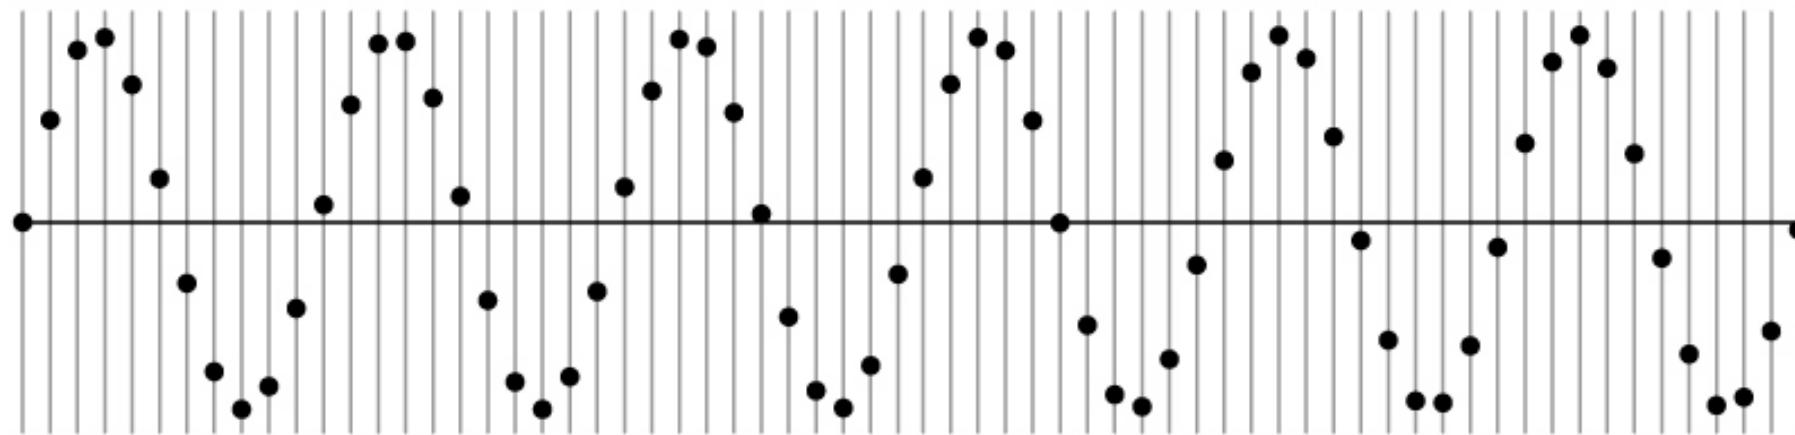
Sampling

Very simple example: a sine wave



Sampling

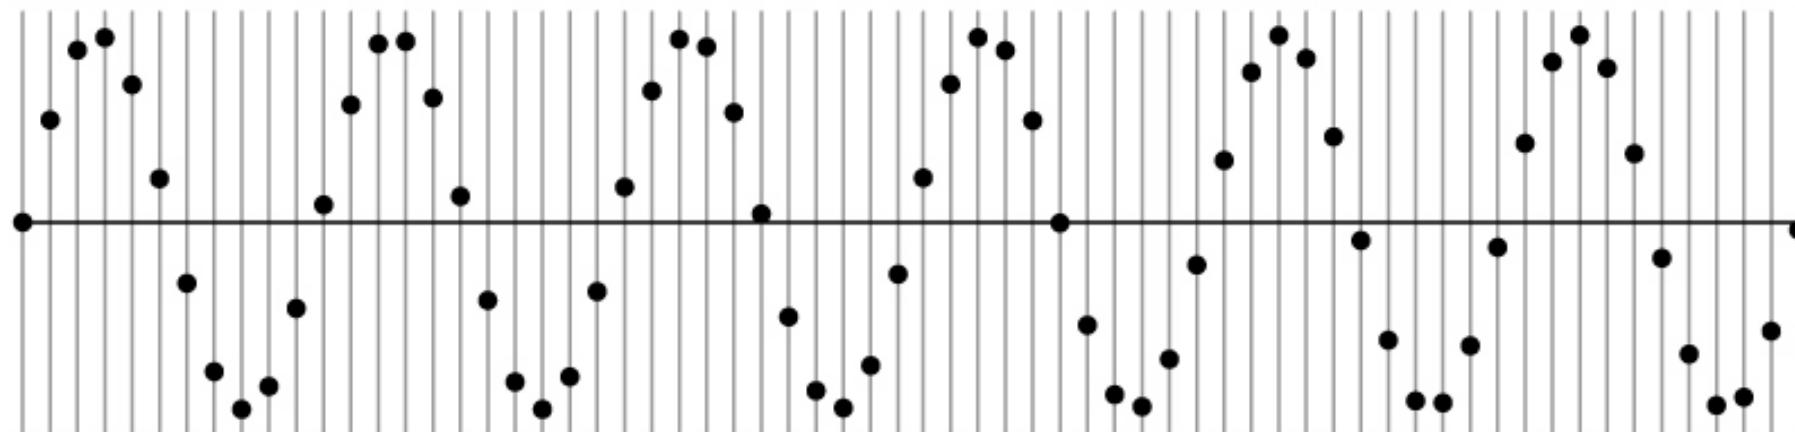
Very simple example: a sine wave



How many samples should I take?
Can I take as *many* samples as I want?

Sampling

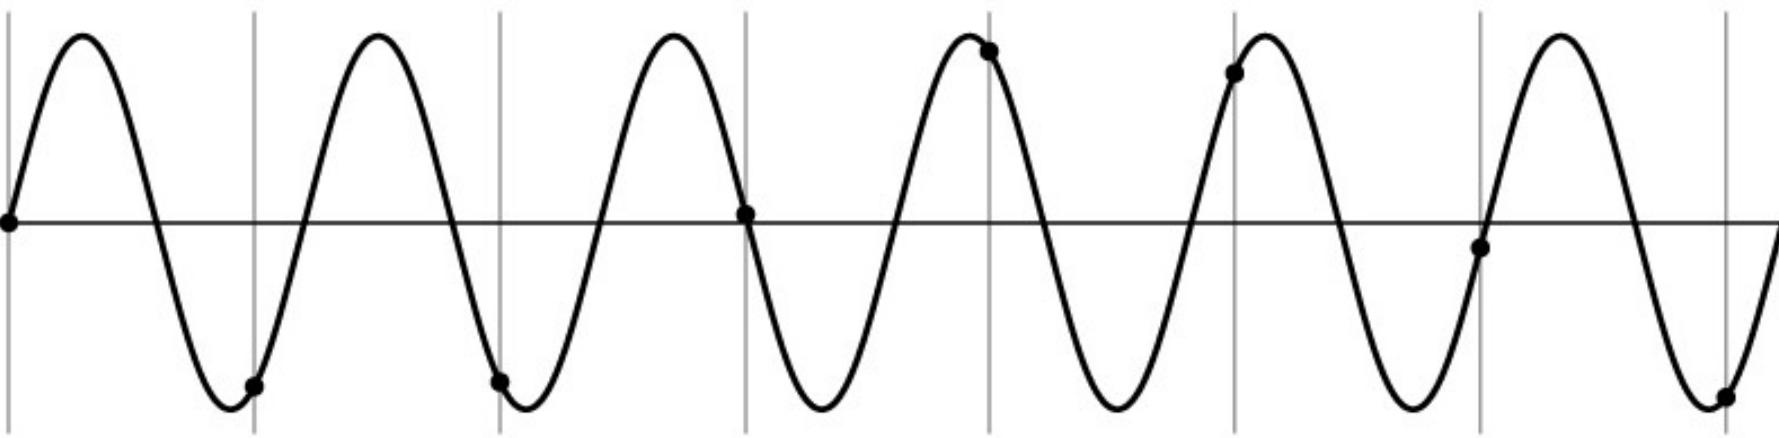
Very simple example: a sine wave



How many samples should I take?
Can I take as **few** samples as I want?

Undersampling

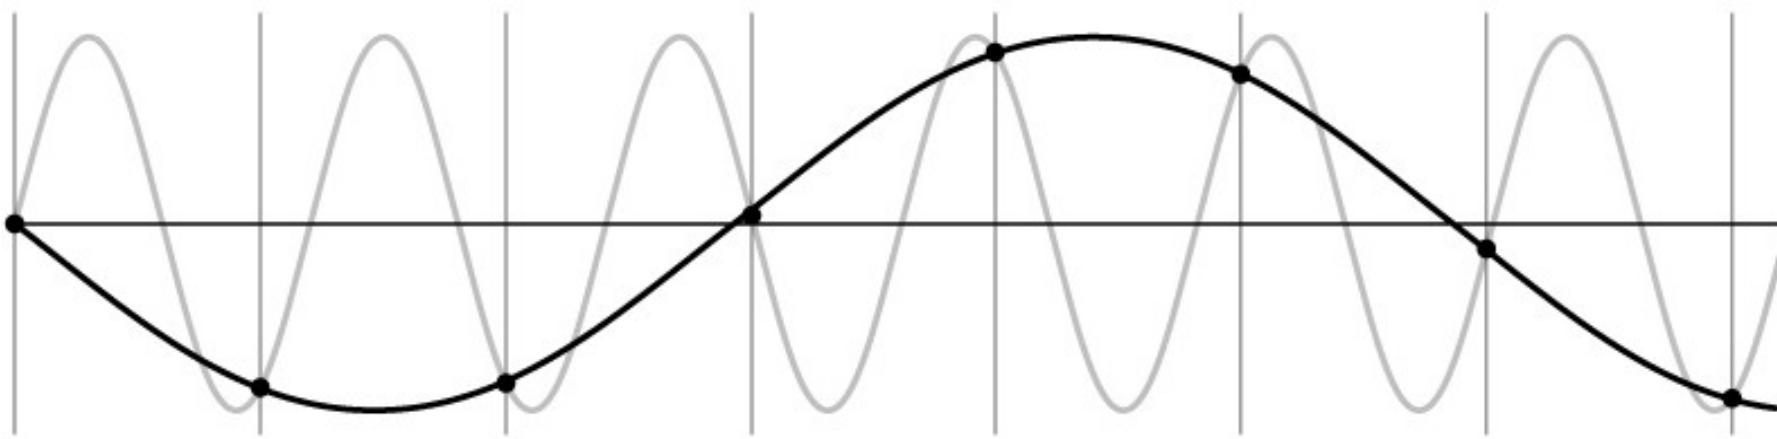
Very simple example: a sine wave



Unsurprising effect: information is lost.

Undersampling

Very simple example: a sine wave

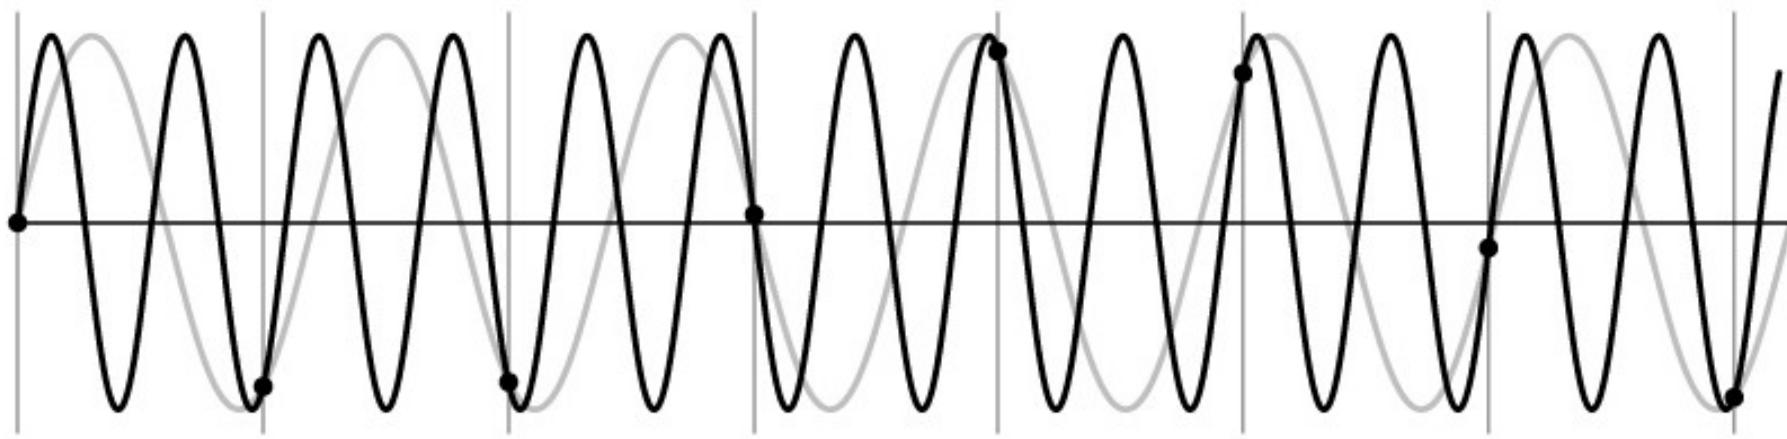


Unsurprising effect: information is lost.

Surprising effect: can confuse the signal with one of *lower* frequency.

Undersampling

Very simple example: a sine wave



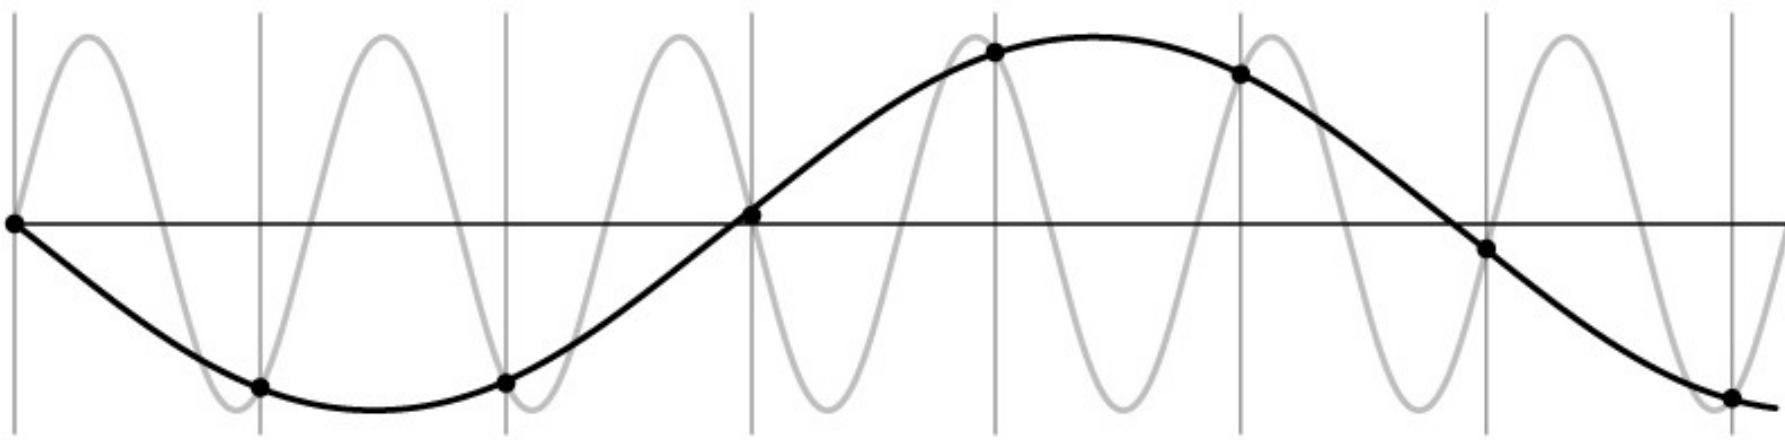
Unsurprising effect: information is lost.

Surprising effect: can confuse the signal with one of *lower* frequency.

Note: we could always confuse the signal with one of *higher* frequency.

Aliasing

Fancy term for: *Undersampling can disguise a signal as one of a lower frequency*

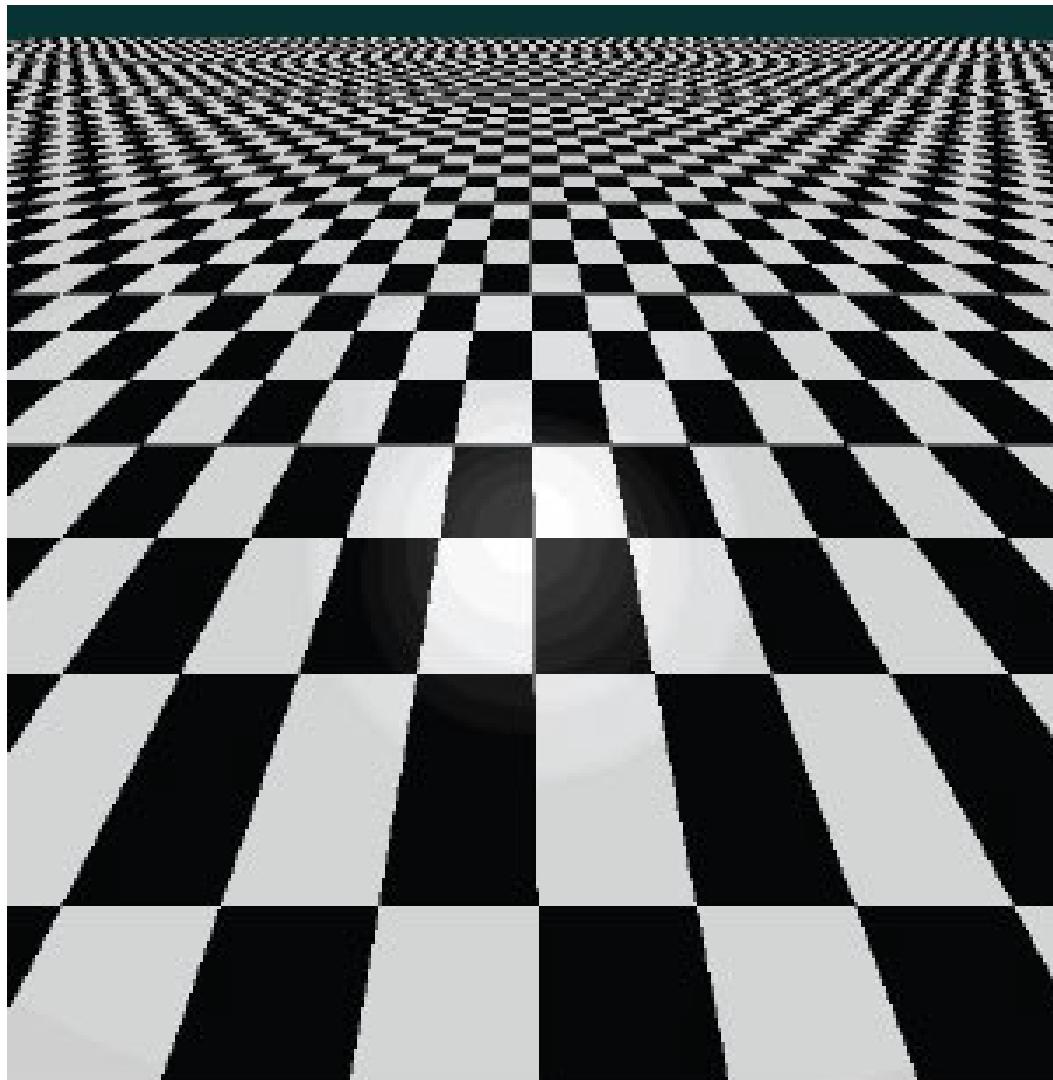


Unsurprising effect: information is lost.

Surprising effect: can confuse the signal with one of *lower* frequency.

Note: we could always confuse the signal with one of *higher* frequency.

Aliasing in textures



Aliasing in photographs

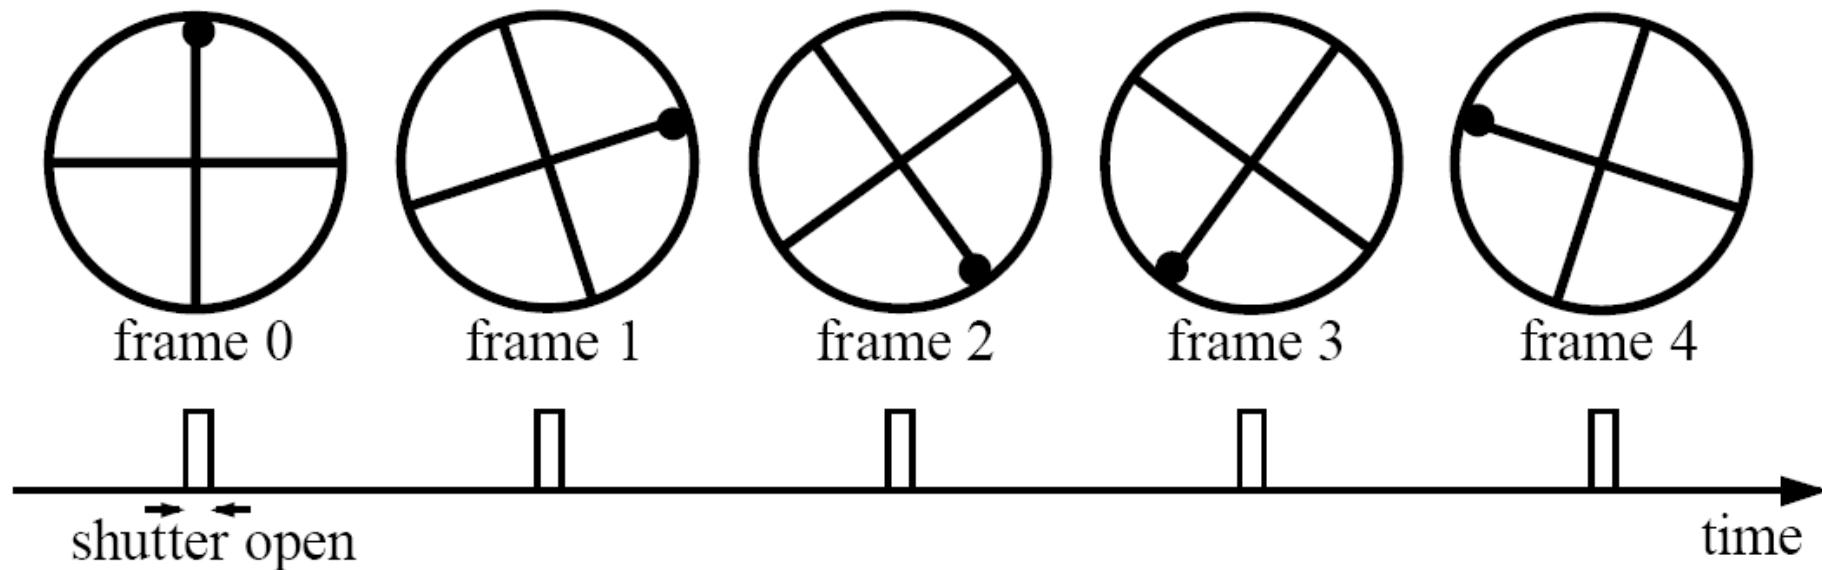
This is also known as “moire”



Temporal aliasing

Imagine a spoked wheel moving to the right (rotating clockwise).
Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):



Without dot, wheel appears to be rotating slowly backwards!
(counterclockwise)

Wagon wheel effect



https://youtu.be/yr3ngmRuGUc?si=0A5t_s5C_y0p39rB



Anti-aliasing

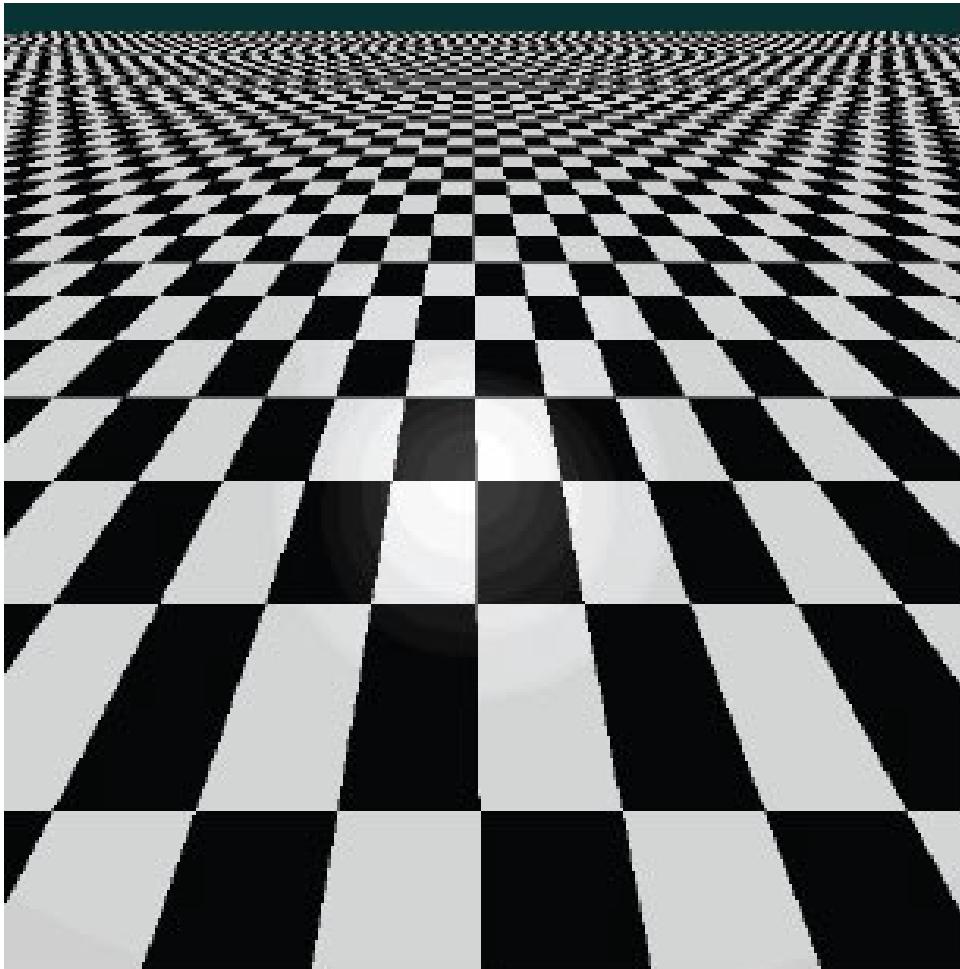
How would you deal with aliasing?

Anti-aliasing

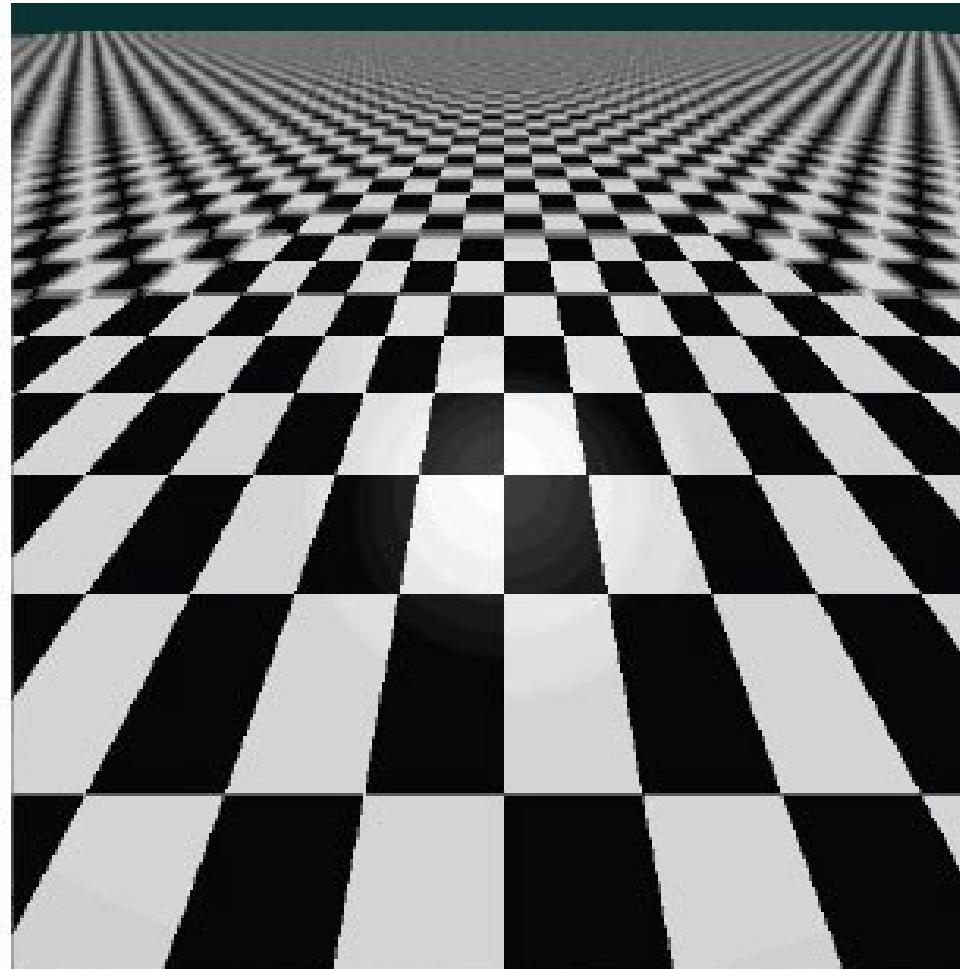
How would you deal with aliasing?

Approach 1: Oversample the signal

Anti-aliasing in textures



aliasing artifacts



anti-aliasing by oversampling

Anti-aliasing

How would you deal with aliasing?

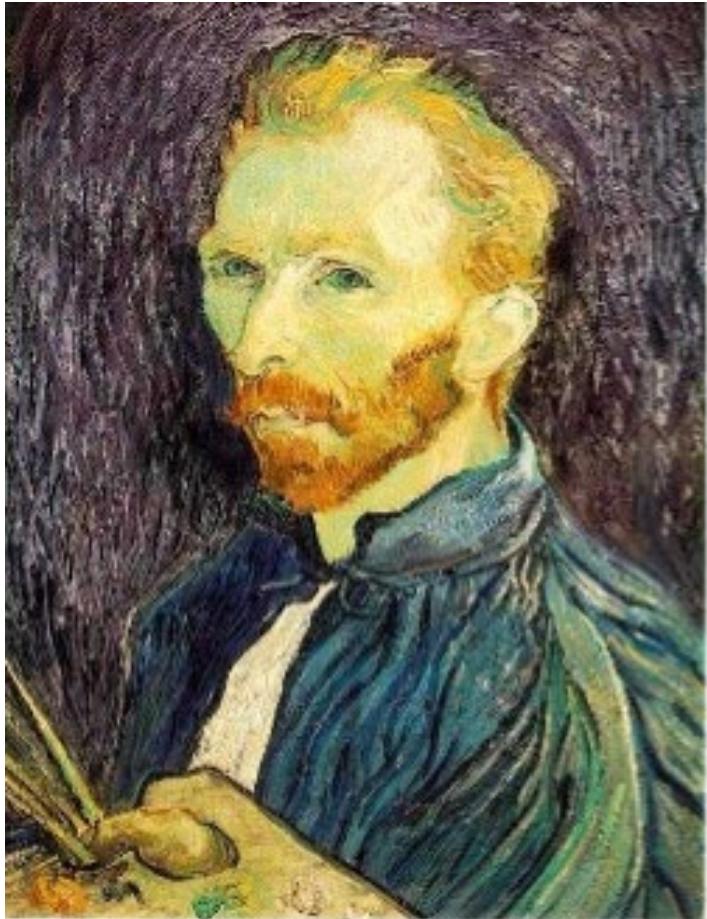
Approach 1: Oversample the signal

Approach 2: Smooth the signal

- Remove some of the detail effects that cause aliasing.
- Lose information, but better than aliasing artifacts.

How would you smooth a signal?

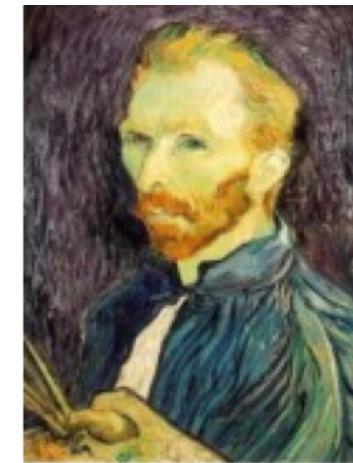
Better image downsampling



1/2

Apply a smoothing filter first, then throw away half the rows and columns

Gaussian filter
delete even rows
delete even columns



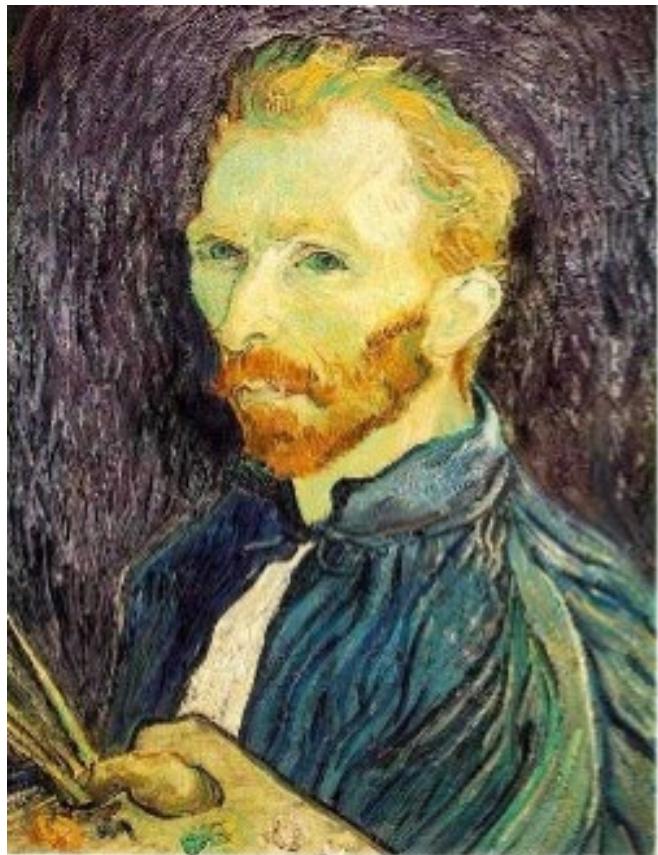
1/4

Gaussian filter
delete even rows
delete even columns

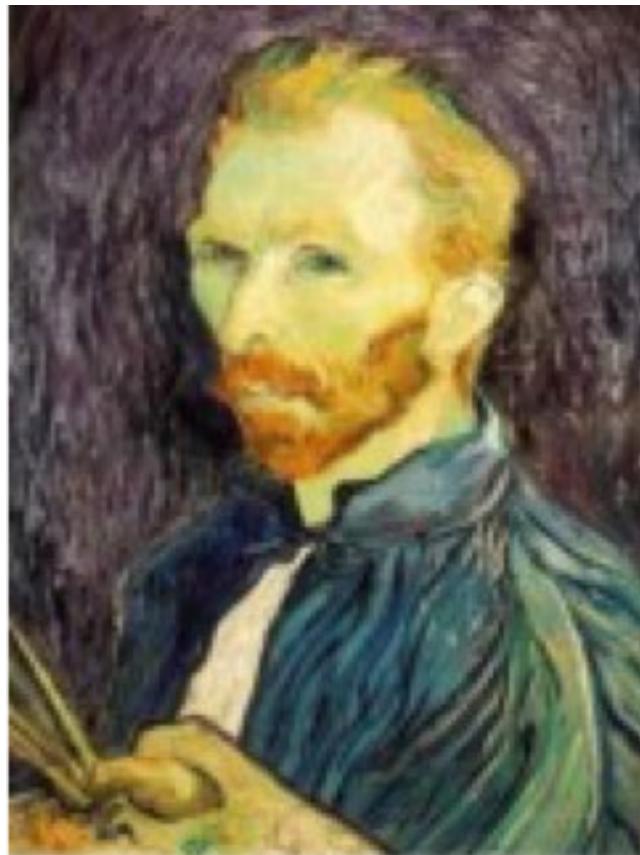


1/8

Better image downsampling



1/2

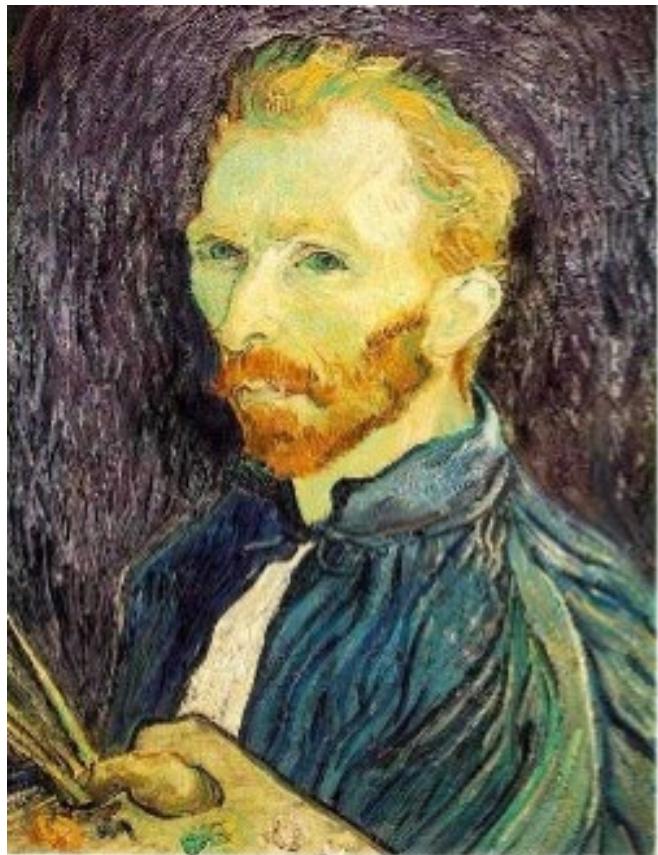


1/4 (2x zoom)

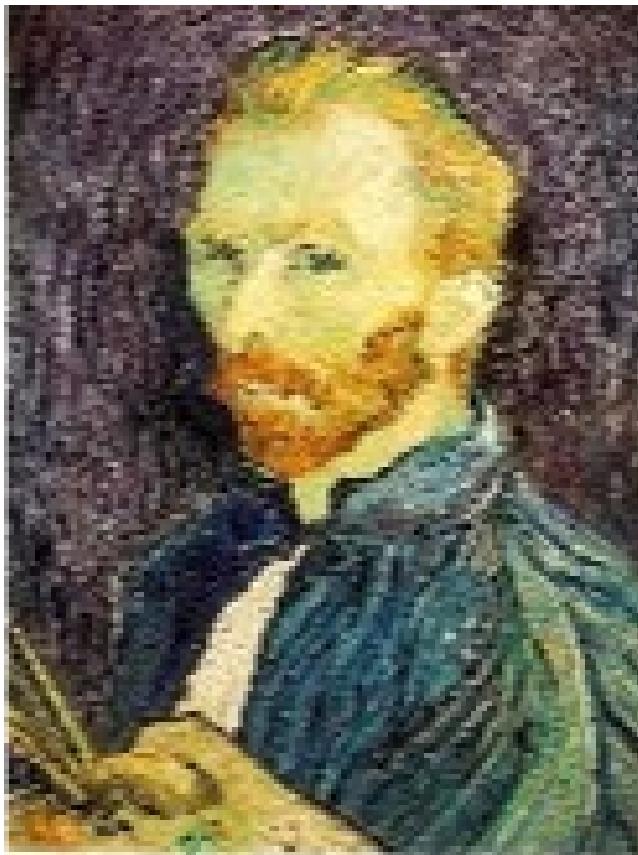


1/8 (4x zoom)

Naïve image downsampling



1/2



1/4 (2x zoom)



1/8 (4x zoom)

Anti-aliasing

Question 1: How much smoothing do I need to do to avoid aliasing?

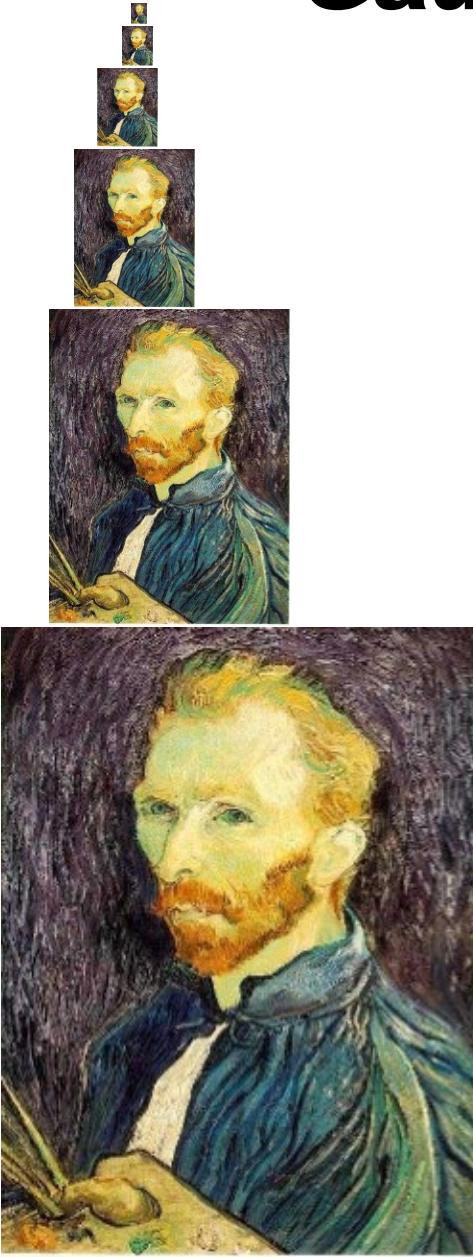
Question 2: How many samples do I need to take to avoid aliasing?

Answer to both: Enough to reach the Nyquist limit

We'll see what this means soon.

Gaussian image pyramid

Gaussian image pyramid

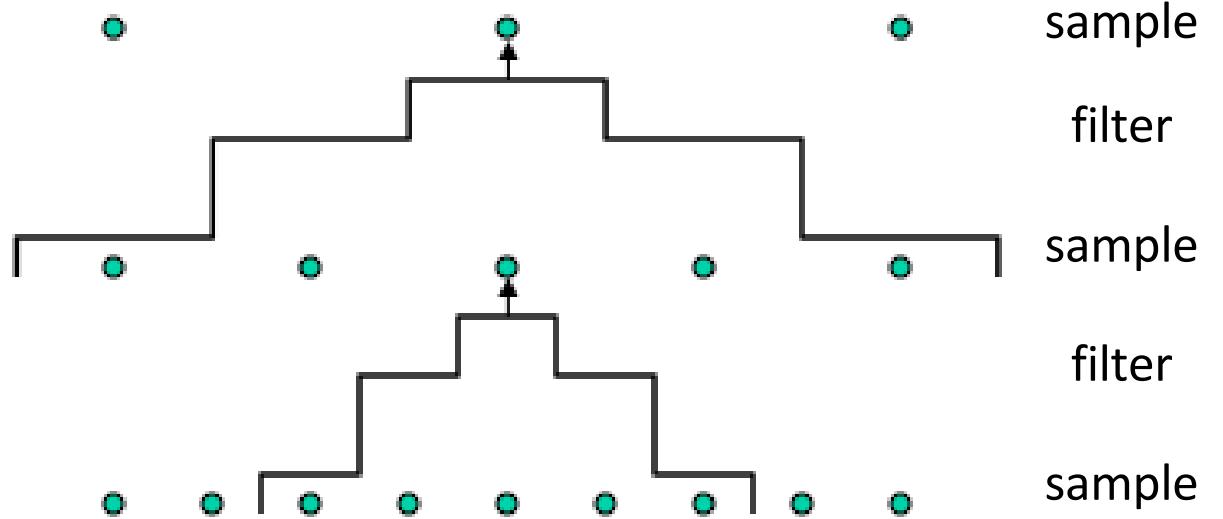


The name of this sequence of subsampled images

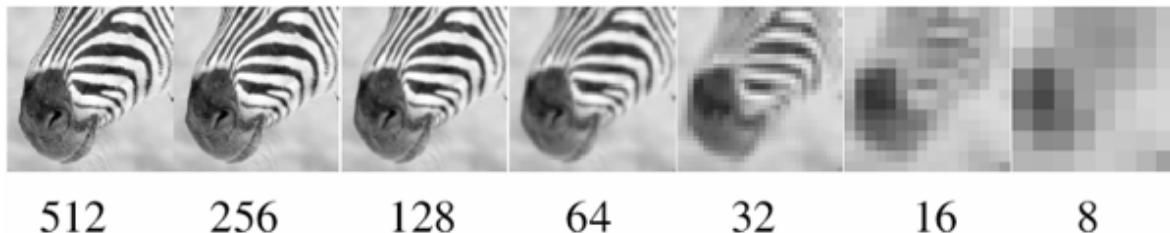
Constructing a Gaussian pyramid

Algorithm

```
repeat:  
    filter  
    subsample  
until min resolution reached
```



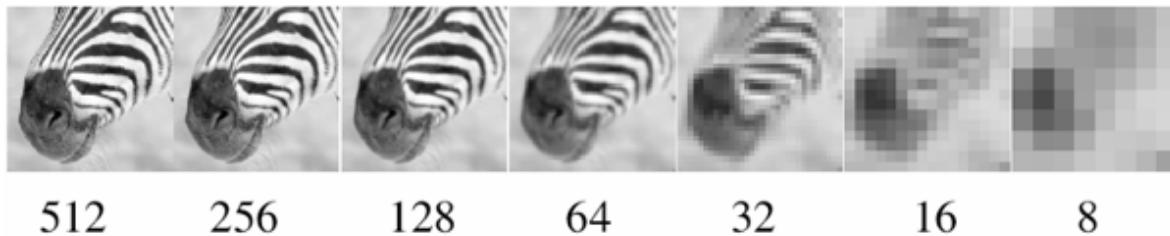
Some properties of the Gaussian pyramid



What happens to the details of the image?



Some properties of the Gaussian pyramid



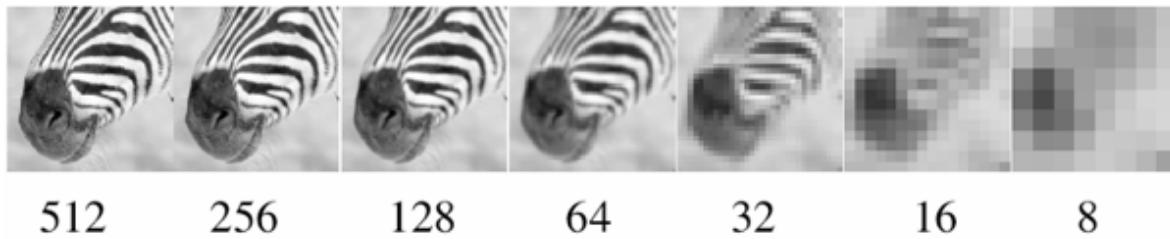
What happens to the details of the image?

- They get smoothed out as we move to higher levels.

What is preserved at the higher levels?



Some properties of the Gaussian pyramid



What happens to the details of the image?

- They get smoothed out as we move to higher levels.

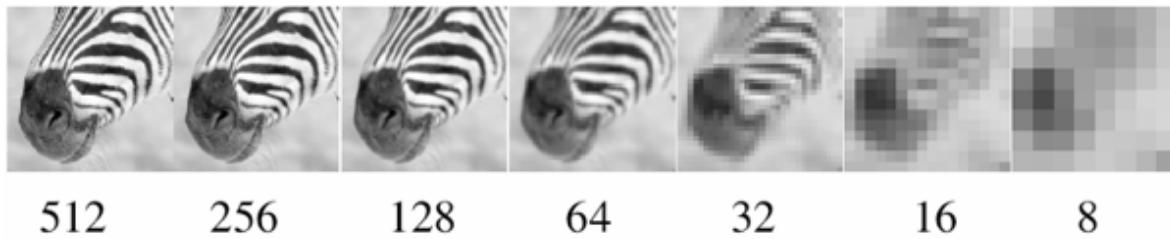


What is preserved at the higher levels?

- Mostly large uniform regions in the original image.

How would you reconstruct the original image from the image at the upper level?

Some properties of the Gaussian pyramid



What happens to the details of the image?

- They get smoothed out as we move to higher levels.

What is preserved at the higher levels?

- Mostly large uniform regions in the original image.

How would you reconstruct the original image from the image at the upper level?

- That's not possible.

Blurring is lossy



level 0



level 1 (before downsampling)

What does the residual look like?

Blurring is lossy



level 0



level 1 (before downsampling)

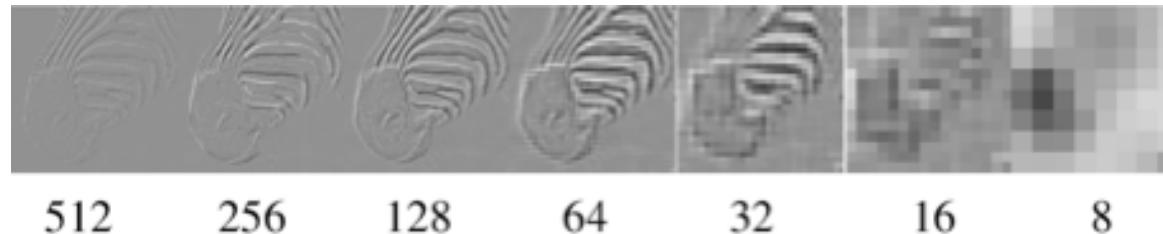


residual

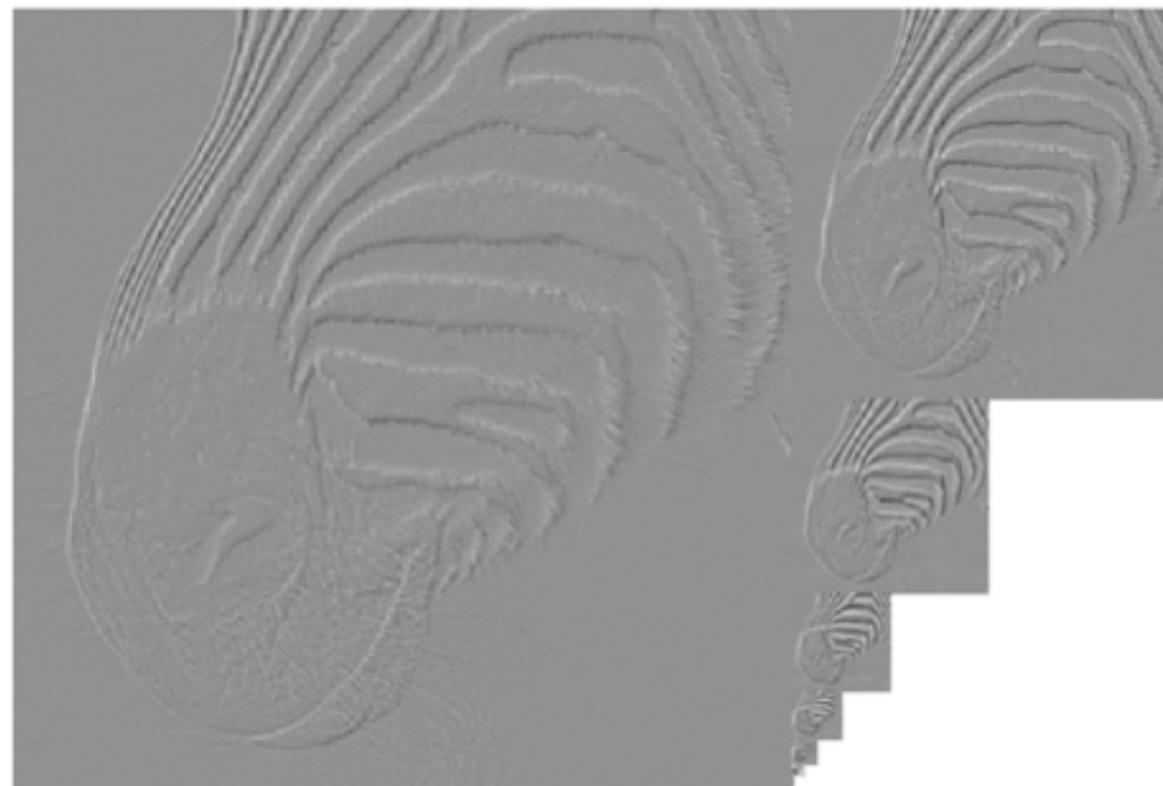
Can we make a pyramid that is lossless?

Laplacian image pyramid

Laplacian image pyramid

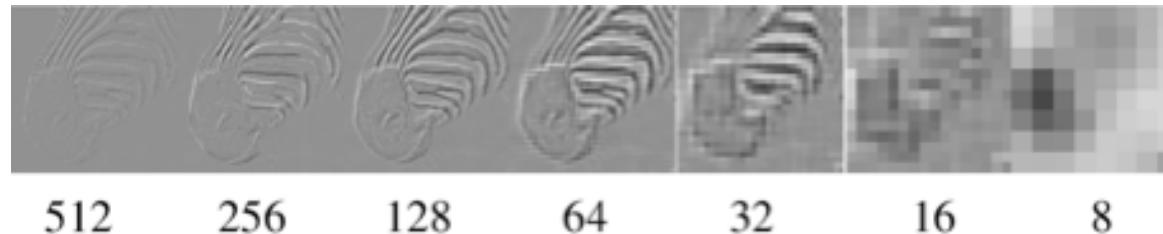


At each level, retain the residuals instead of the blurred images themselves.

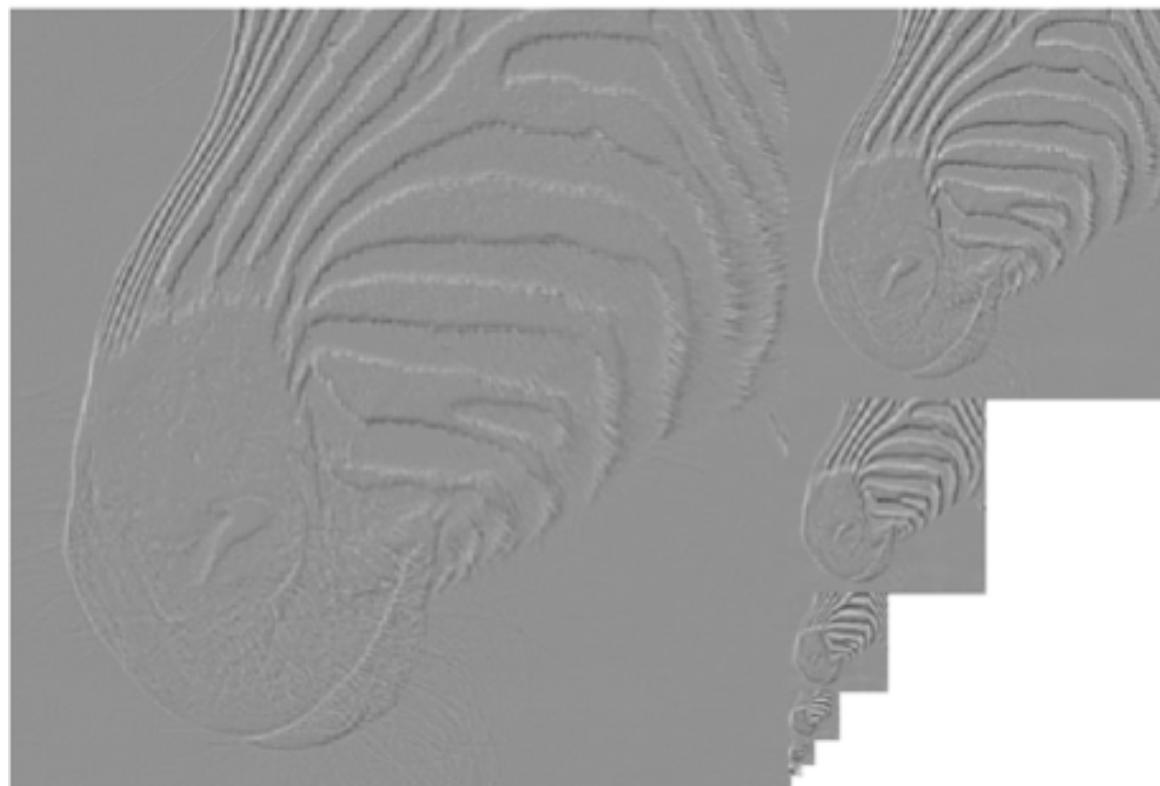


Can we reconstruct the original image using the pyramid?

Laplacian image pyramid



At each level, retain the residuals instead of the blurred images themselves.



Can we reconstruct the original image using the pyramid?

- Yes we can!

What do we need to store to be able to reconstruct the original image?

Let's start by looking at just one level



level 0



level 1 (upsampled)



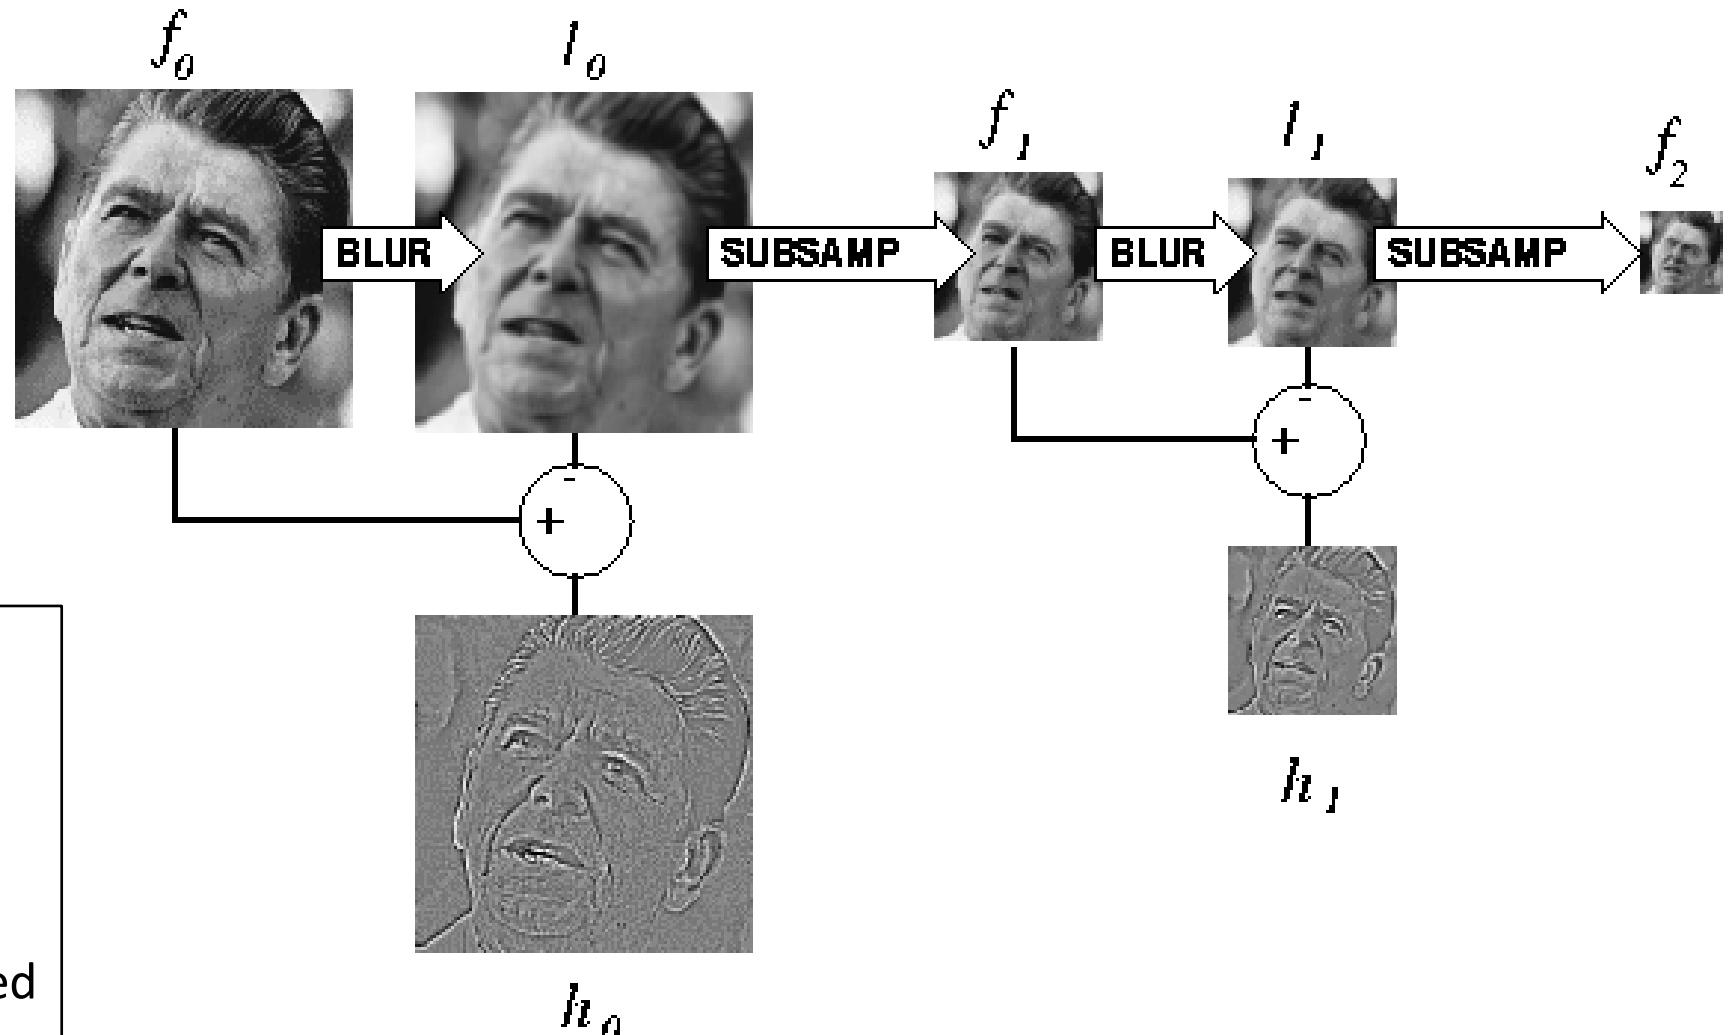
residual

Does this mean we need to store both residuals and the blurred copies of the original?

Constructing a Laplacian pyramid

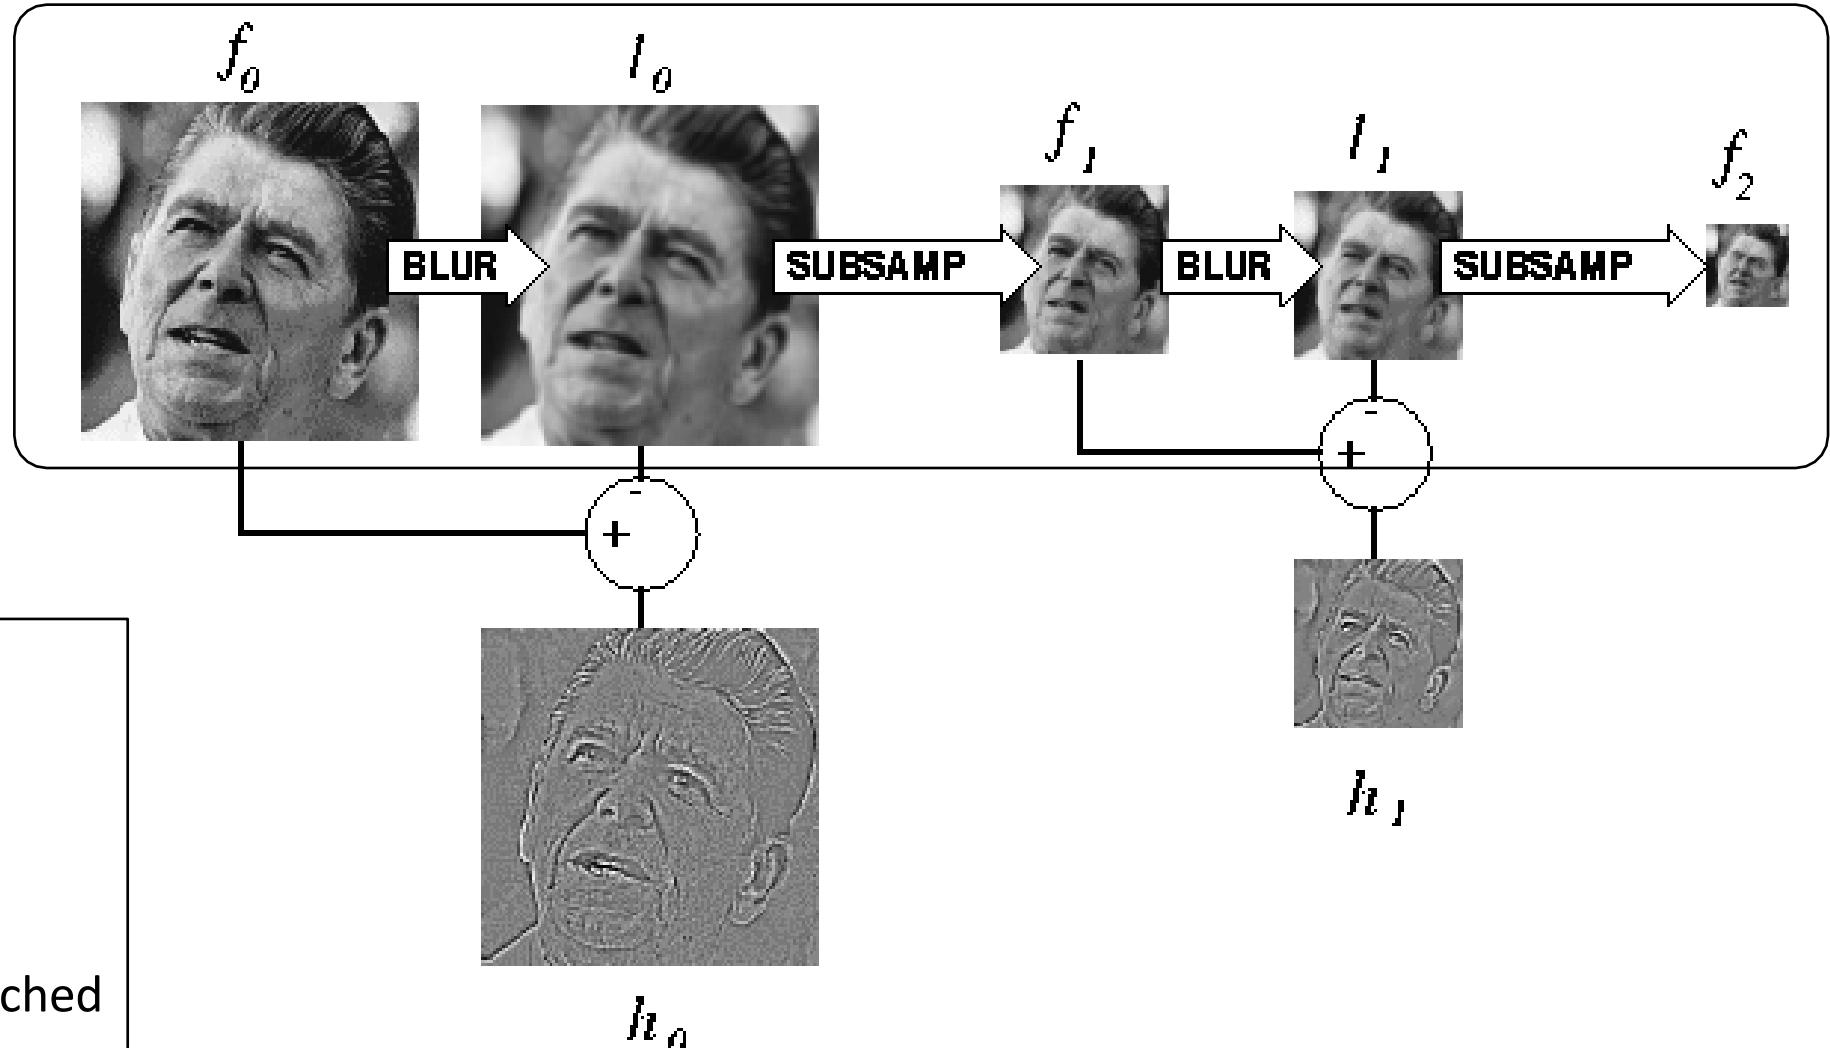
Algorithm

```
repeat:  
    filter  
    compute residual  
    subsample  
  
until min resolution reached
```



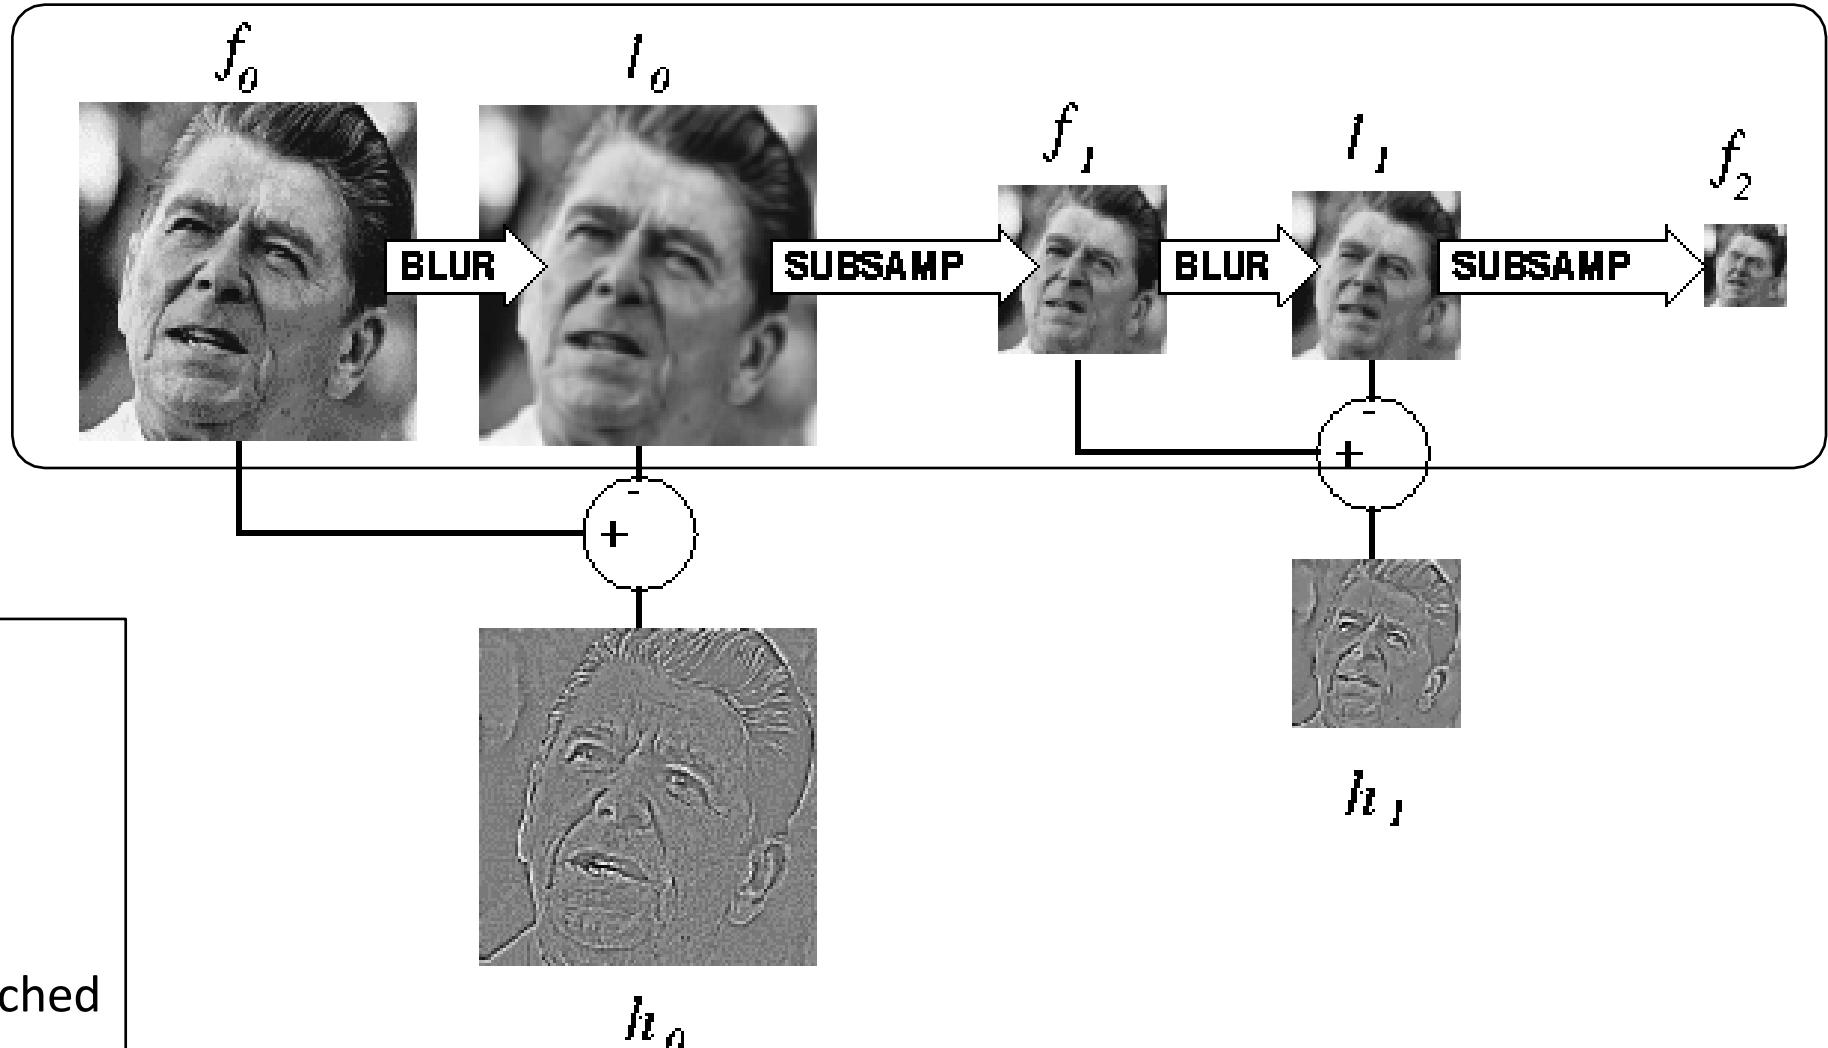
Constructing a Laplacian pyramid

What is this part?



Constructing a Laplacian pyramid

It's a Gaussian pyramid.



Algorithm

```
repeat:  
    filter  
    compute residual  
    subsample  
until min resolution reached
```

What do we need to construct the original image?



What do we need to construct the original image?



h_1

(1) residuals



h_o



What do we need to construct the original image?

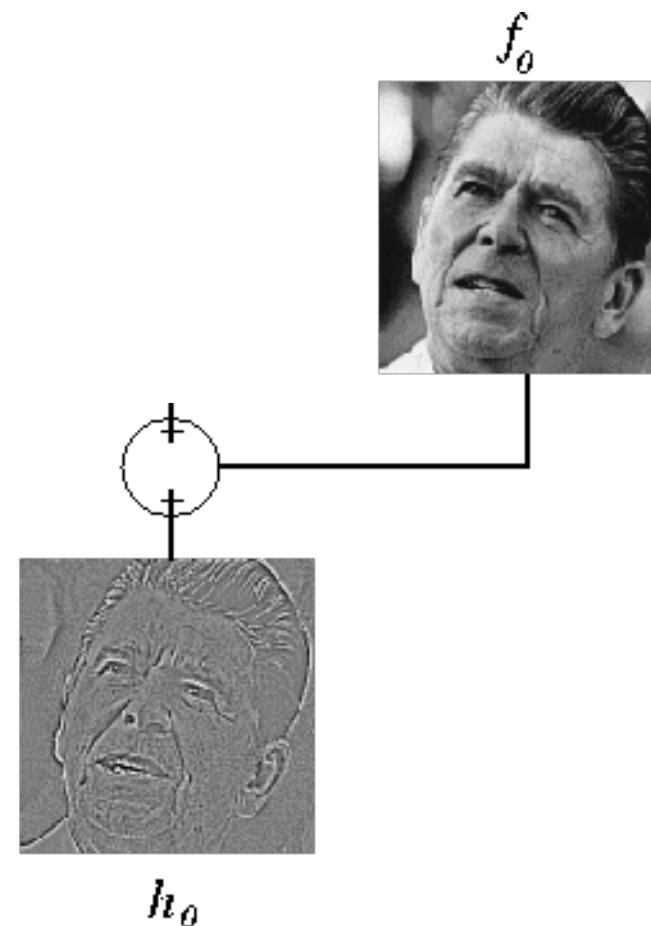
(2) smallest
image

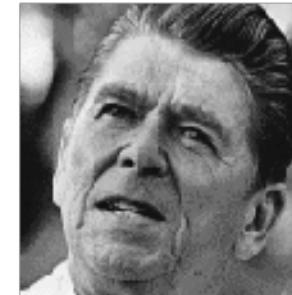
$$f_2$$




$$h_1$$


(1) residuals

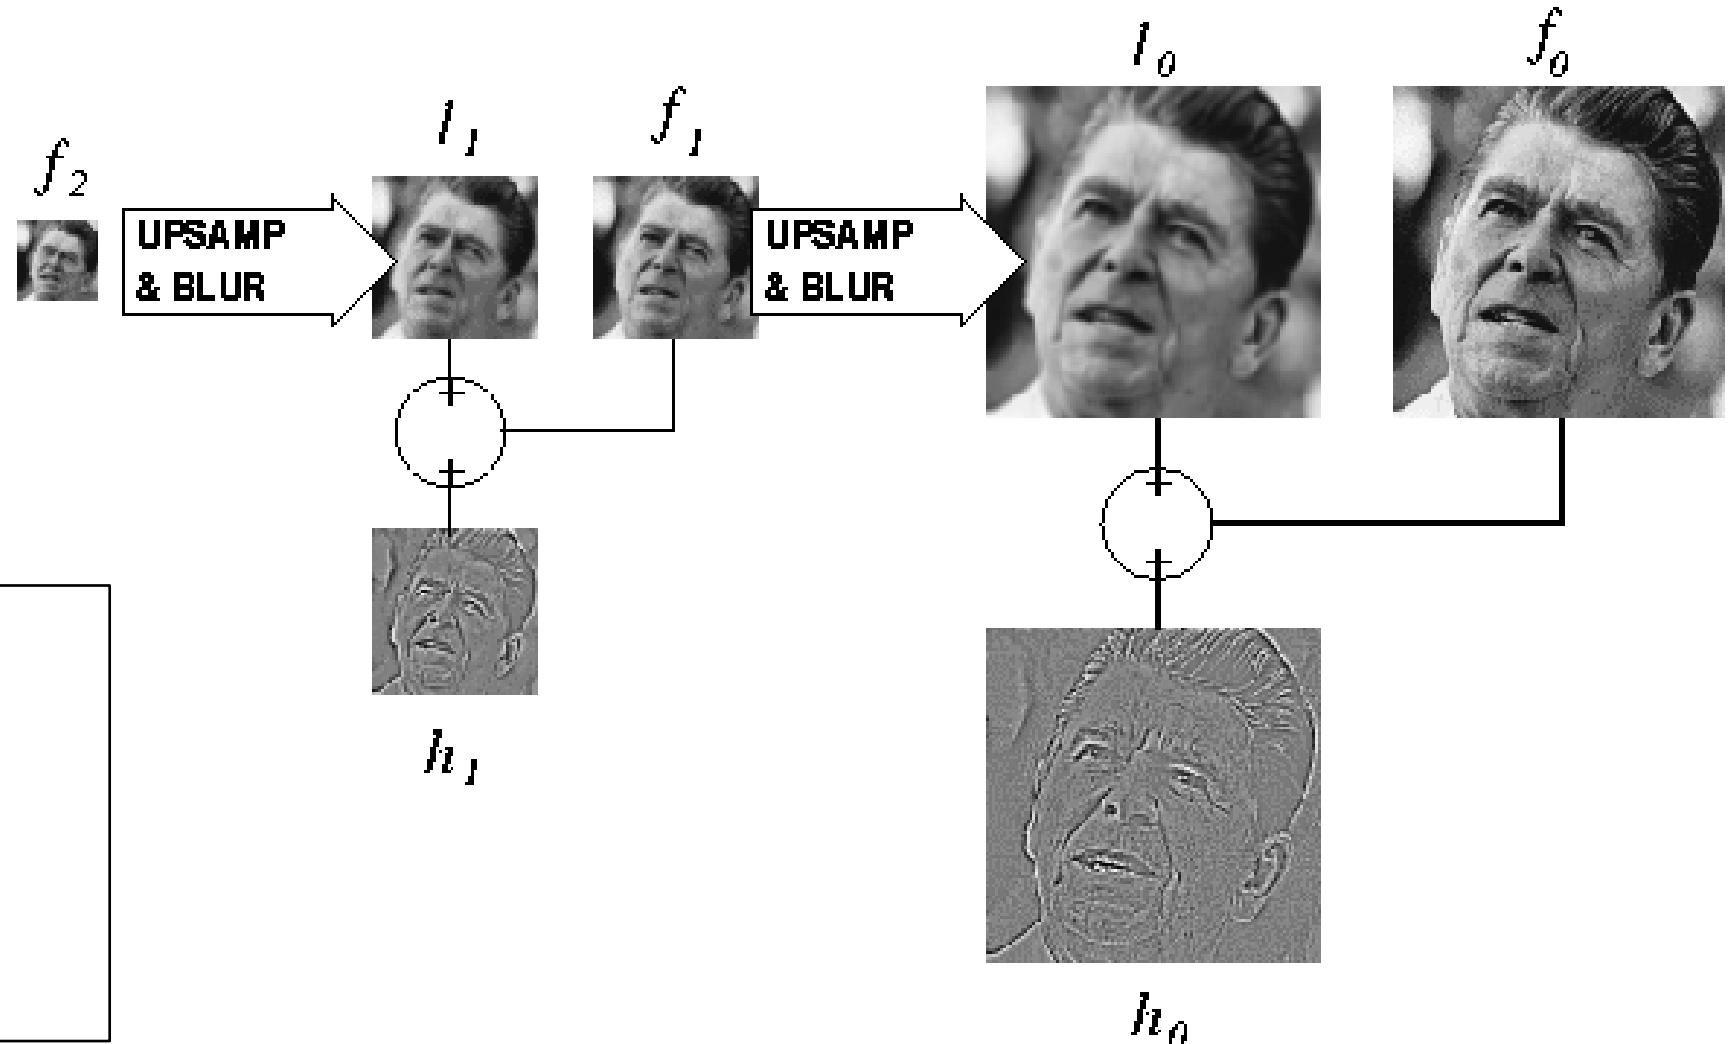


$$h_0$$


Reconstructing the original image

Algorithm

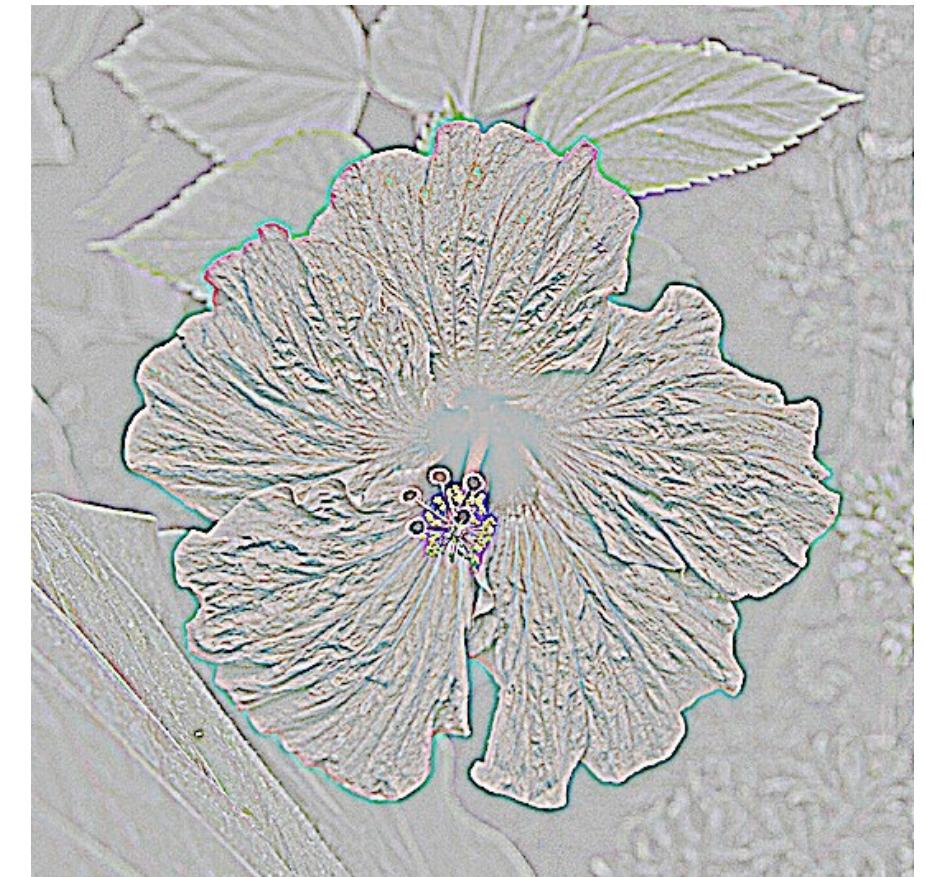
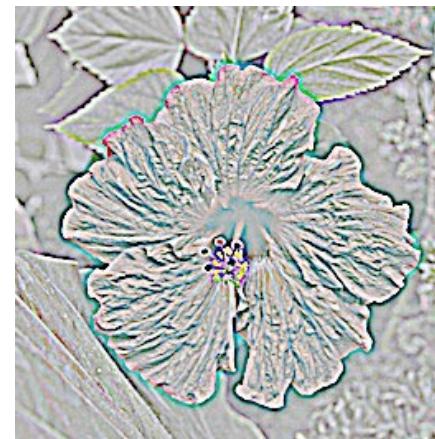
```
repeat:  
    upsample  
    sum with residual  
until orig resolution  
reached
```



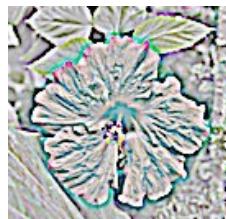
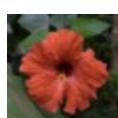
Gaussian vs Laplacian Pyramid



Shown in opposite
order for space.

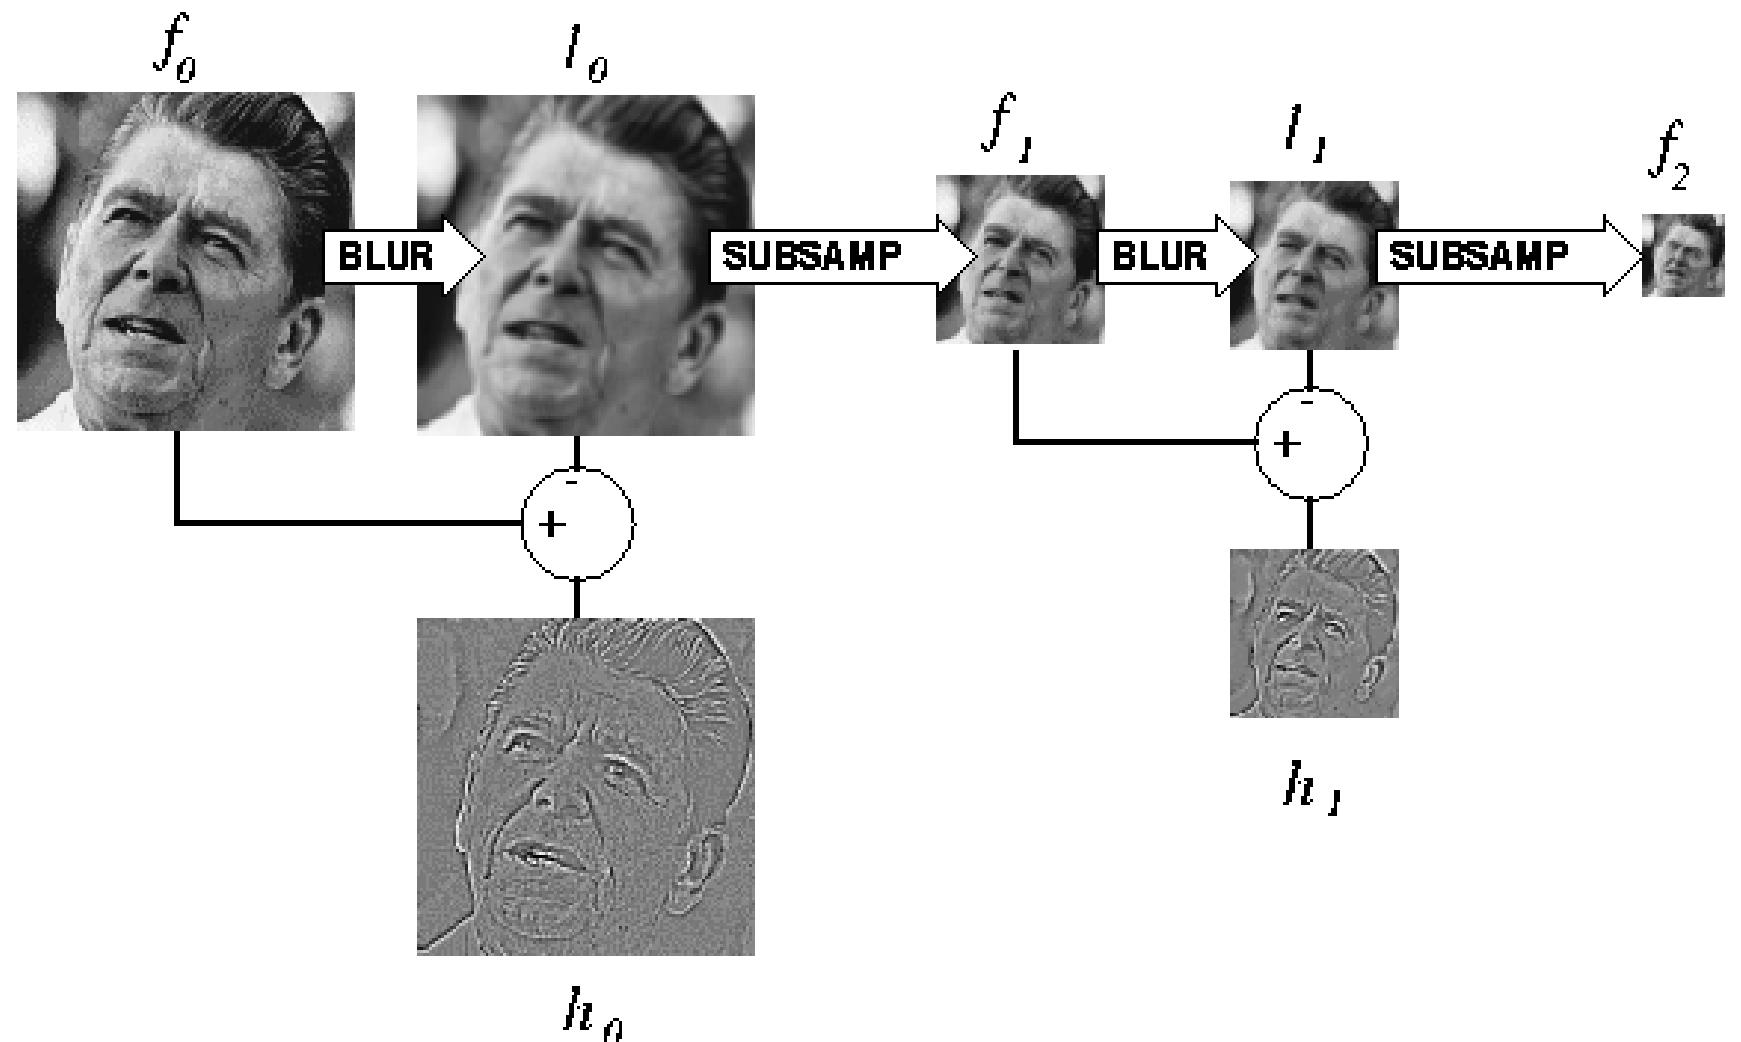


Which one takes
more space to store?



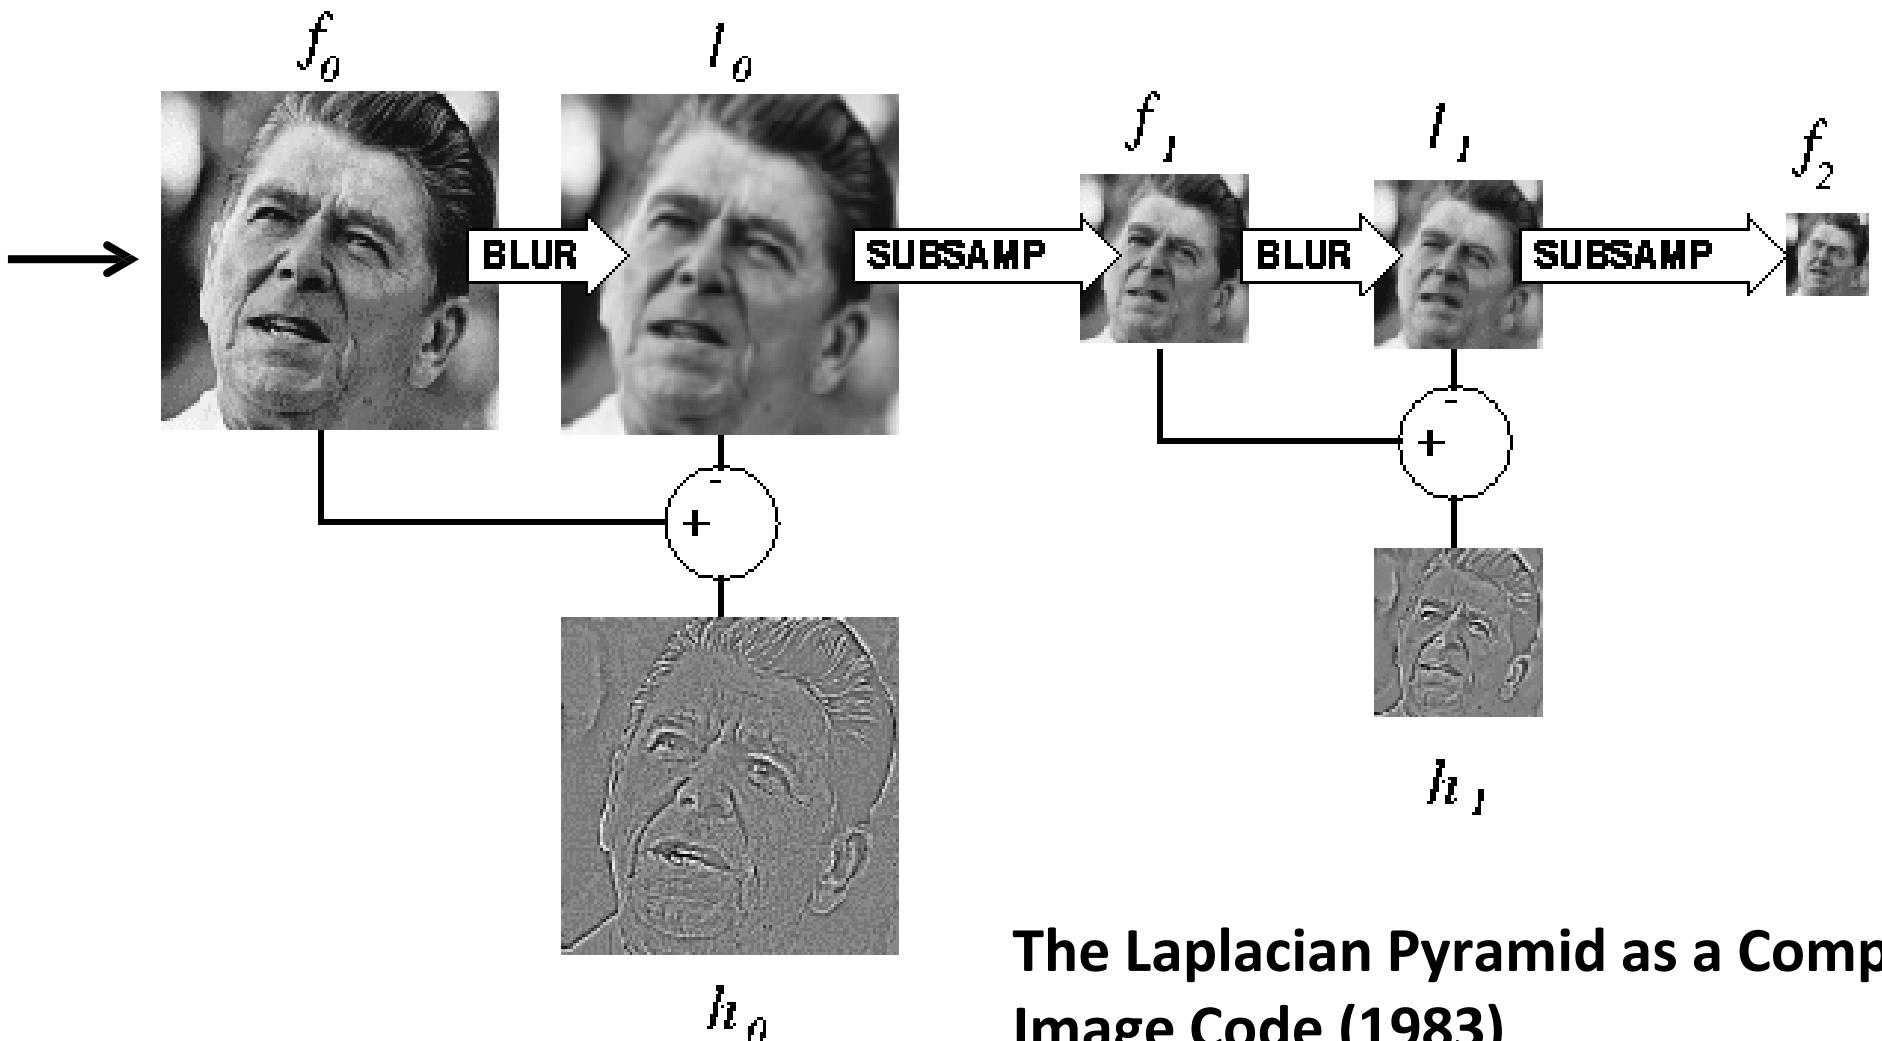
Kaveh Fathian

Why Reagan?



Why Reagan?

Ronald Reagan was President when the Laplacian pyramid was invented



The Laplacian Pyramid as a Compact Image Code (1983)

Peter J. Burt , Edward H. Adelson

Still used extensively



foreground details enhanced, background details reduced



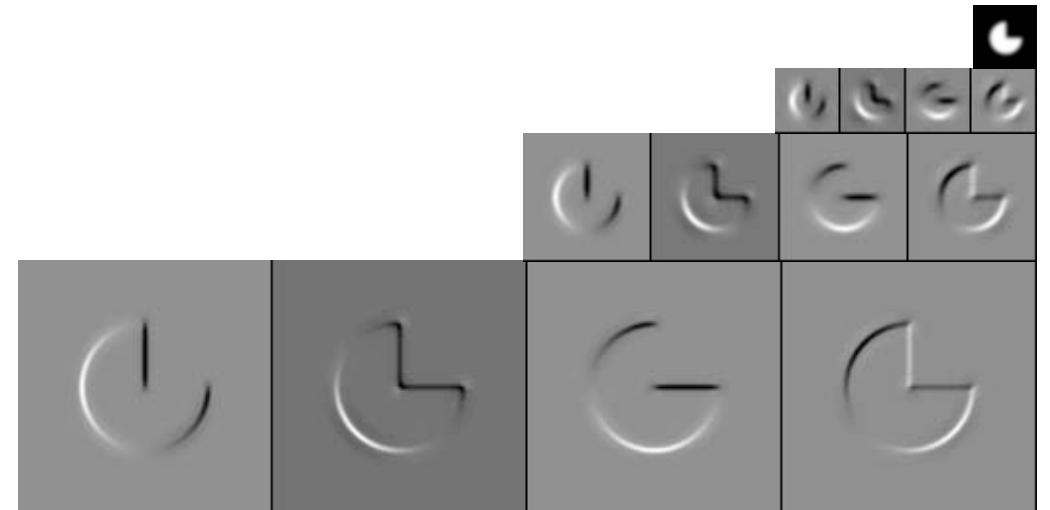
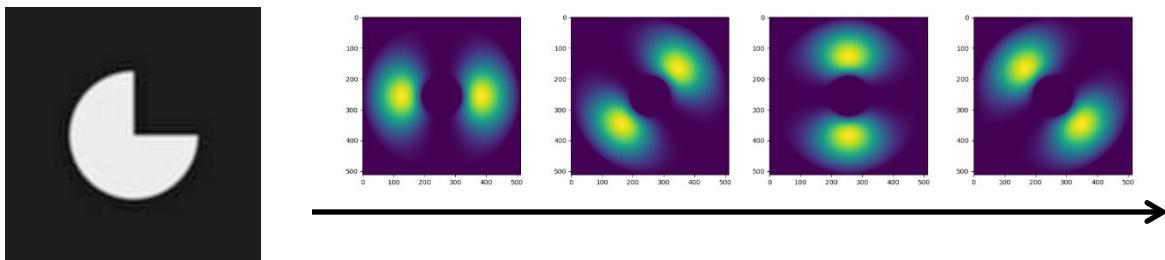
input image



user-provided mask

Other types of pyramids

Steerable pyramid: At each level keep multiple versions, one for each direction



Wavelets: Huge area in image processing



What are image pyramids used for?

image compression



multi-scale texture mapping

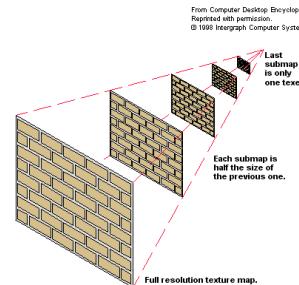
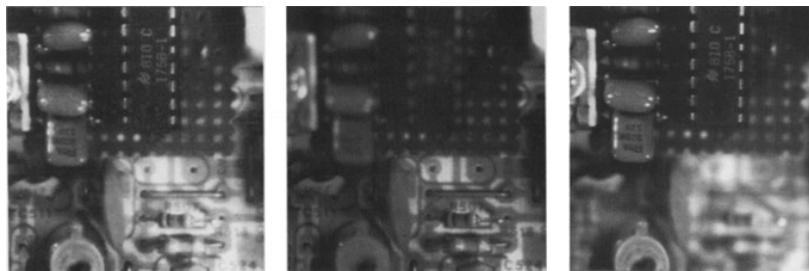


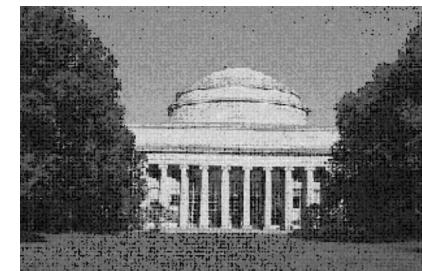
image blending



focal stack compositing



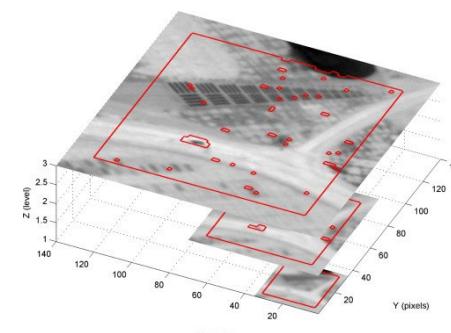
denoising



multi-scale detection



multi-scale registration



Fourier series

Some history

Who is this guy?



What is he famous for?



Jean Baptiste Joseph Fourier
(1768-1830)

What is he famous for?



Jean Baptiste Joseph Fourier
(1768-1830)

The Fourier series claim (1807):

'Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.'

... and apparently also for the discovery
of the greenhouse effect

Is this claim true?



Jean Baptiste Joseph Fourier
(1768-1830)

The Fourier series claim (1807):

'Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.'

Is this claim true?



Jean Baptiste Joseph Fourier
(1768-1830)

The Fourier series claim (1807):

'Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.'

Well, almost.

- The theorem requires additional conditions.
- Close enough to be named after him.
- Very surprising result at the time.

Is this claim true?



Jean Baptiste Joseph Fourier
(1768-1830)

The Fourier series claim (1807):

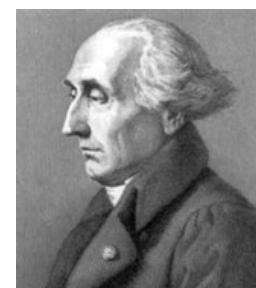
'Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.'

Well, almost.

- The theorem requires additional conditions.
- Close enough to be named after him.
- Very surprising result at the time.



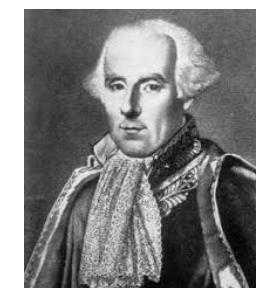
Malus



Lagrange



Legendre



Laplace

The committee examining his paper had expressed skepticism, in part due to not so rigorous proofs

Basic building block

$$A \sin(\omega x + \phi)$$

Fourier's claim: Add enough of these to get any periodic signal you want!

Basic building block

$$A \sin(\omega x + \phi)$$

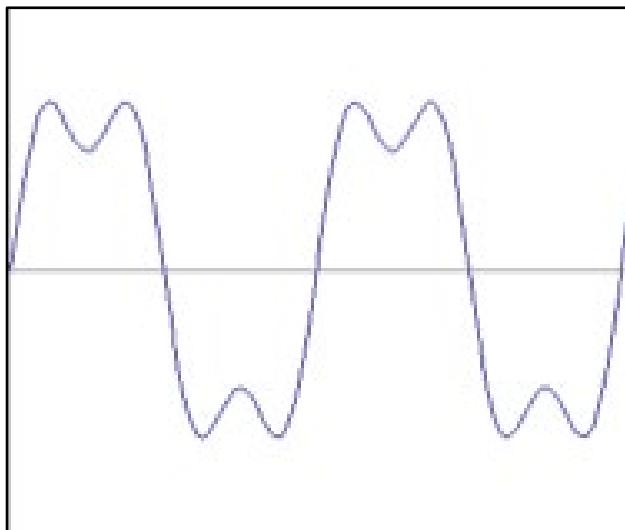
Diagram illustrating the components of the basic building block:

- amplitude (points to the coefficient A)
- sinusoid (points to the \sin function)
- angular frequency (points to the coefficient ω)
- variable (points to the variable x)
- phase (points to the phase shift term ϕ)

Fourier's claim: Add enough of these to get any periodic signal you want!

Examples

How would you generate this function?



=

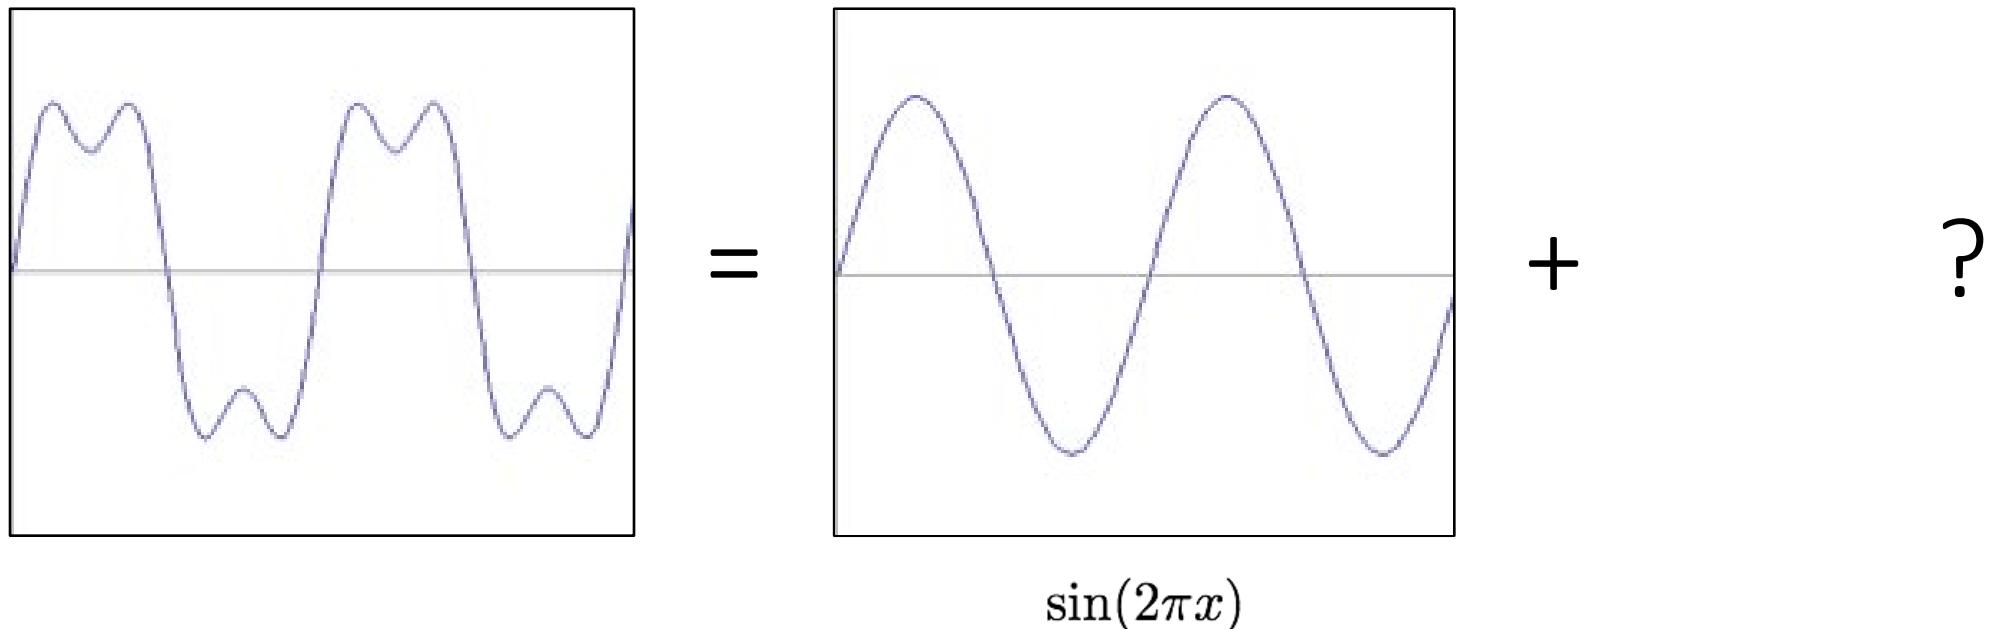
?

+

?

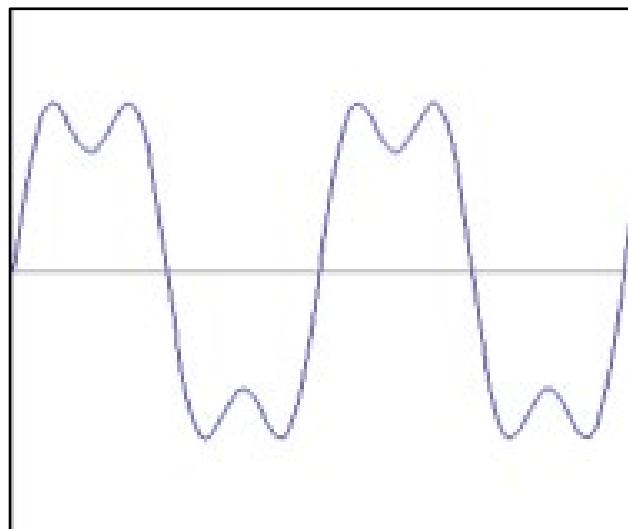
Examples

How would you generate this function?

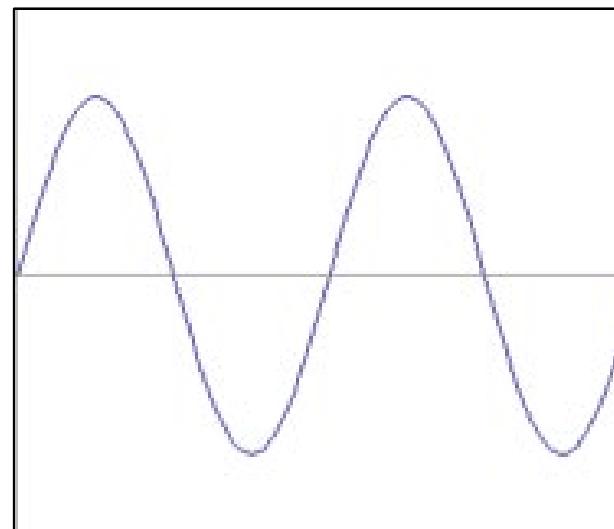


Examples

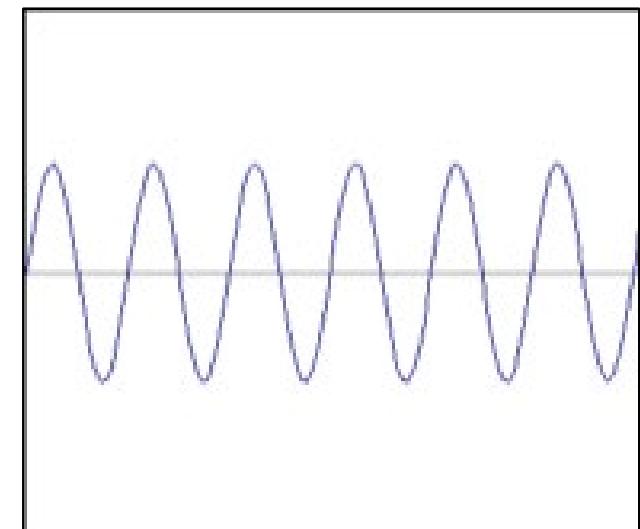
How would you generate this function?



=



+



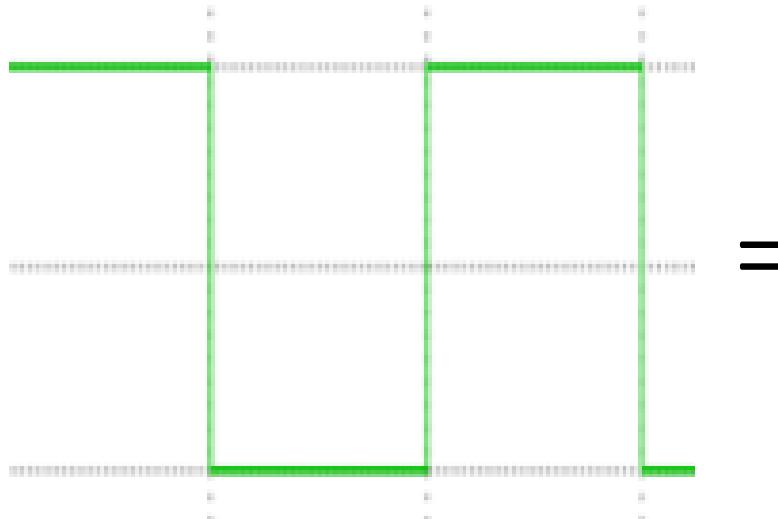
$$f(x) = \sin(2\pi x) + \frac{1}{3} \sin(2\pi 3x)$$

$$\sin(2\pi x)$$

$$\frac{1}{3} \sin(2\pi 3x)$$

Examples

How would you generate this function?



=

?

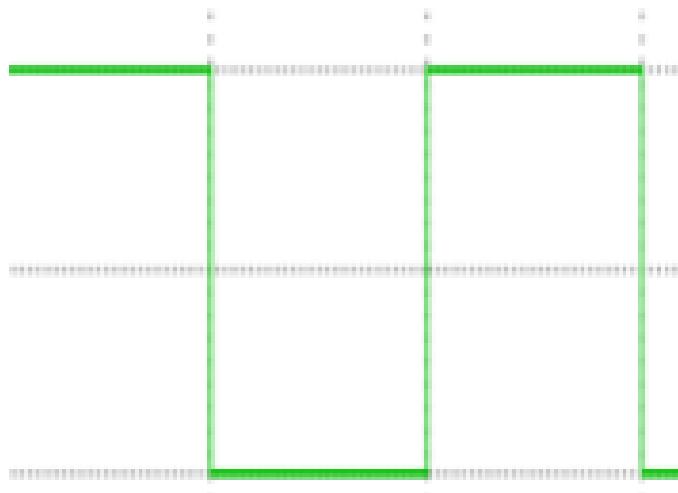
+

?

square wave

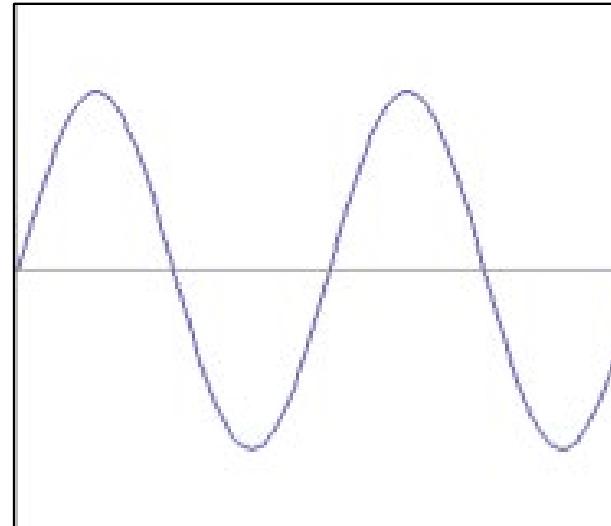
Examples

How would you generate this function?

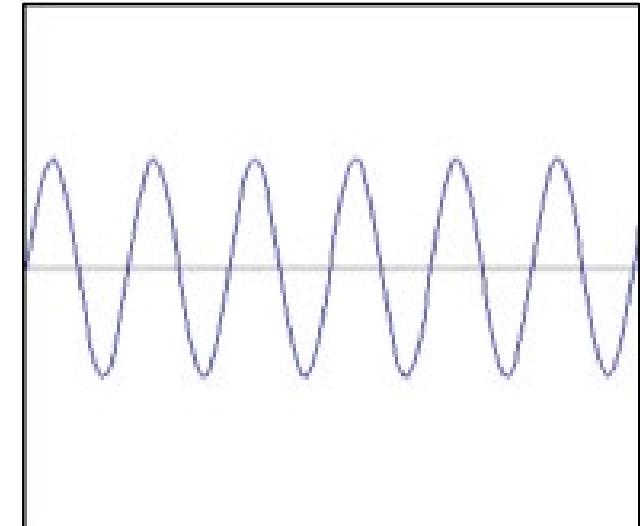


square wave

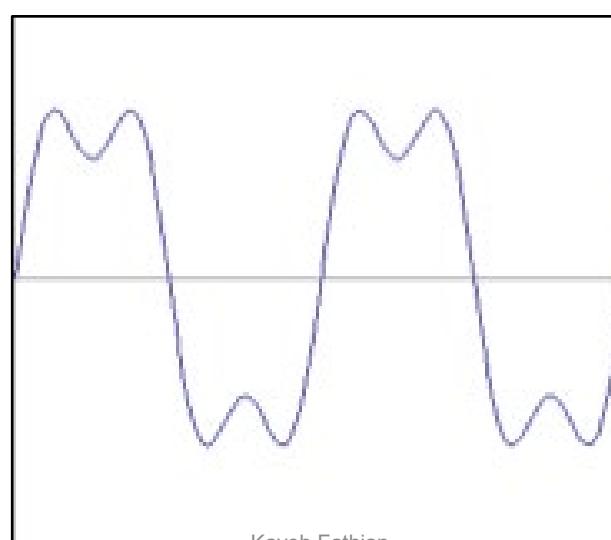
\approx



+

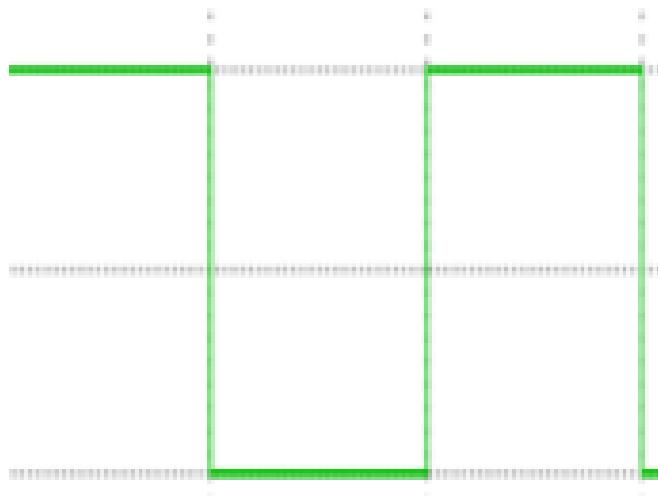


$=$



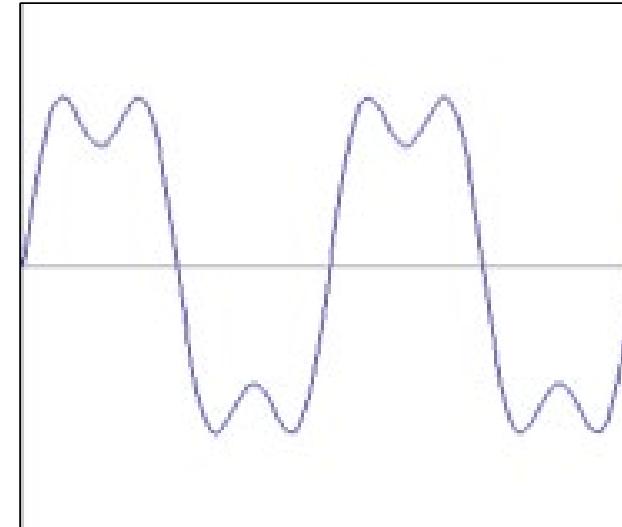
Examples

How would you generate this function?

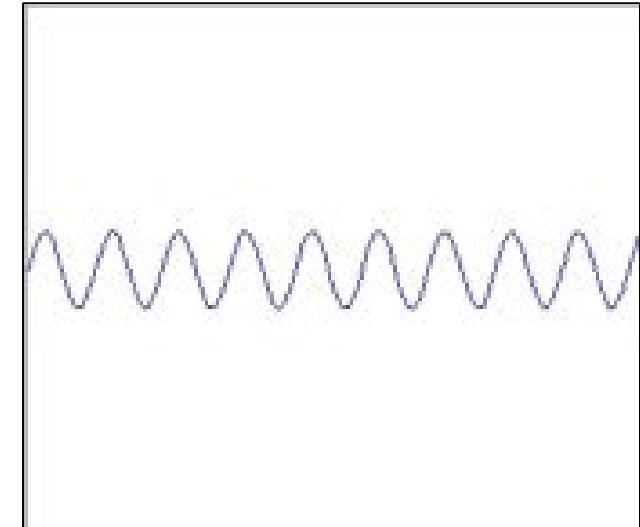


square wave

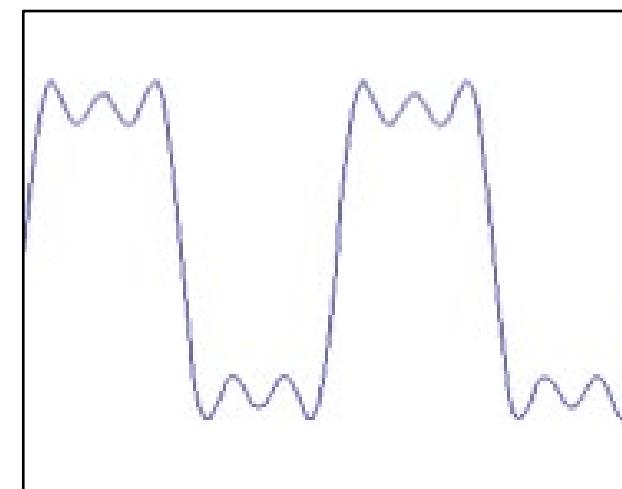
\approx



+

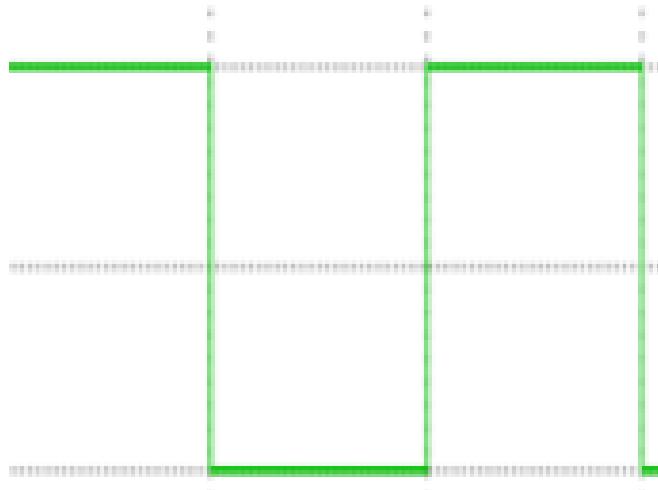


$=$



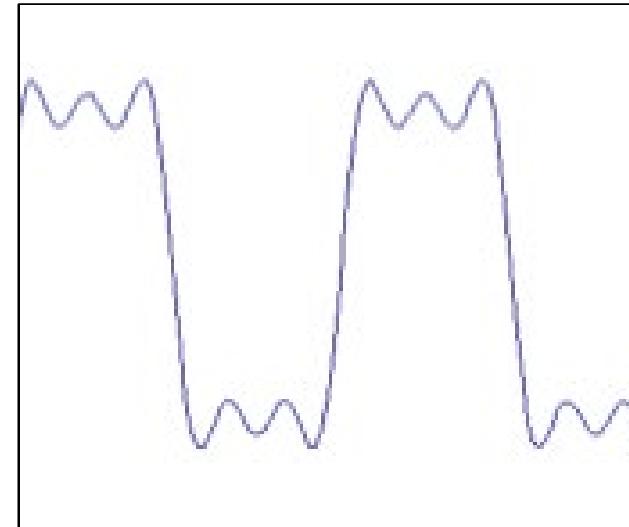
Examples

How would you generate this function?

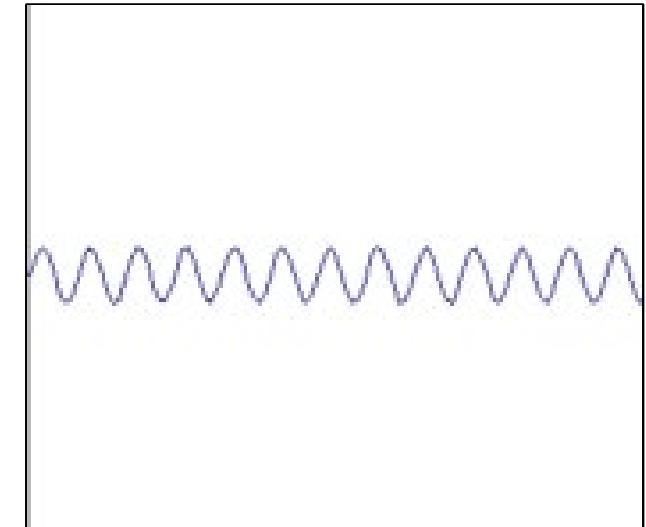


square wave

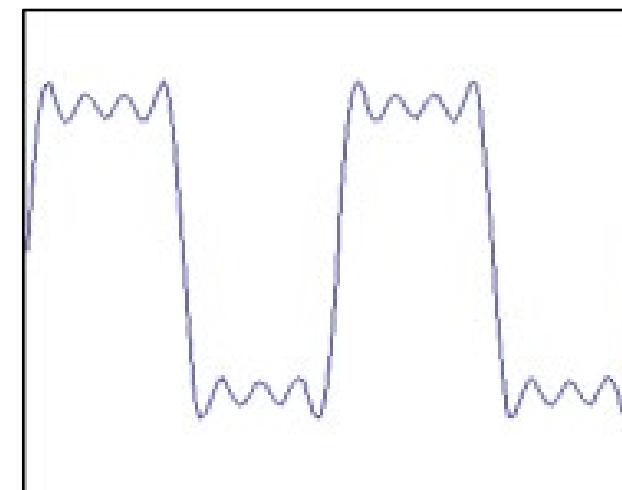
\approx



+

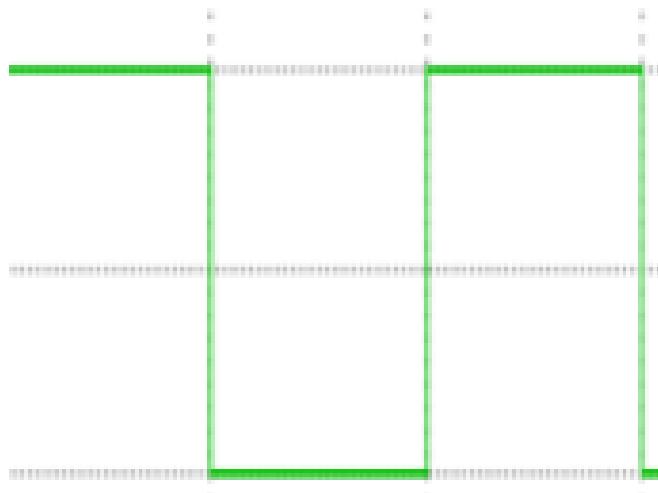


$=$



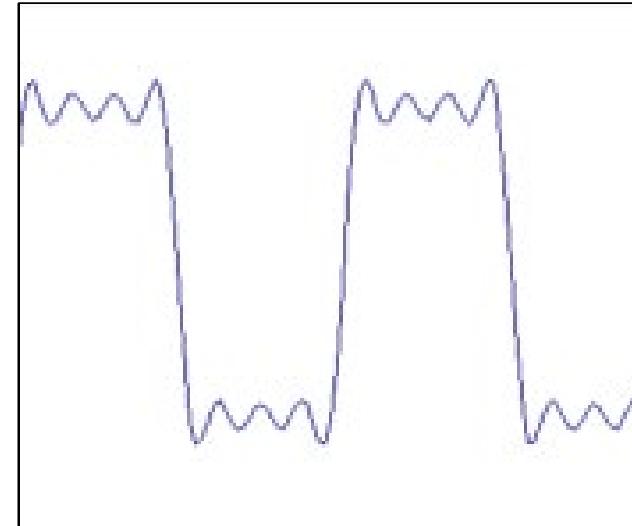
Examples

How would you generate this function?

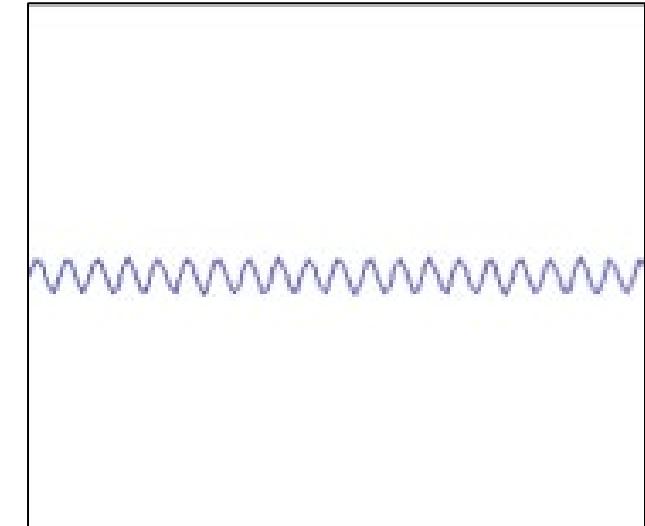


square wave

\approx



+

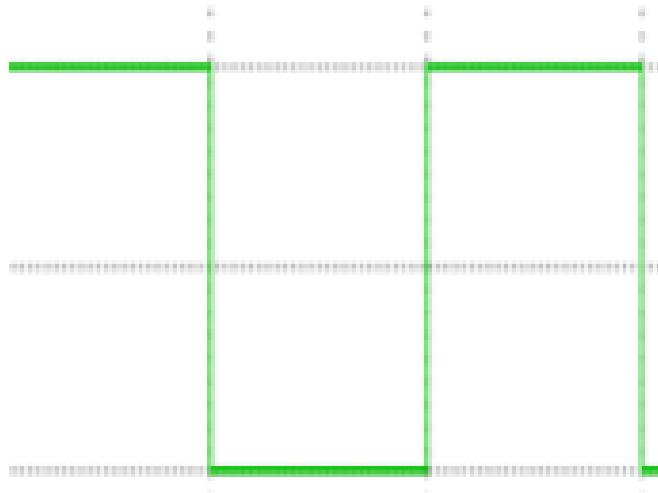


$=$



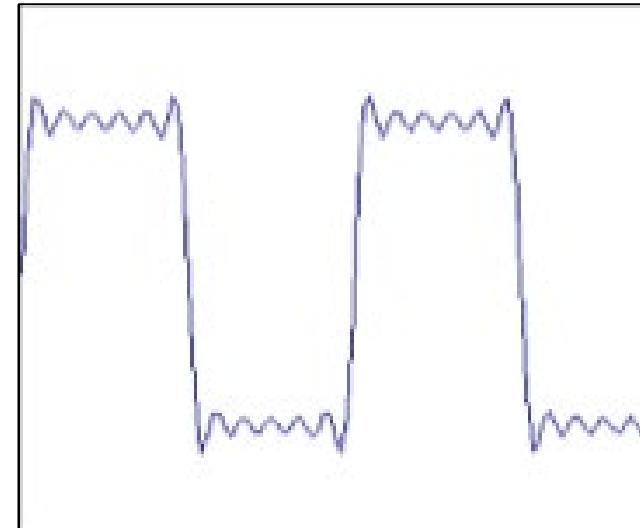
Examples

How would you generate this function?

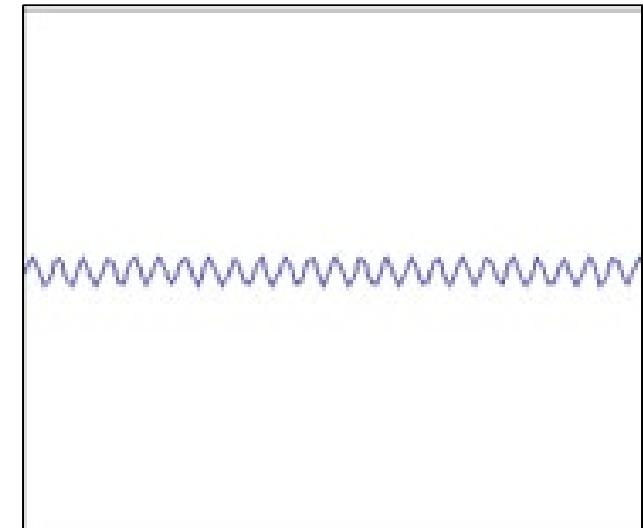


square wave

\approx



+

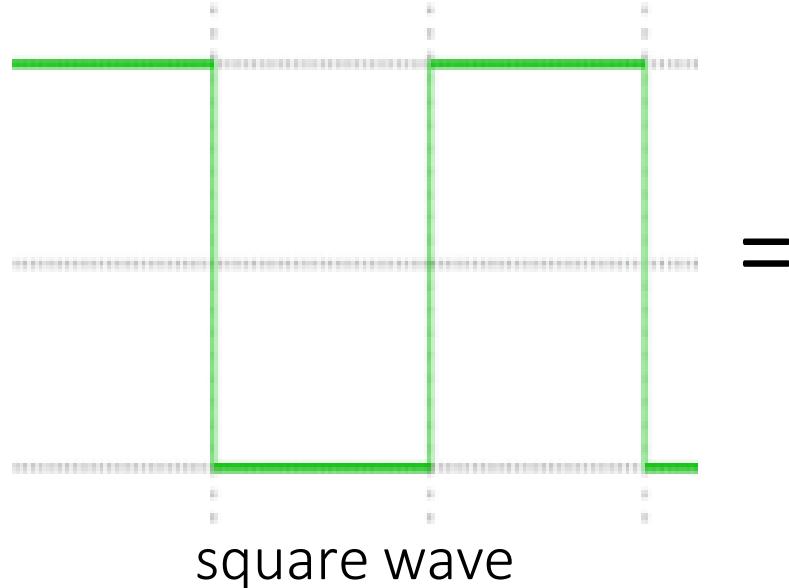


$=$



How would you express
this mathematically?

Examples



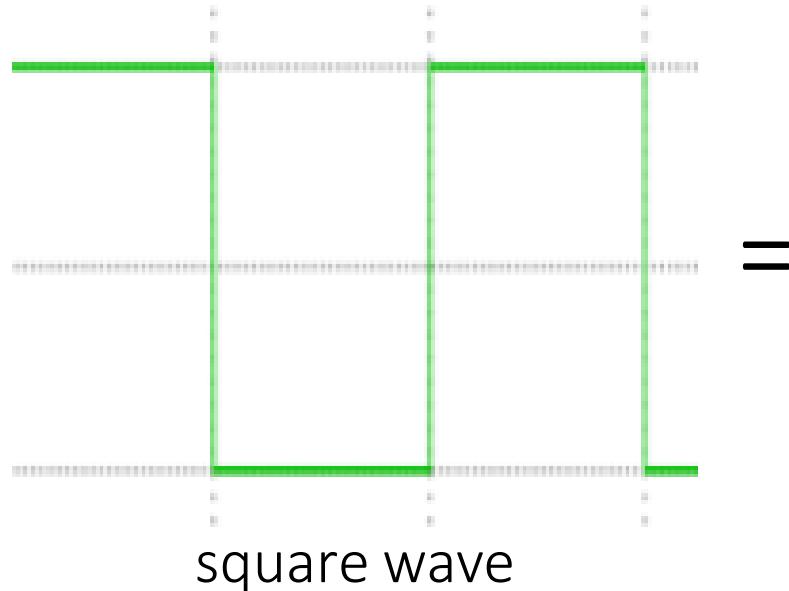
=

$$A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kx)$$

infinite sum of sine waves

How would you visualize this in the frequency domain?

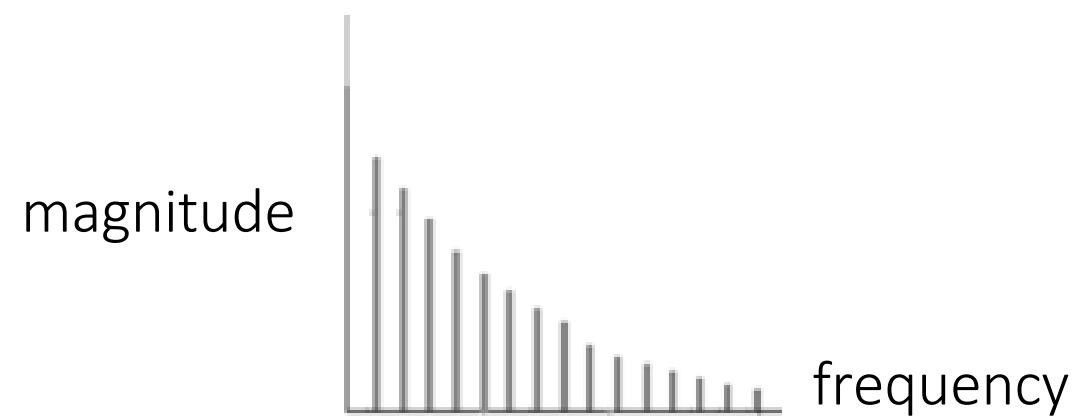
Examples



=

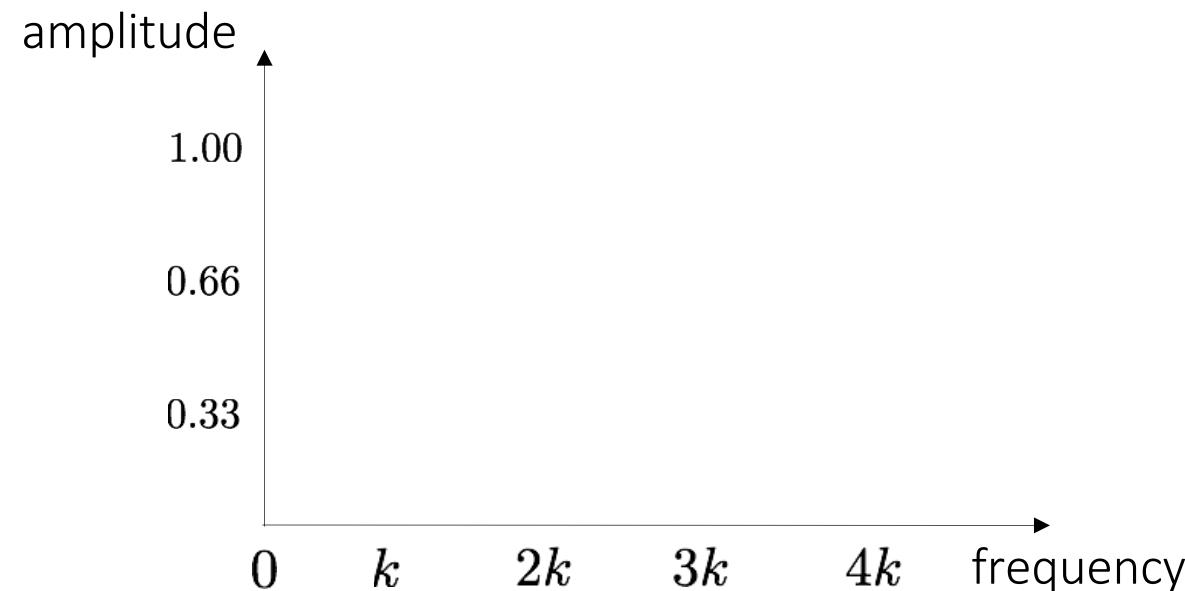
$$A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kx)$$

infinite sum of sine waves

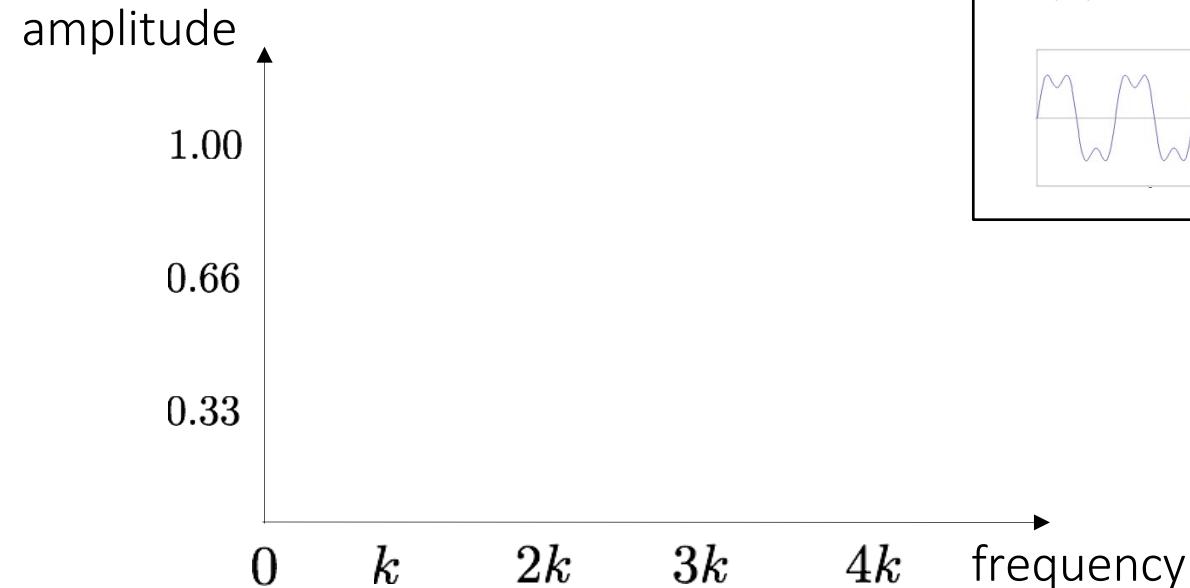


Frequency domain

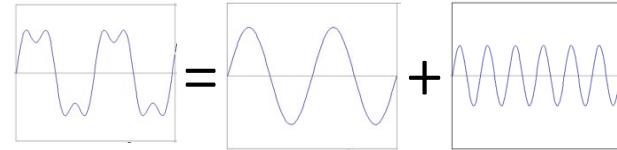
Visualizing the frequency spectrum



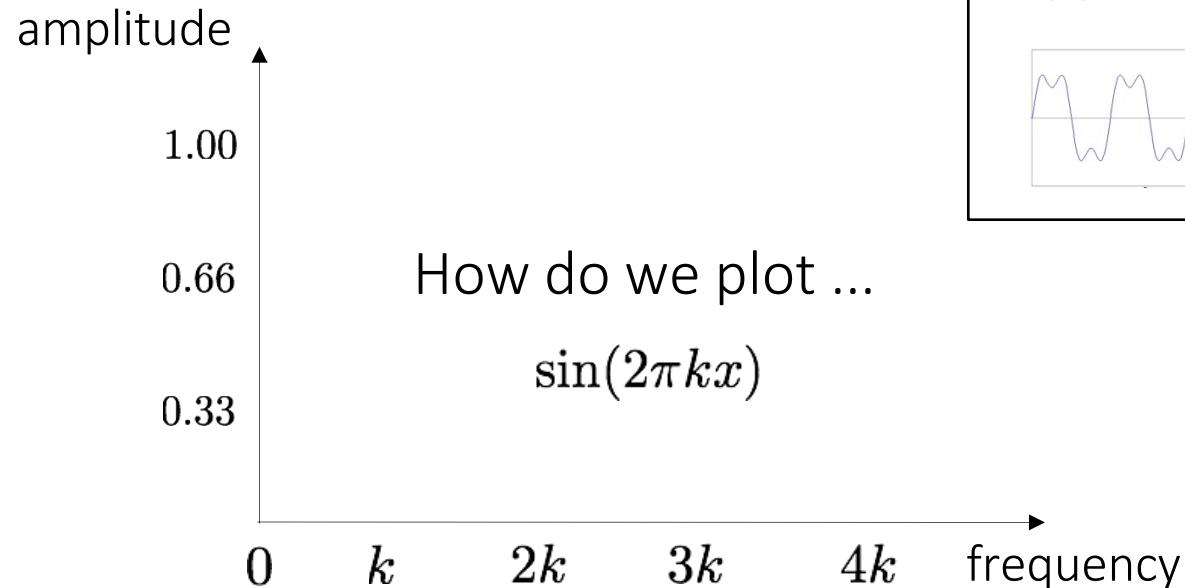
Visualizing the frequency spectrum



$$f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$$

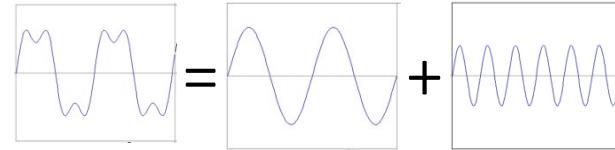


Visualizing the frequency spectrum

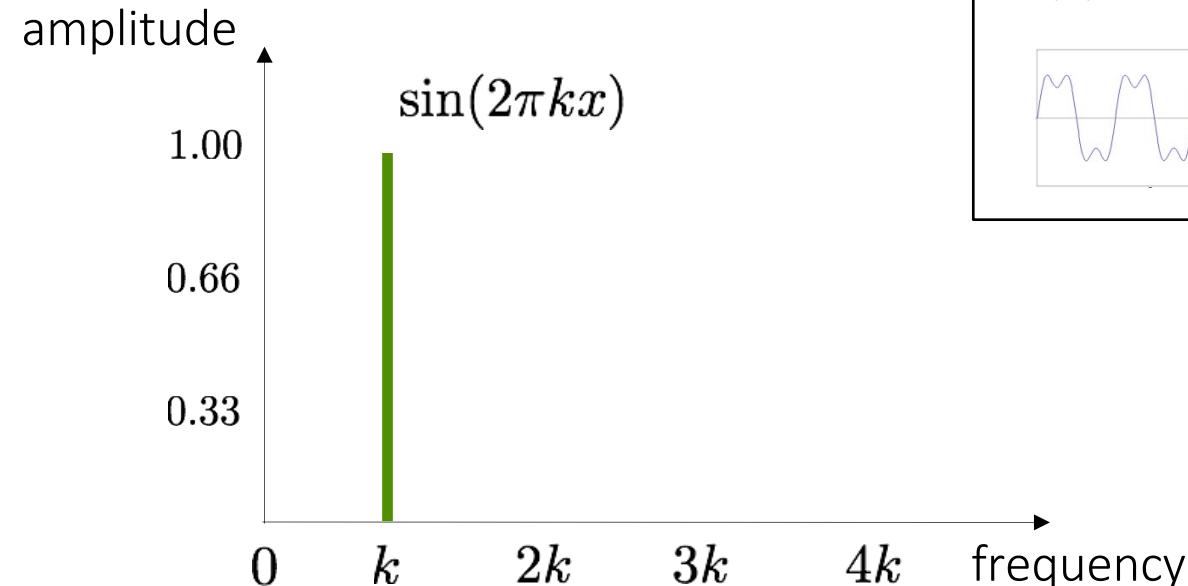


Recall the temporal domain visualization

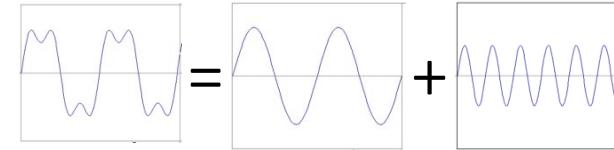
$$f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$$



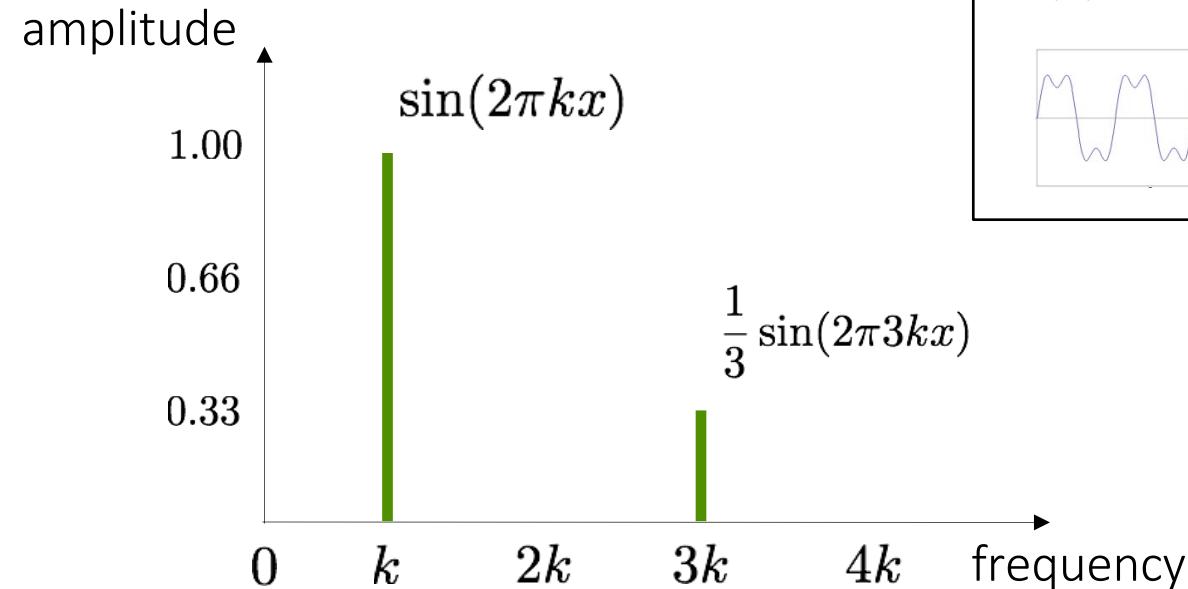
Visualizing the frequency spectrum



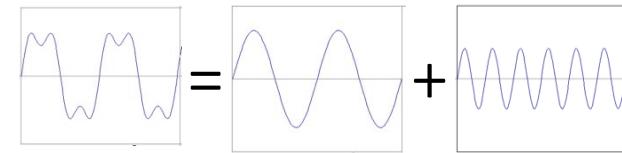
$$f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$$



Visualizing the frequency spectrum

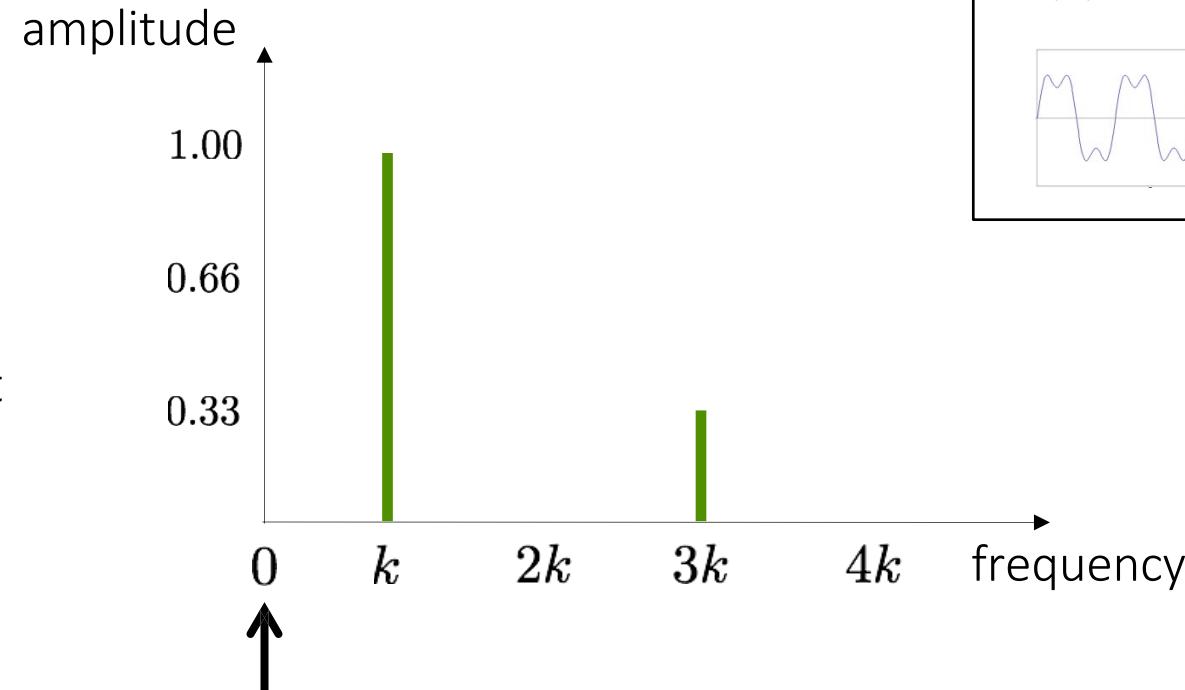


$$f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$$



Visualizing the frequency spectrum

not visualizing the
symmetric negative part



What is at zero
frequency?

Recall the temporal domain visualization

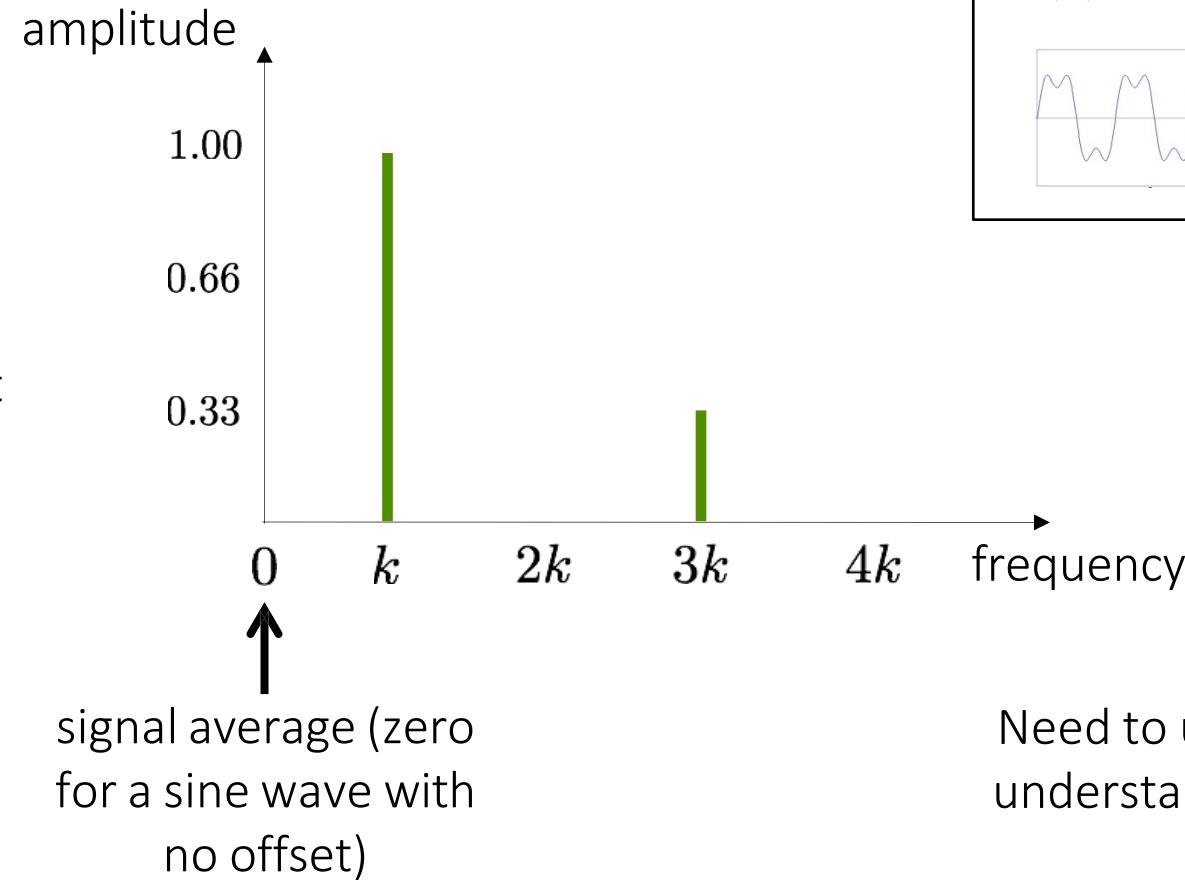
$$f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$$

The equation $f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$ is shown in a box. Below it, three plots illustrate the decomposition of the function. The first plot shows a blue sine wave. The second plot shows a blue sine wave with a lower frequency. The third plot shows a blue sine wave with a higher frequency. An equals sign and a plus sign are placed between the first two plots, indicating the sum of the two components.

Need to understand this to
understand the 2D version!

Visualizing the frequency spectrum

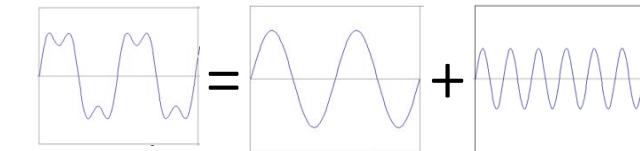
not visualizing the
symmetric negative part



Recall the temporal domain visualization

$$f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$$

The equation $f(x) = \sin(2\pi kx) + \frac{1}{3} \sin(2\pi 3kx)$ is shown in a box. Below the equation, there are three separate plots: a blue sine wave, a red sine wave, and their sum, which is a purple wave. The blue and red waves are centered around zero, while the purple wave has an overall positive offset.

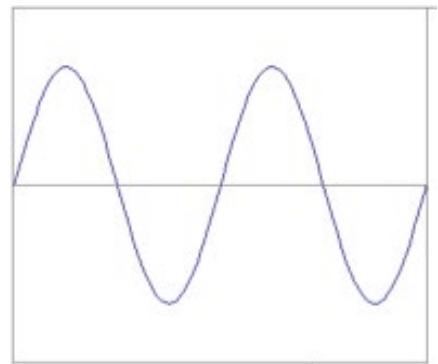


Need to understand this to
understand the 2D version!

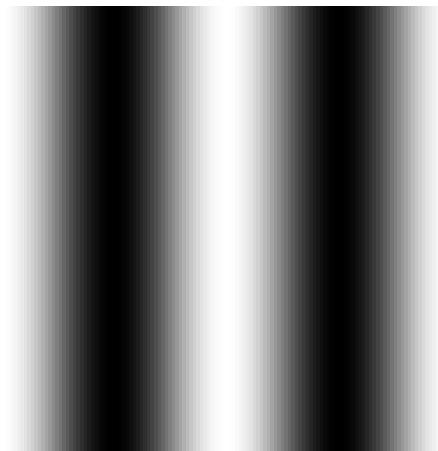
Examples

Spatial domain visualization

1D



2D



Frequency domain visualization

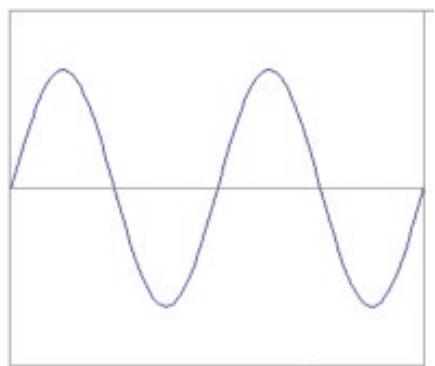


?

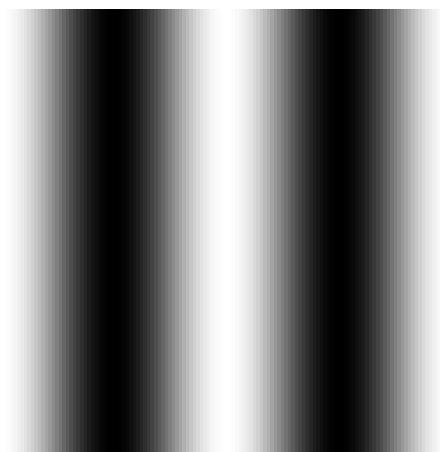
Examples

Spatial domain visualization

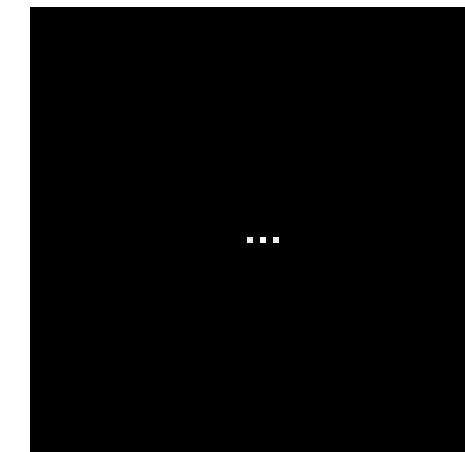
1D



2D



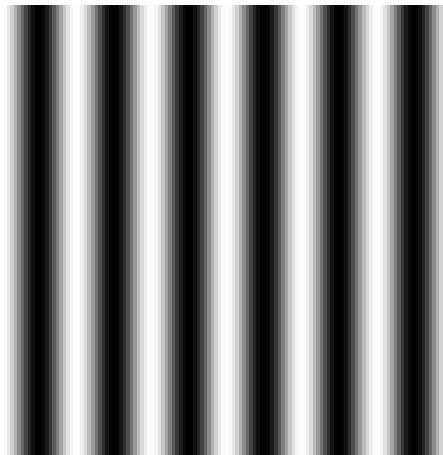
Frequency domain visualization



What do the three dots
correspond to?

Examples

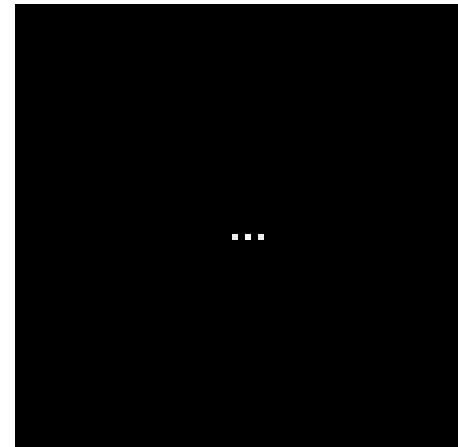
Spatial domain visualization



Frequency domain visualization

?

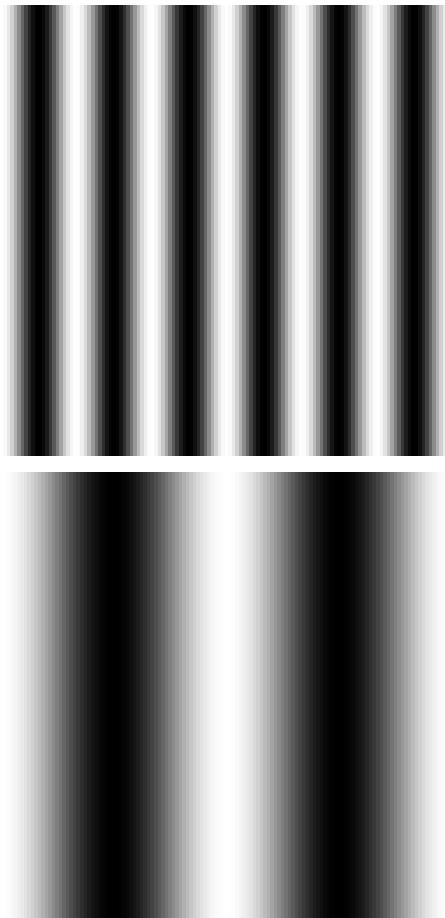
k_y



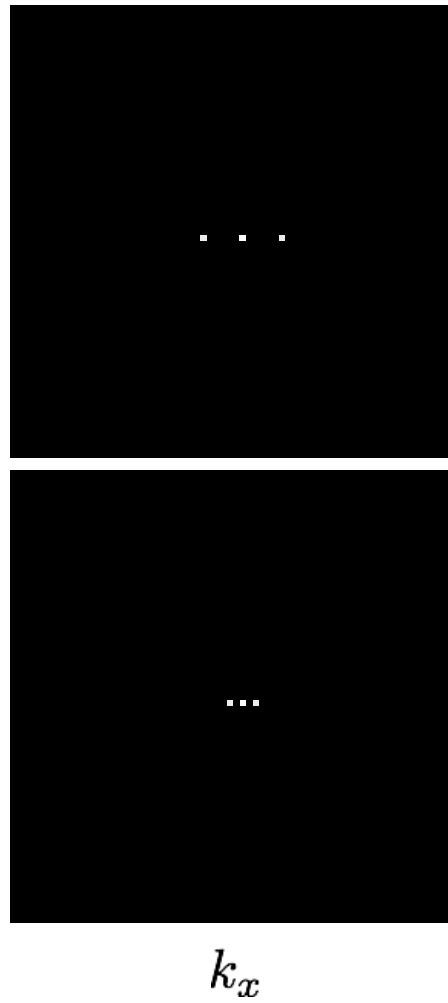
k_x

Examples

Spatial domain visualization

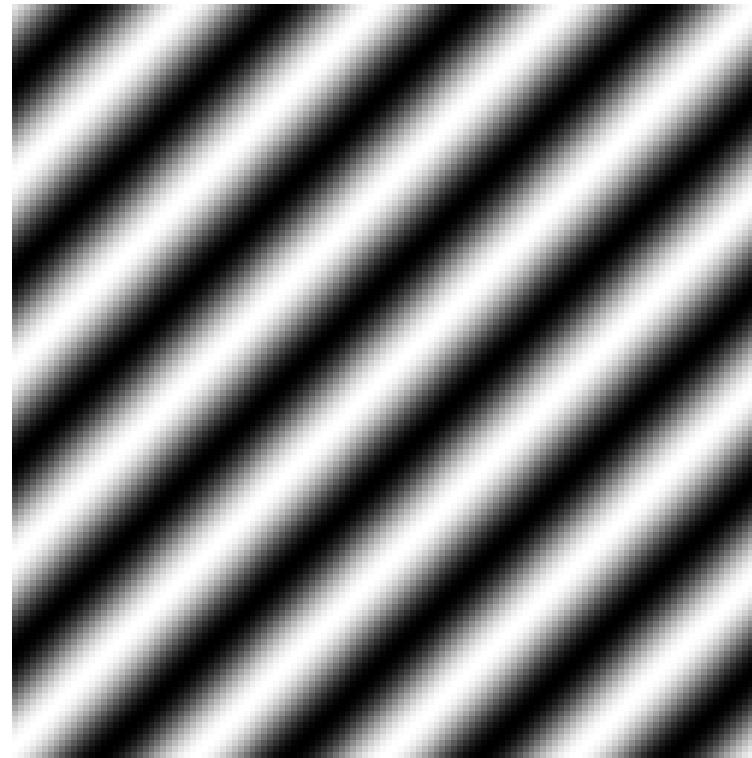


Frequency domain visualization



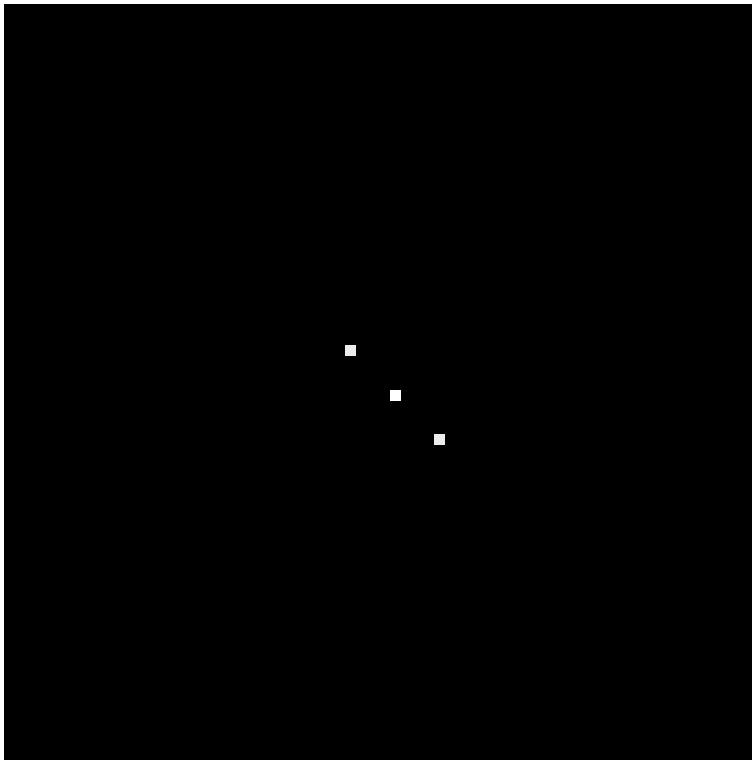
Examples

How would you generate this image with sine waves?



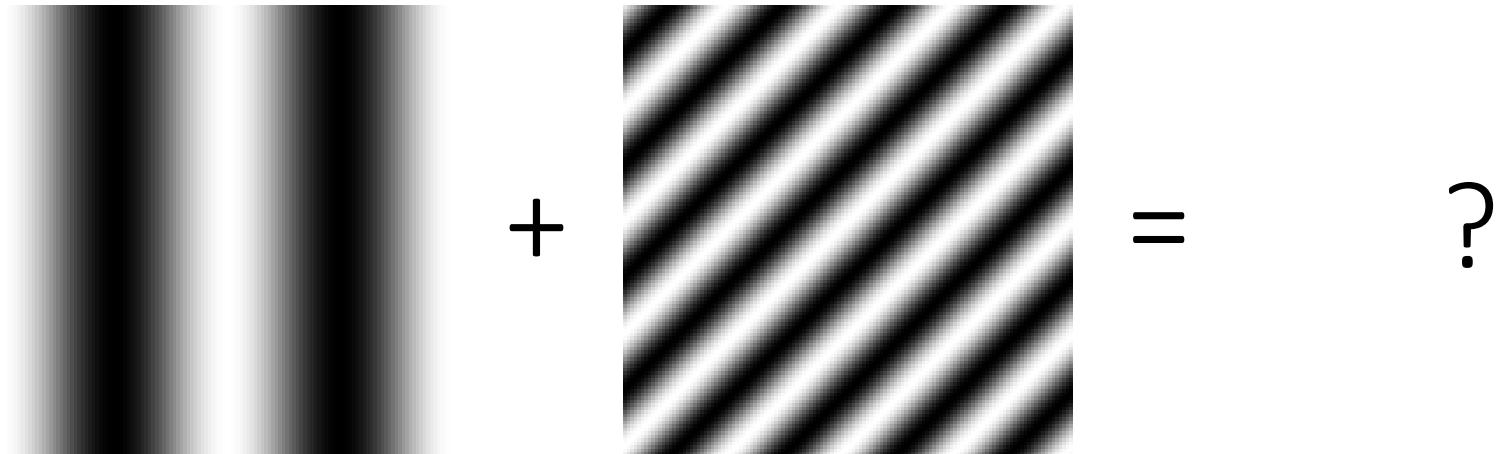
Examples

How would you generate this image with sine waves?

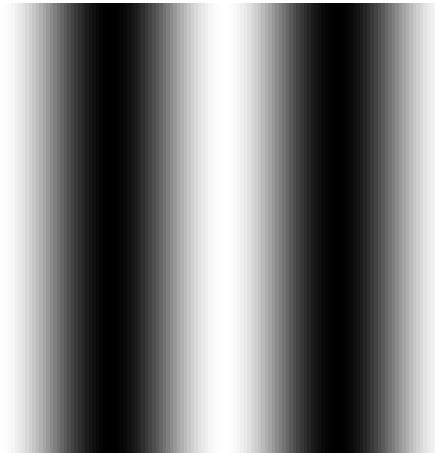


Has both an x and
y components

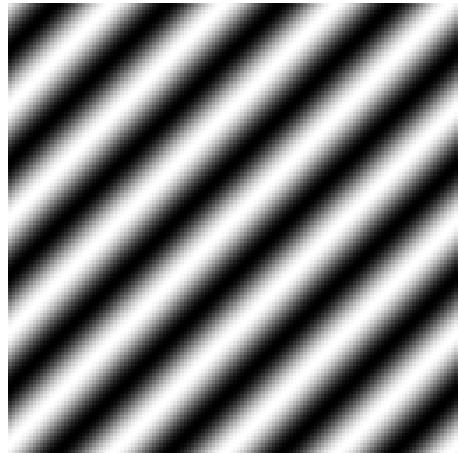
Examples



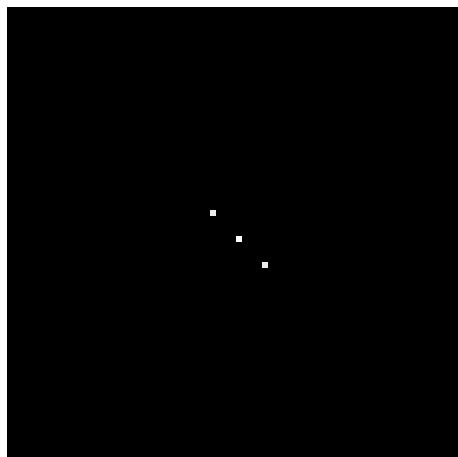
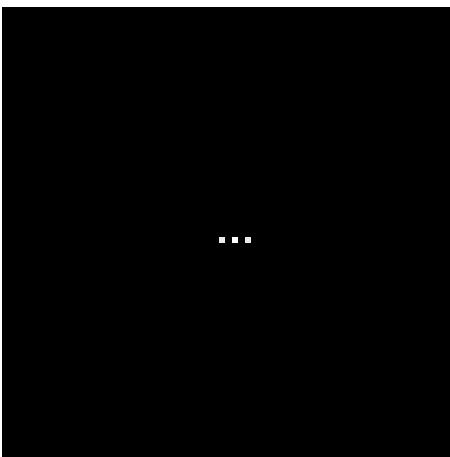
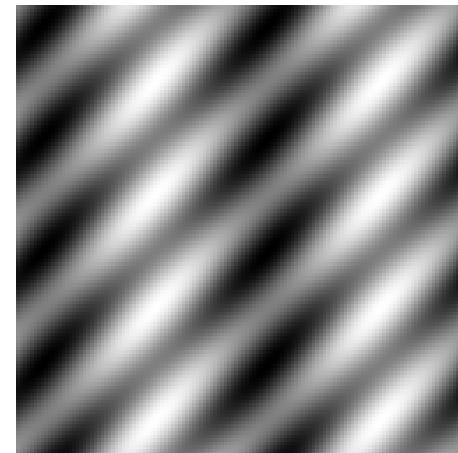
Examples



+

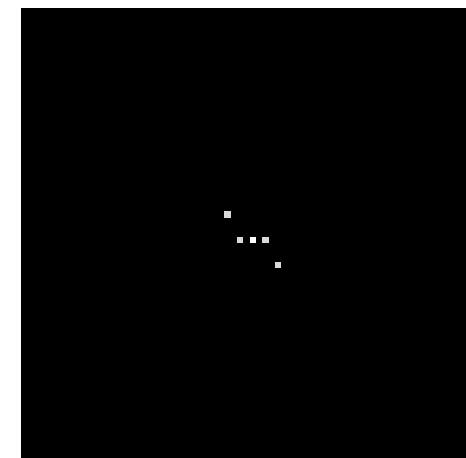
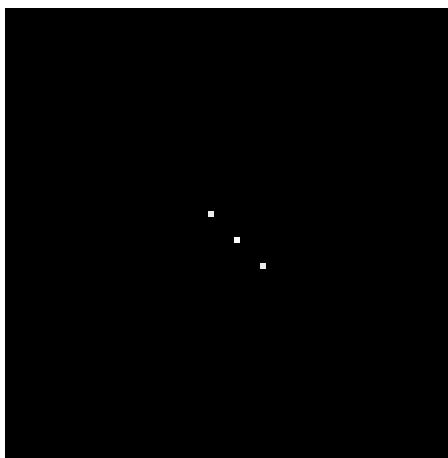
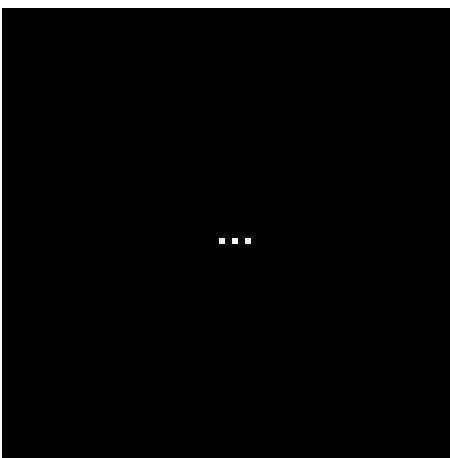
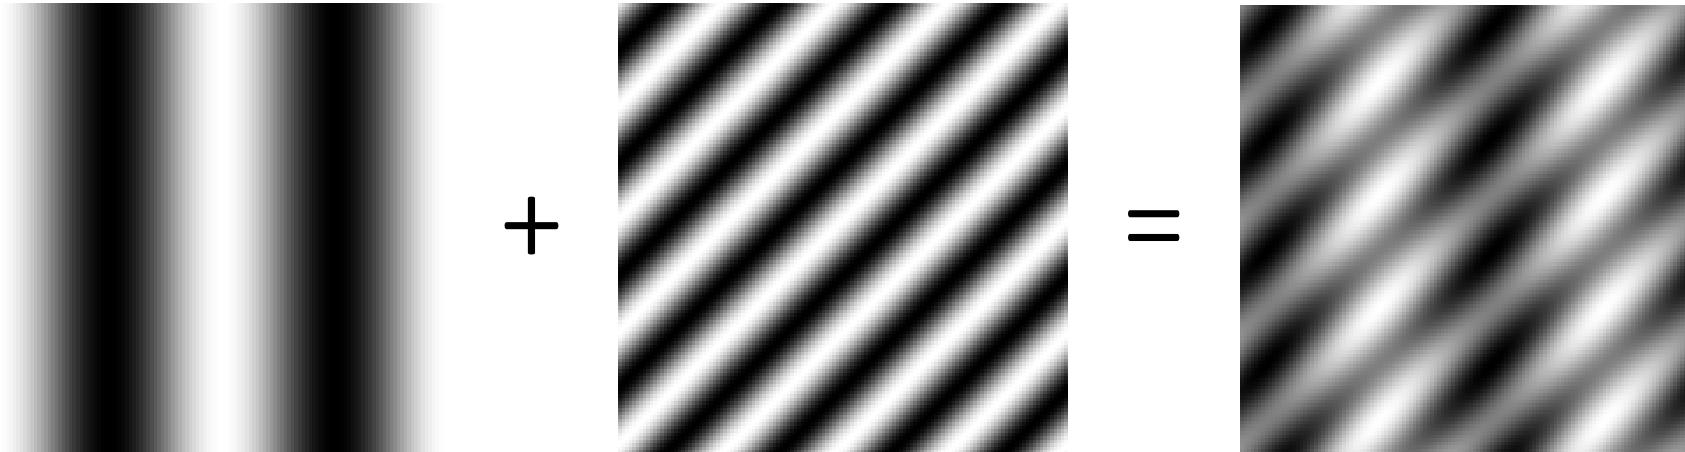


=



?

Examples



Basic building block

$$A \sin(\omega x + \phi)$$

Diagram illustrating the components of the basic building block:

- amplitude (pointing to A)
- sinusoid (pointing to \sin)
- angular frequency (pointing to ω)
- variable (pointing to x)
- phase (pointing to ϕ)

What about non-periodic signals?

Fourier's claim: Add enough of these to get any periodic signal you want!

Fourier Transform

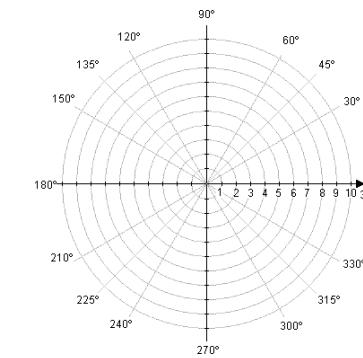
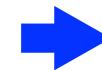
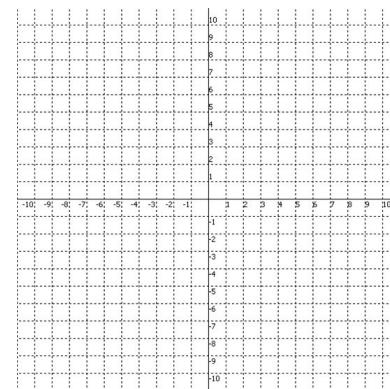
Recalling some basics

Complex numbers have two parts:

rectangular
coordinates

$$R + jI$$

what's this? what's this?



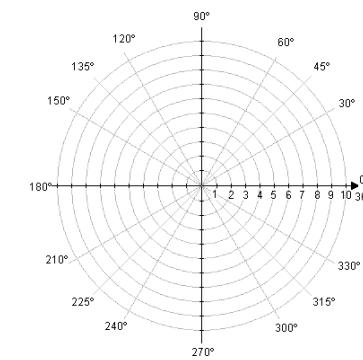
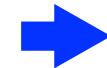
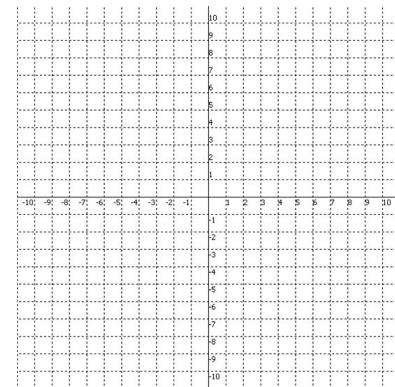
Recalling some basics

Complex numbers have two parts:

rectangular
coordinates

$$R + jI$$

real imaginary



Recalling some basics

Complex numbers have two parts:

rectangular
coordinates

$$R + jI$$

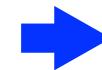
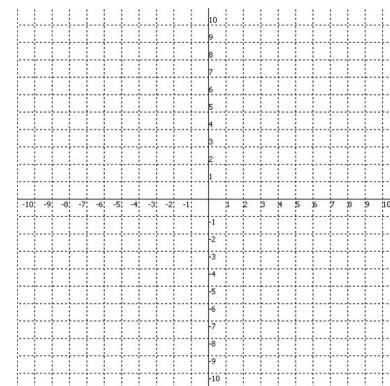
real imaginary

Alternative reparameterization:

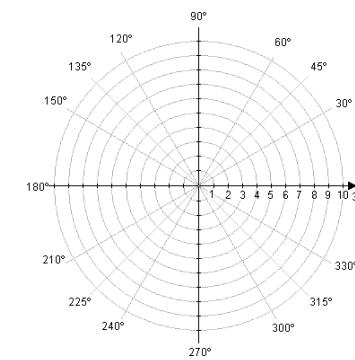
polar
coordinates

$$r(\cos \theta + j \sin \theta)$$

how do we compute these?



polar transform



Recalling some basics

Complex numbers have two parts:

rectangular
coordinates

$$R + jI$$

real imaginary

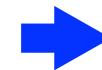
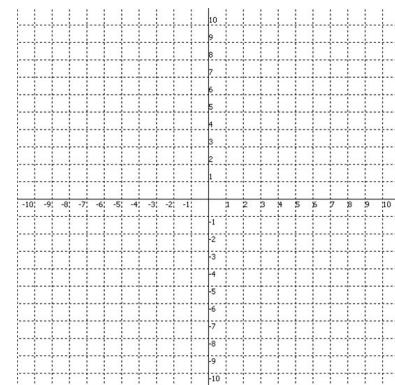
Alternative reparameterization:

polar
coordinates

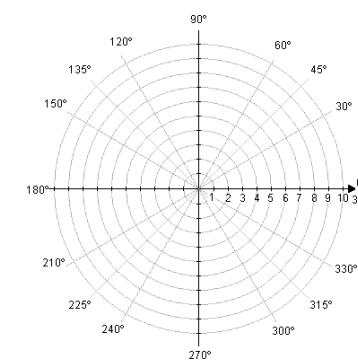
$$r(\cos \theta + j \sin \theta)$$

polar transform

$$\theta = \tan^{-1}\left(\frac{I}{R}\right) \quad r = \sqrt{R^2 + I^2}$$



polar transform



Recalling some basics

Complex numbers have two parts:

rectangular
coordinates

$$R + jI$$

real imaginary

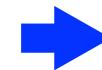
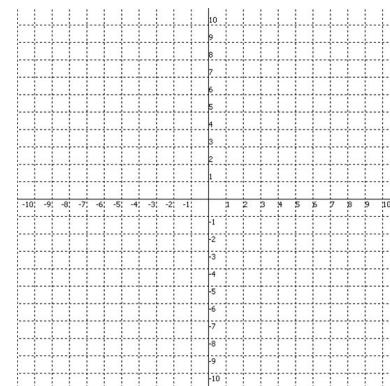
Alternative reparameterization:

polar
coordinates

$$r(\cos \theta + j \sin \theta)$$

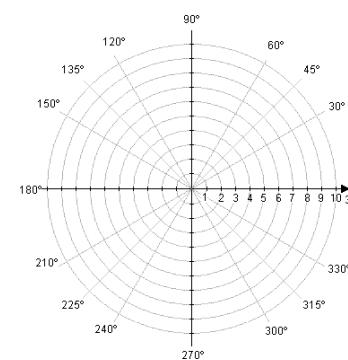
polar transform

$$\theta = \tan^{-1}\left(\frac{I}{R}\right) \quad r = \sqrt{R^2 + I^2}$$



polar transform

How do you write
these in exponential
form?



Recalling some basics

Complex numbers have two parts:

rectangular
coordinates

$$R + jI$$

real imaginary

Alternative reparameterization:

polar
coordinates

$$r(\cos \theta + j \sin \theta)$$

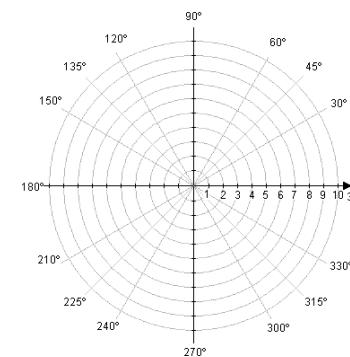
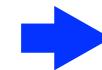
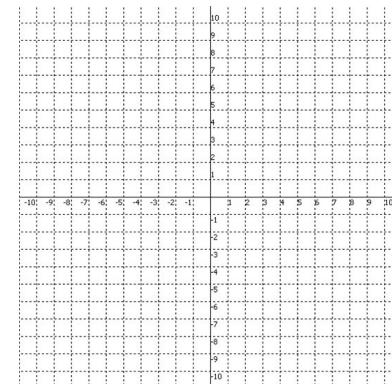
polar transform

$$\theta = \tan^{-1}\left(\frac{I}{R}\right) \quad r = \sqrt{R^2 + I^2}$$

or
equivalently

$$re^{j\theta}$$

how did we get this?



exponential
form

Recalling some basics

Complex numbers have two parts:

rectangular
coordinates

$$R + jI$$

real imaginary

Alternative reparameterization:

polar
coordinates

$$r(\cos \theta + j \sin \theta)$$

polar transform

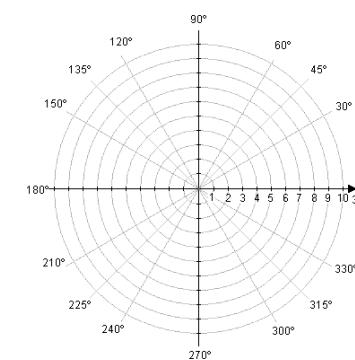
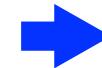
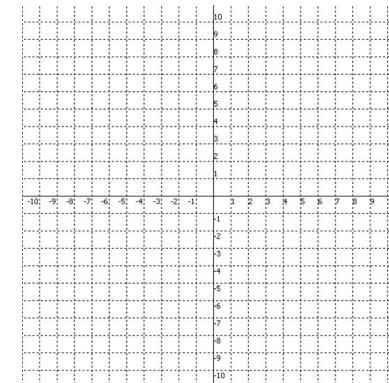
$$\theta = \tan^{-1}\left(\frac{I}{R}\right) \quad r = \sqrt{R^2 + I^2}$$

or
equivalently

$$re^{j\theta}$$

Euler's formula

$$e^{j\theta} = \cos \theta + j \sin \theta$$



exponential
form

This will help us understand the Fourier transform equations

Fourier transform

continuous

$$F(k) = \int_{-\infty}^{\infty} f(x)e^{-j2\pi kx}dx$$

inverse Fourier transform

discrete

$$F(k) = \frac{1}{N} \sum_{x=0}^{N-1} f(x)e^{-j2\pi kx/N}$$

$k = 0, 1, 2, \dots, N - 1$

$$f(x) = \sum_{k=0}^{N-1} F(k)e^{j2\pi kx/N}$$

$x = 0, 1, 2, \dots, N - 1$

Where is the connection to the 'summation of sine waves' idea?

Fourier transform

continuous

$$F(k) = \int_{-\infty}^{\infty} f(x)e^{-j2\pi kx}dx$$

inverse Fourier transform

$$f(x) = \int_{-\infty}^{\infty} F(k)e^{j2\pi kx}dk$$

discrete

$$F(k) = \frac{1}{N} \sum_{x=0}^{N-1} f(x)e^{-j2\pi kx/N}$$
$$k = 0, 1, 2, \dots, N-1$$

$$f(x) = \sum_{k=0}^{N-1} F(k)e^{j2\pi kx/N}$$
$$x = 0, 1, 2, \dots, N-1$$

Where is the connection to the 'summation of sine waves' idea?

Fourier transform

Where is the connection to the ‘summation of sine waves’ idea?

$$f(x) = \sum_{k=0}^{N-1} F(k) e^{j2\pi kx/N}$$

Euler's formula

$$e^{j\theta} = \cos \theta + j \sin \theta$$

sum over frequencies

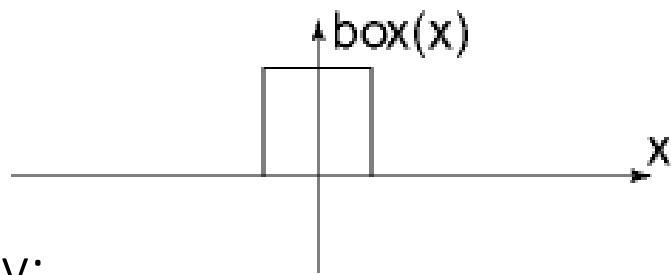
$$f(x) = \sum_{k=0}^{N-1} F(k) \left\{ \cos(2\pi kx) + j \sin(2\pi kx) \right\}$$

scaling parameter

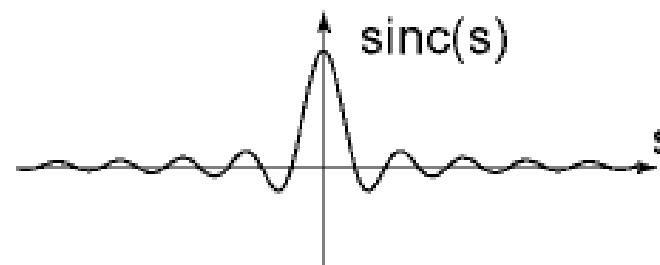
wave components

Fourier transform pairs

spatial domain

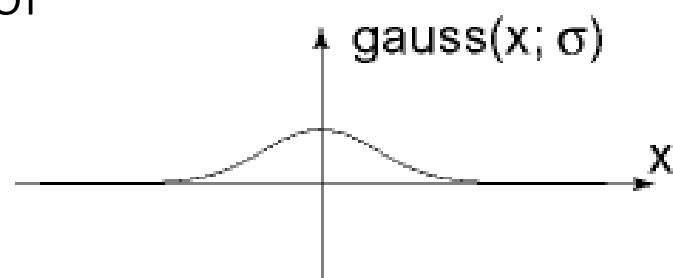


frequency domain

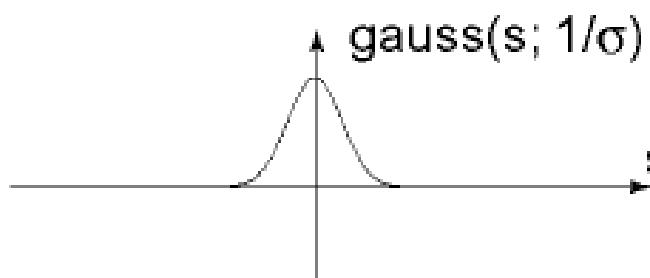


Note the symmetry:
duality property of
Fourier transform

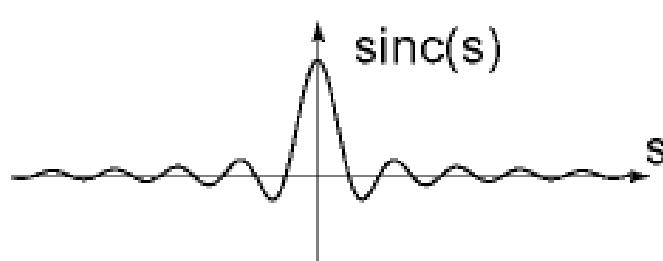
gauss($x; \sigma$)



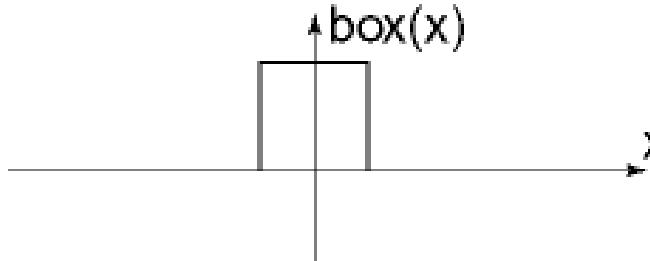
gauss($s; 1/\sigma$)



sinc(s)



box(x)



Fourier transform for images

1D Discrete Fourier Transform (DFT) transforms a signal $f[n]$ into $F[u]$ as:

$$F[u] = \sum_{n=0}^{N-1} f[n] \exp\left(-2\pi j \frac{un}{N}\right)$$

2D Discrete Fourier Transform (DFT) transforms an image $f[n,m]$ into $F[u,v]$ as:

$$F[u, v] = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} f[n, m] \exp\left(-2\pi j \left(\frac{un}{N} + \frac{vm}{M}\right)\right)$$

Computing the discrete Fourier transform (DFT)

$F(k) = \frac{1}{N} \sum_{x=0}^{N-1} f(x)e^{-j2\pi kx/N}$ is just a matrix multiplication:

$$\mathbf{F} = \mathbf{W}\mathbf{f}$$

$$\begin{bmatrix} F(0) \\ F(1) \\ F(2) \\ F(3) \\ \vdots \\ F(N-1) \end{bmatrix} = \begin{bmatrix} W^0 & W^0 & W^0 & W^0 & \dots & W^0 \\ W^0 & W^1 & W^2 & W^3 & \dots & W^{N-1} \\ W^0 & W^2 & W^4 & W^6 & \dots & W^{N-2} \\ W^0 & W^3 & W^6 & W^9 & \dots & W^{N-3} \\ \vdots & \vdots & & & \ddots & \vdots \\ W^0 & W^{N-1} & W^{N-2} & W^{N-3} & \dots & W^1 \end{bmatrix} \begin{bmatrix} f(0) \\ f(1) \\ f(2) \\ f(3) \\ \vdots \\ f(N-1) \end{bmatrix} \quad W = e^{-j2\pi/N}$$

In practice this is implemented using the *fast Fourier transform* (FFT) algorithm.

Visualizing the image Fourier transform

$$F[u, v] = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} f[n, m] \exp \left(-2\pi j \left(\frac{un}{N} + \frac{vm}{M} \right) \right)$$

The values of $F [u,v]$ are complex.

Using the real and imaginary components:

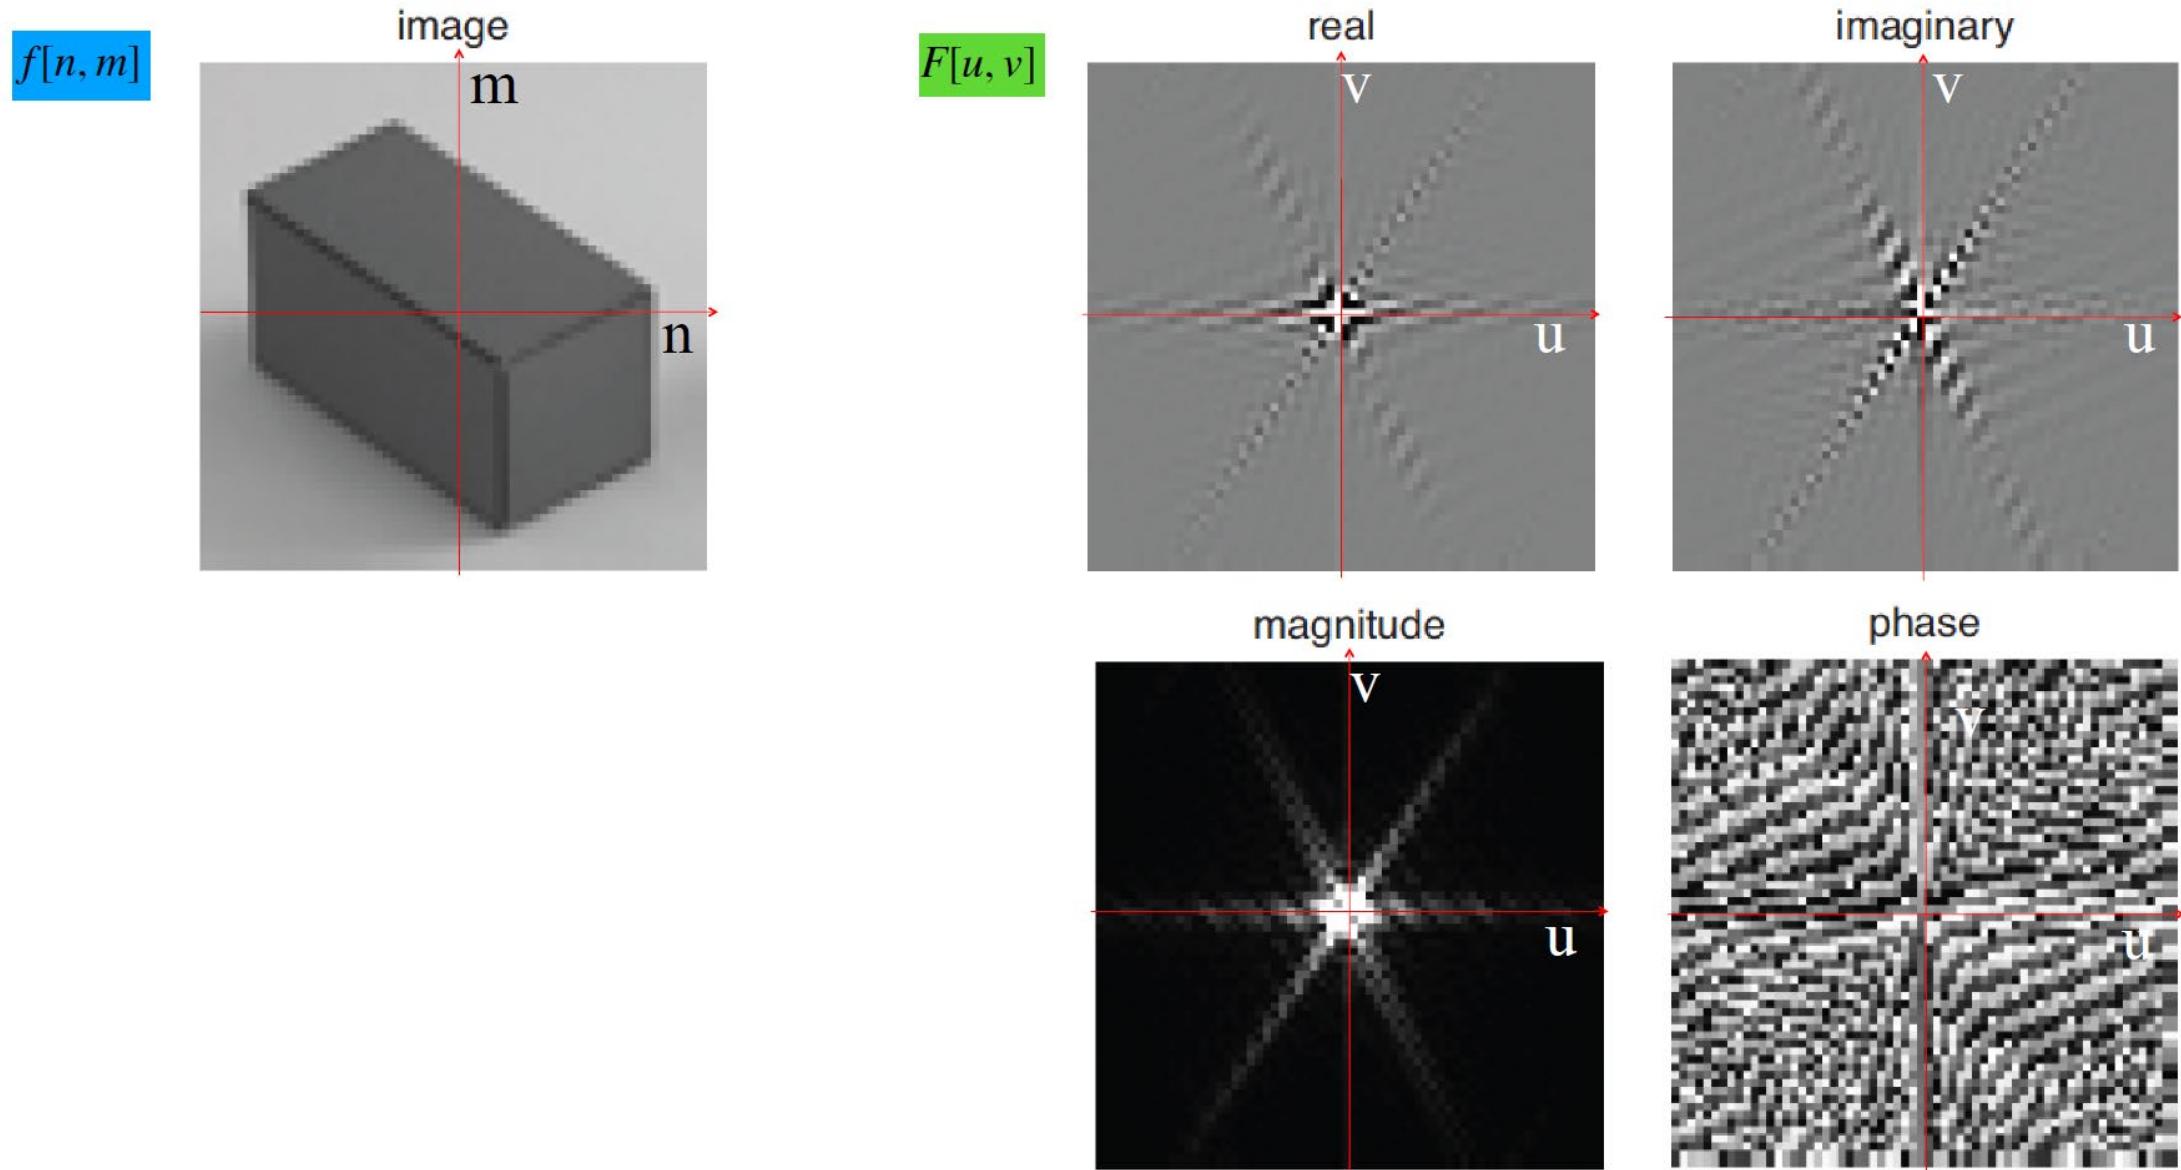
$$F [u, v] = Re \{F [u, v]\} + j Imag \{F [u, v]\}$$

Or using a polar decomposition:

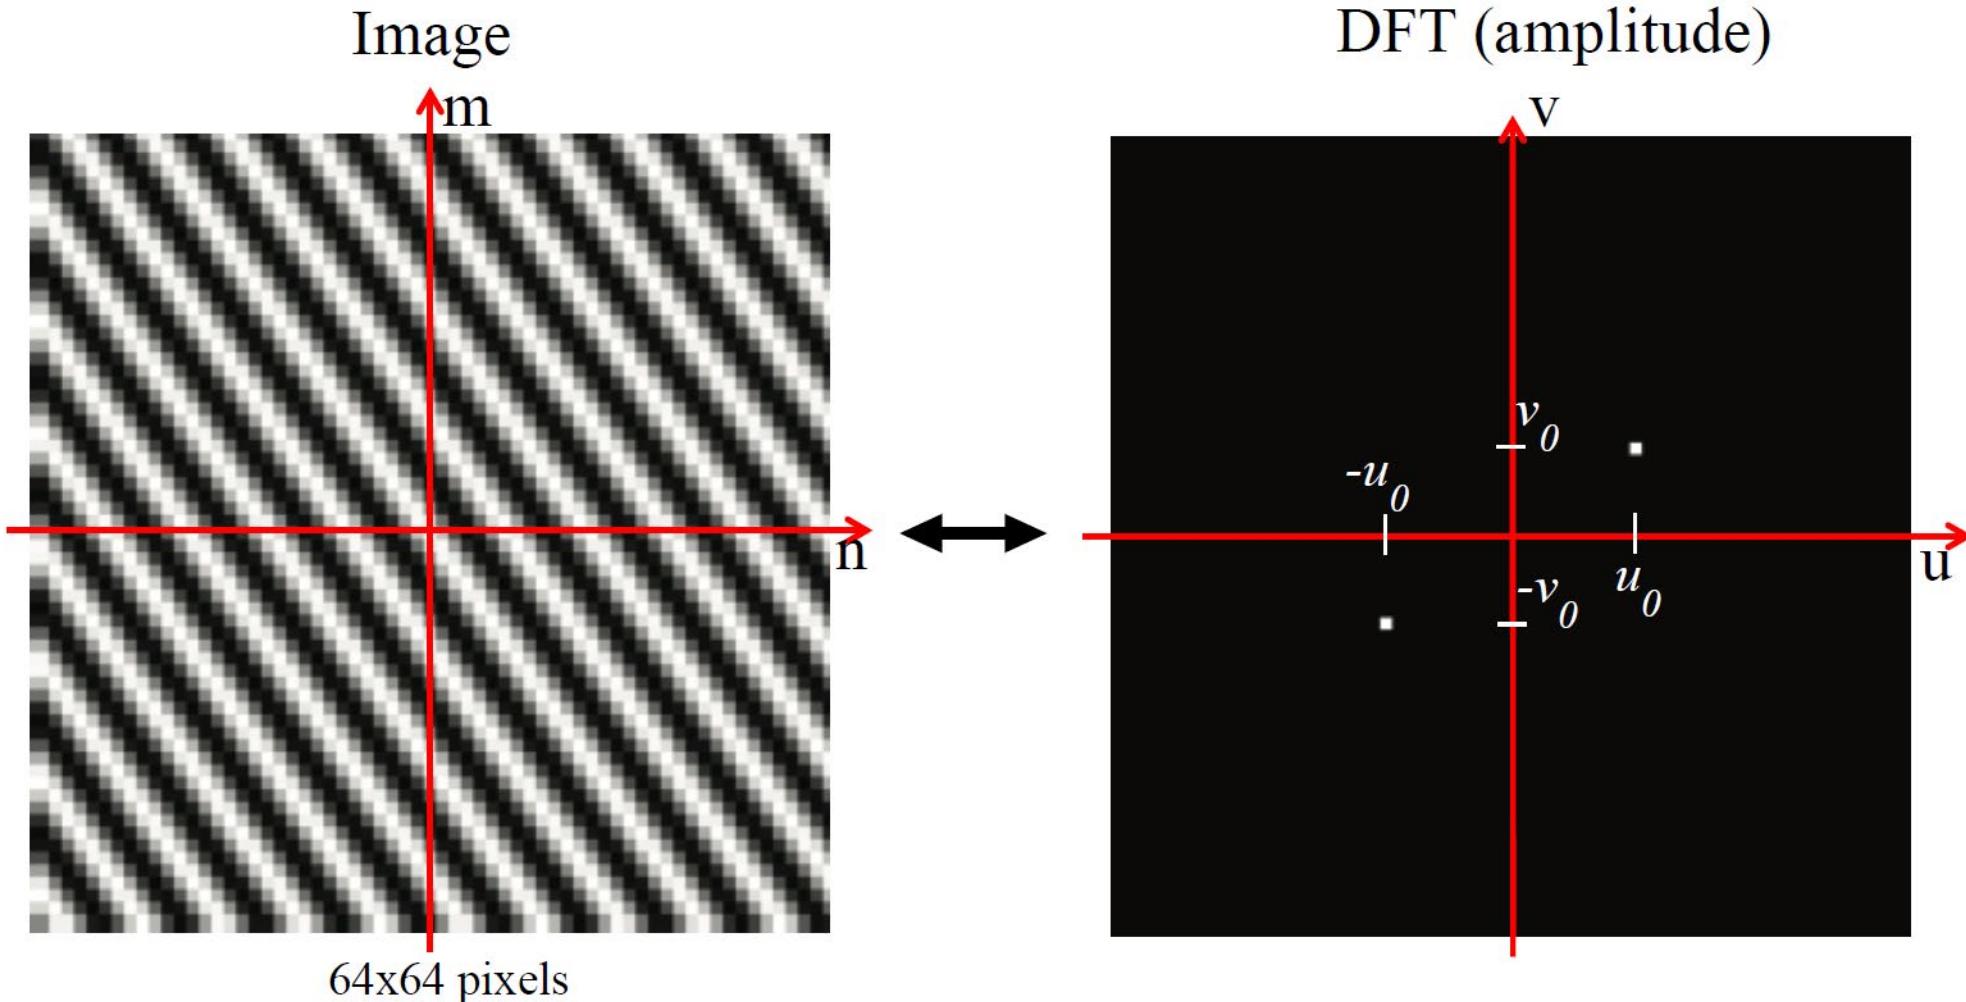
$$F [u, v] = A [u, v] \exp (j \theta [u, v])$$

Amplitude Phase

Visualizing the image Fourier transform

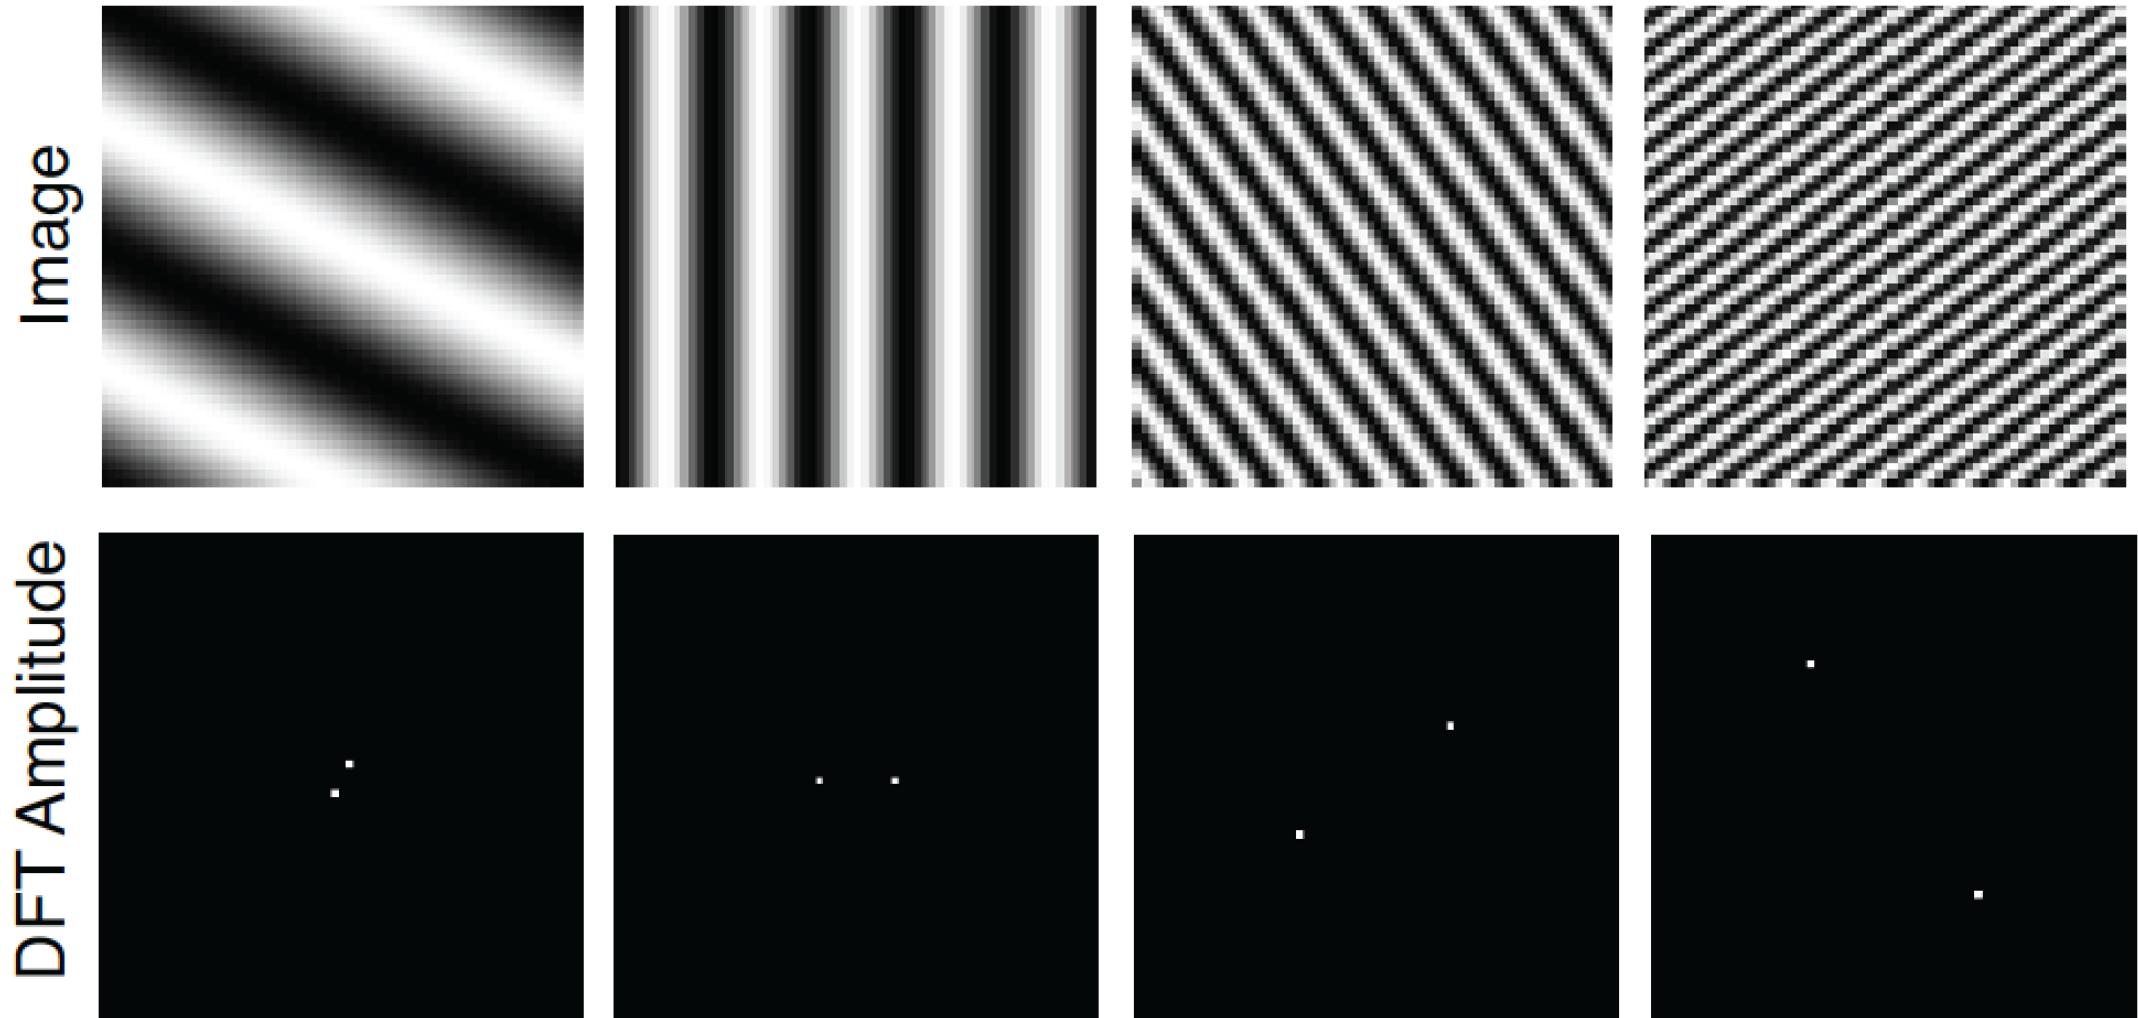


Visualizing the image Fourier transform



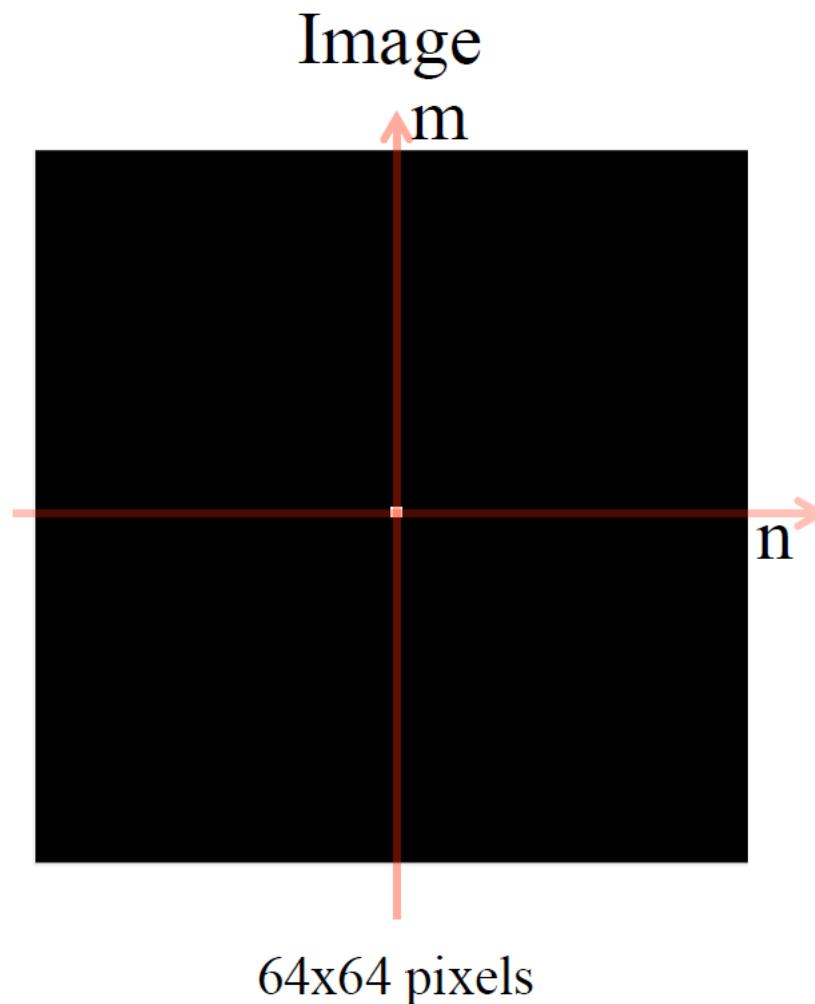
$$\cos \left(2\pi \left(\frac{u_0 n}{N} + \frac{v_0 m}{M} \right) \right) \quad \longleftrightarrow \quad \frac{1}{2} (\delta [u - u_0, v - v_0] + \delta [u + u_0, v + v_0])$$

Simple Fourier transforms

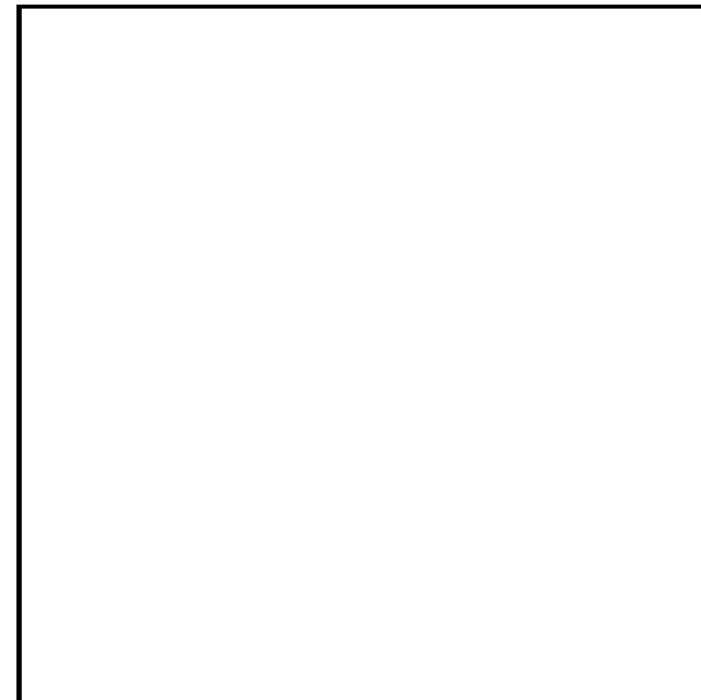


Images are 64x64 pixels. The wave is a cosine, therefore DFT phase is zero.

Some important Fourier transforms



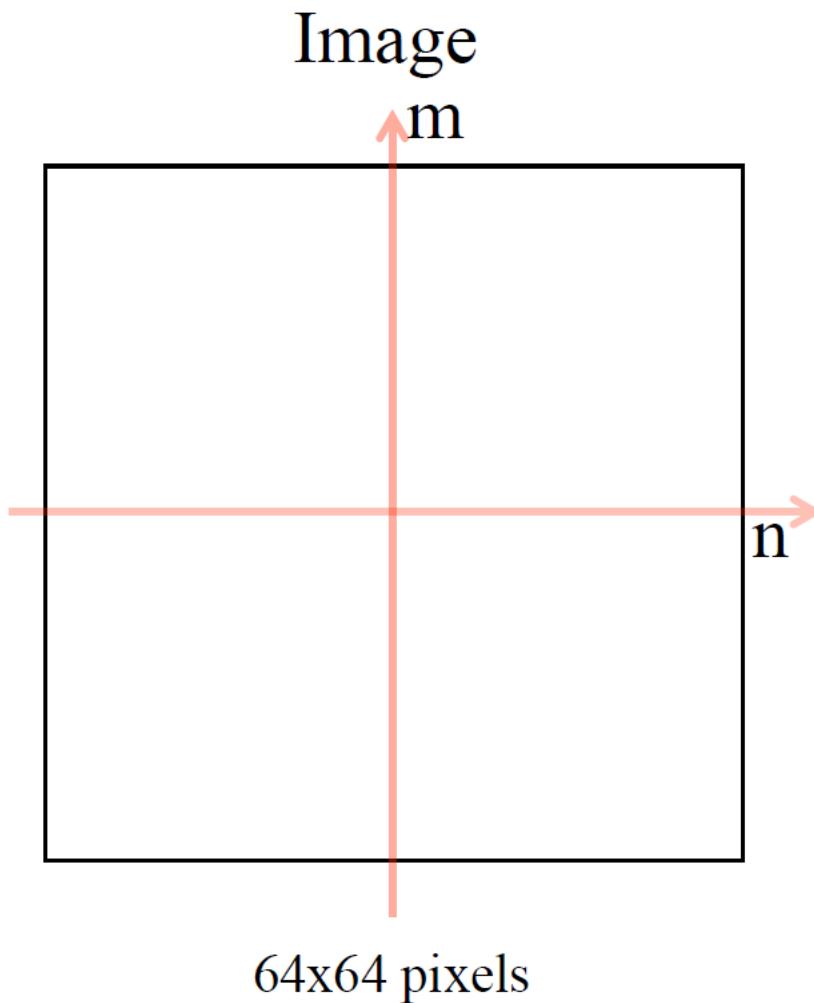
Magnitude DFT



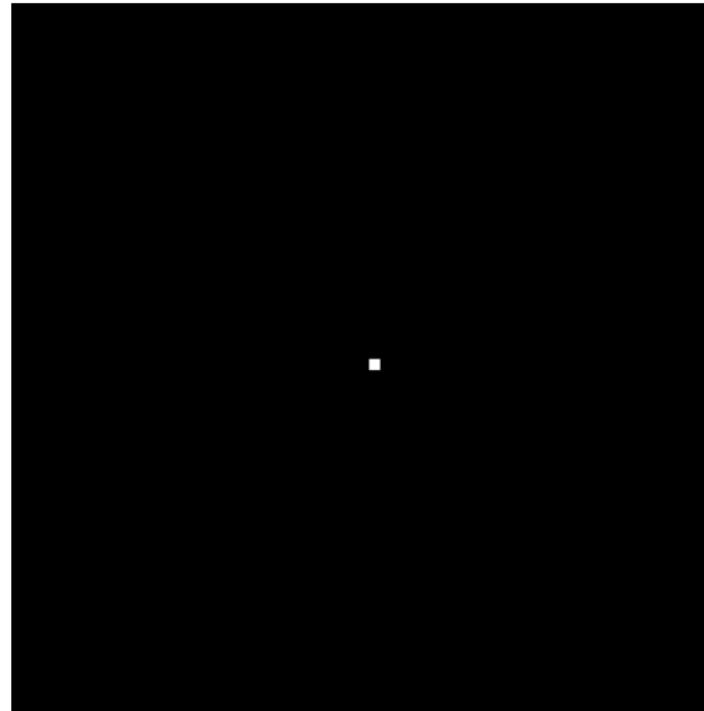
Phase DFT



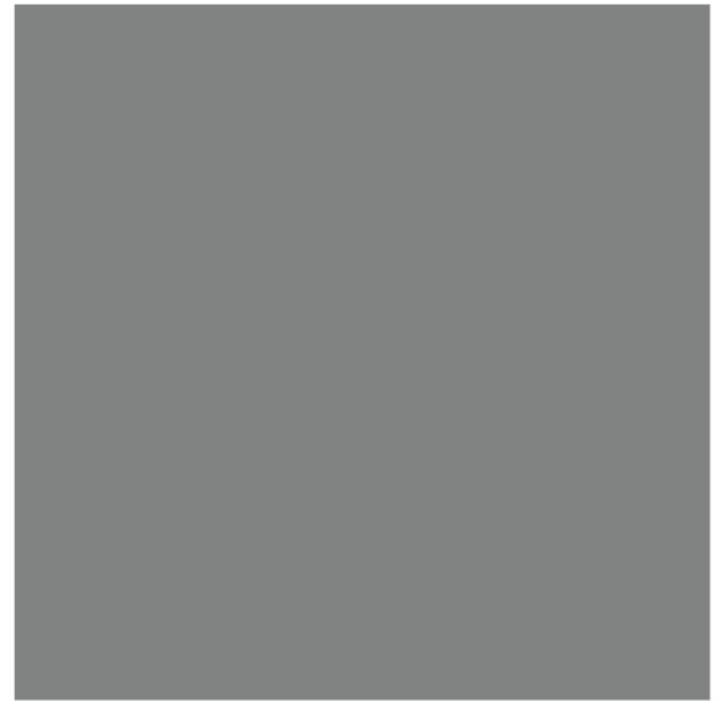
Some important Fourier transforms



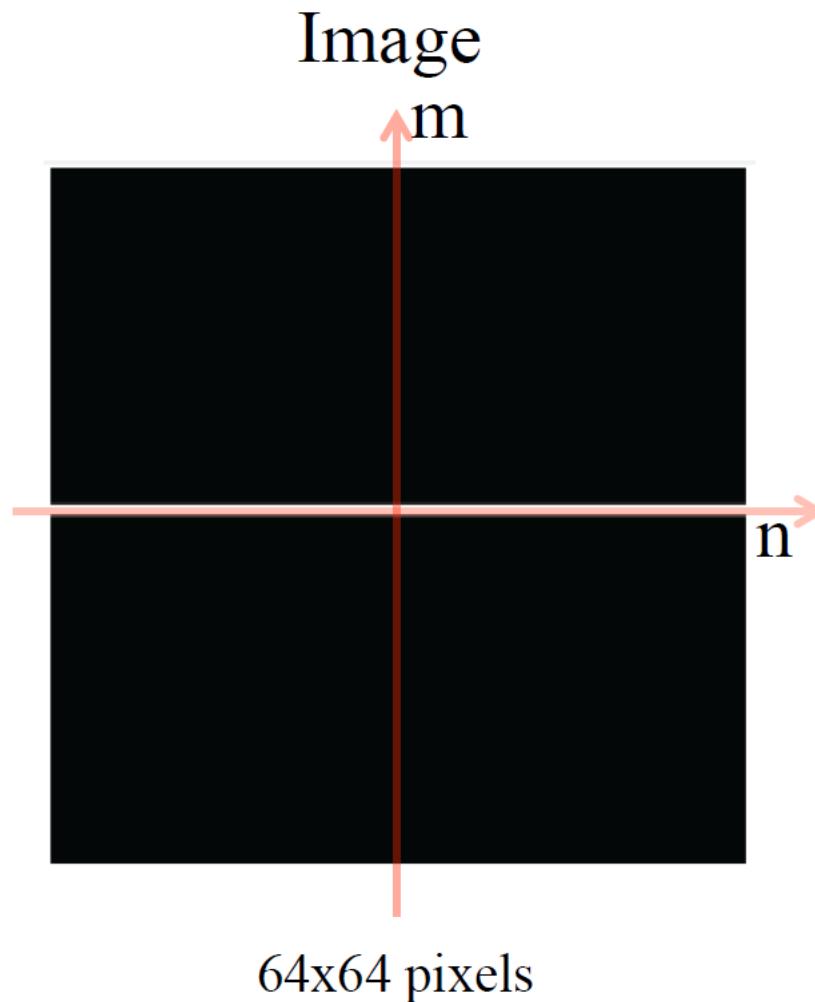
Magnitude DFT



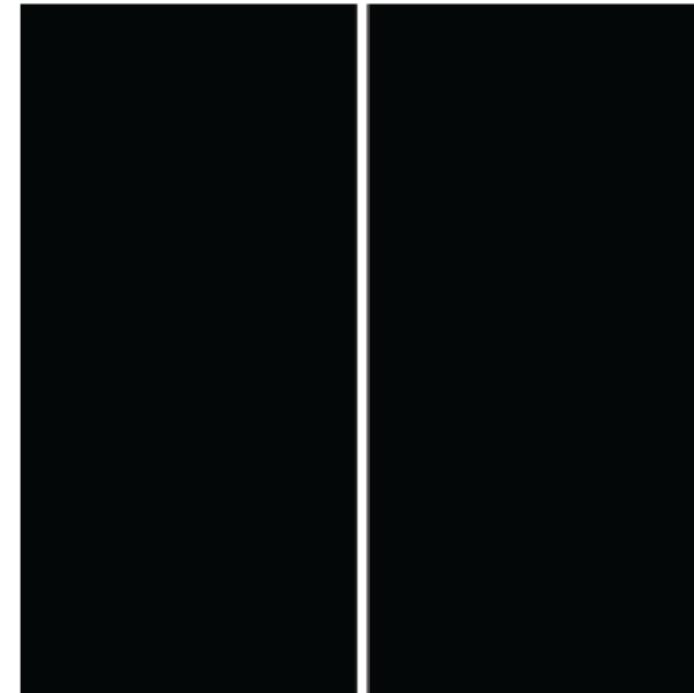
Phase DFT



Some important Fourier transforms



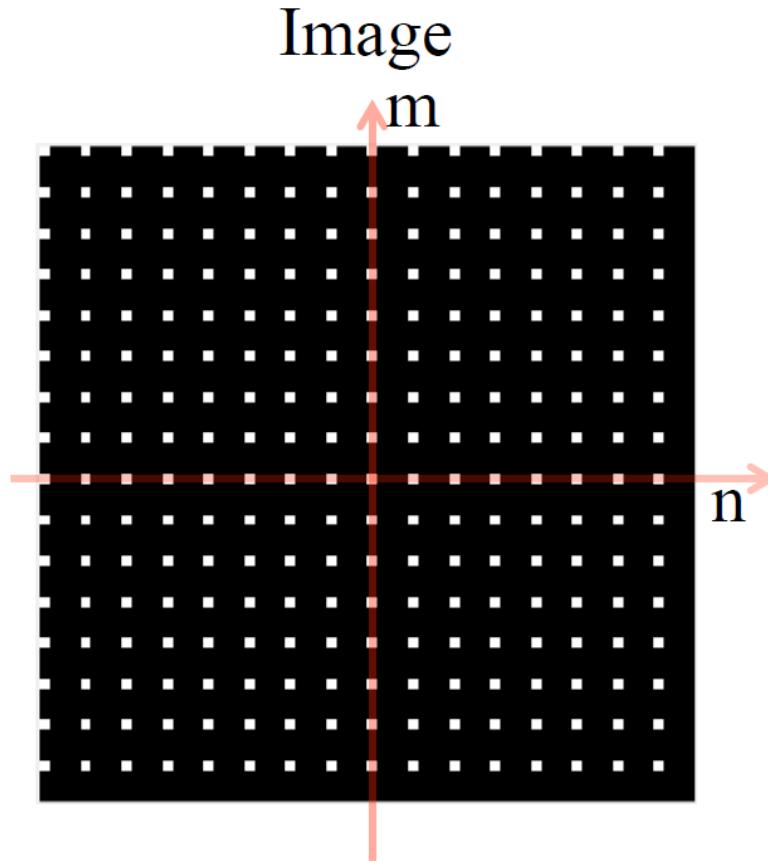
Magnitude DFT



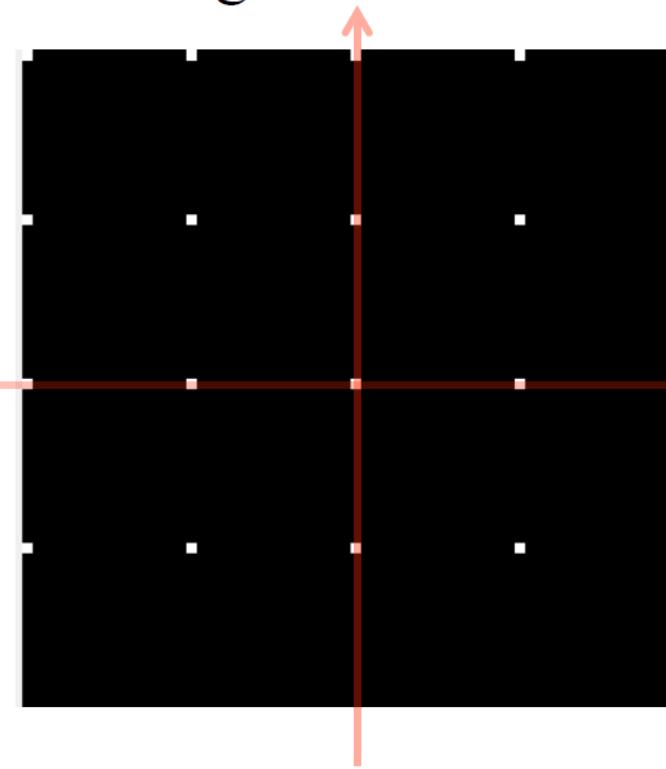
Phase DFT



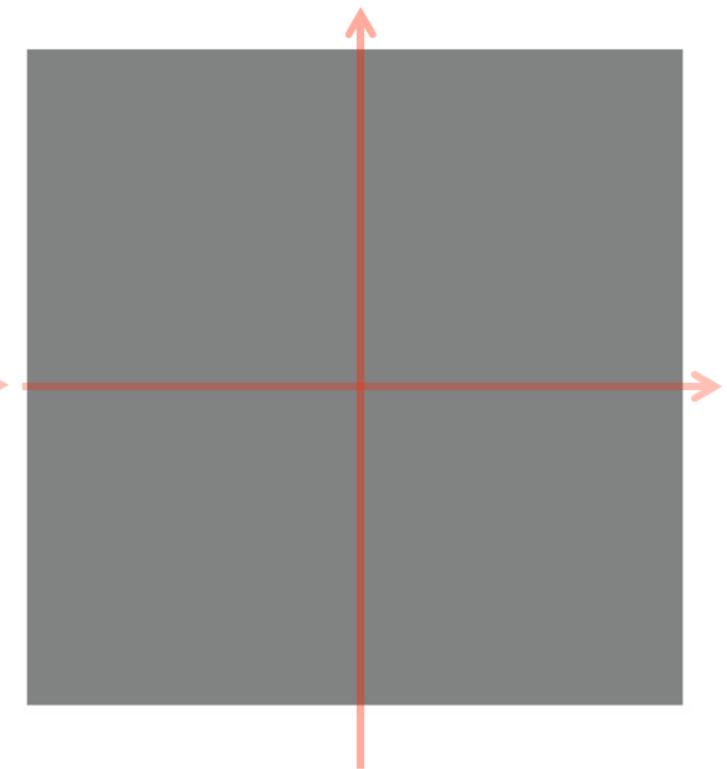
Some important Fourier transforms



Magnitude DFT



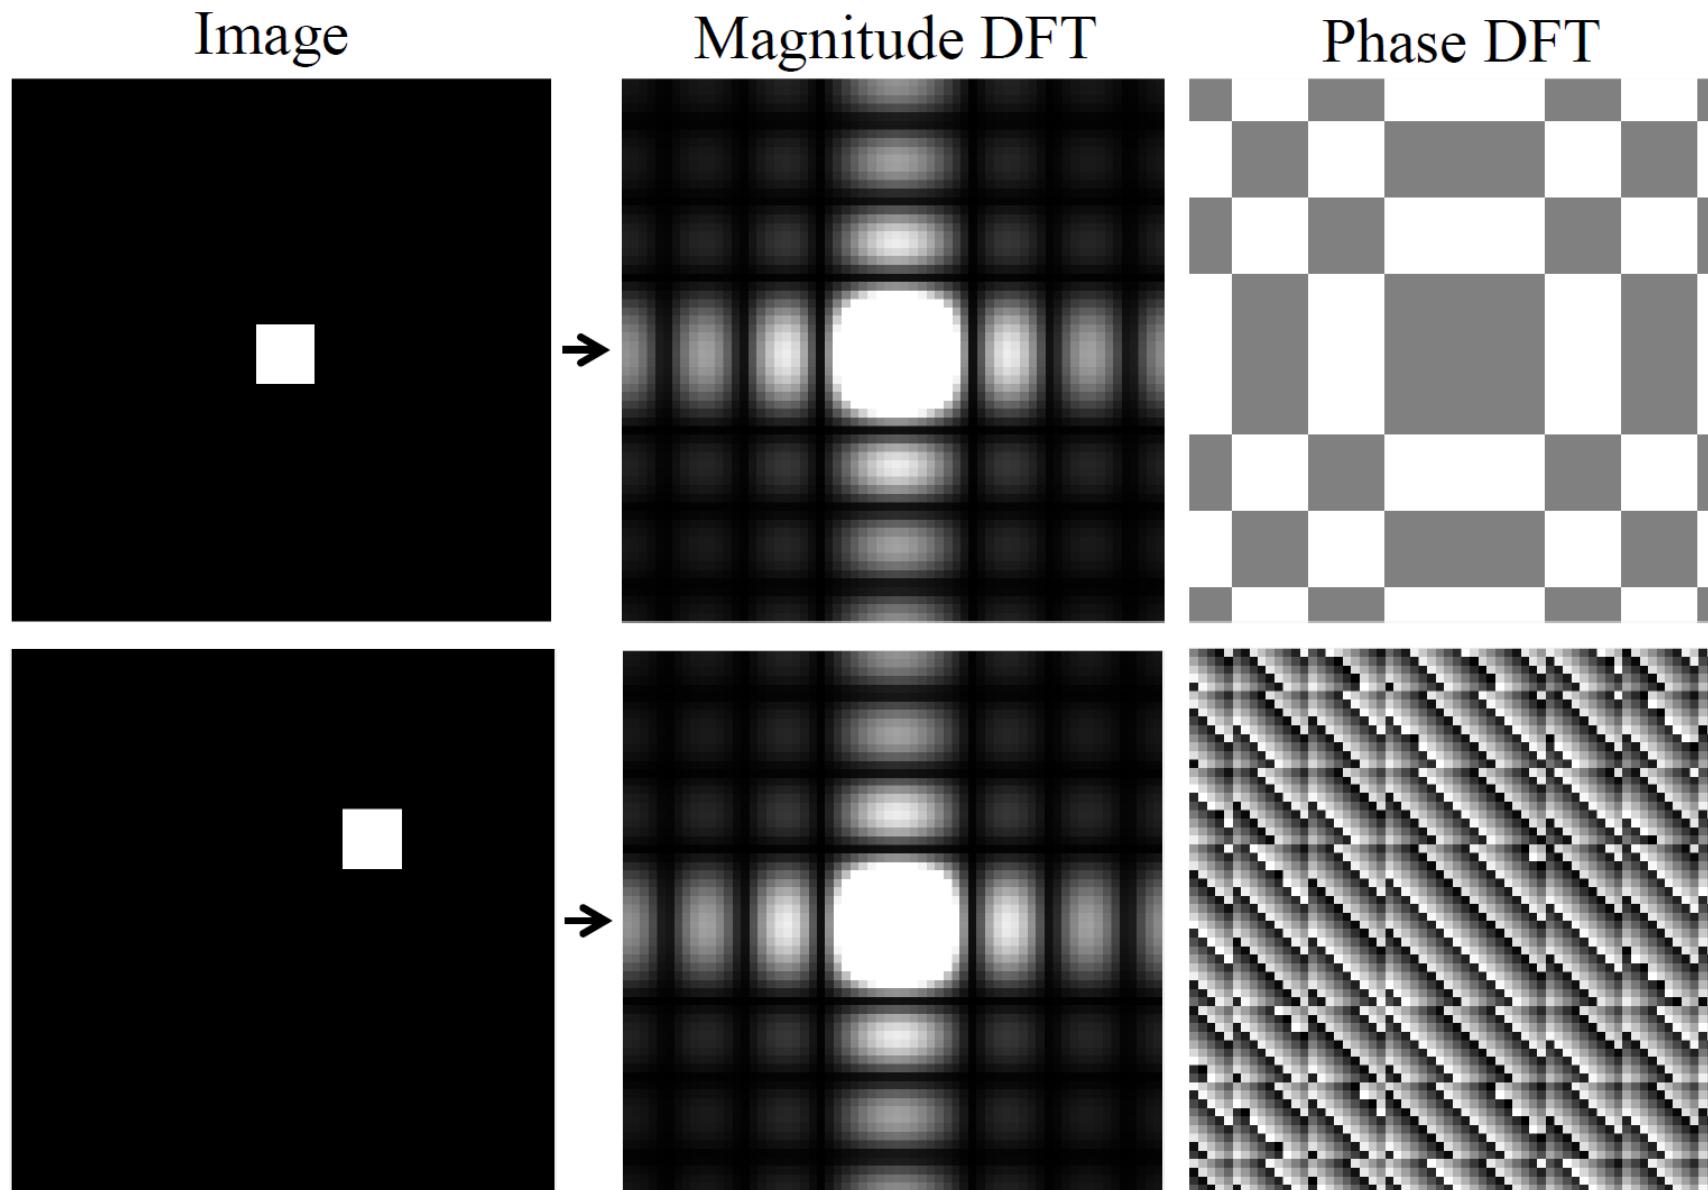
Phase DFT



$$\delta_k [n, m] = \sum_{s=0}^{N/k-1} \sum_{r=0}^{M/k-1} \delta [n - sk, m - rk] \quad \leftrightarrow \quad \Delta_k [u, v] = \frac{NM}{k^2} \sum_{s=0}^{k-1} \sum_{r=0}^{k-1} \delta \left[u - s \frac{N}{k}, v - r \frac{M}{k} \right]$$

delta train, delta comb, impulse train

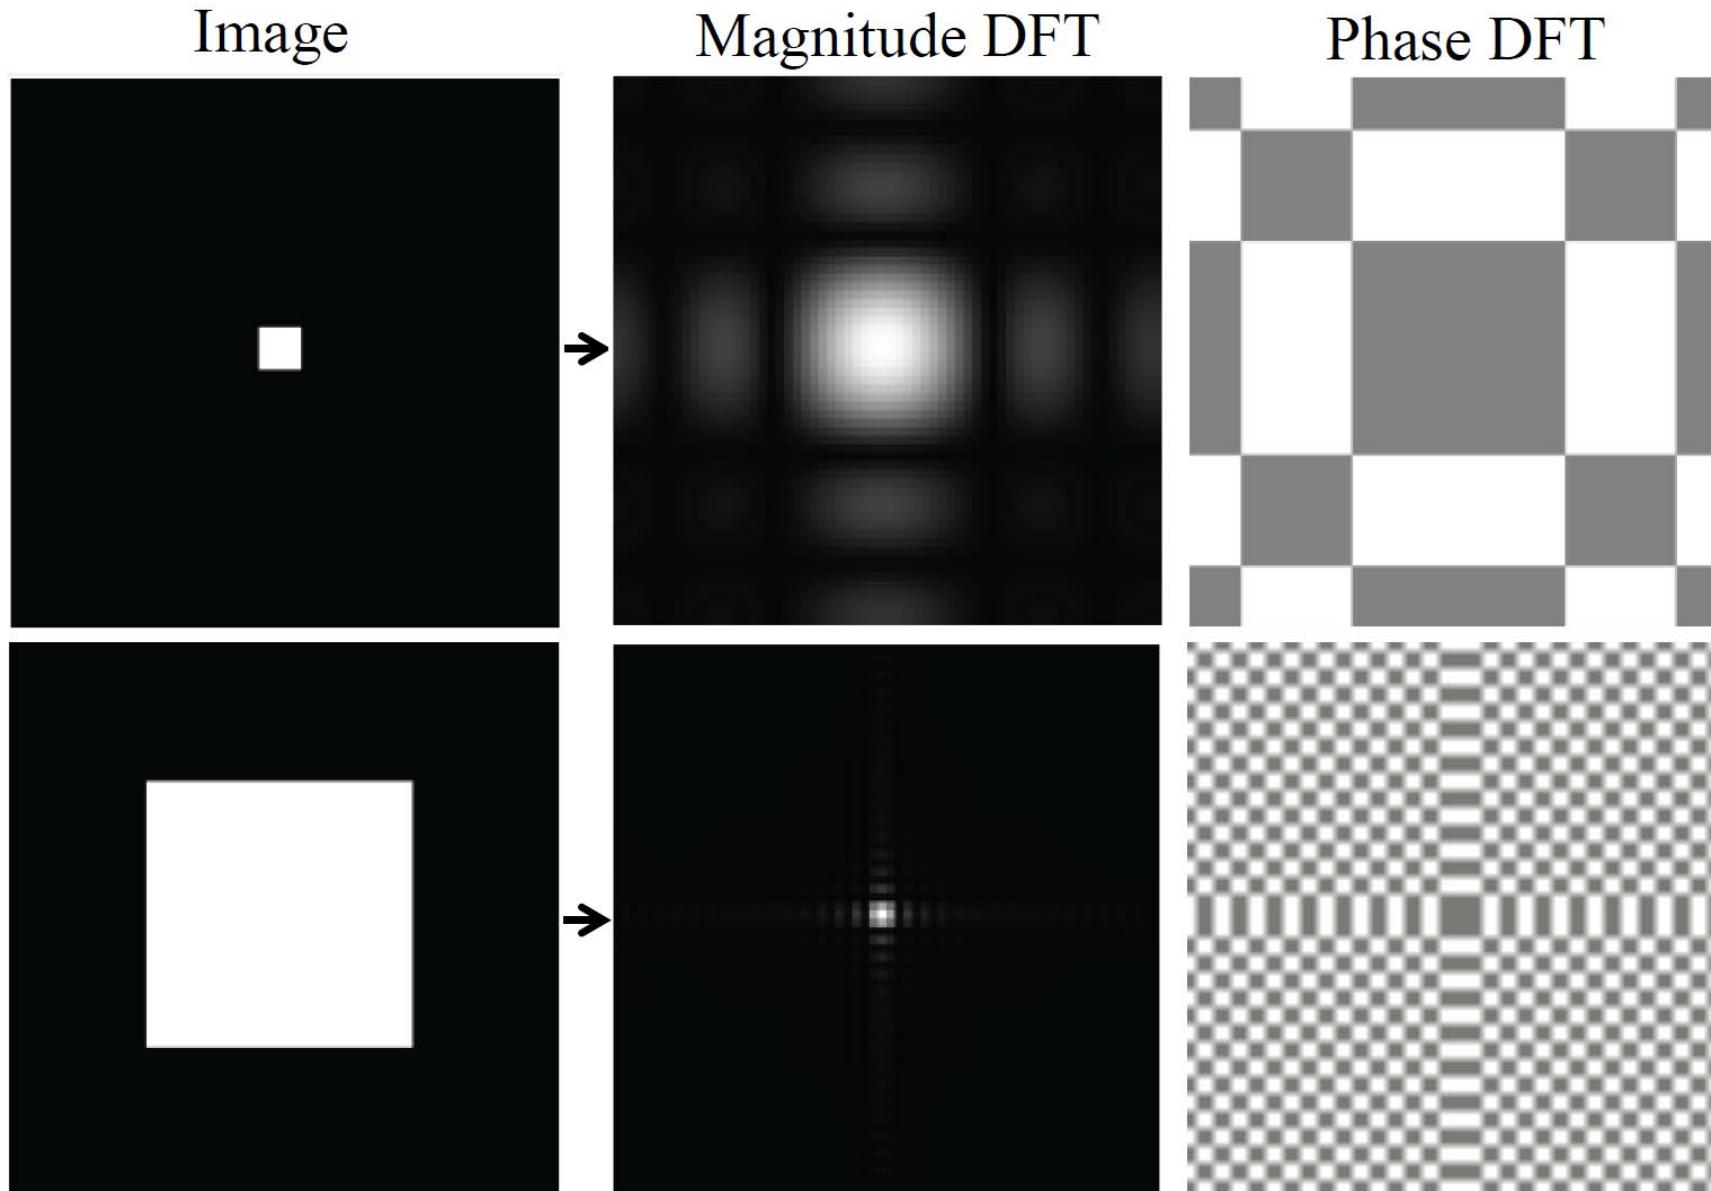
Some important Fourier transforms



Translation

Shifts of an image only produce changes on the phase of the DFT.

Some important Fourier transforms

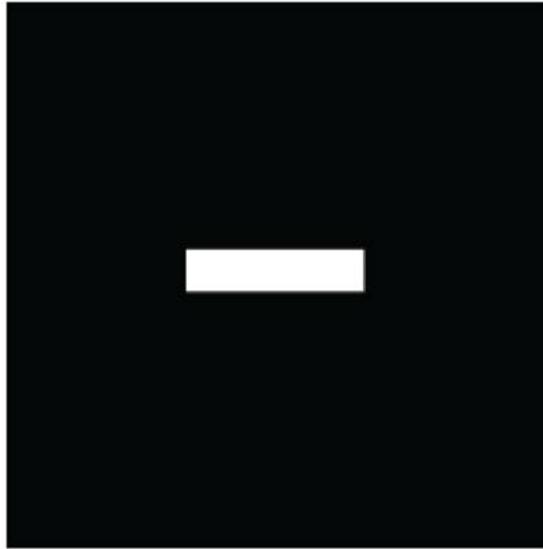


Scale

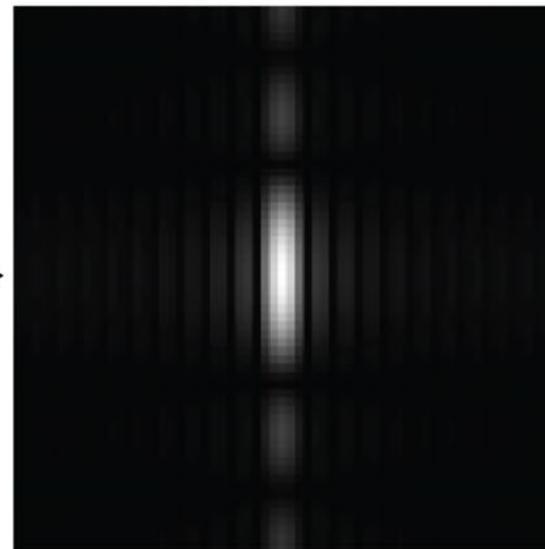
Small image details produce content in high spatial frequencies

Some important Fourier transforms

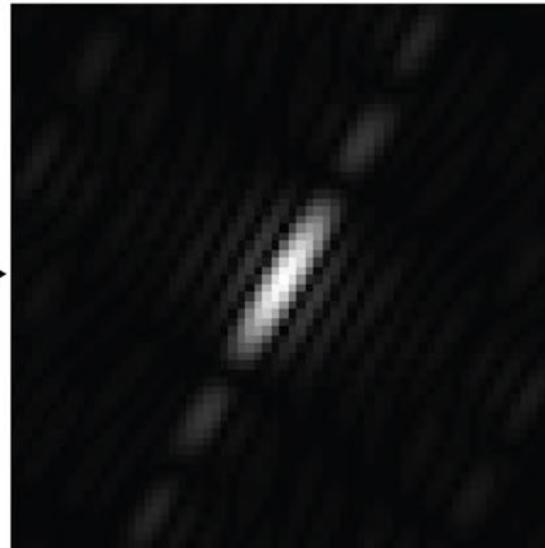
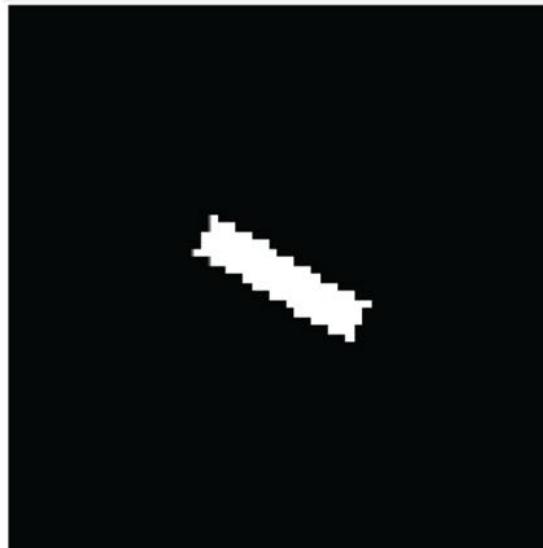
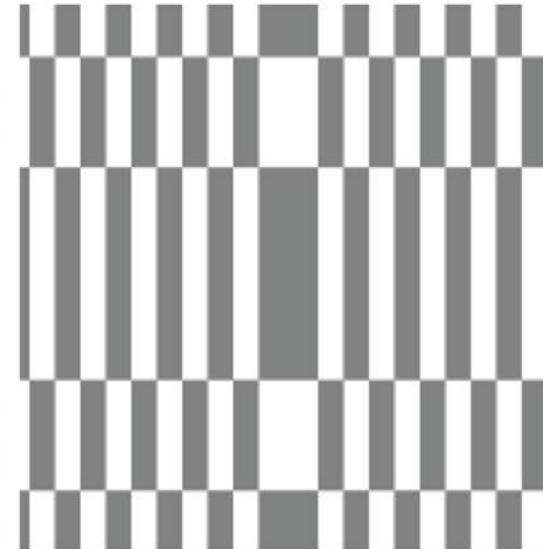
Image



Magnitude DFT



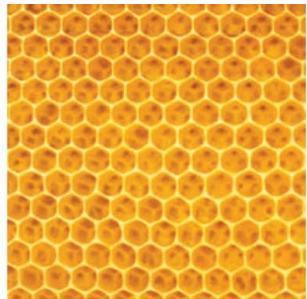
Phase DFT



Orientation

A line transforms to a line oriented perpendicularly to the first.

The DFT Game: find the right pairs



a)

b)

c)

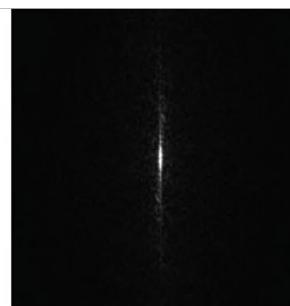
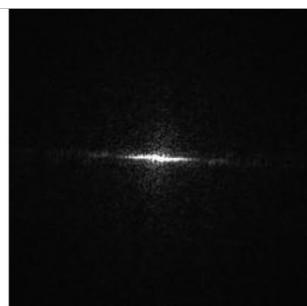
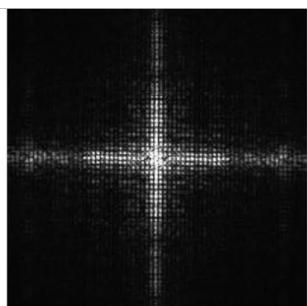
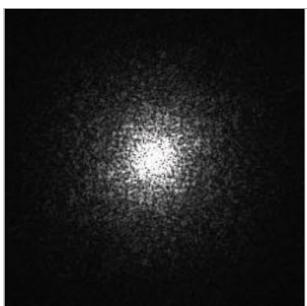
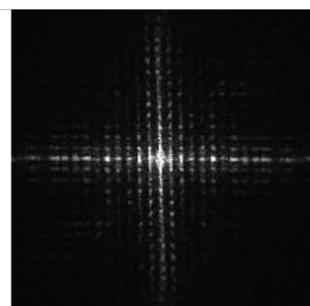
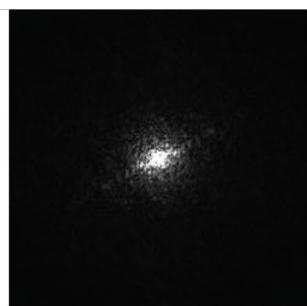
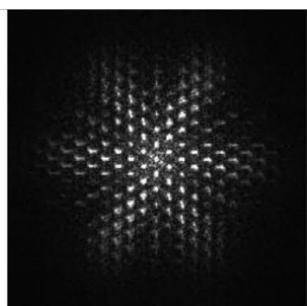
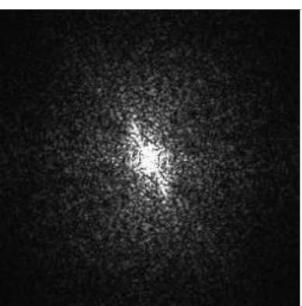
d)

e)

f)

g)

h)



1)

2)

3)

4)

5)

6)

7)

8)

The inverse Discrete Fourier transform

2D Discrete Fourier Transform (DFT) transforms an image $f[n,m]$ into $F[u,v]$ as:

$$F[u, v] = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} f[n, m] \exp\left(-2\pi j \left(\frac{un}{N} + \frac{vm}{M}\right)\right)$$

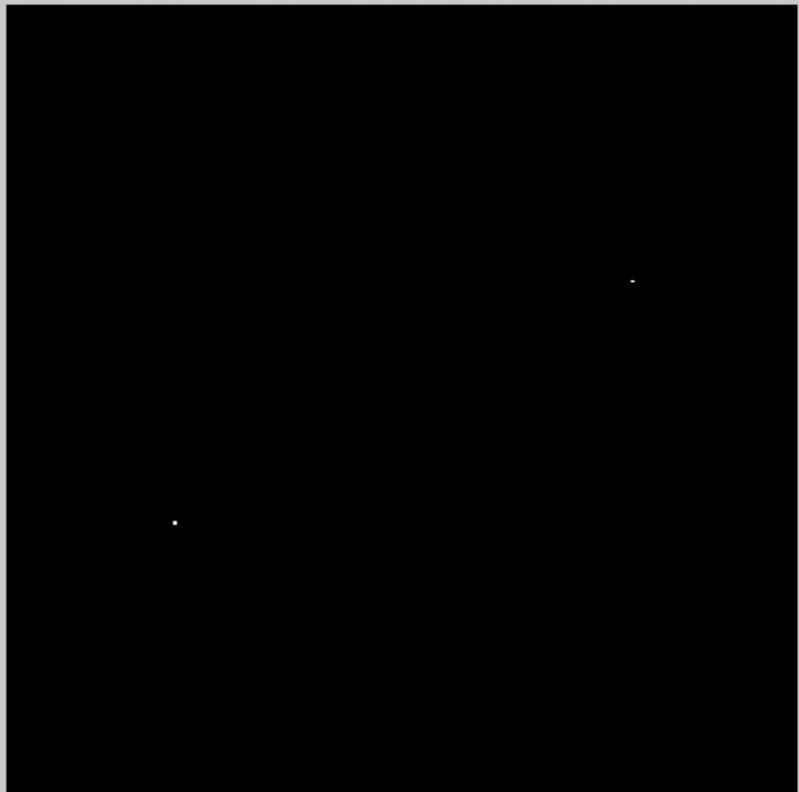
The inverse of the 2D DFT is:

$$f[n, m] = \frac{1}{NM} \sum_{u=0}^{N-1} \sum_{v=0}^{M-1} F[u, v] \exp\left(+2\pi j \left(\frac{un}{N} + \frac{vm}{M}\right)\right)$$

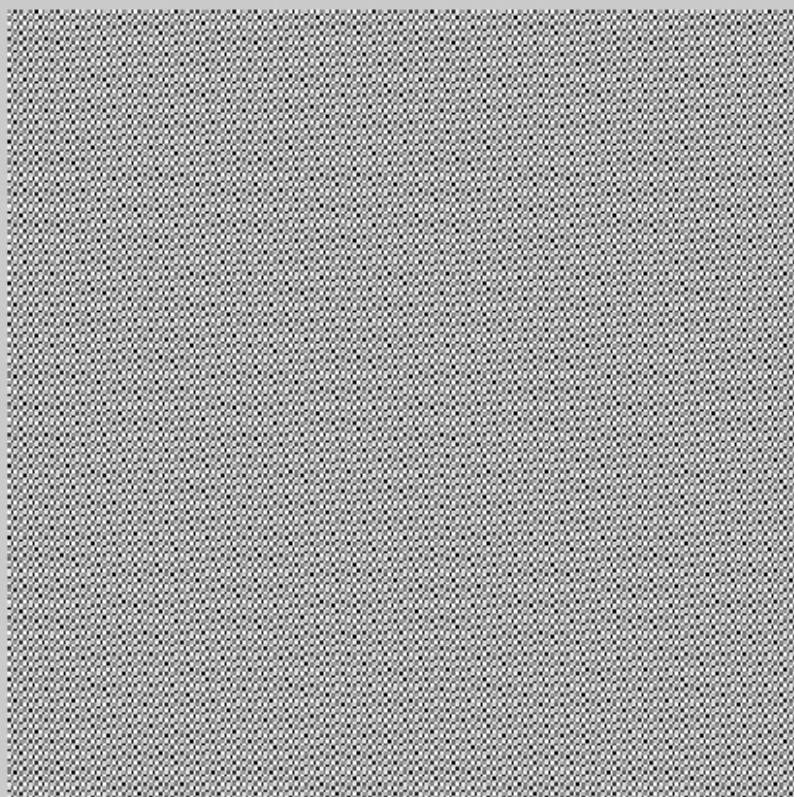
How does summing waves ends up giving back a picture?

2

2



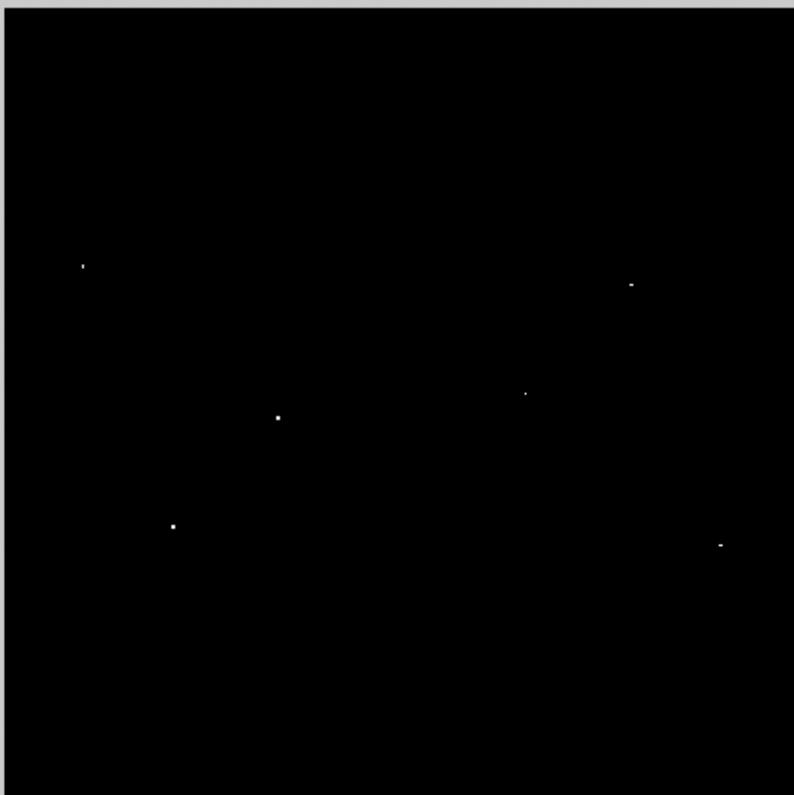
#1: Range [0, 1]
Dims [256, 256]



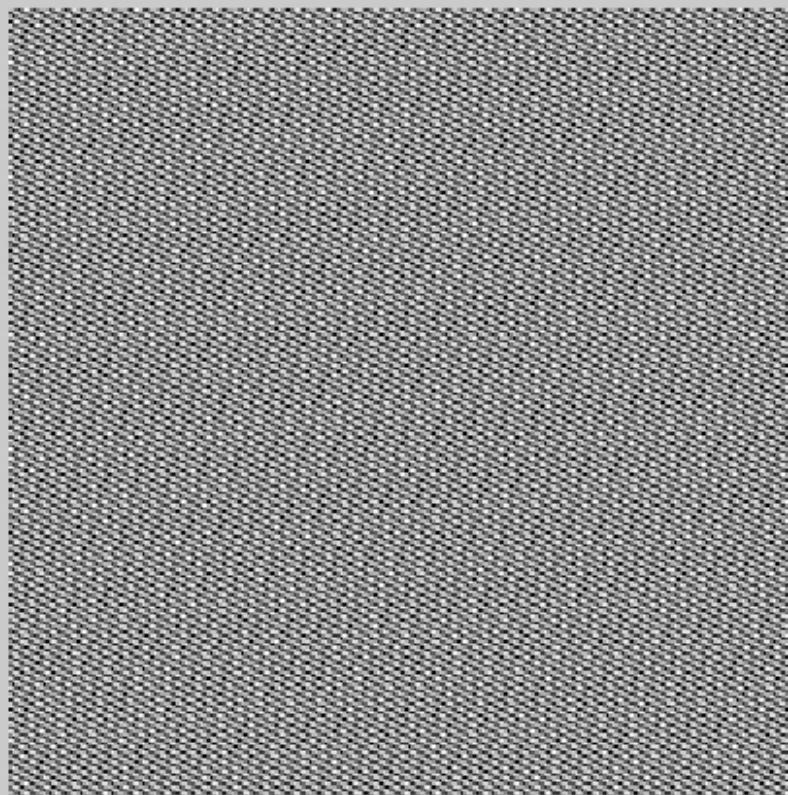
#2: Range [0.000109, 0.0267]
Dims [256, 256]

6

6

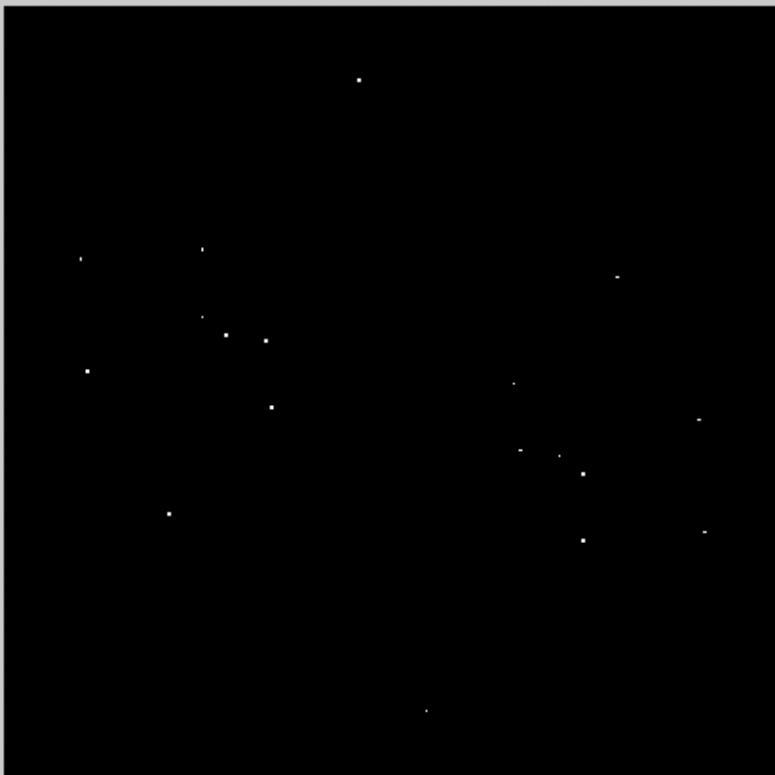


#1: Range [0, 1]
Dims [256, 256]

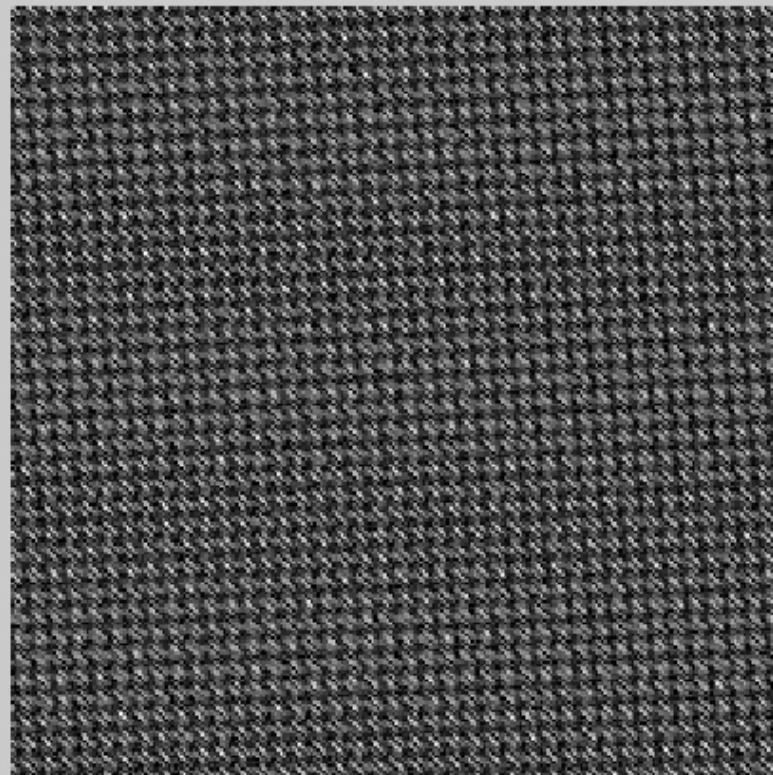


#2: Range [1.89e-007, 0.226]
Dims [256, 256]

18



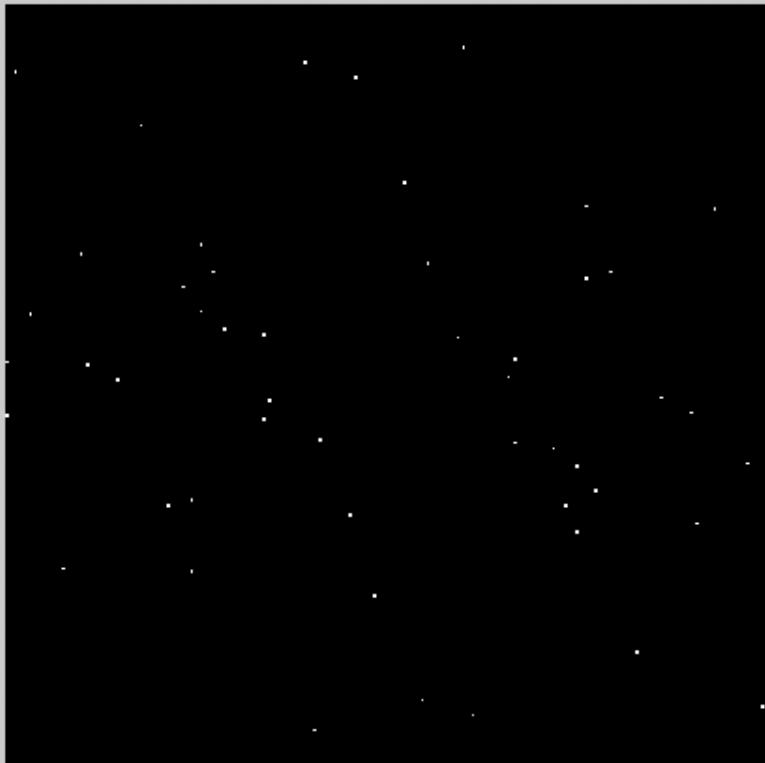
#1: Range [0, 1]
Dims [256, 256]



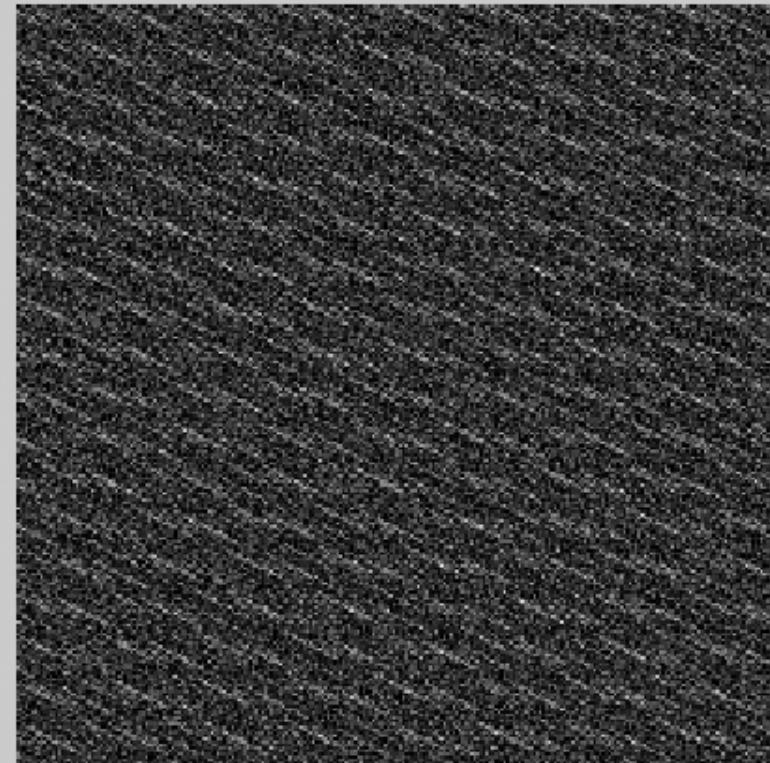
#2: Range [4.79e-007, 0.503]
Dims [256, 256]

50

50



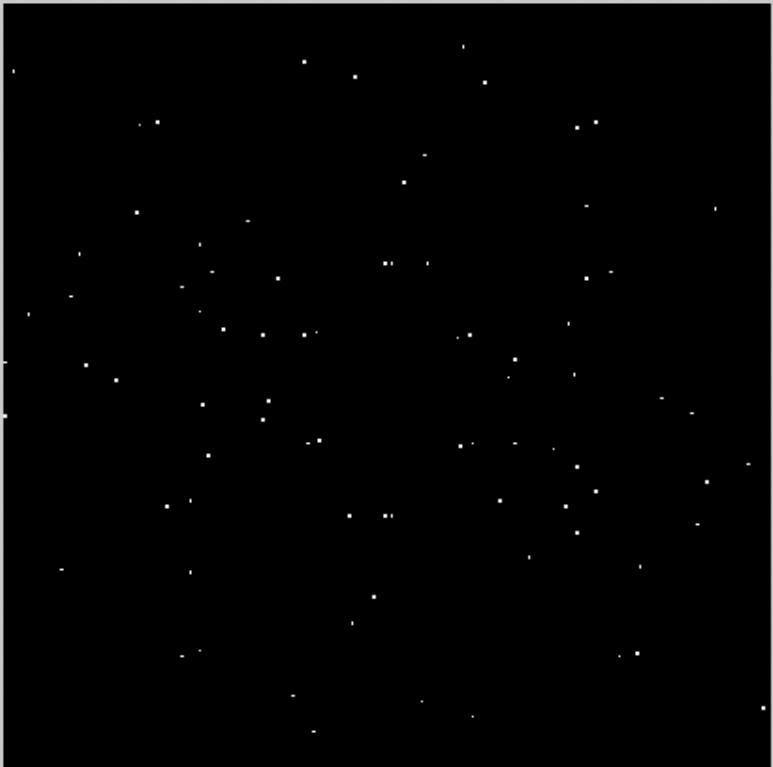
#1: Range [0, 1]
Dims [256, 256]



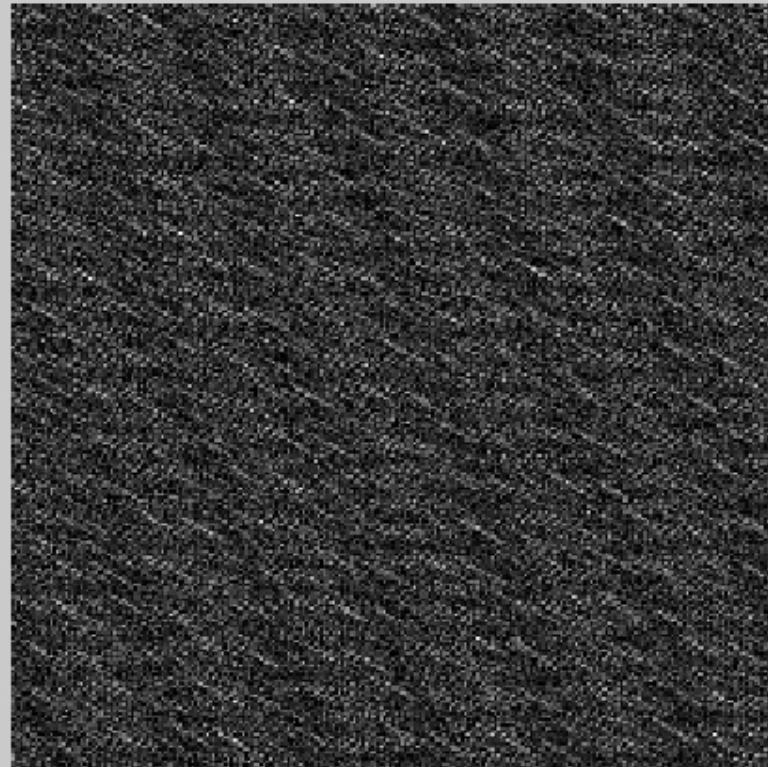
#2: Range [8.5e-006, 1.7]
Dims [256, 256]

82

82



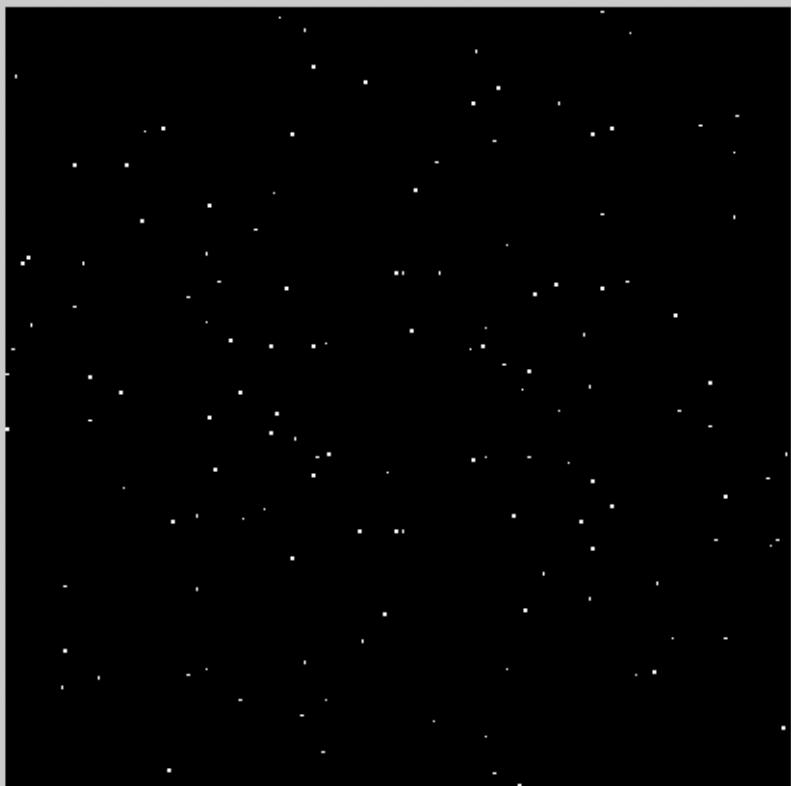
#1: Range [0, 1]
Dims [256, 256]



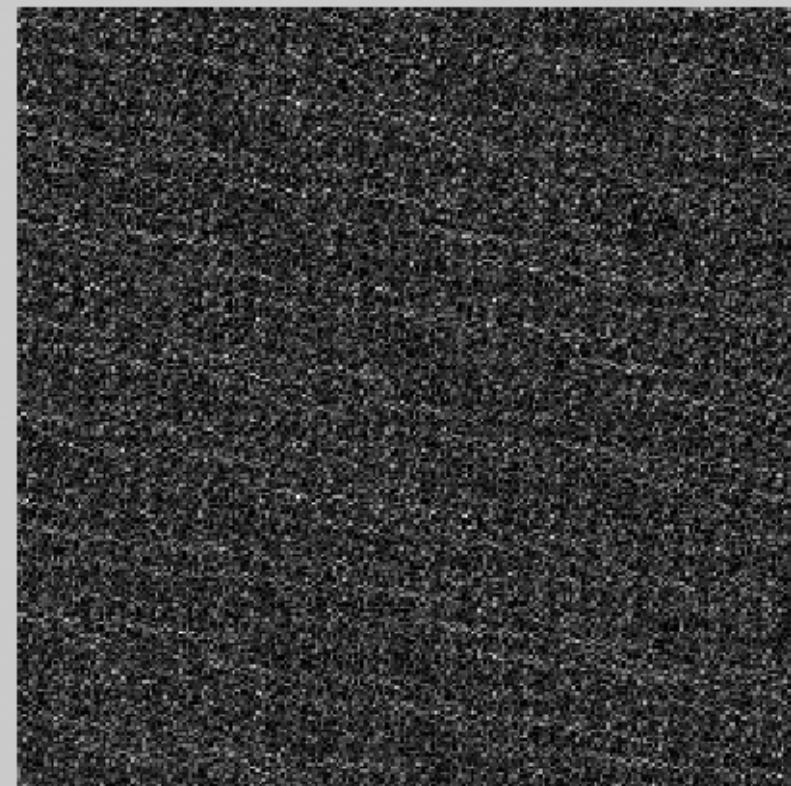
#2: Range [3.85e-007, 2.21]
Dims [256, 256]

136

136



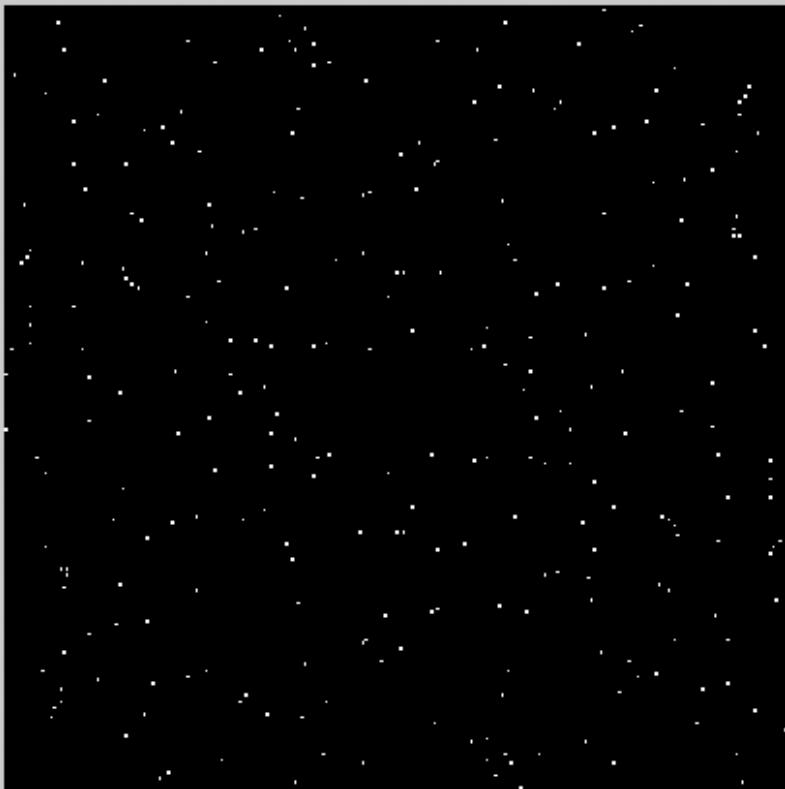
#1: Range [0, 1]
Dims [256, 256]



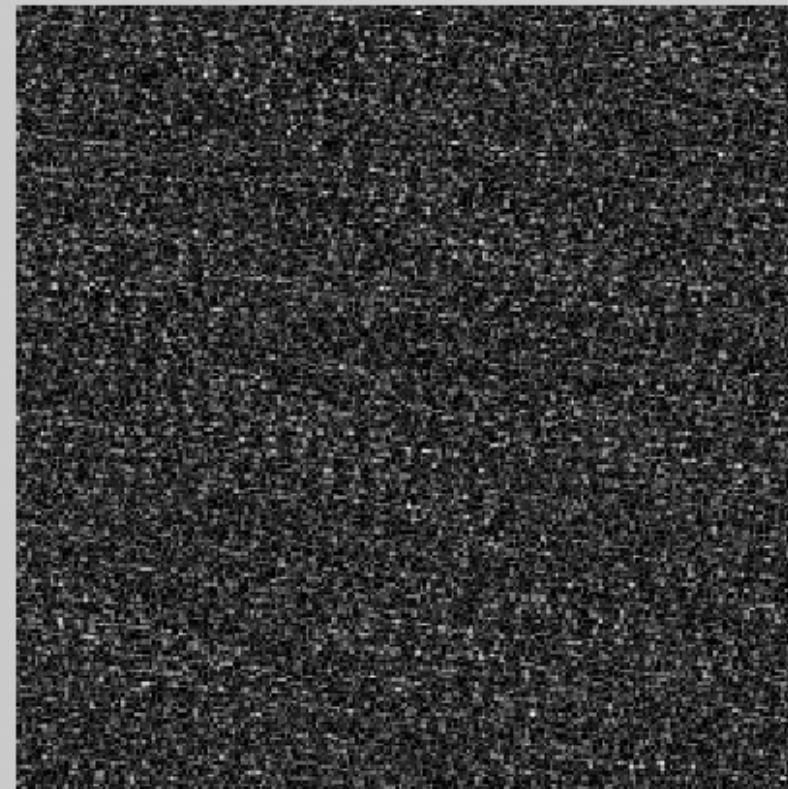
#2: Range [8.25e-006, 3.48]
Dims [256, 256]

282

282



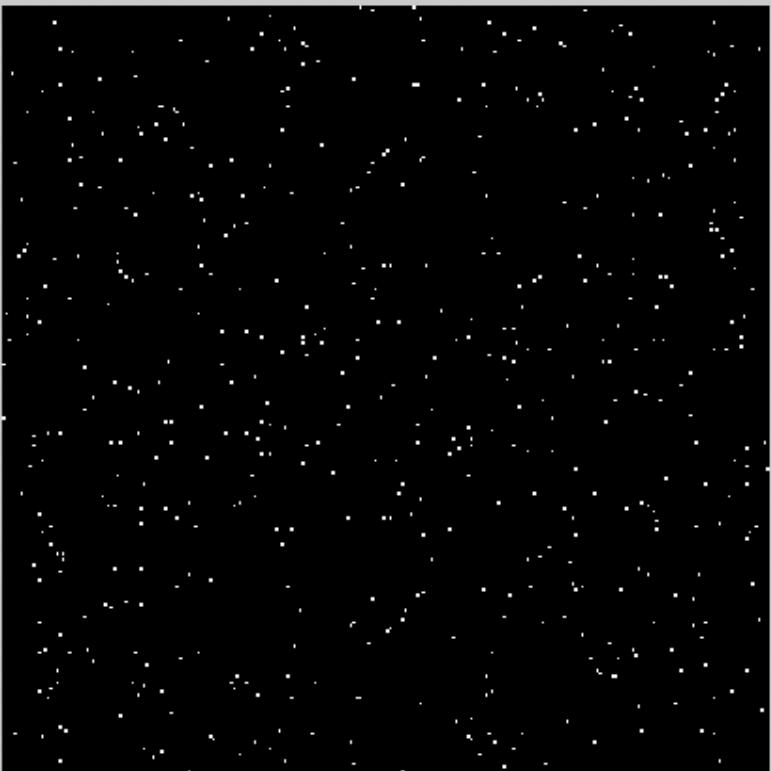
#1: Range [0, 1]
Dims [256, 256]



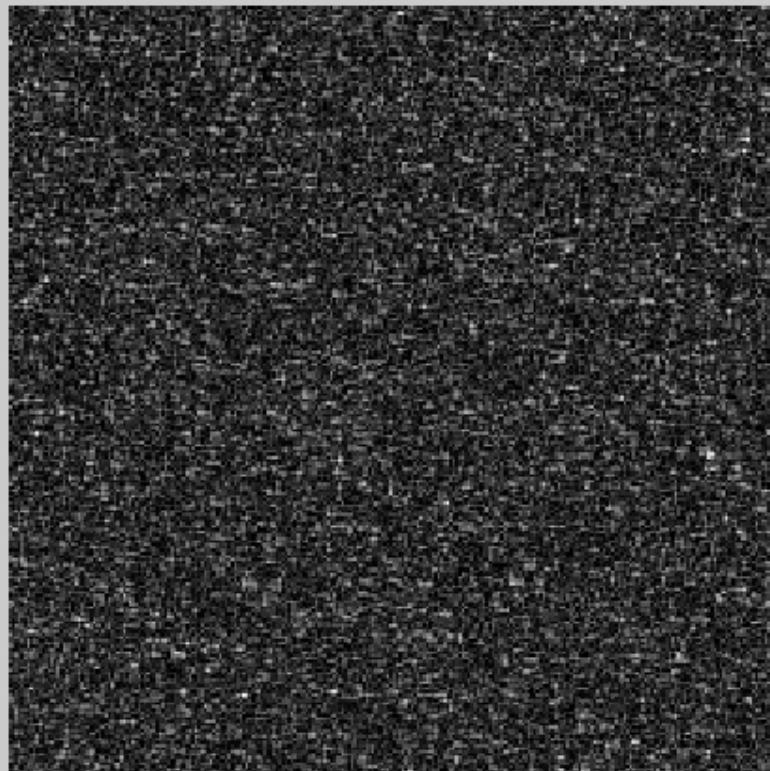
#2: Range [1.39e-005, 5.88]
Dims [256, 256]

538

538



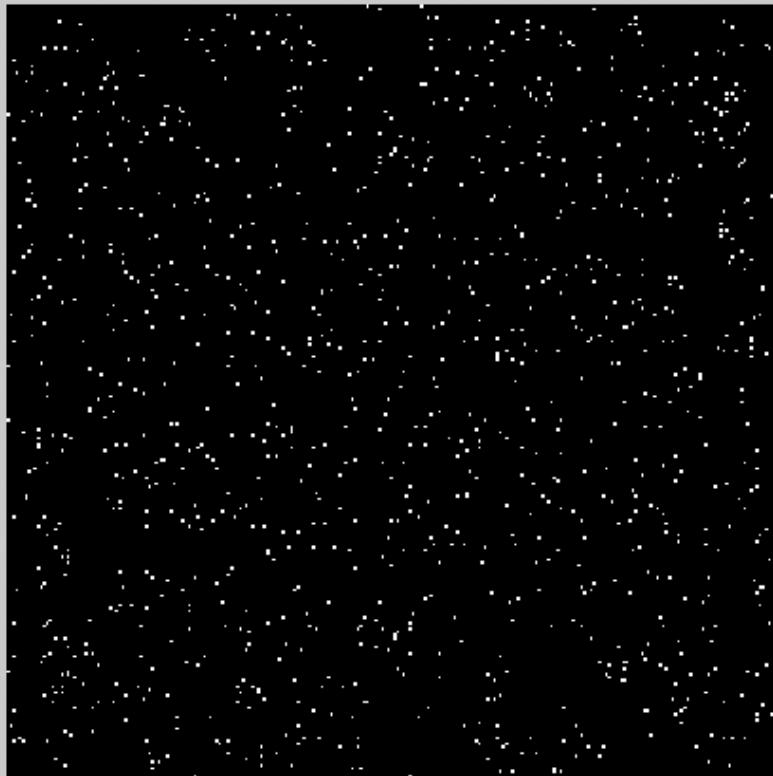
#1: Range [0, 1]
Dims [256, 256]



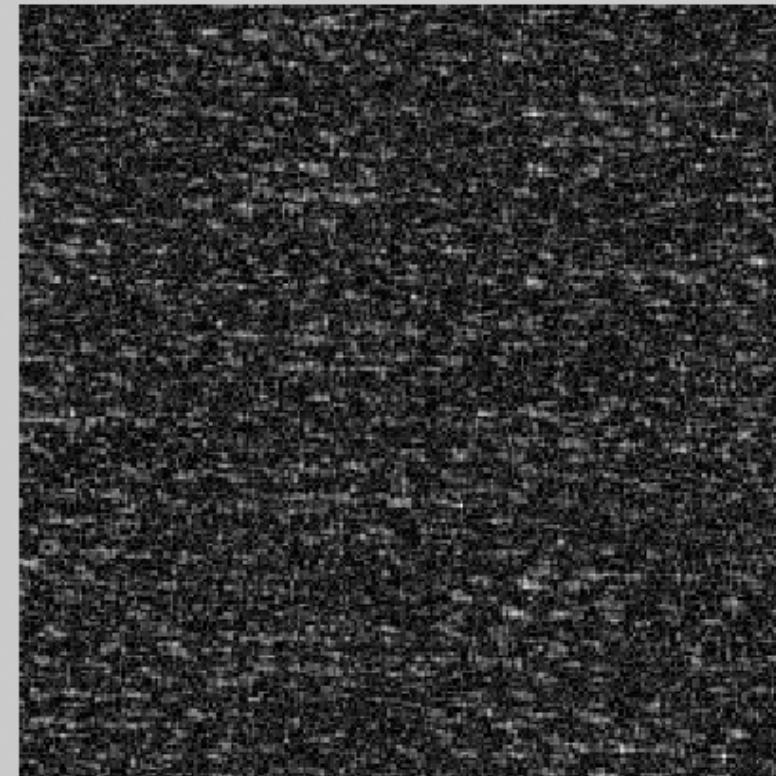
#2: Range [6.17e-006, 8.4]
Dims [256, 256]

1088

1088



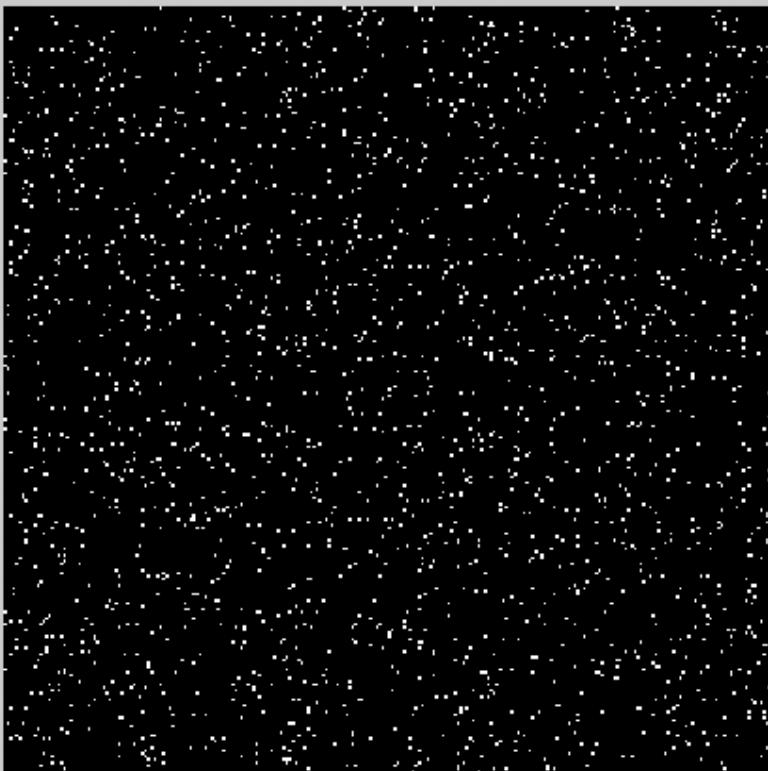
#1: Range [0, 1]
Dims [256, 256]



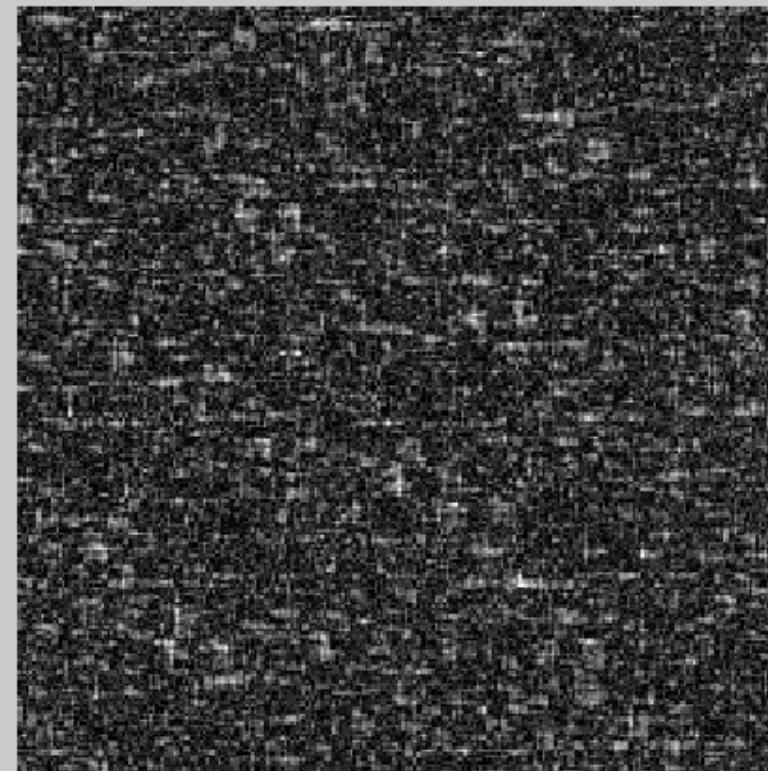
#2: Range [9.99e-005, 15]
Dims [256, 256]

2094

2094



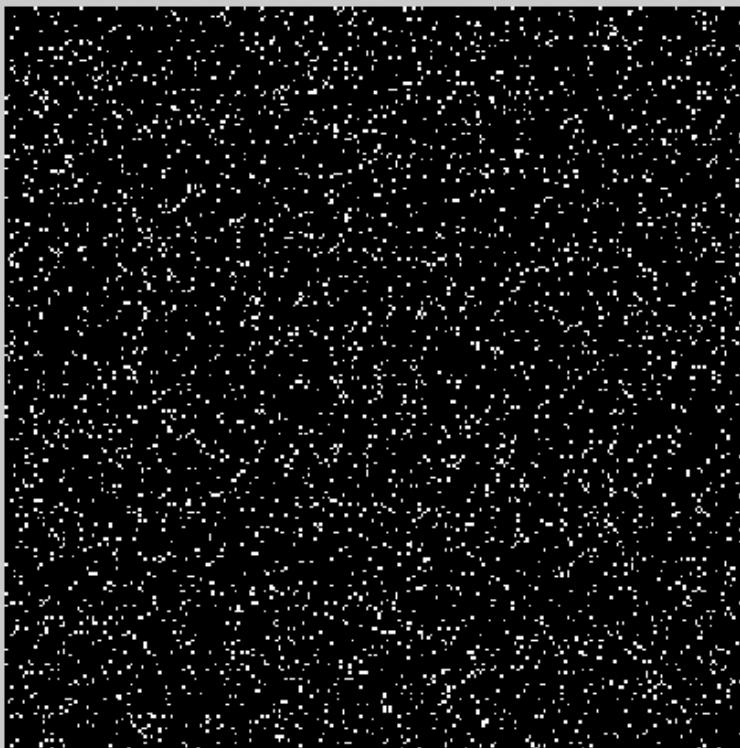
#1: Range [0, 1]
Dims [256, 256]



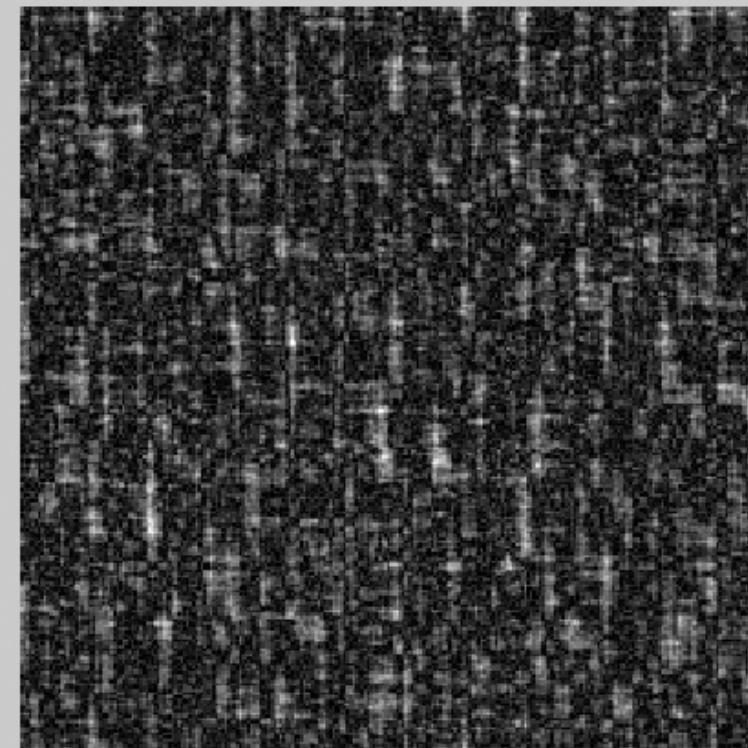
#2: Range [8.7e-005, 19]
Dims [256, 256]

4052

4052



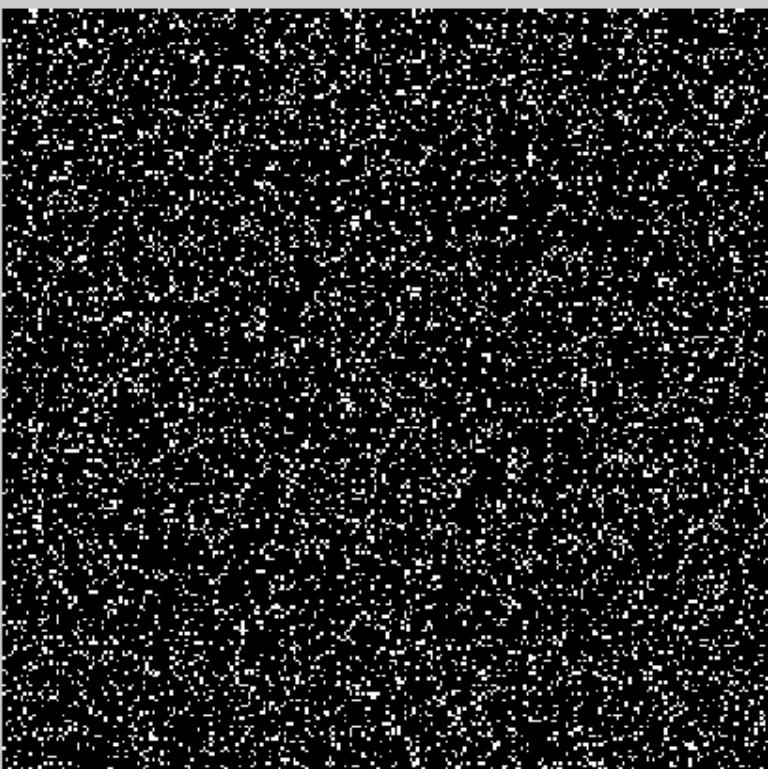
#1: Range [0, 1]
Dims [256, 256]



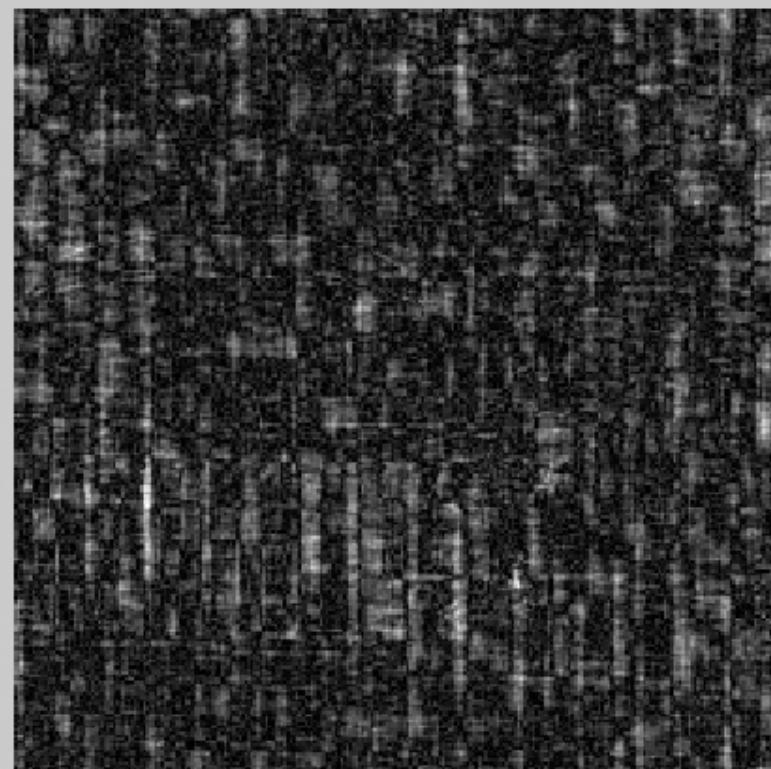
#2: Range [0.000556, 37.7]
Dims [256, 256]

8056.

8056



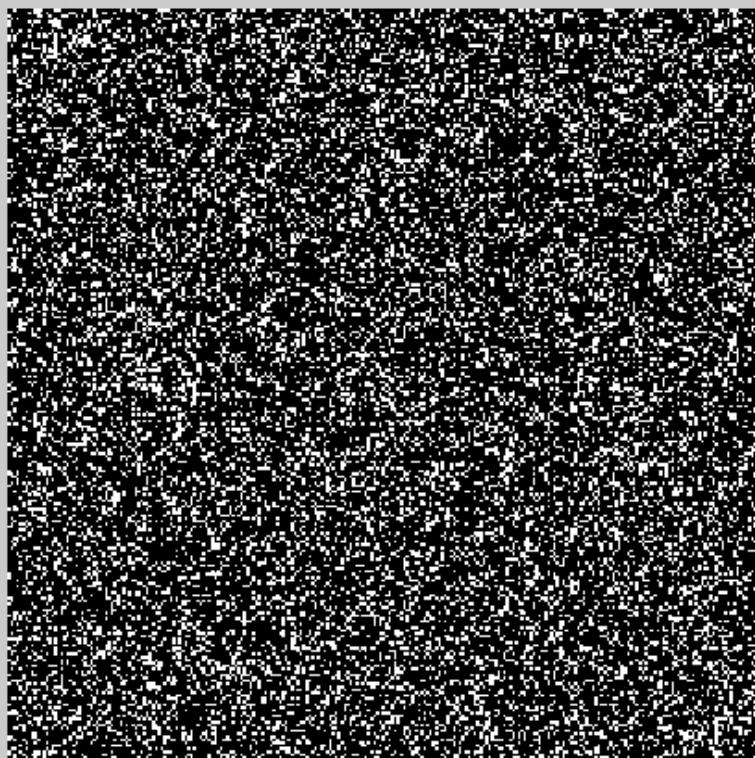
#1: Range [0, 1]
Dims [256, 256]



#2: Range [0.00032, 64.5]
Dims [256, 256]

15366

15366



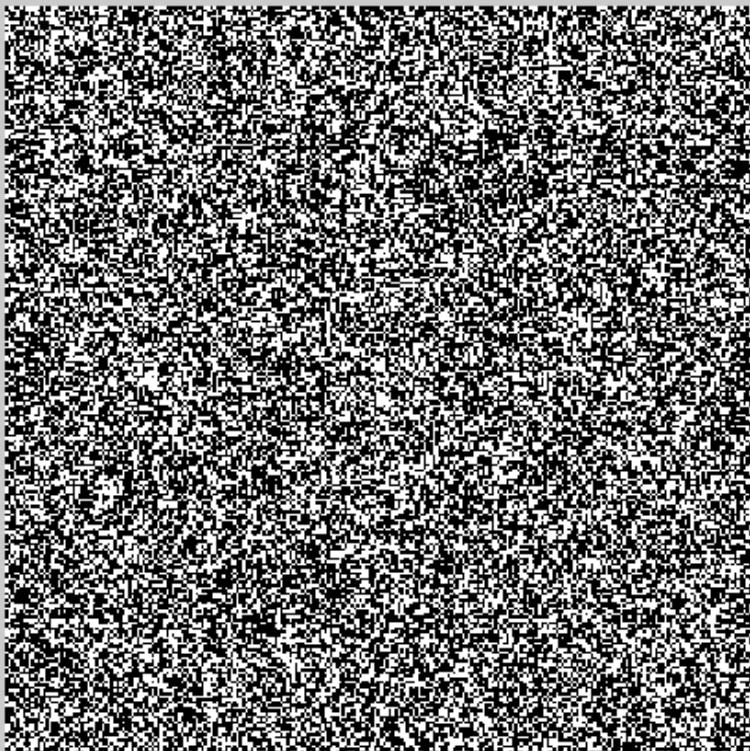
#1: Range [0, 1]
Dims [256, 256]



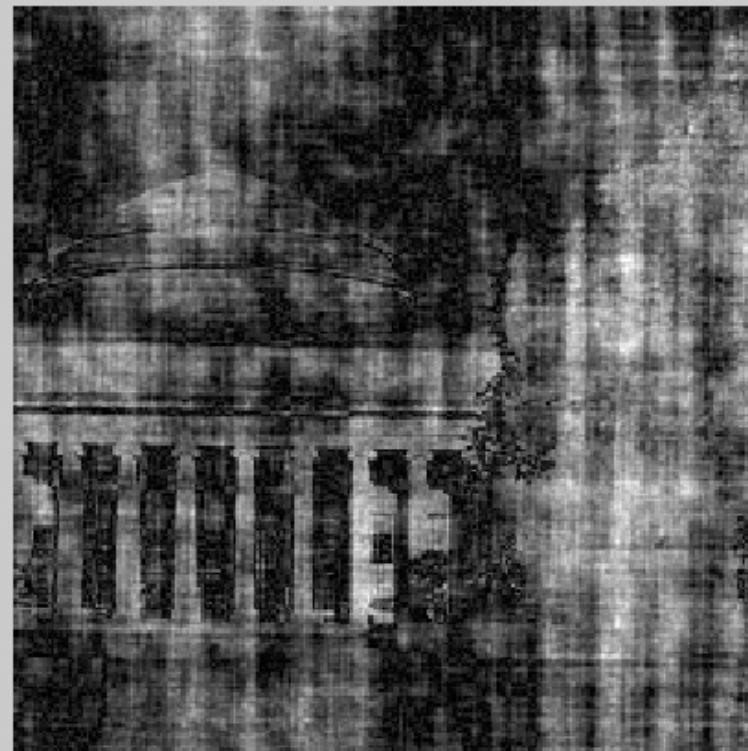
#2: Range [0.000231, 91.1]
Dims [256, 256]

28743

28743



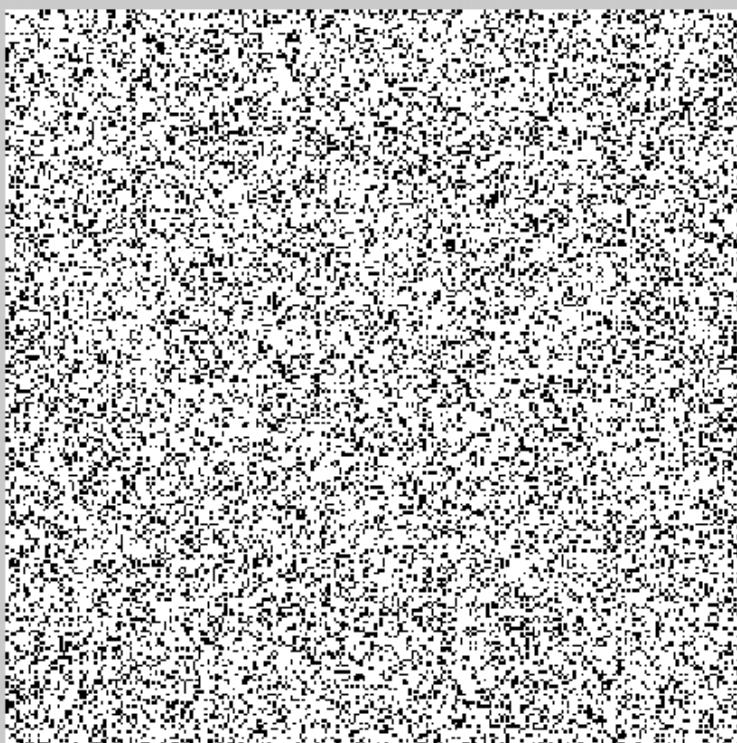
#1: Range [0, 1]
Dims [256, 256]



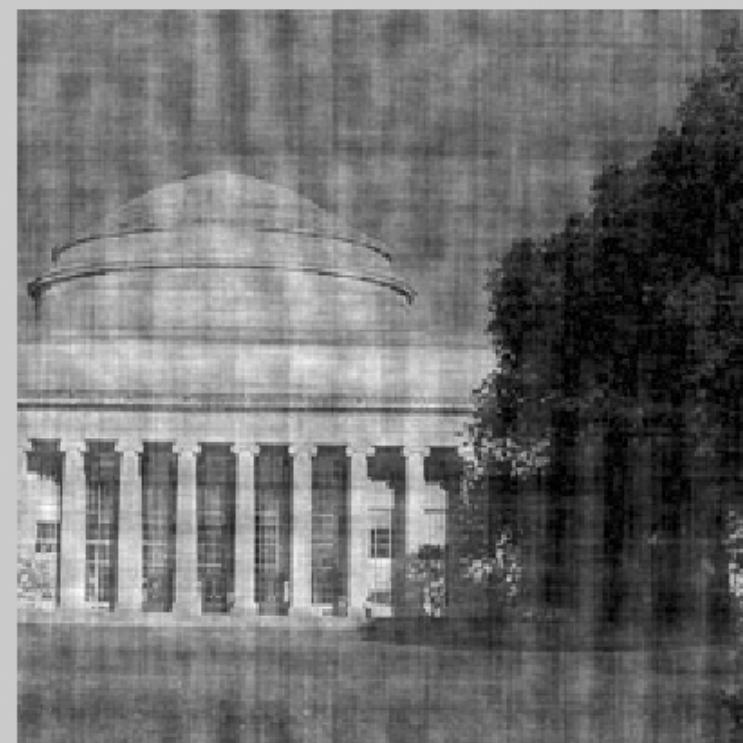
#2: Range [0.00109, 146]
Dims [256, 256]

49190.

49190



#1: Range [0, 1]
Dims [256, 256]



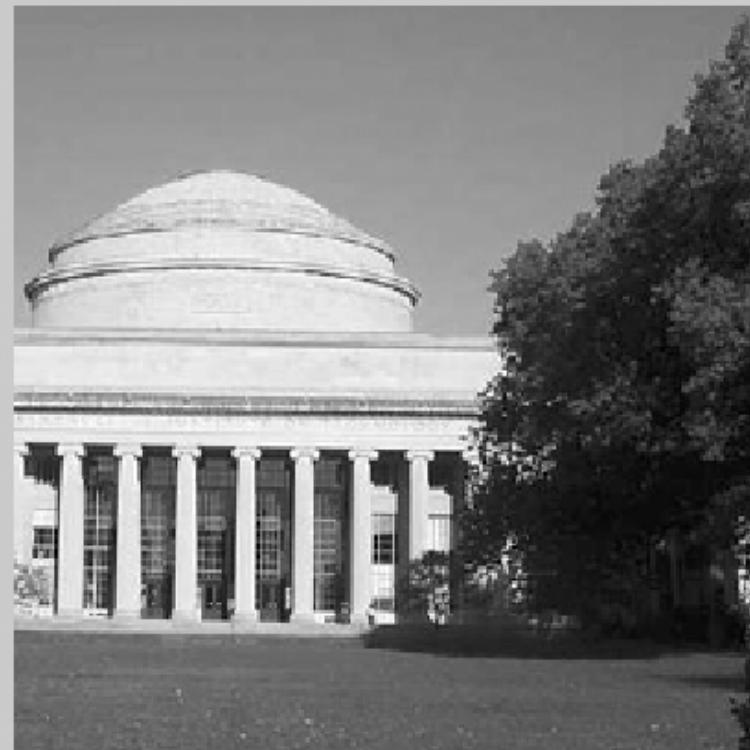
#2: Range [0.00758, 294]
Dims [256, 256]

65536.

65536.

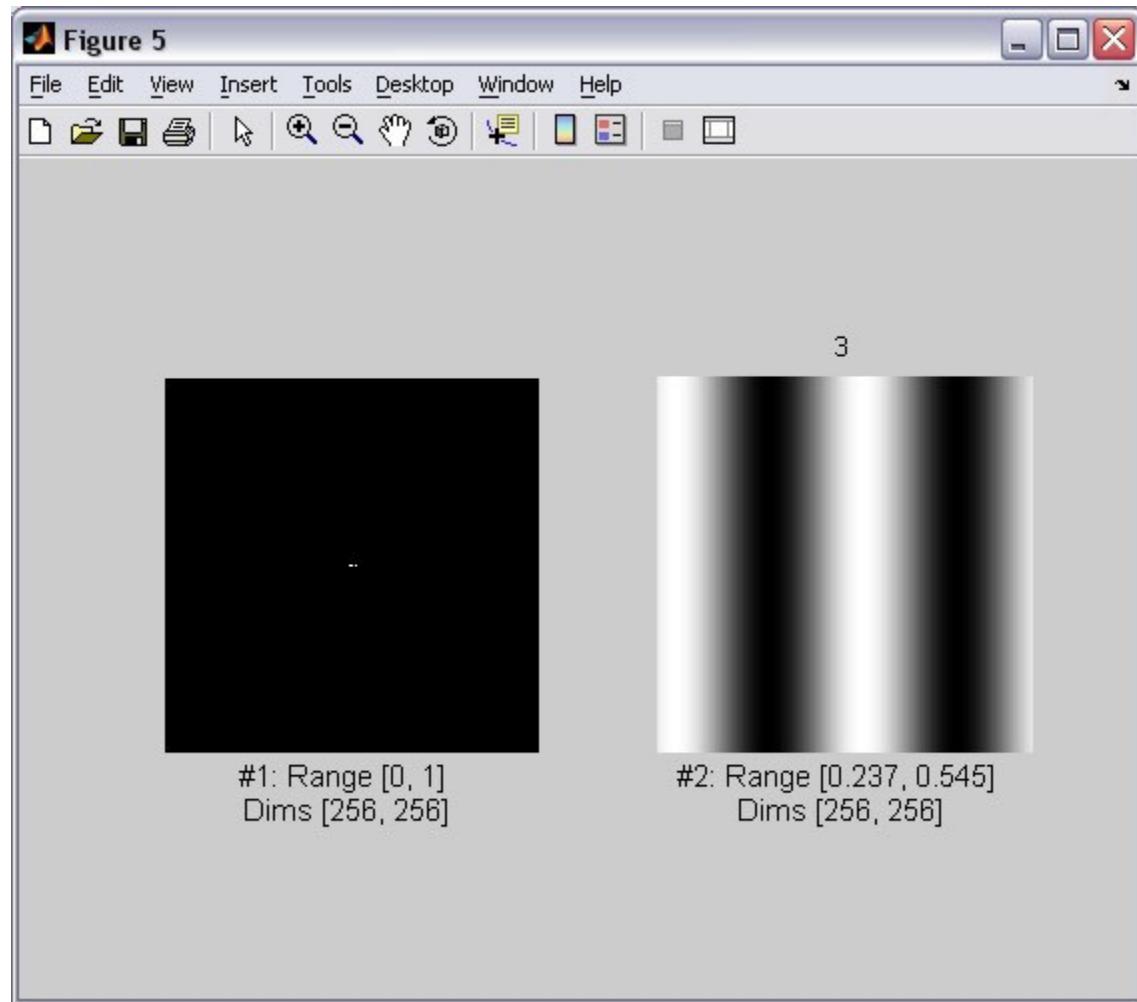


#1: Range [0.5, 1.5]
Dims [256, 256]



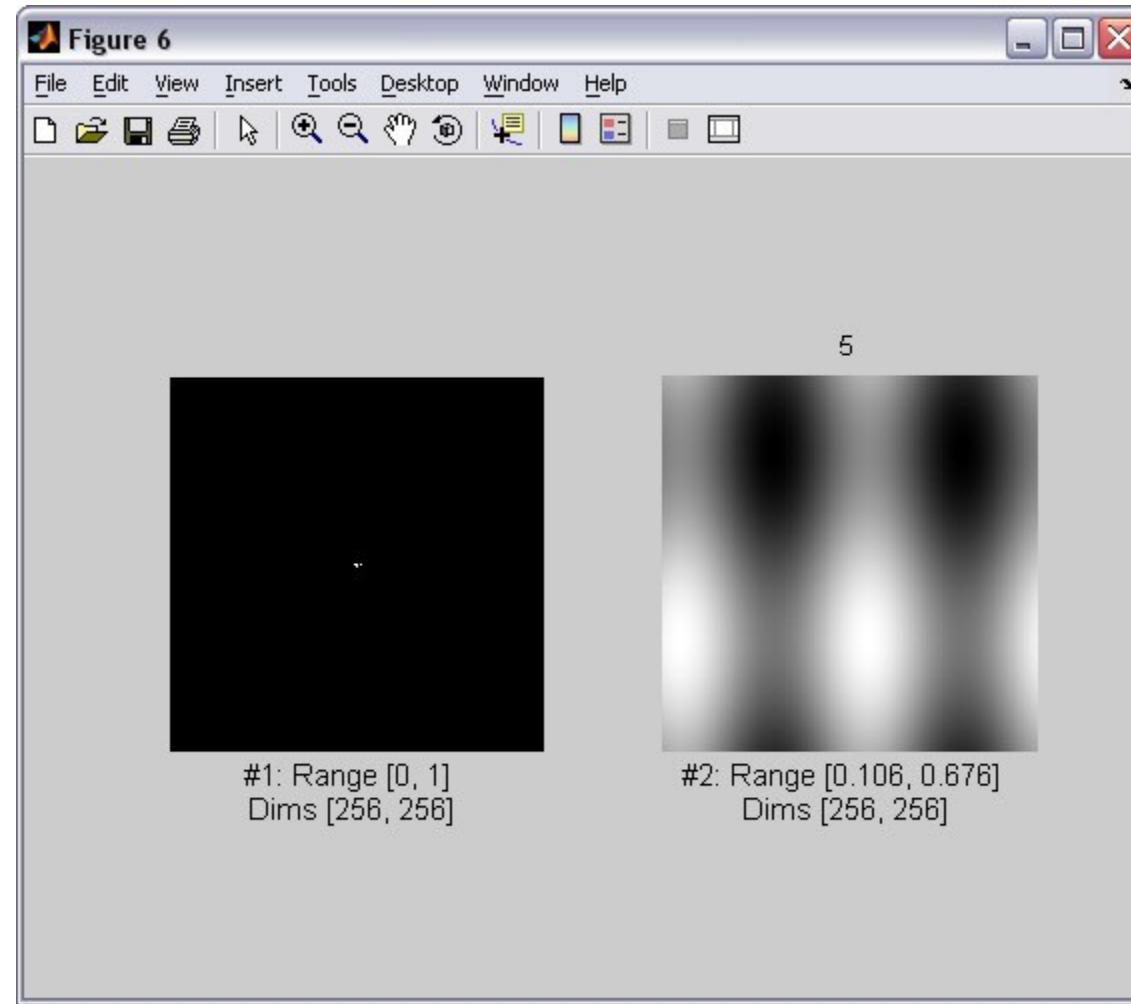
#2: Range [4.43e-015, 255]
Dims [256, 256]

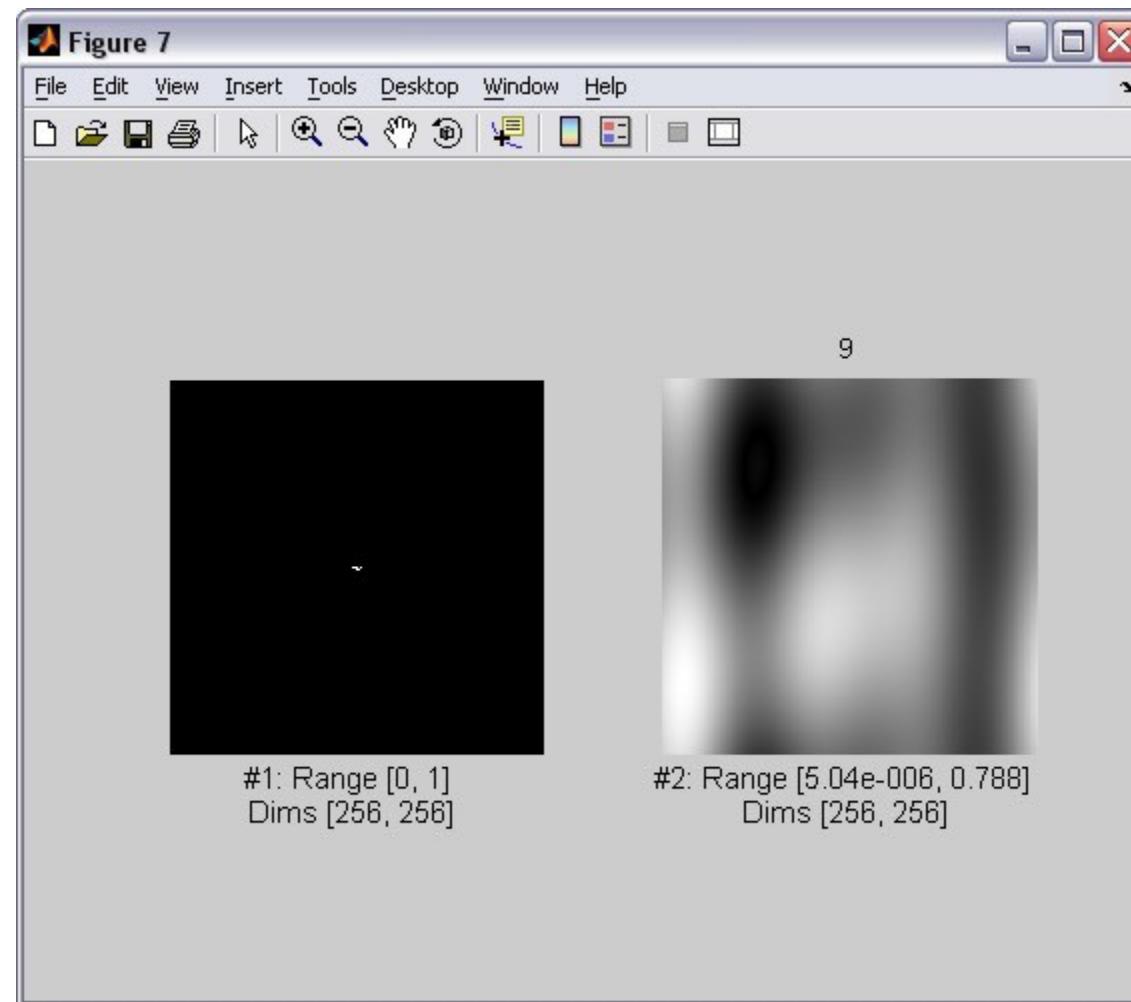
3

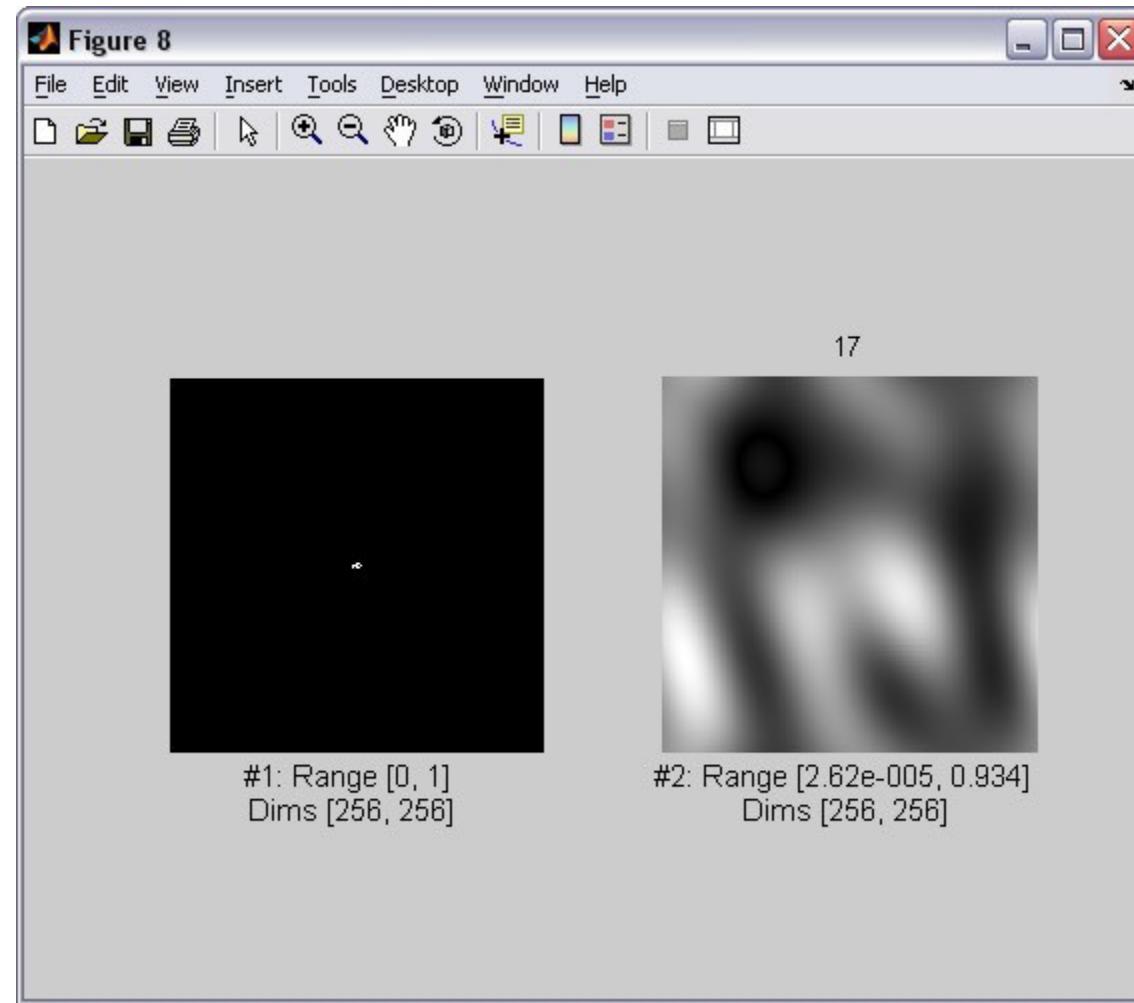


Now, an analogous sequence of images, but selecting Fourier components in descending order of magnitude.

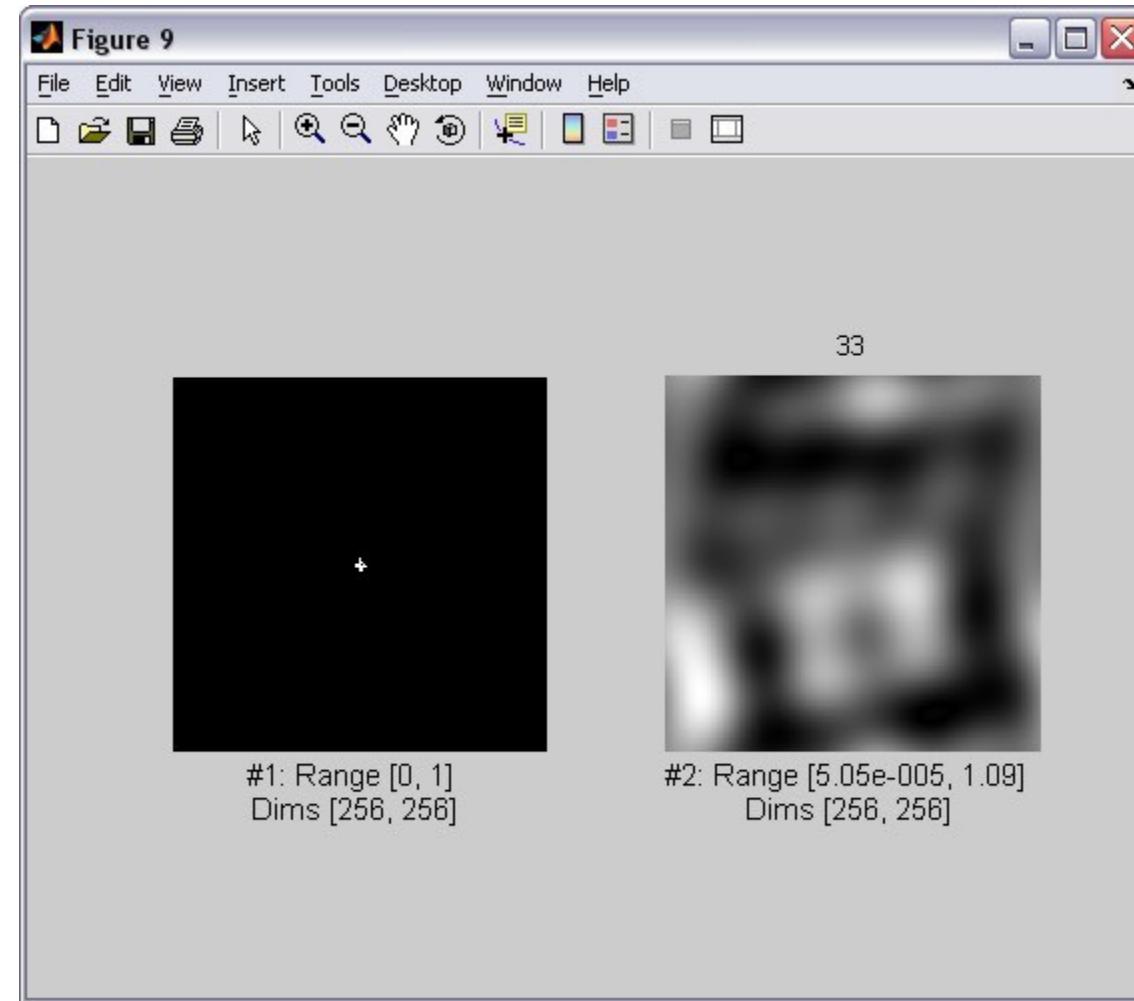
5



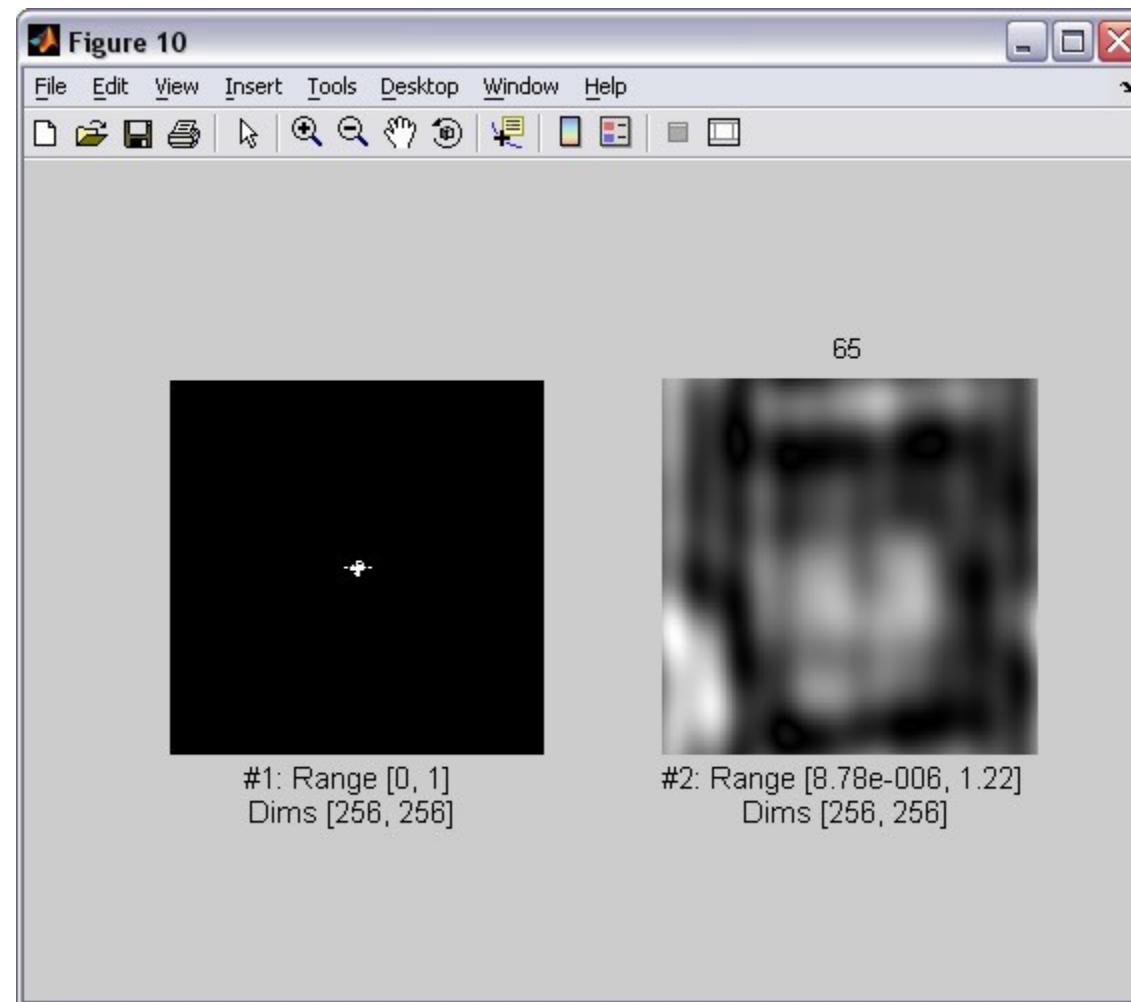




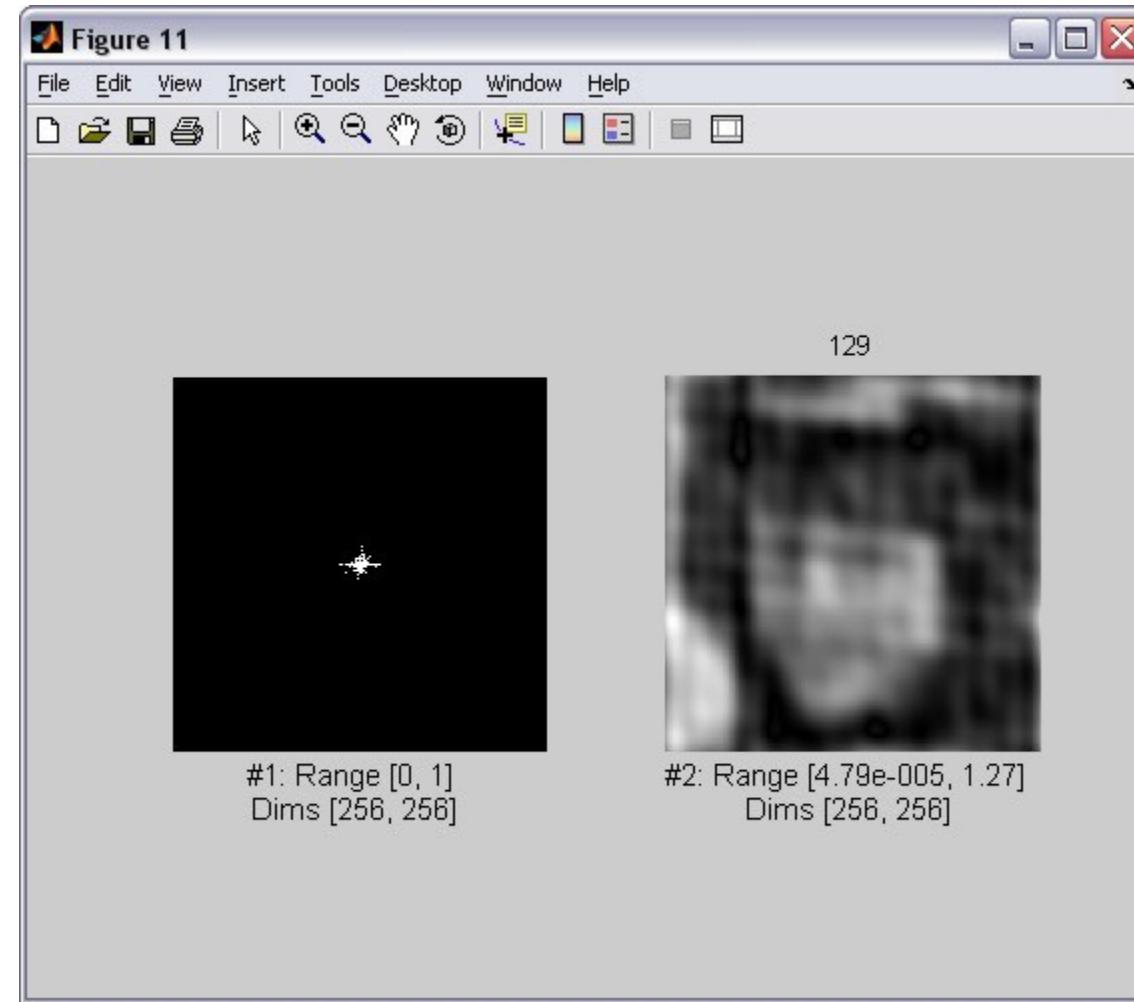
33



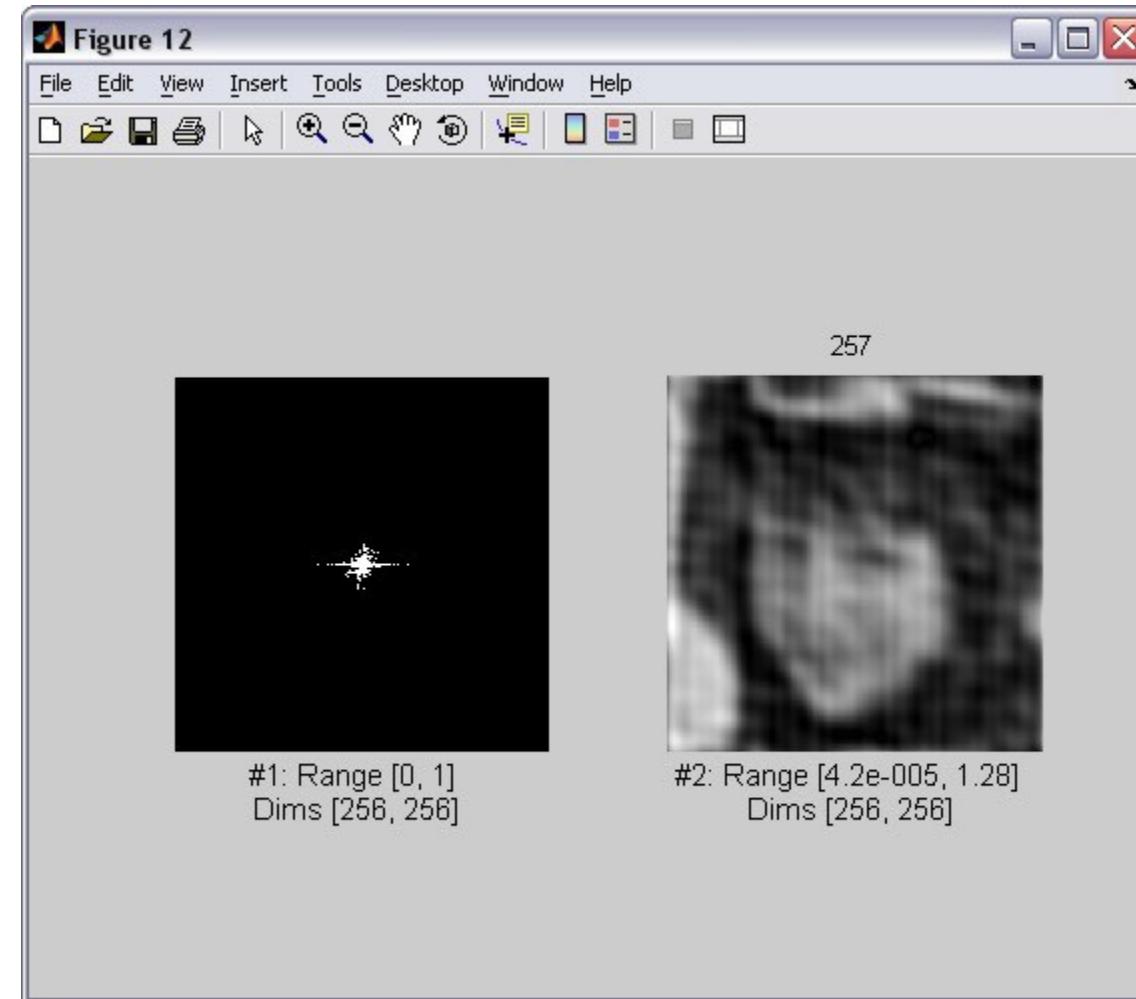
65



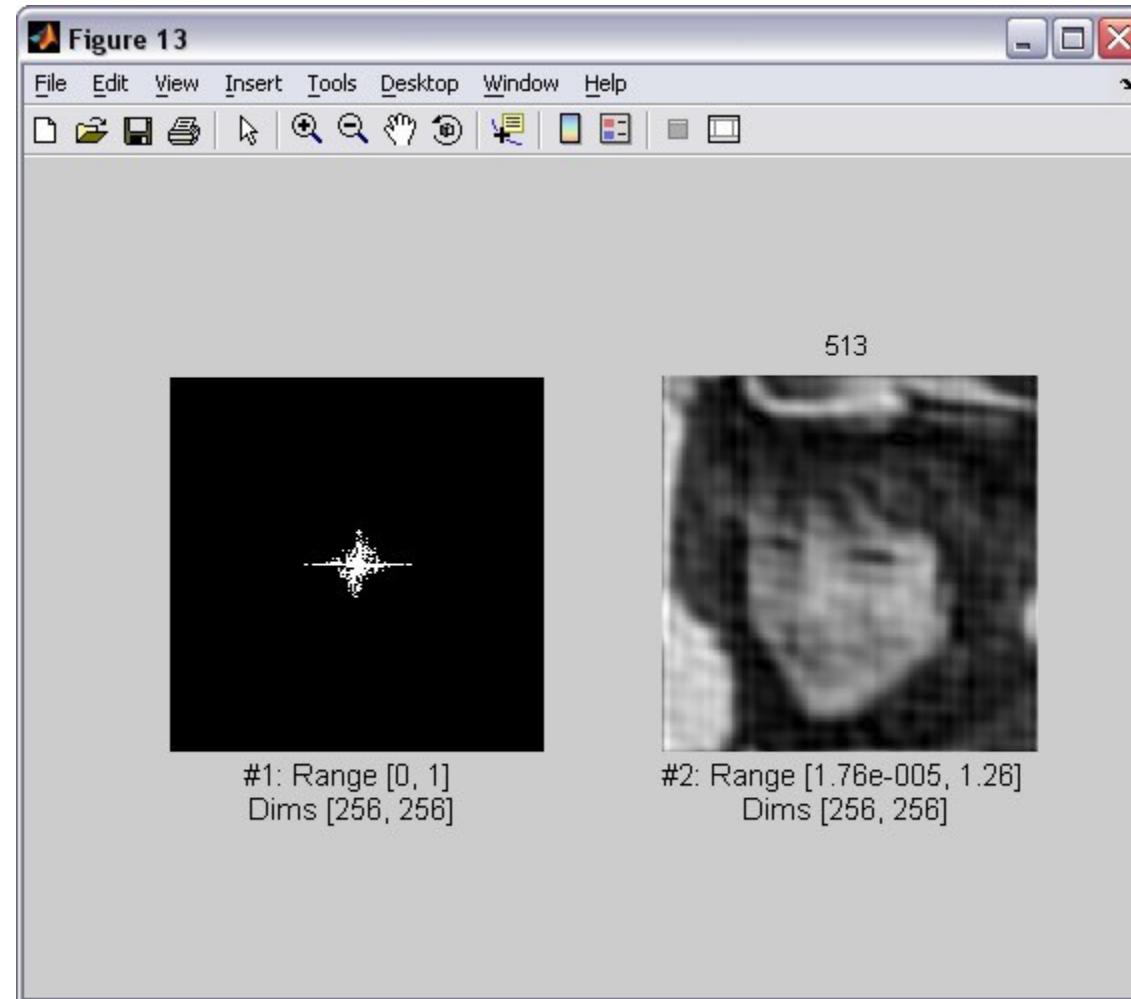
129



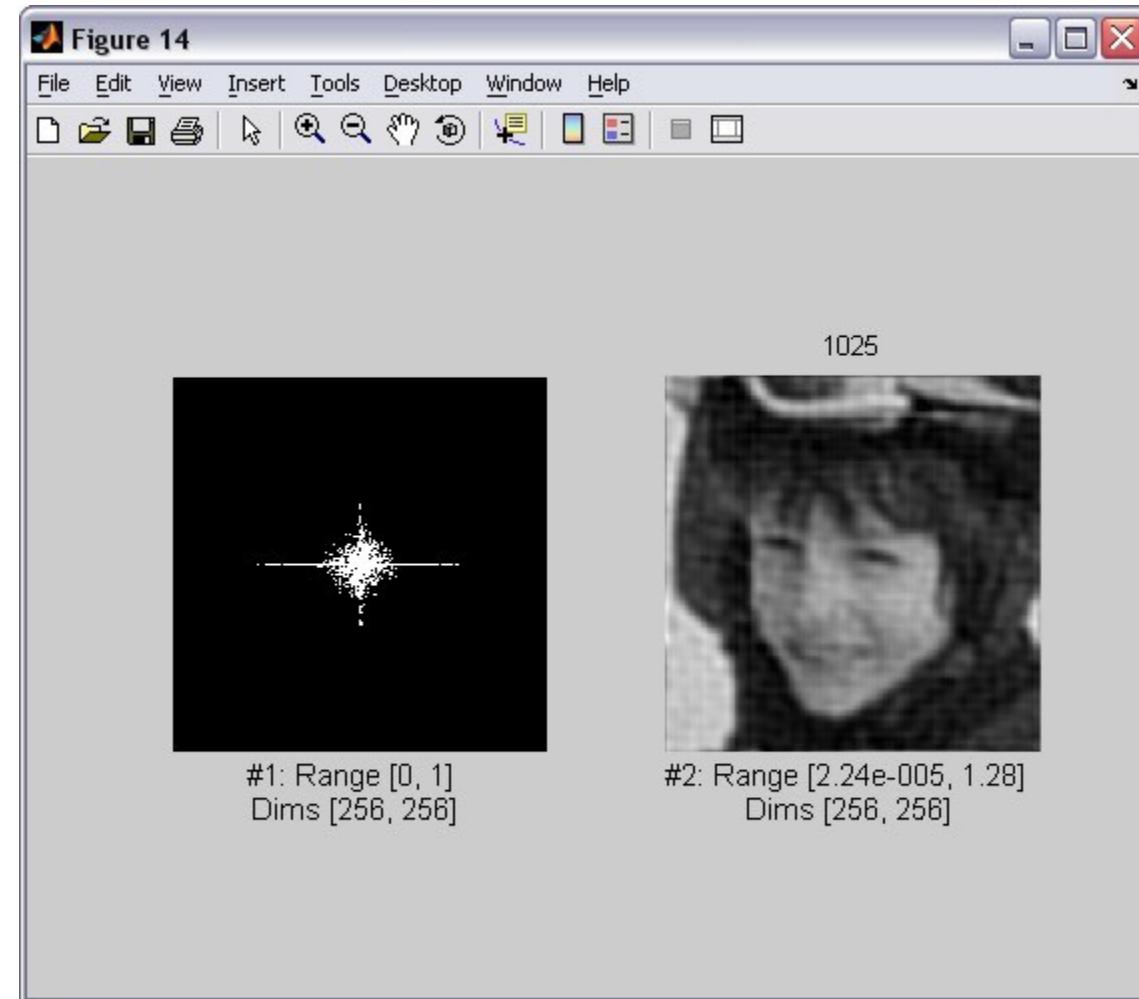
257



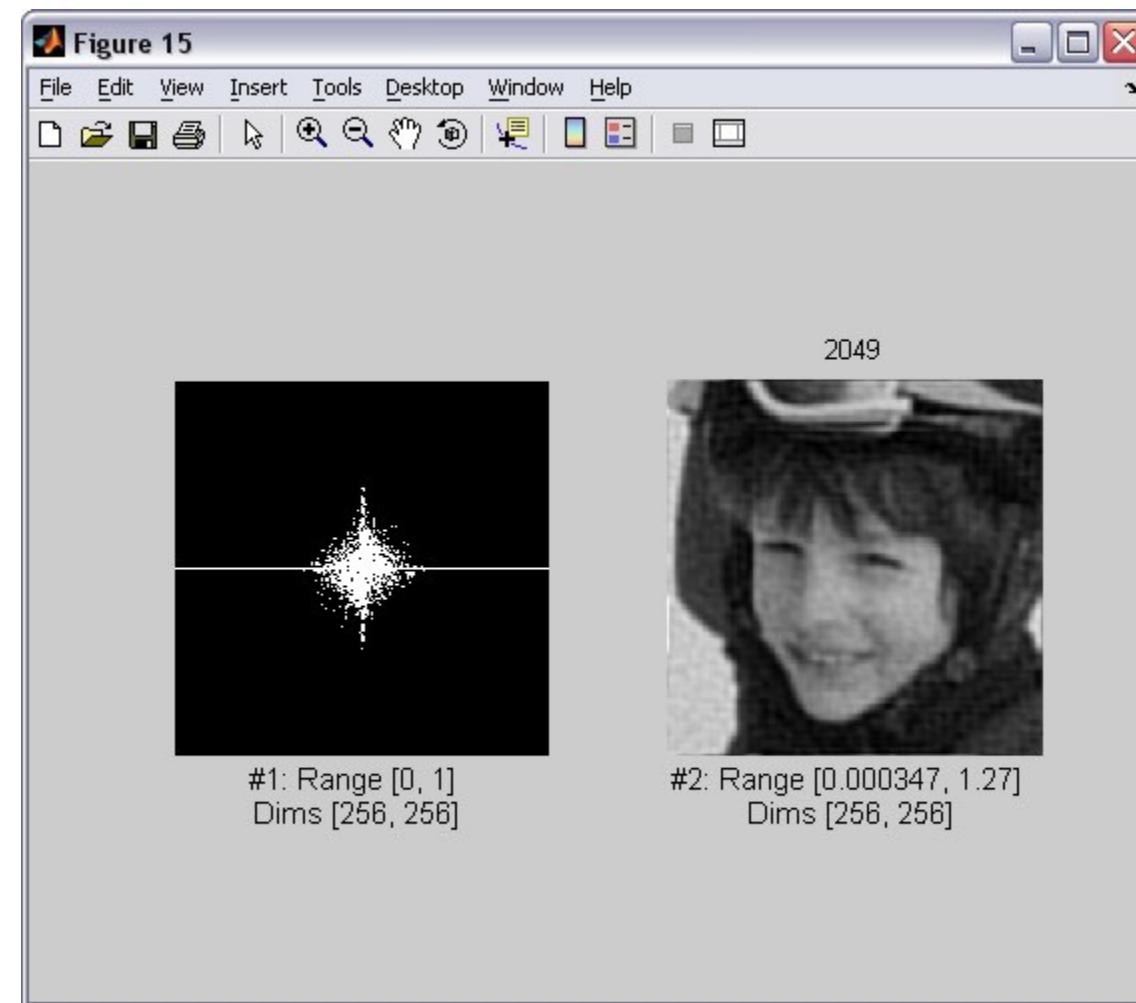
513



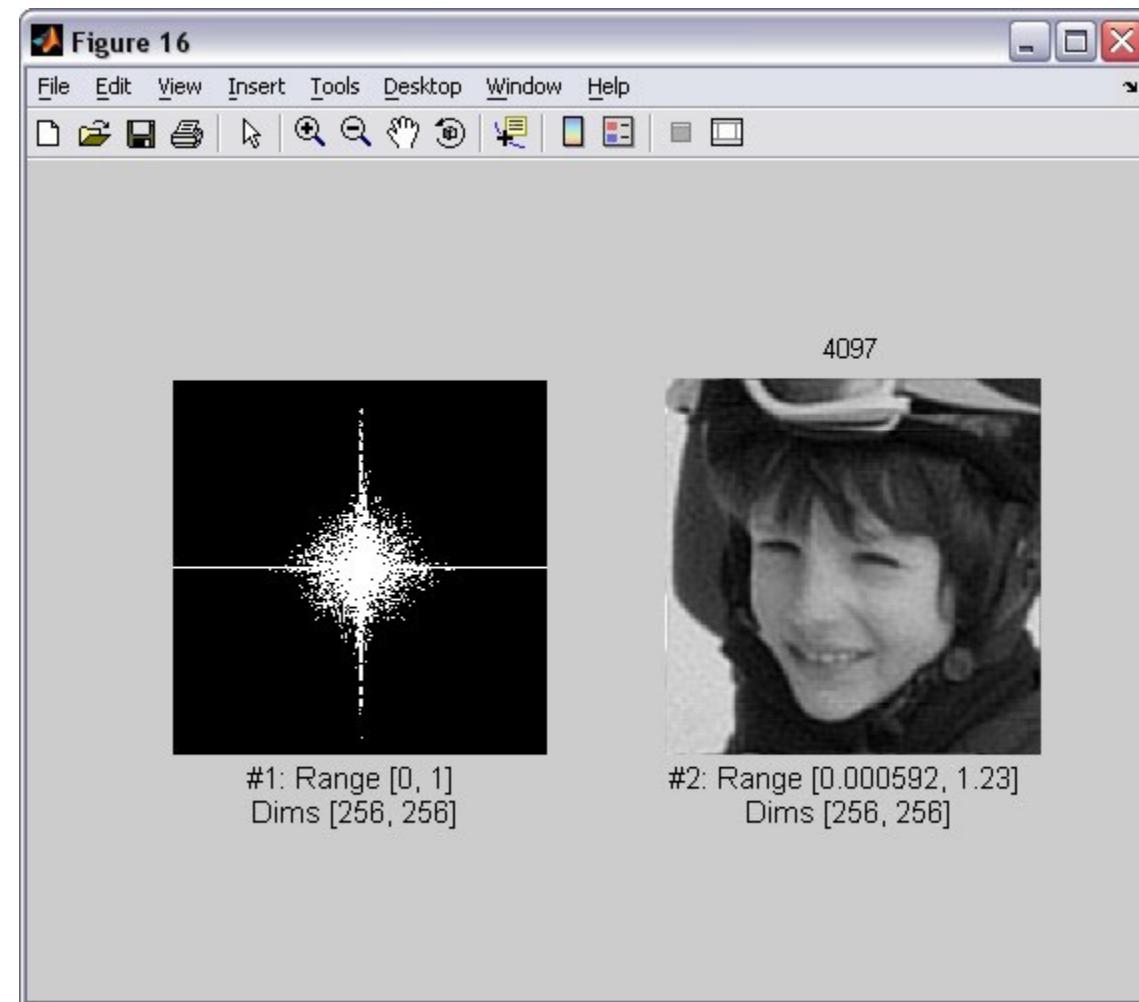
1025



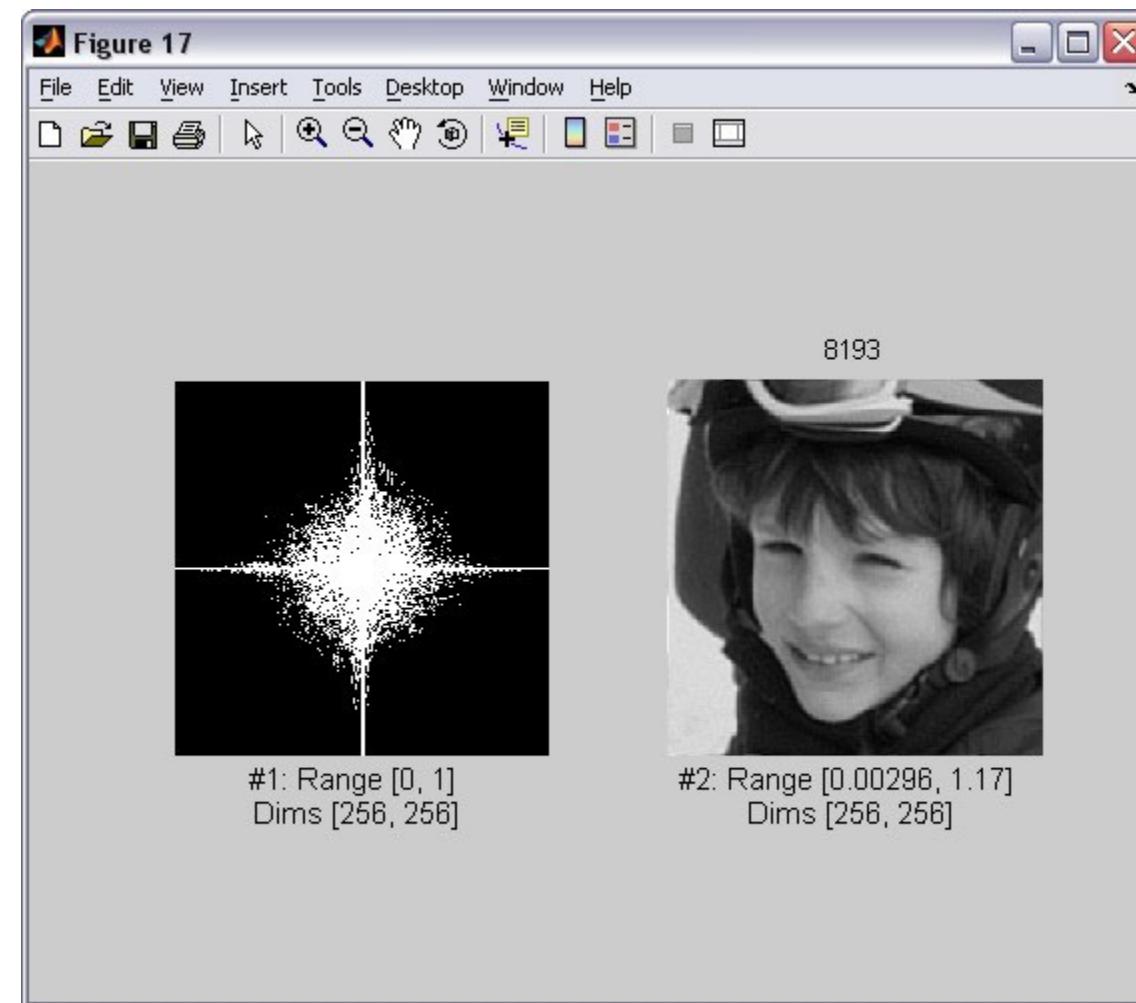
2049



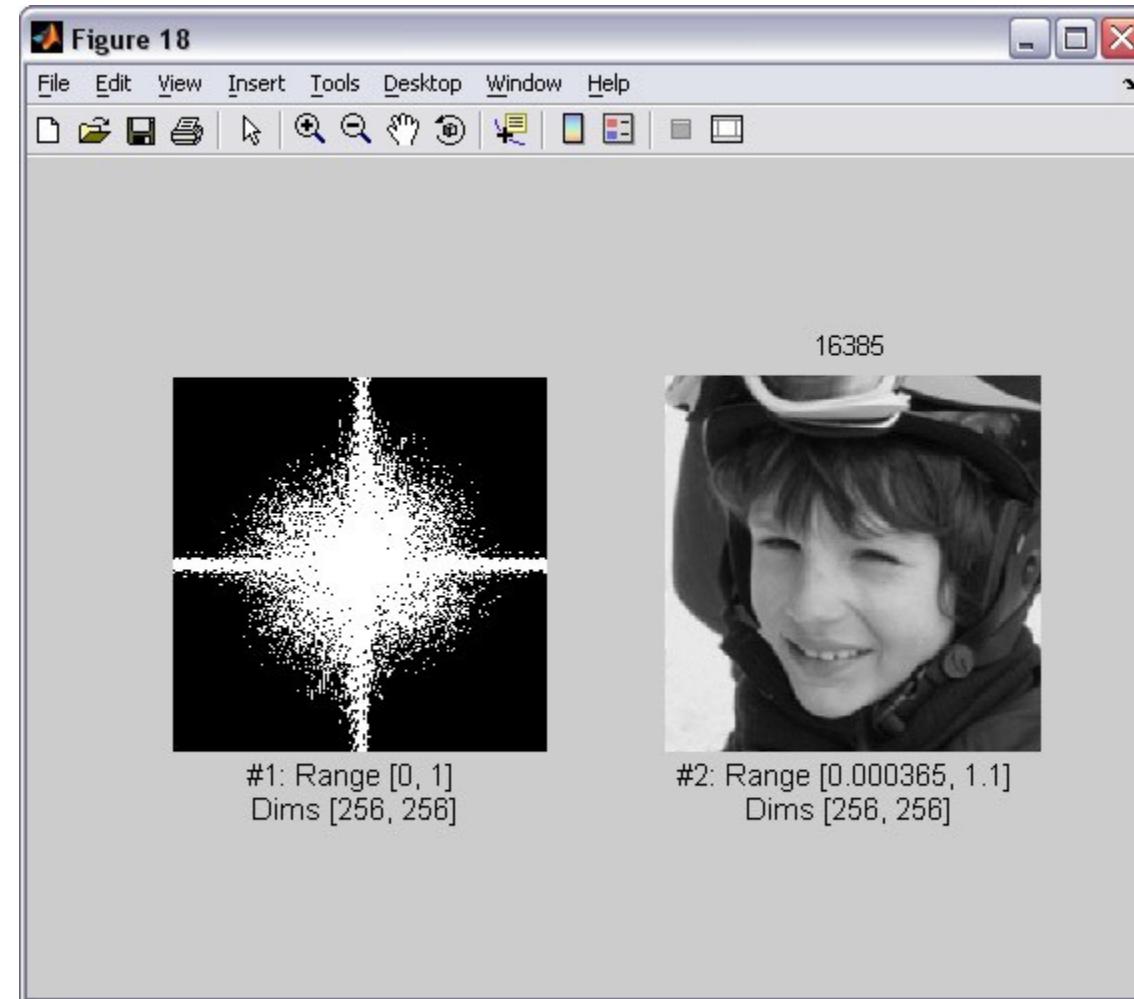
4097



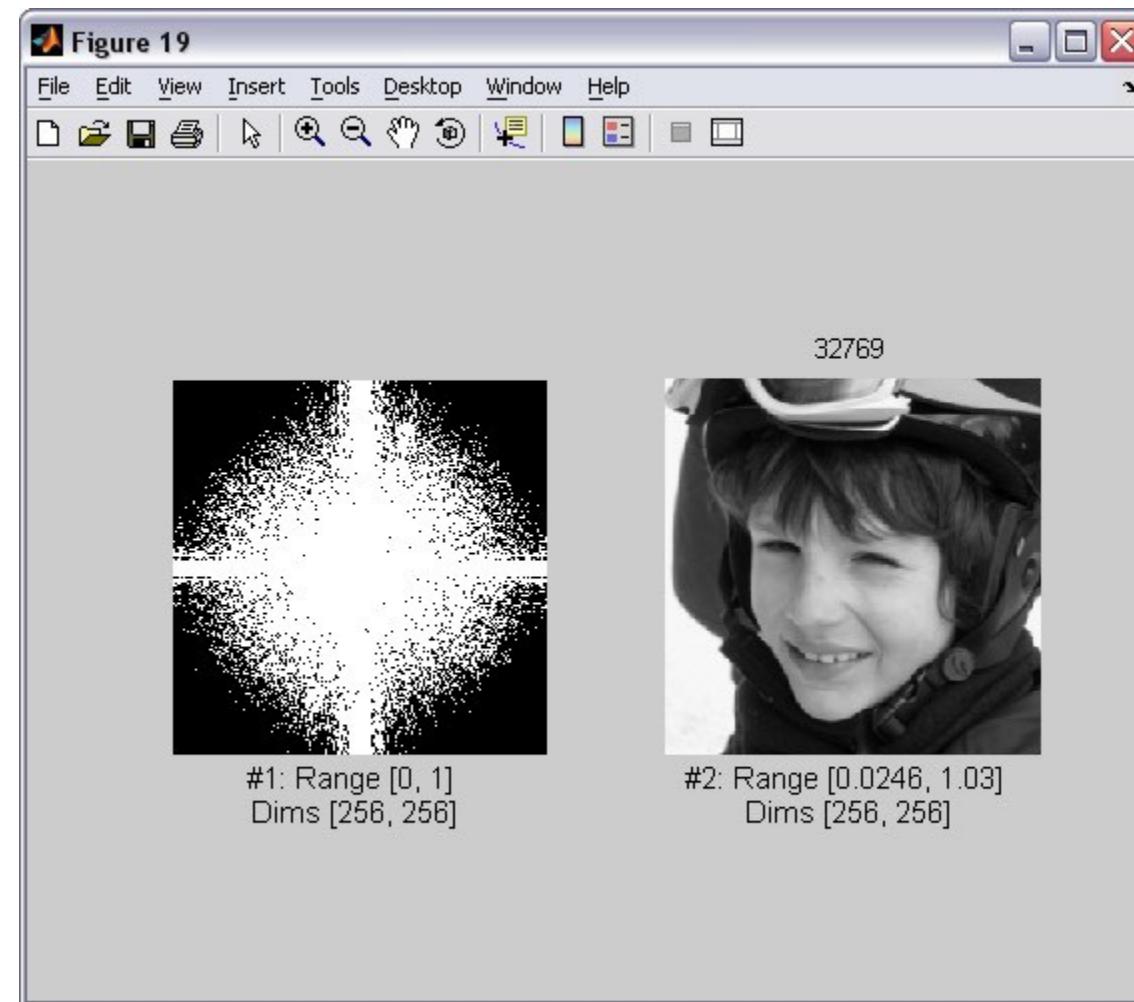
8193



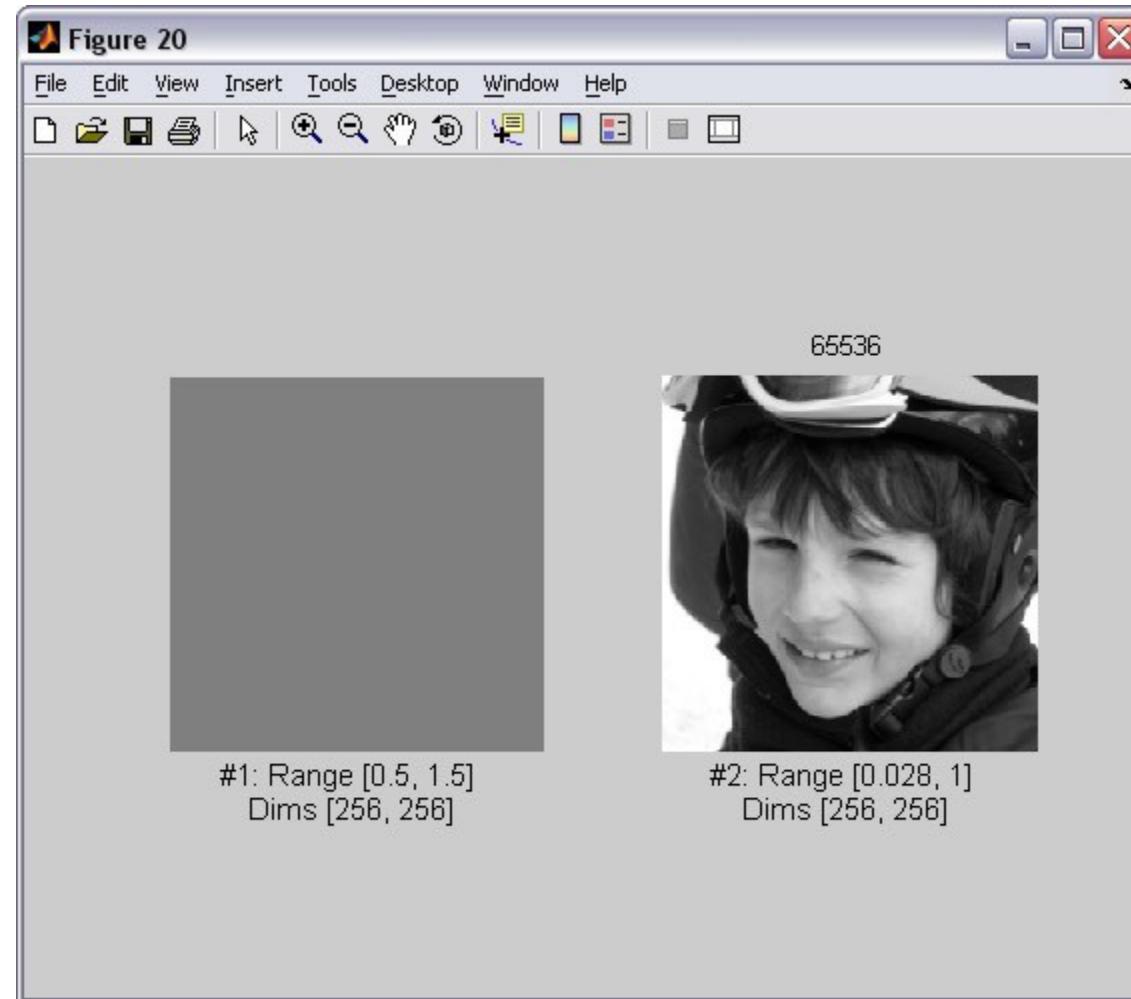
16385



32769

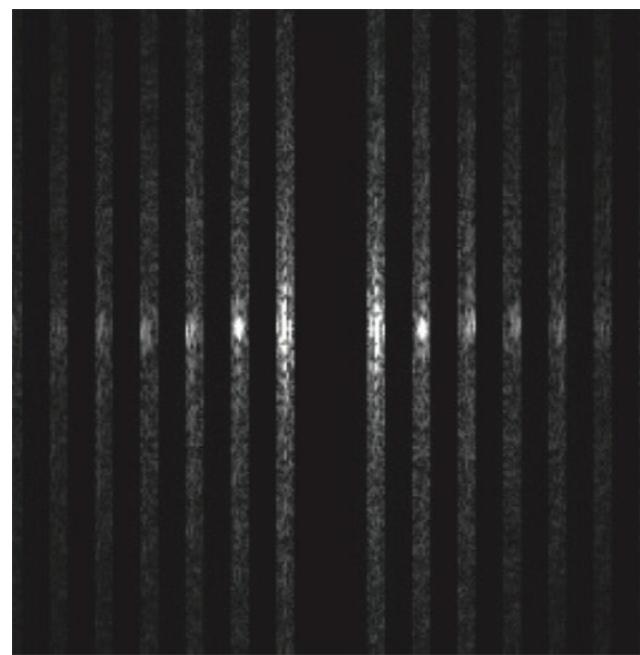
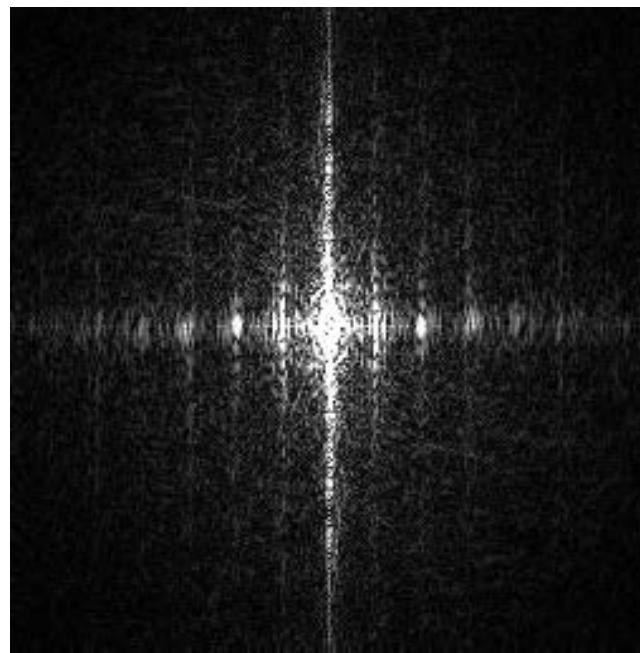


65536

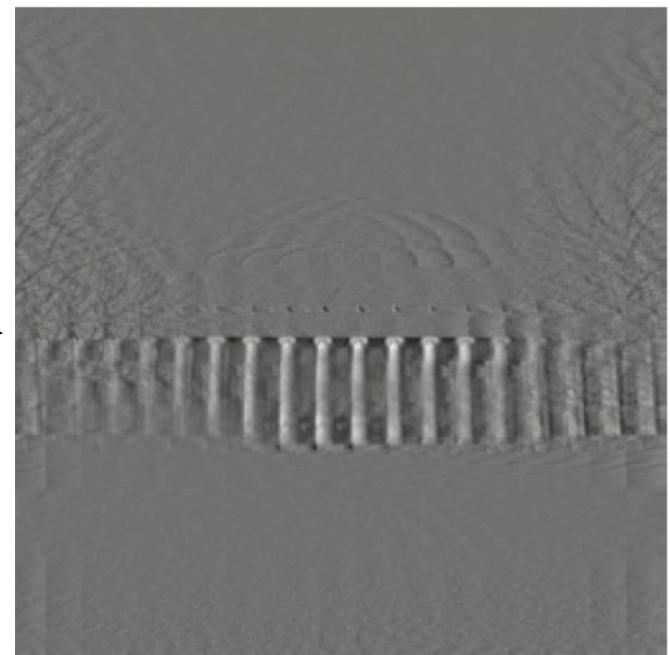




DFT
→

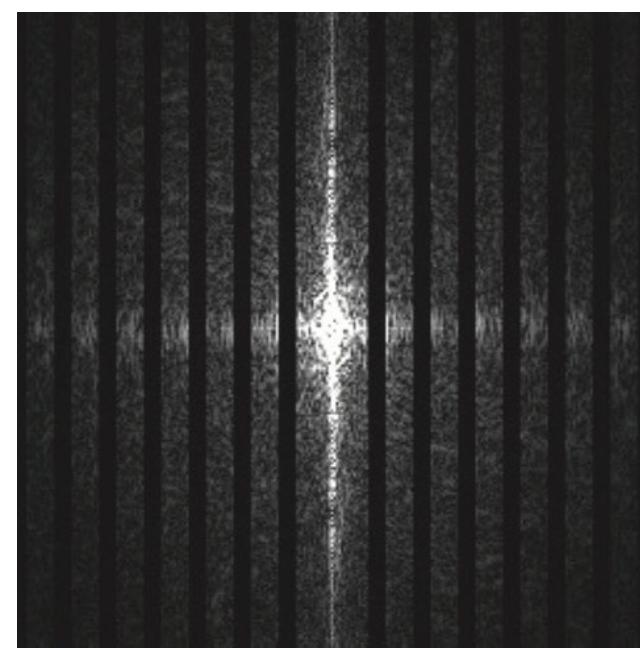
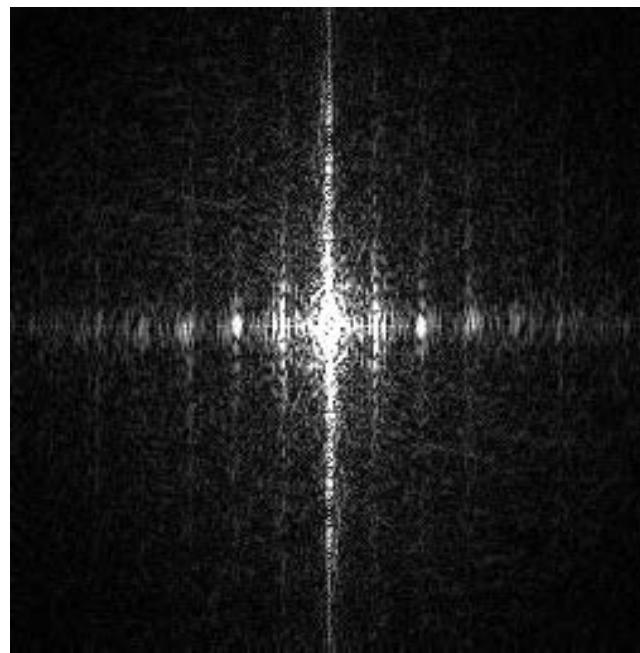


DFT^{-1}
→





DFT
→



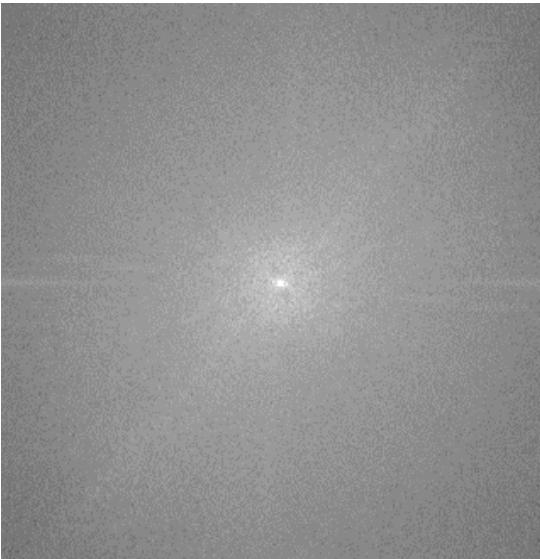
DFT^{-1}
→



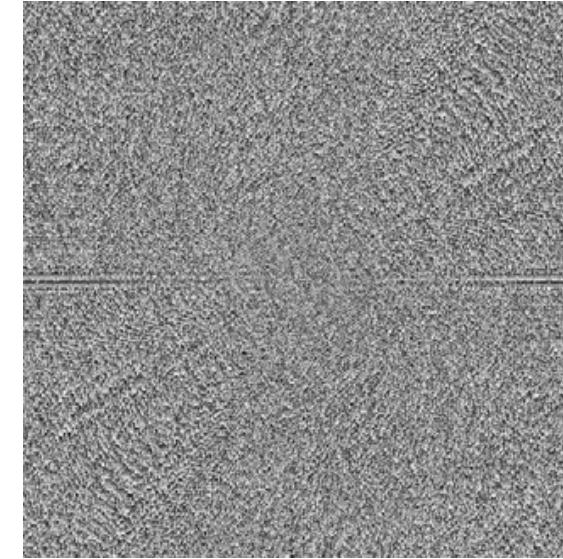
Fourier transforms of natural images



original

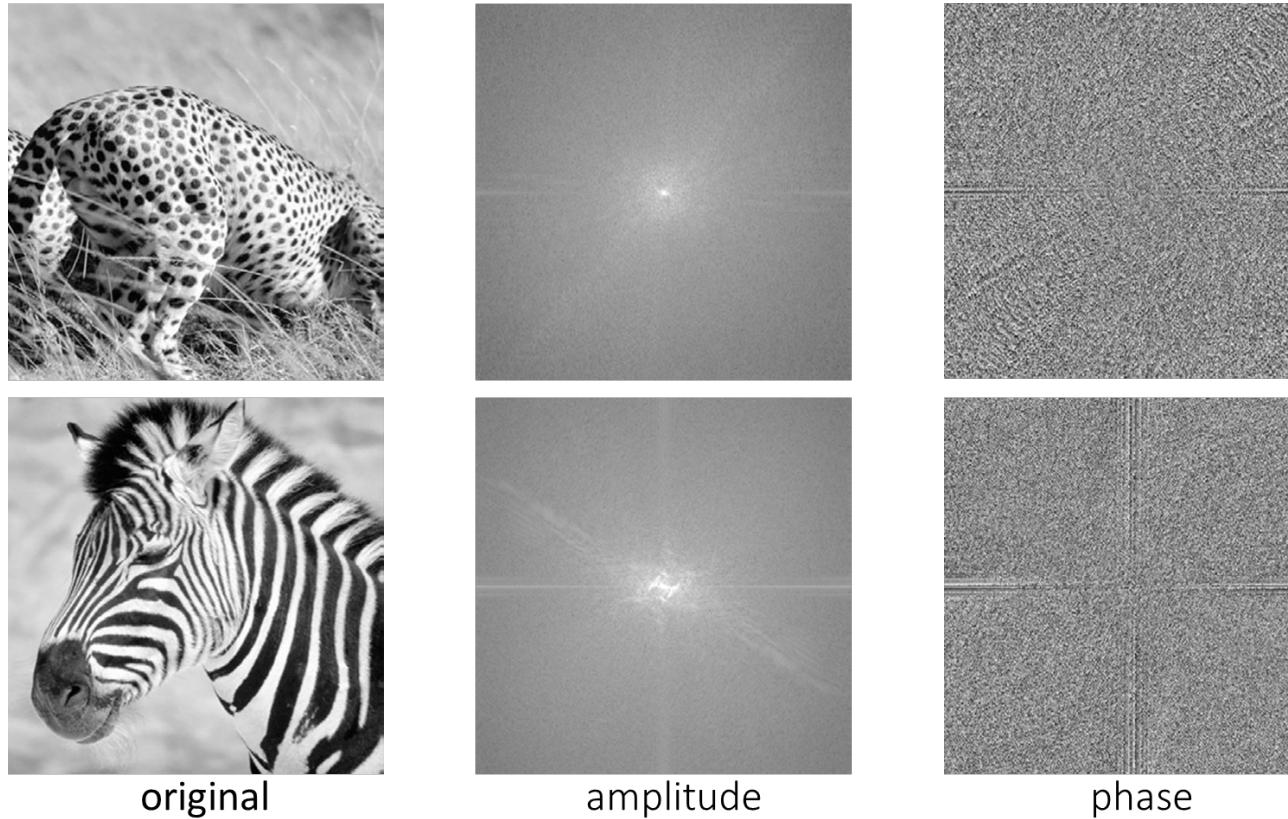


amplitude



phase

Fourier transforms of natural images

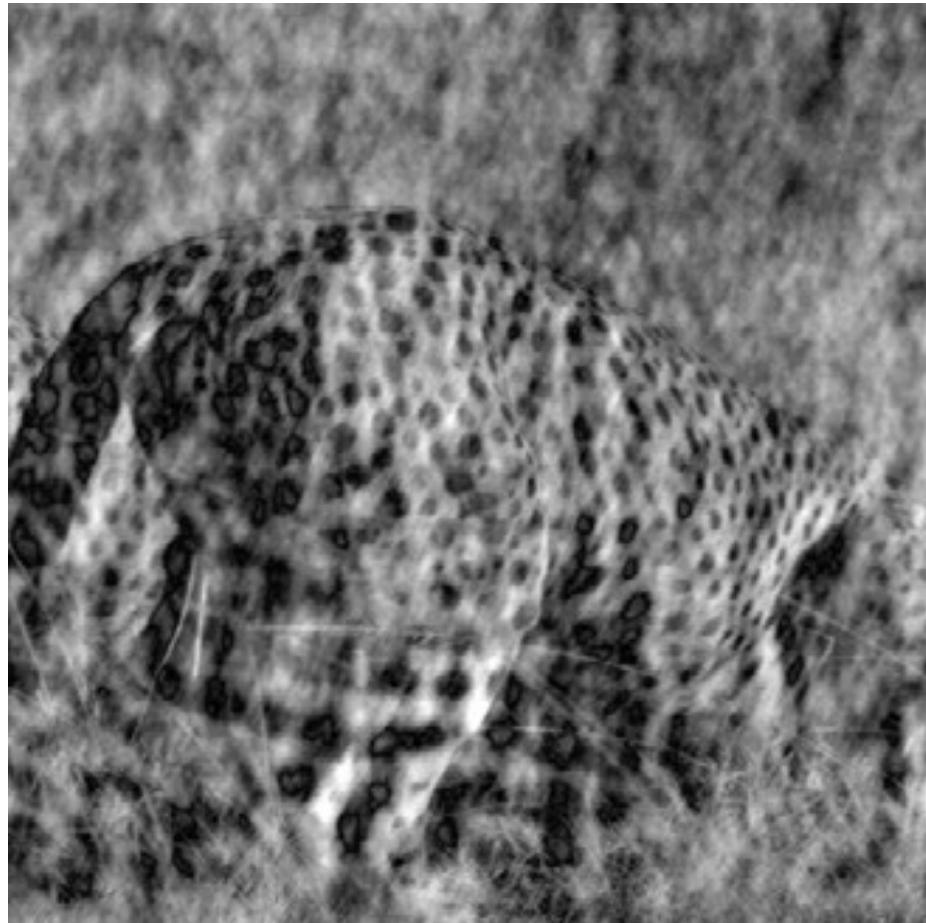


Curious fact:

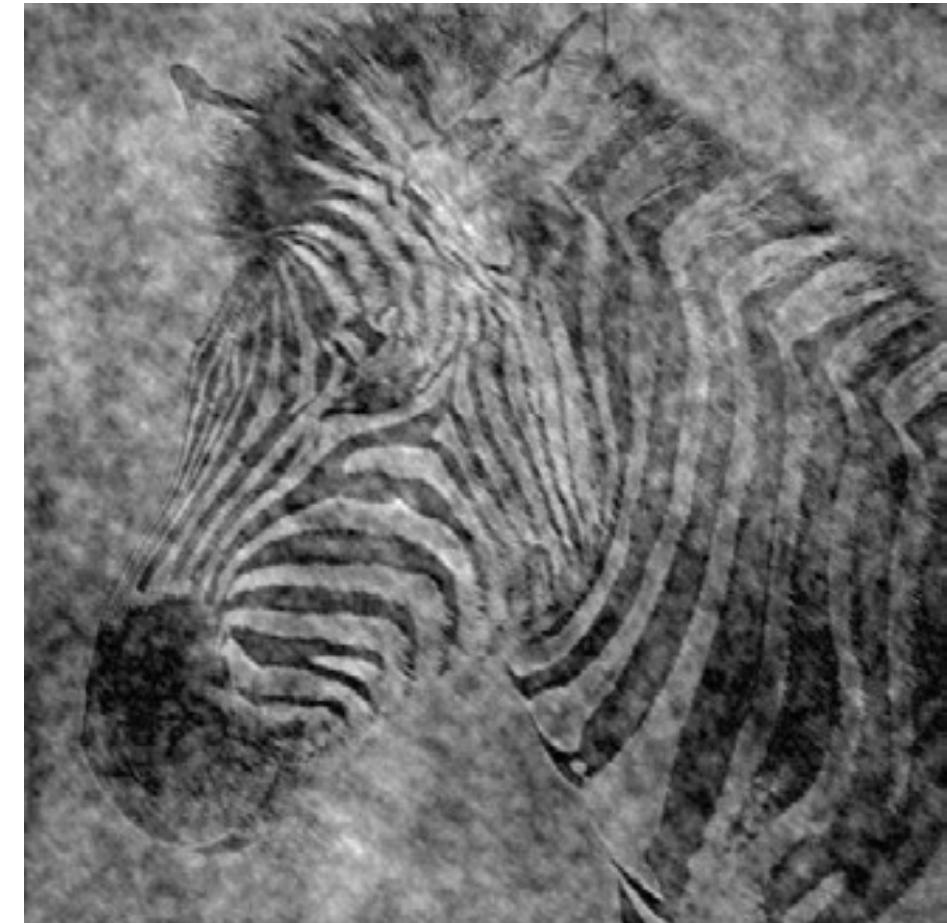
- All natural images have about the same magnitude transform
- Hence, phase seems to matter, but magnitude largely doesn't

Fourier transforms of natural images

Image phase matters!



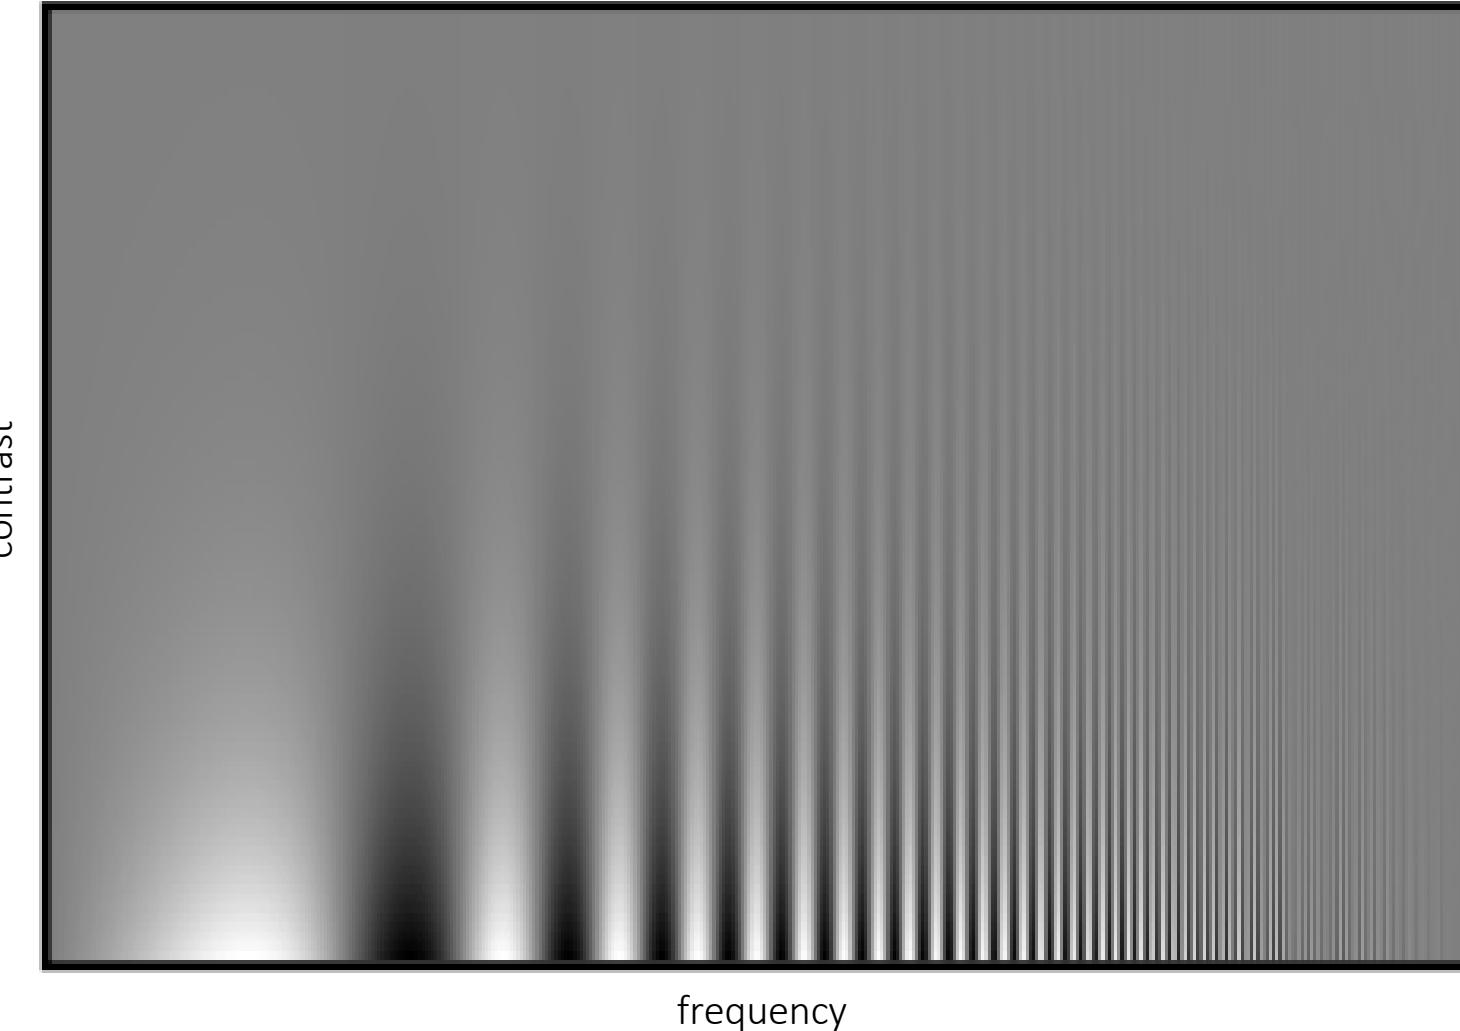
cheetah phase with zebra amplitude



zebra phase with cheetah amplitude

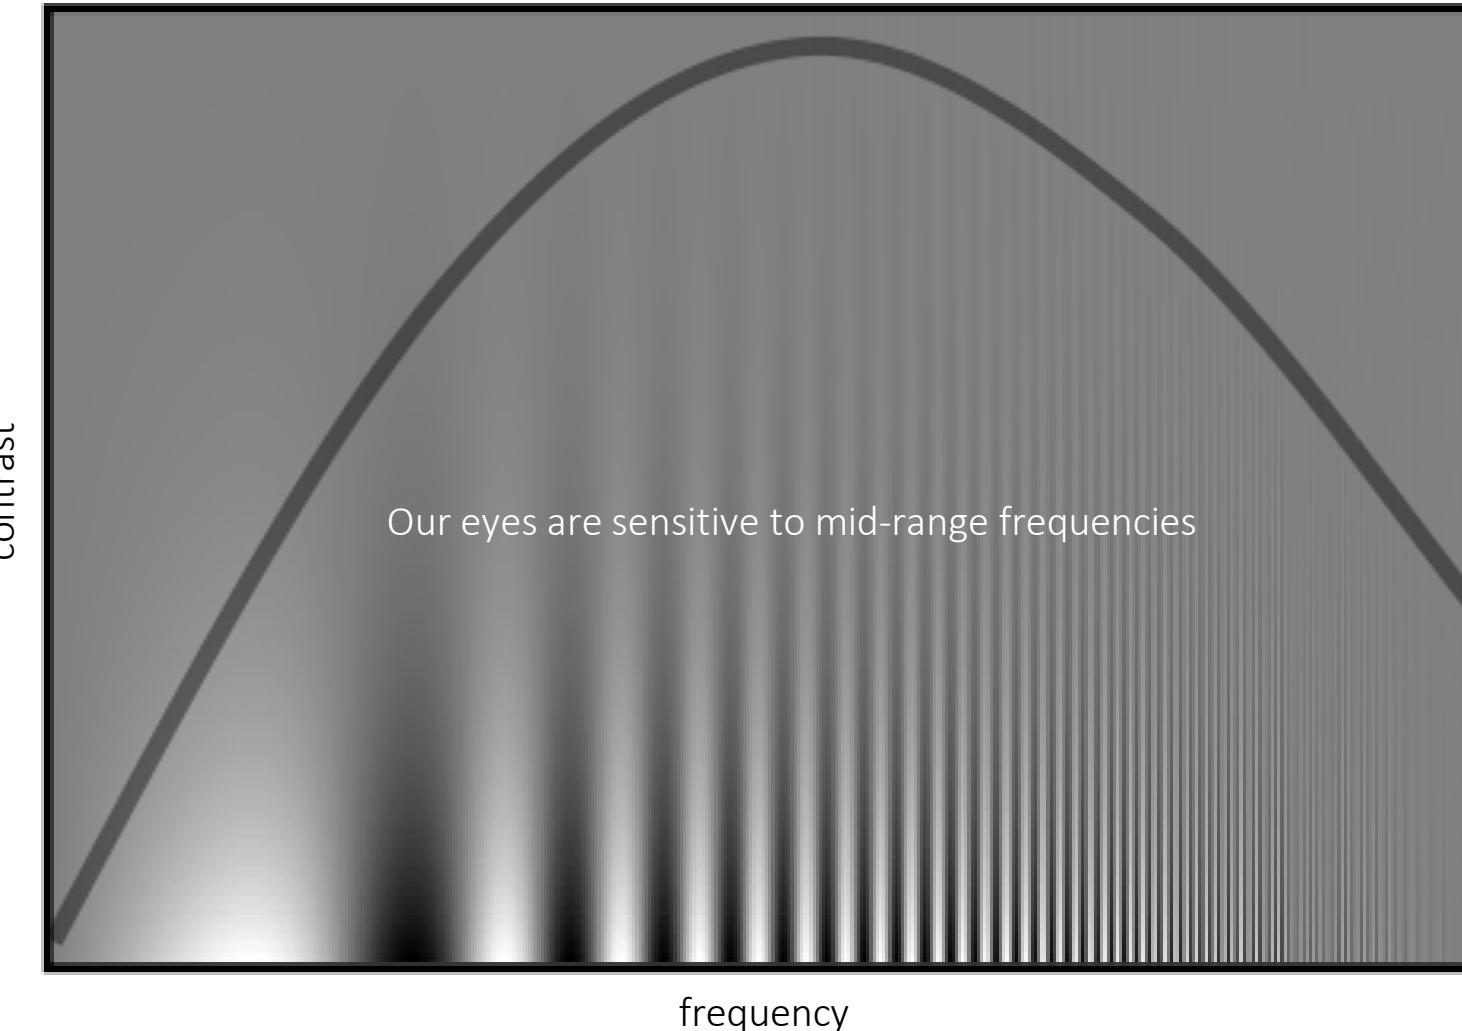
Variable frequency sensitivity

Experiment: Where do you see the stripes?



Variable frequency sensitivity

Campbell-Robson contrast sensitivity curve



- Early processing in humans filters for various orientations and scales of frequency
- Perceptual cues in the mid frequencies dominate perception

Frequency-domain filtering

The convolution theorem

Why do we care about all this?

The Fourier transform of the convolution of two functions is the product of their Fourier transforms:

$$\mathcal{F}\{g * h\} = \mathcal{F}\{g\}\mathcal{F}\{h\}$$

The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms:

$$\mathcal{F}^{-1}\{gh\} = \mathcal{F}^{-1}\{g\} * \mathcal{F}^{-1}\{h\}$$

Convolution in spatial domain is equivalent to multiplication in frequency domain!

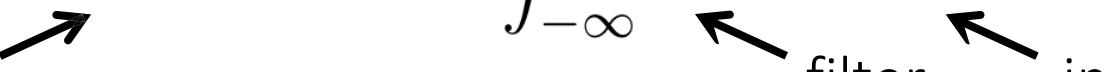
Convolution for 1D continuous signals

What do we use convolution for?

Definition of linear shift-invariant filtering as convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

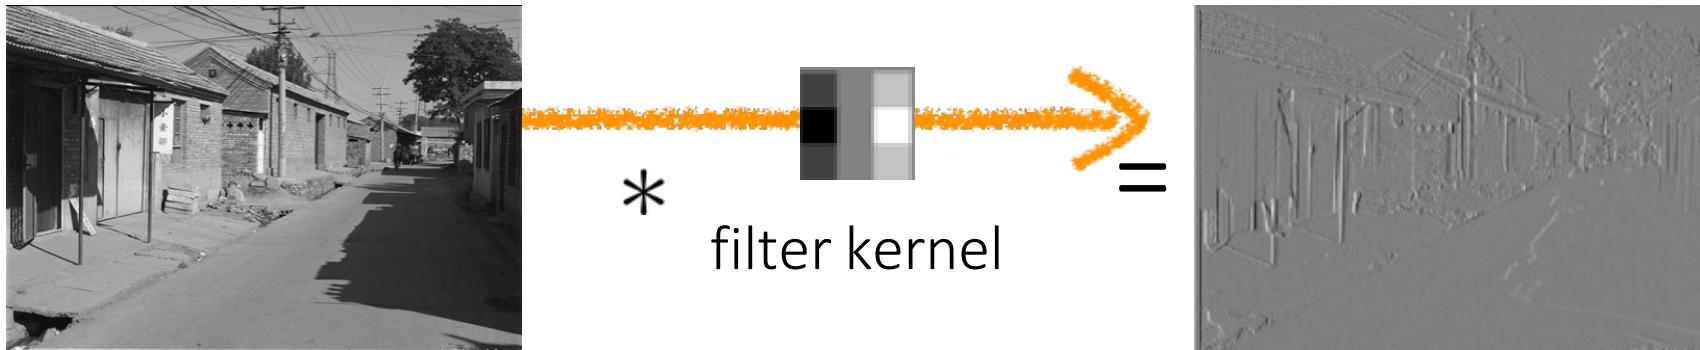
filtered signal filter input signal



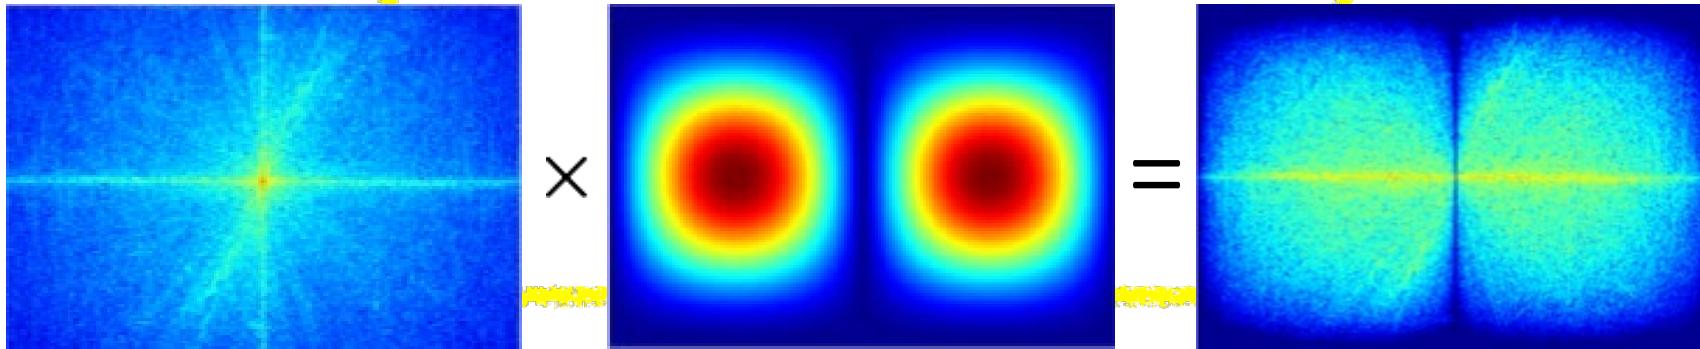
Using the convolution theorem, we can interpret and implement all types of linear shift-invariant filtering as multiplication in frequency domain.

Why implement convolution in frequency domain?

Spatial domain filtering



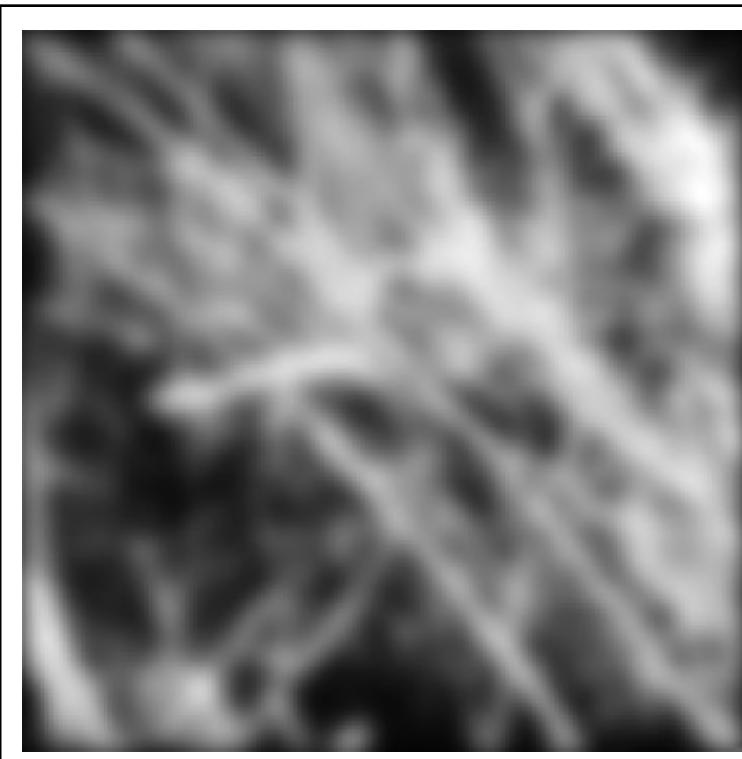
Fourier transform



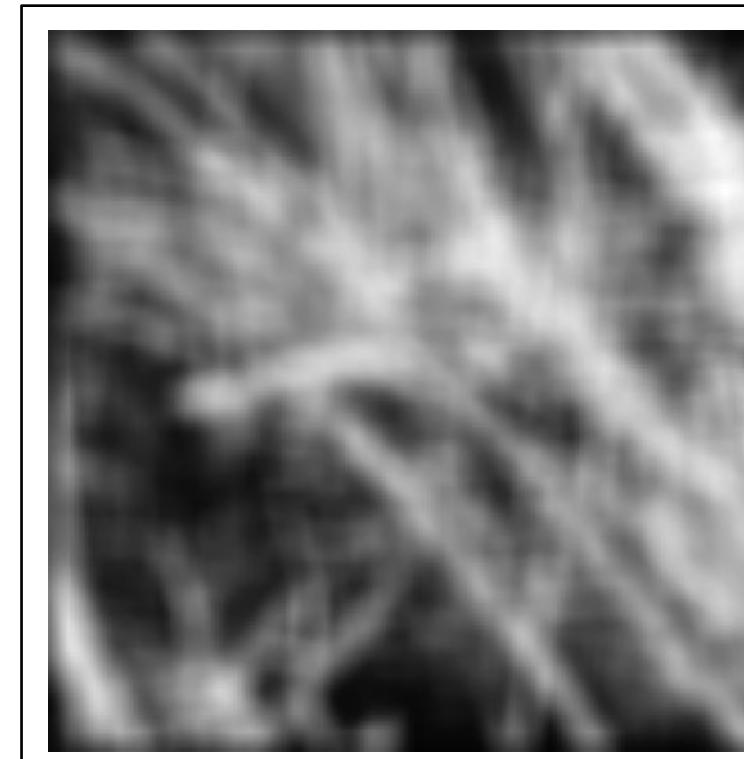
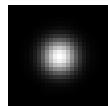
Frequency domain filtering

Revisiting blurring

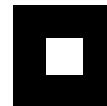
Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?



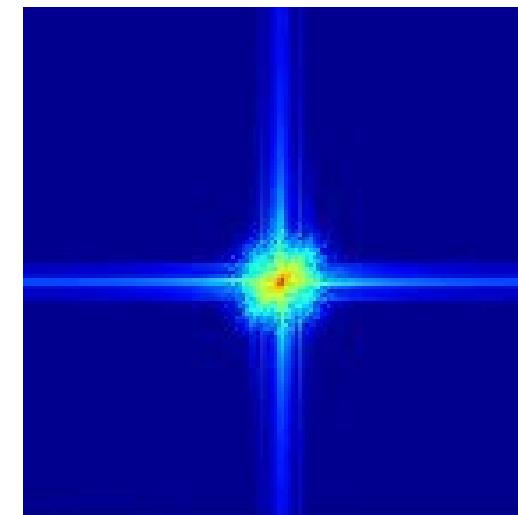
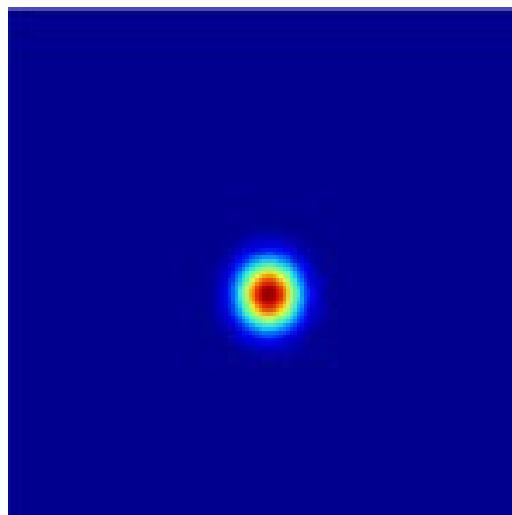
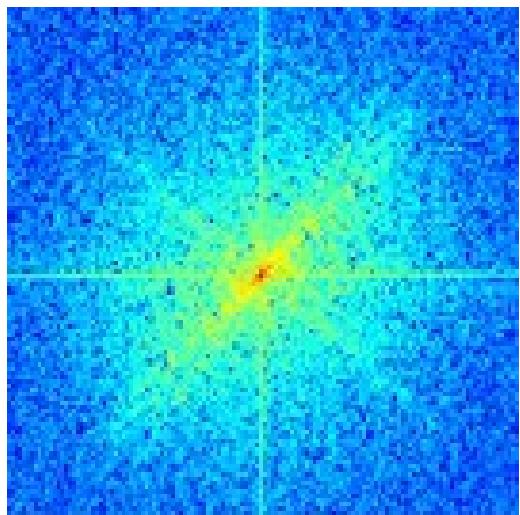
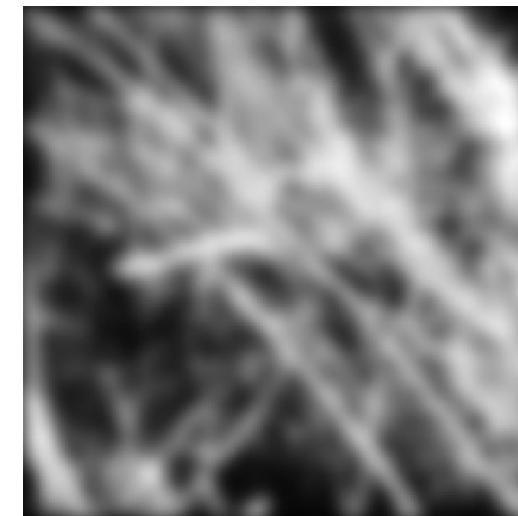
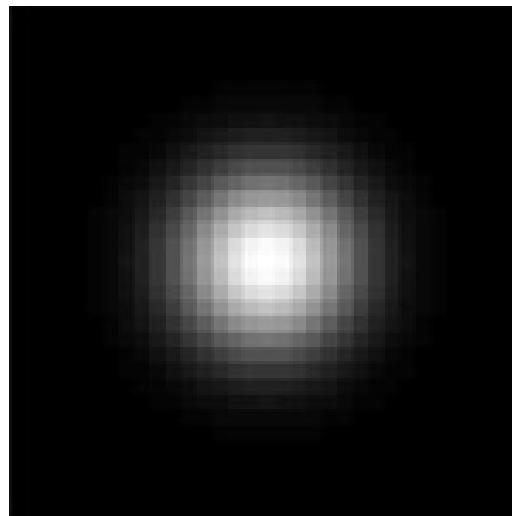
Gaussian
filter



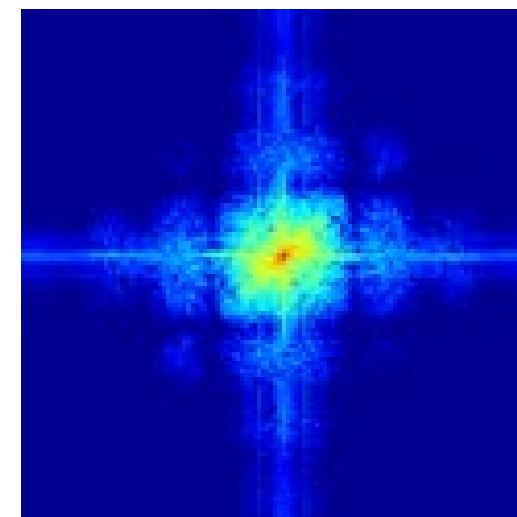
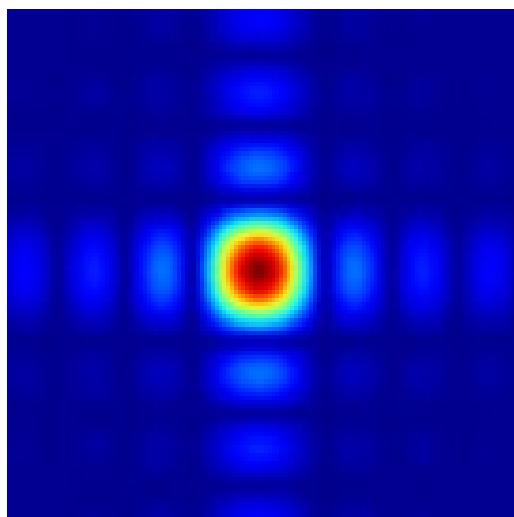
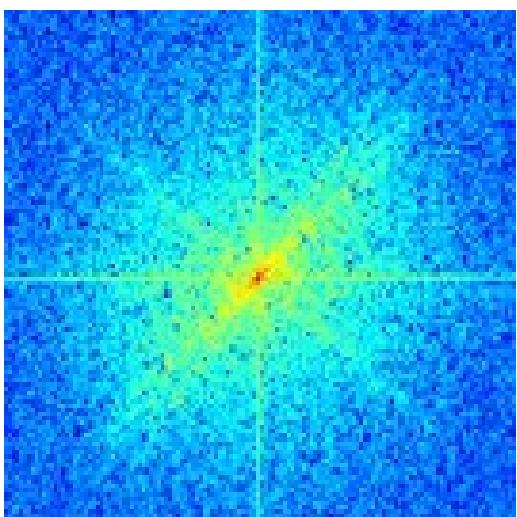
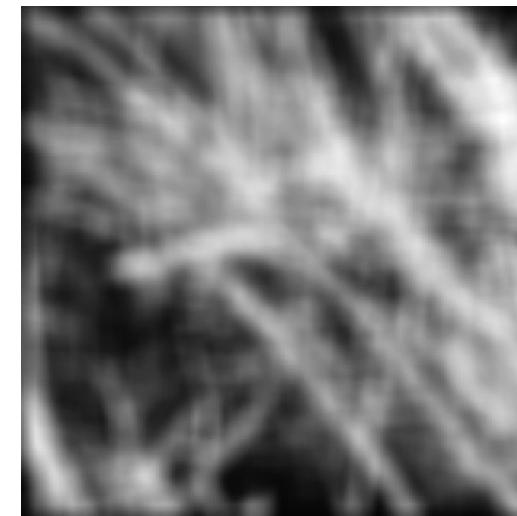
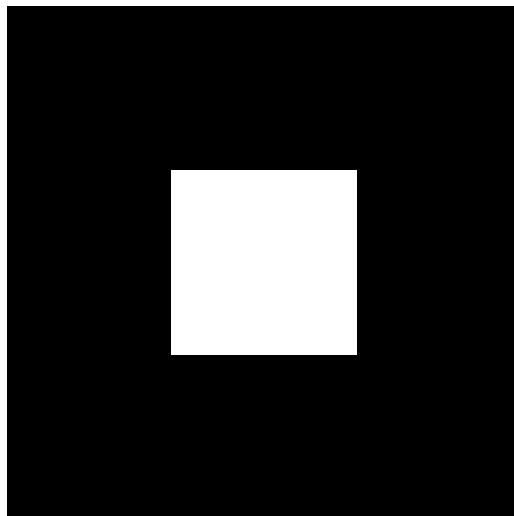
Box
filter



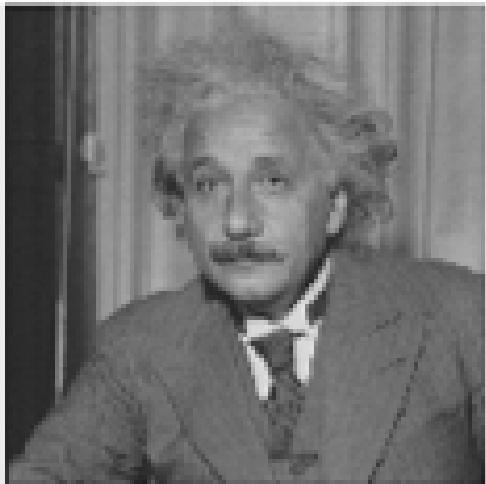
Gaussian blur



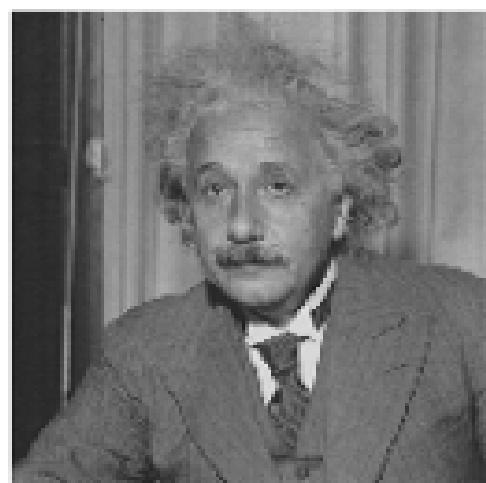
Box blur



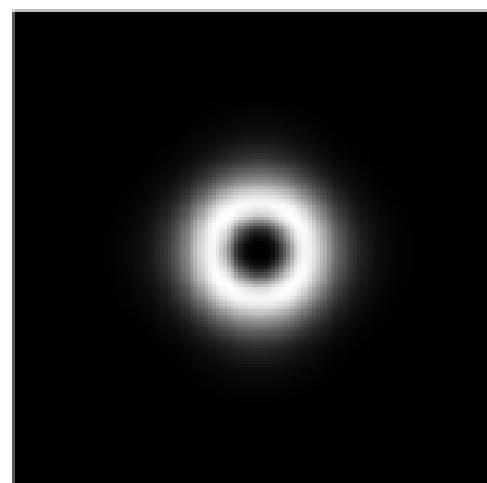
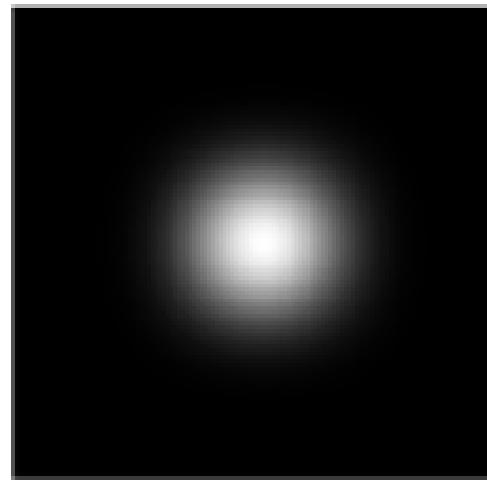
More filtering examples



?

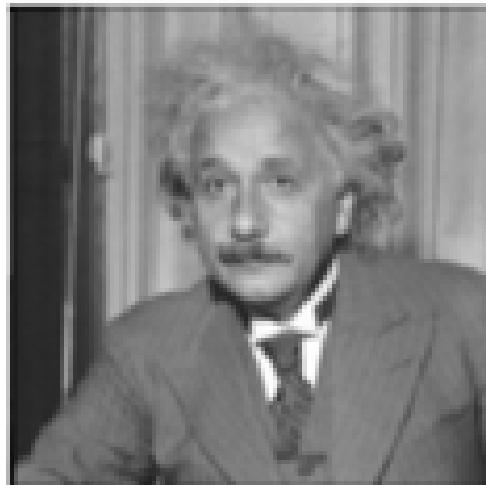
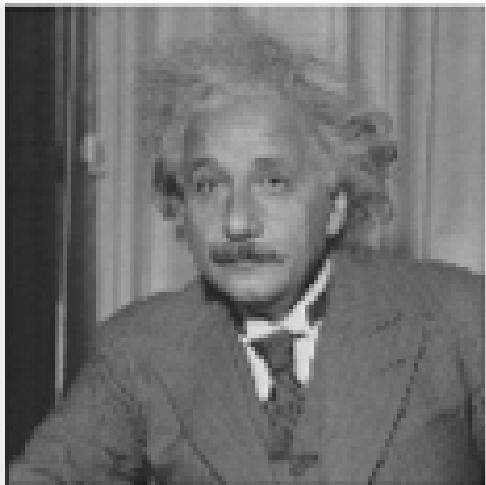


?

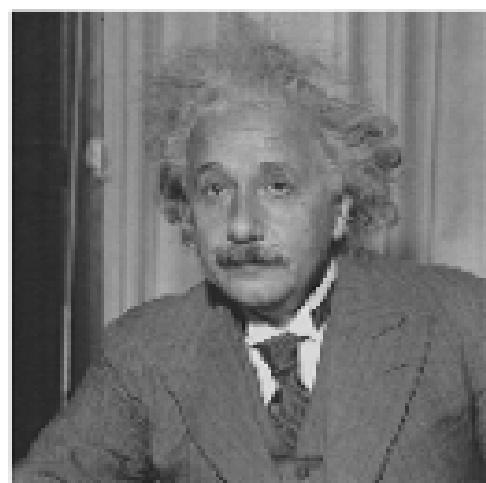
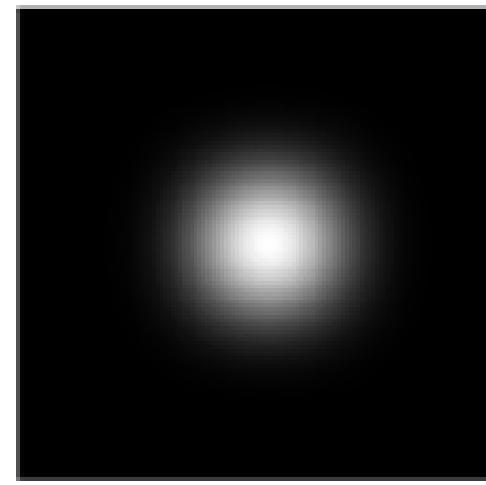


filters shown
in frequency-
domain

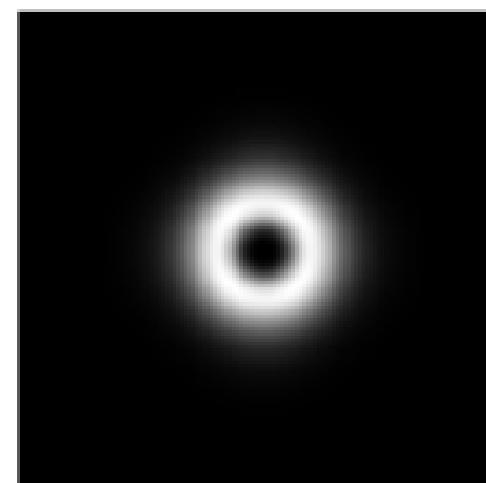
More filtering examples



low-pass



band-pass



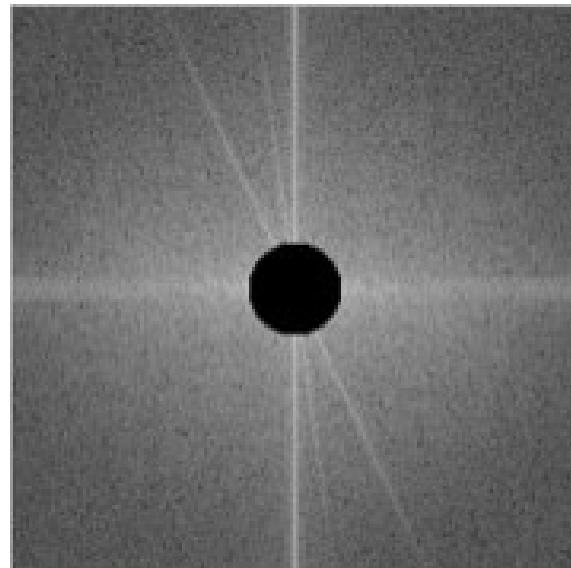
filters shown
in frequency-
domain

More filtering examples

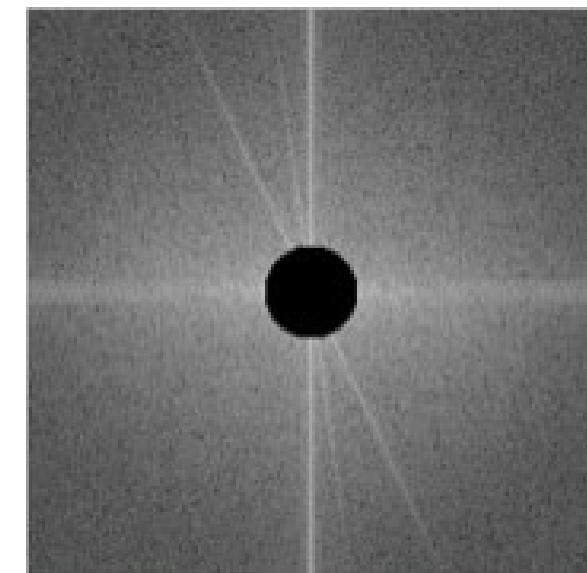
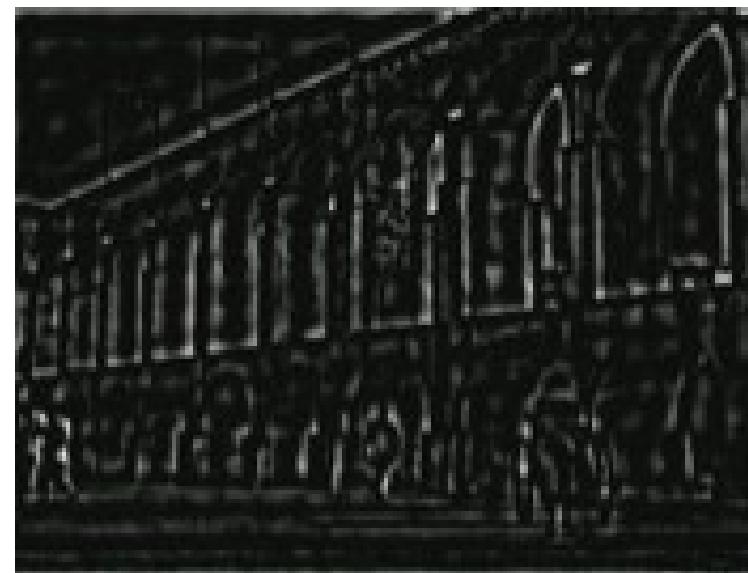


?

high-pass



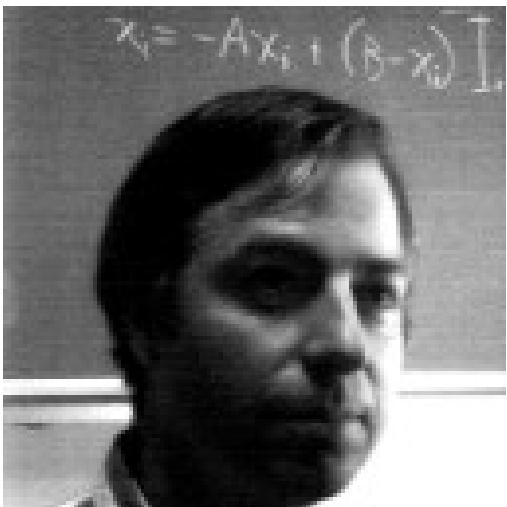
More filtering examples



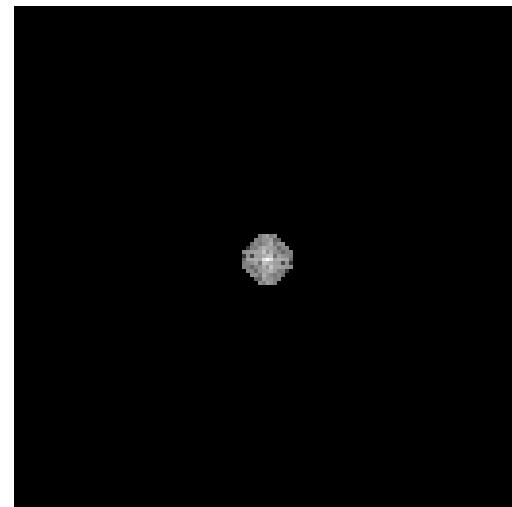
high-pass

More filtering examples

original image

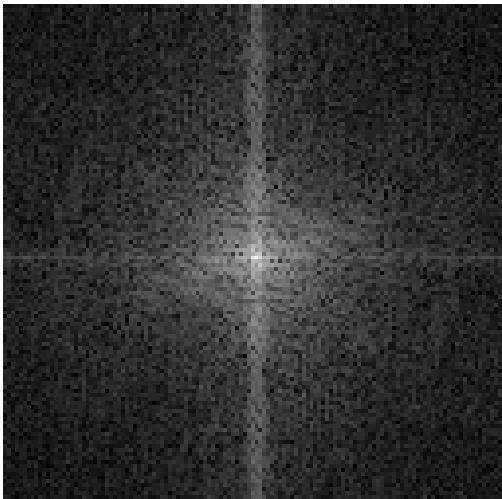


low-pass filter



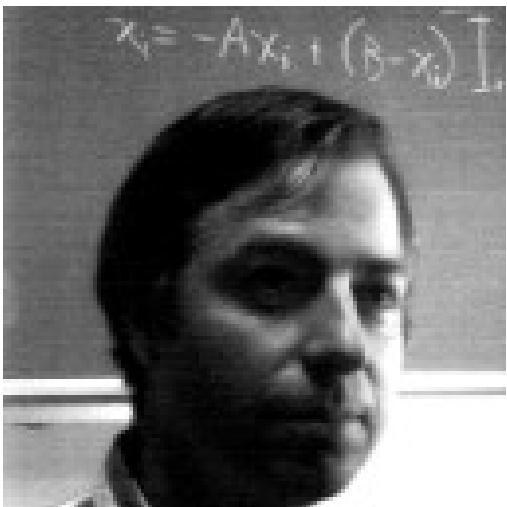
?

frequency magnitude

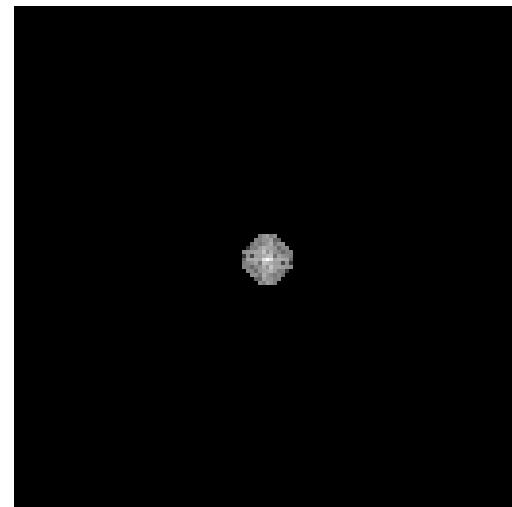


More filtering examples

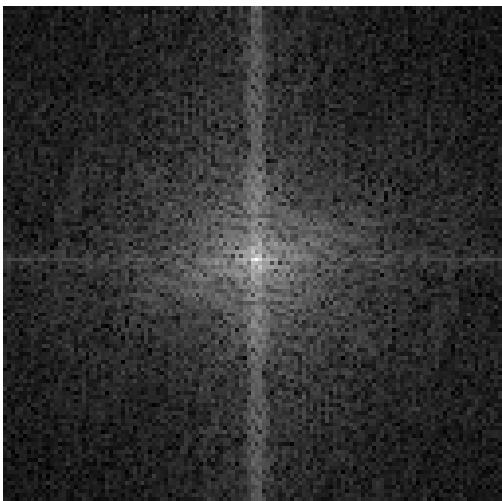
original image



low-pass filter

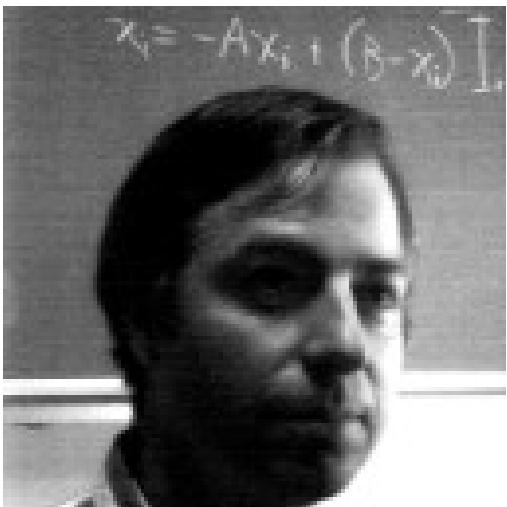


frequency magnitude

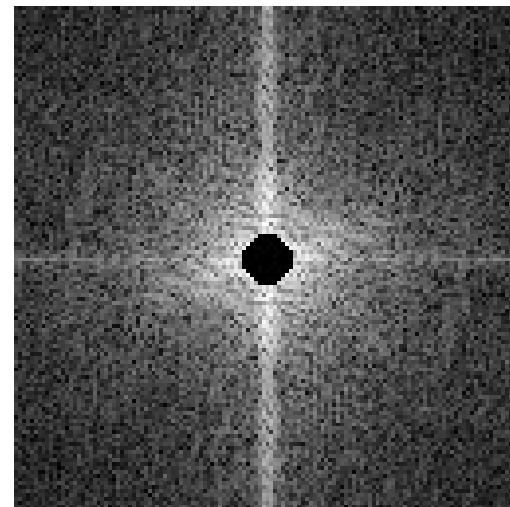


More filtering examples

original image

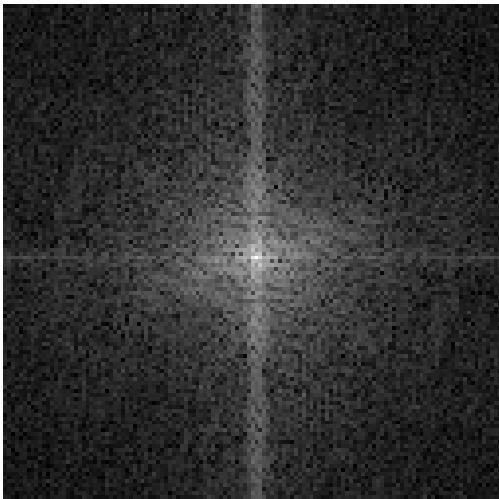


high-pass filter



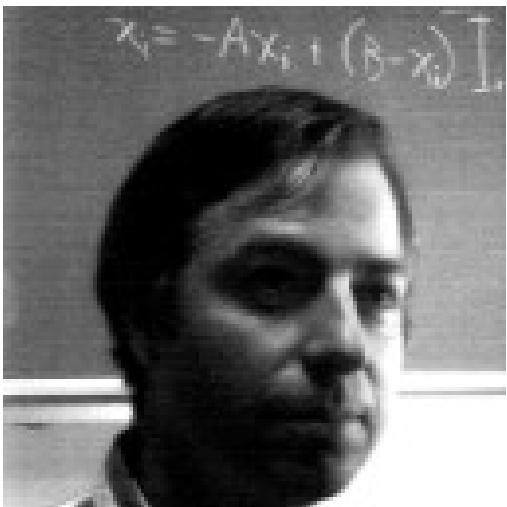
?

frequency magnitude

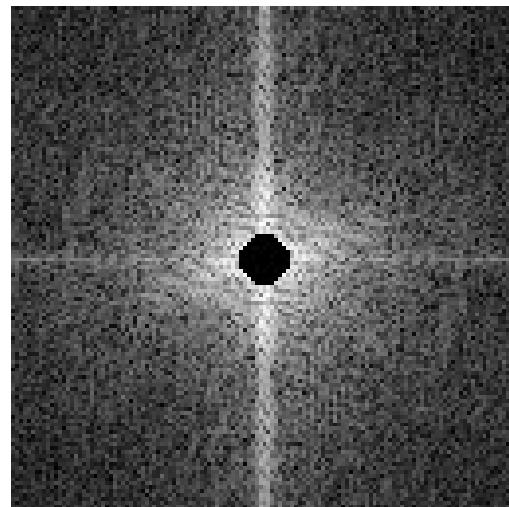


More filtering examples

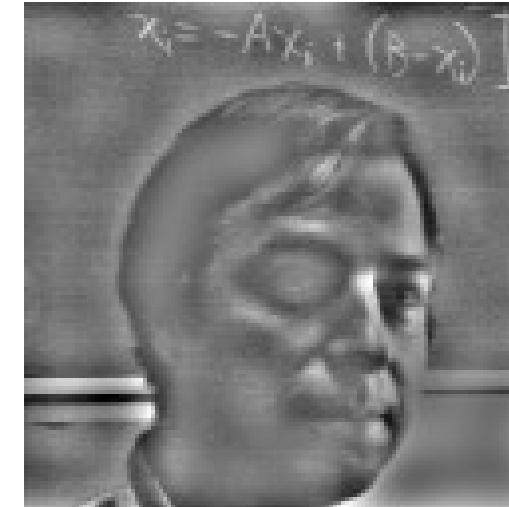
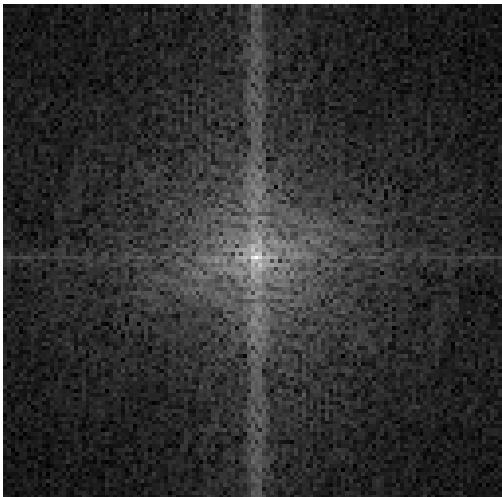
original image



high-pass filter

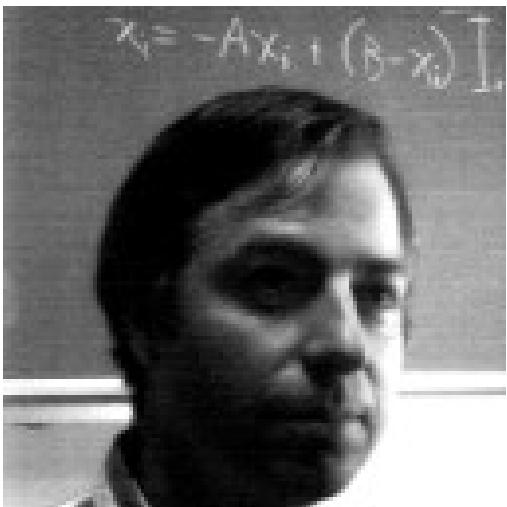


frequency magnitude

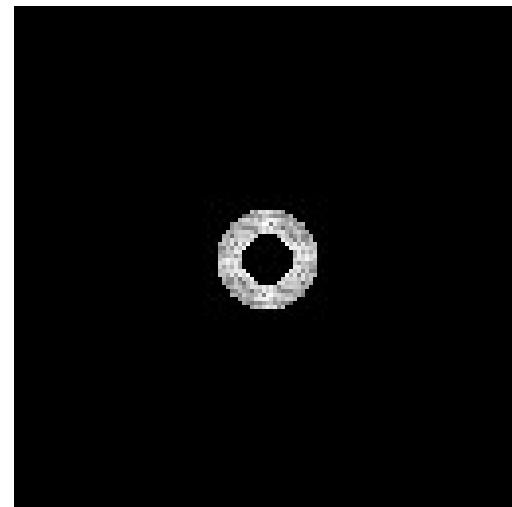


More filtering examples

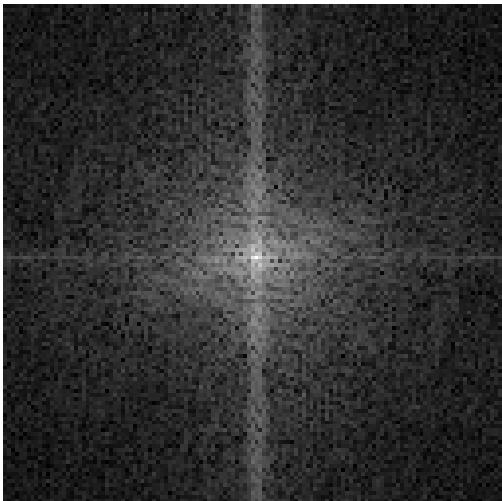
original image



mid-pass filter

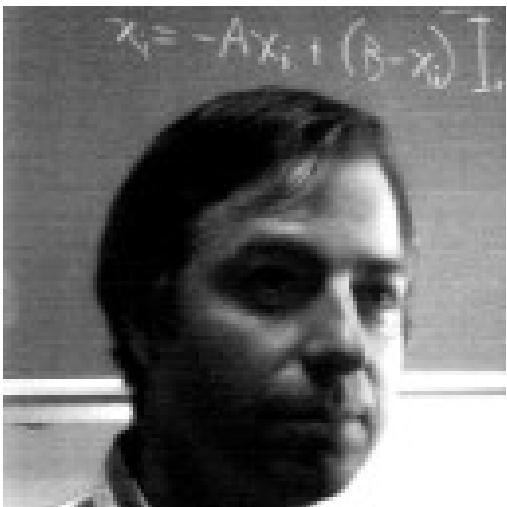


frequency magnitude

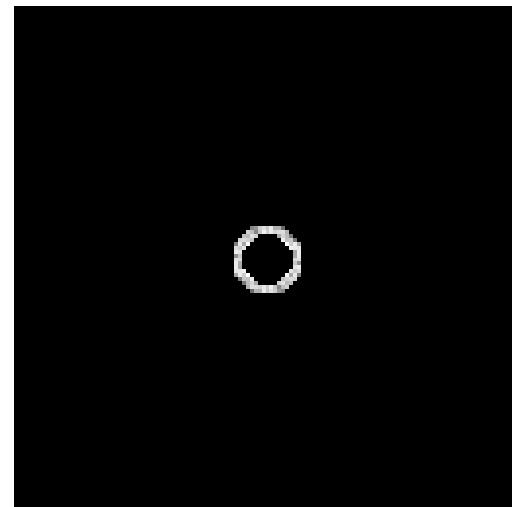


More filtering examples

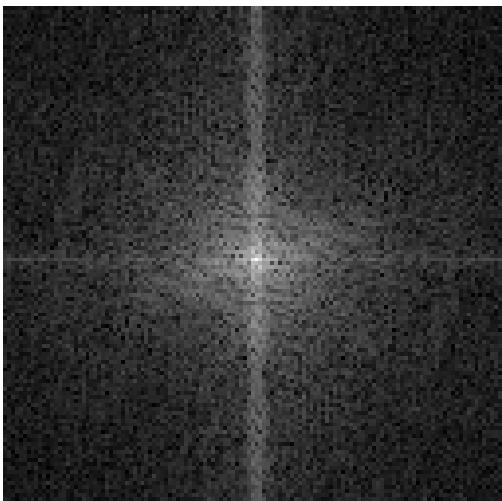
original image



mid-pass filter

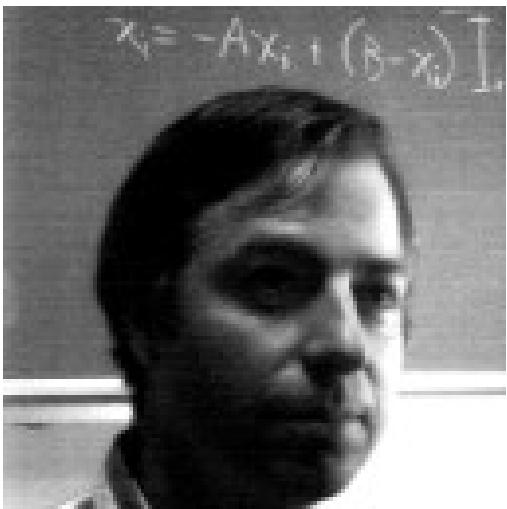


frequency magnitude

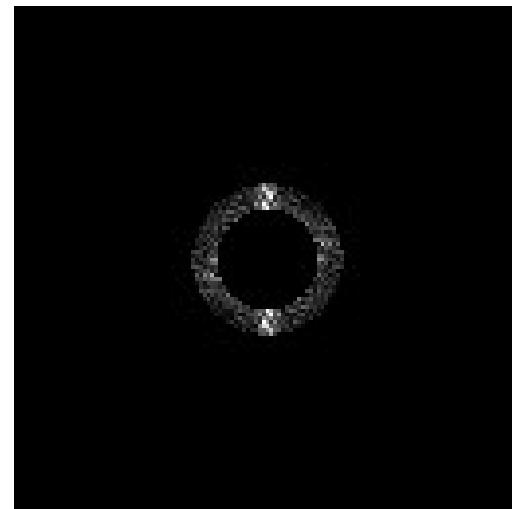


More filtering examples

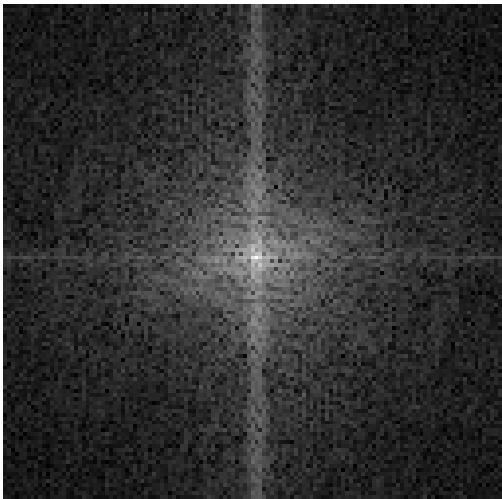
original image



mid-pass filter

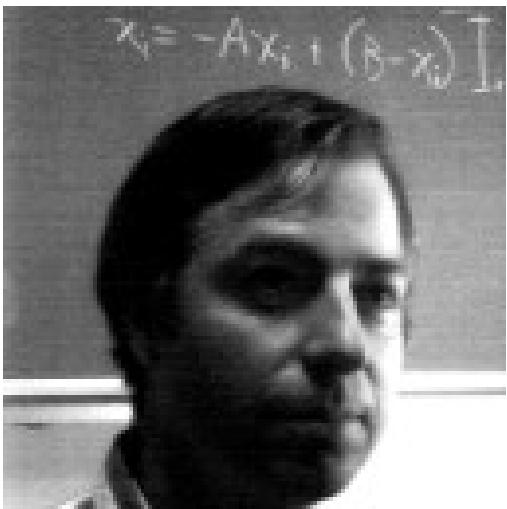


frequency magnitude

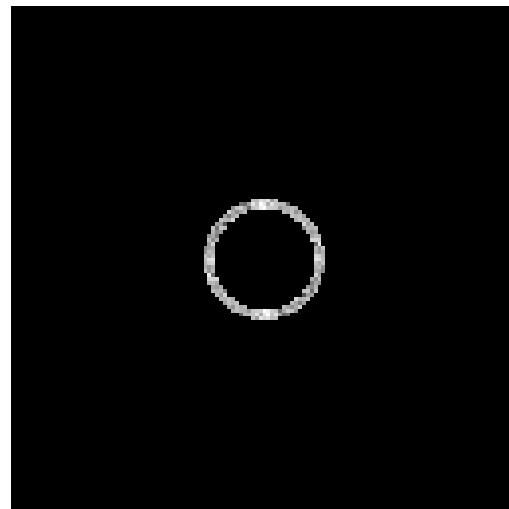


More filtering examples

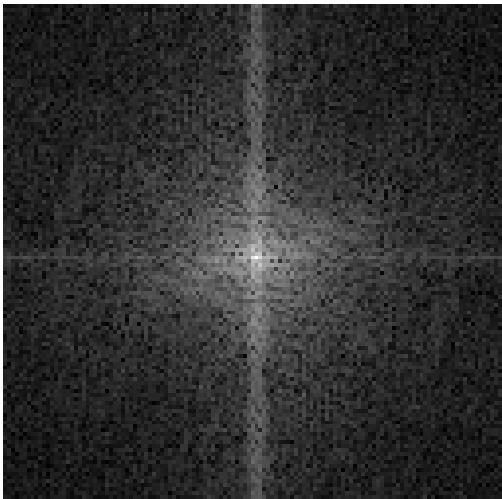
original image



mid-pass filter



frequency magnitude



Revisiting sampling

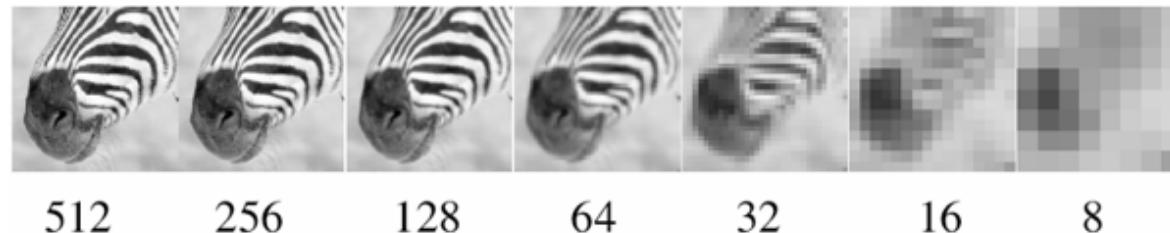
The Nyquist-Shannon sampling theorem

A continuous signal can be perfectly reconstructed from its discrete version using linear interpolation, if sampling occurred with frequency:

$$f_s \geq 2f_{\max} \quad \leftarrow \quad \text{This is called the Nyquist frequency}$$

Equivalent reformulation: When downsampling, aliasing does not occur if samples are taken at the Nyquist frequency or higher.

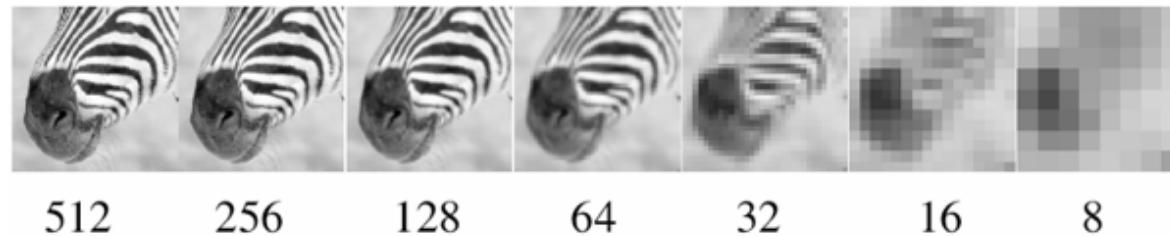
Gaussian pyramid



How does the Nyquist-Shannon theorem relate to the Gaussian pyramid?



Gaussian pyramid



How does the Nyquist-Shannon theorem relate to the Gaussian pyramid?

- Gaussian blurring is low-pass filtering.
- By blurring we try to sufficiently decrease the Nyquist frequency to avoid aliasing.

How large should the Gauss blur we use be?

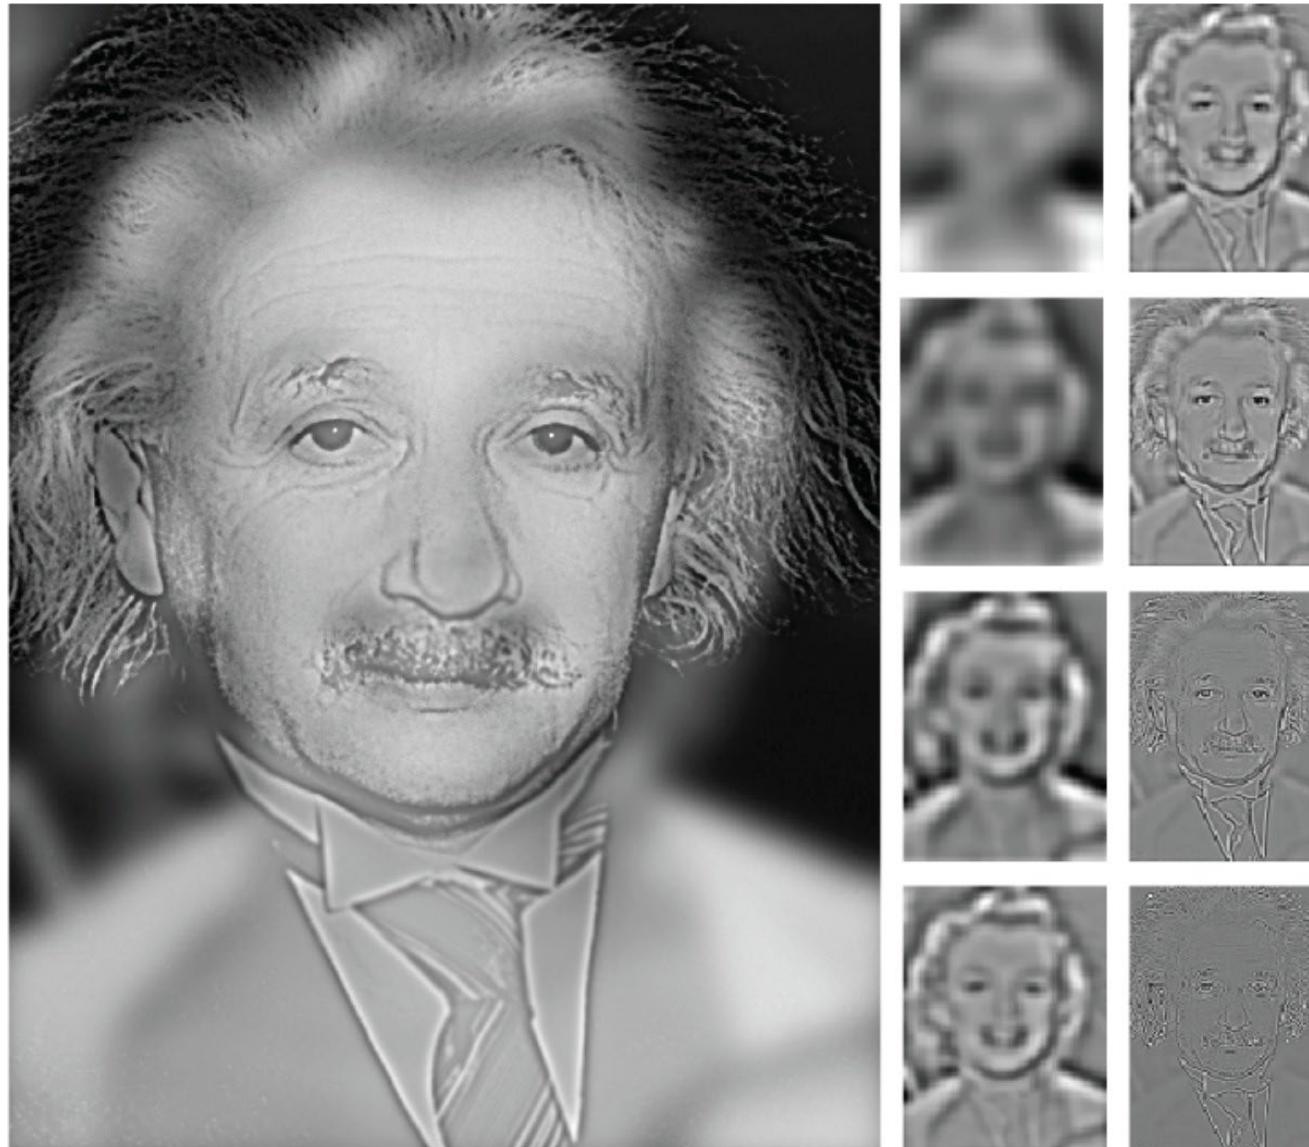
Frequency-domain filtering in human vision



“Hybrid image”

Aude Oliva and Philippe Schyns

Frequency-domain filtering in human vision



http://cvcl.mit.edu/hybrid_gallery/gallery.html